

## Talos Vulnerability Report

TALOS-2020-0996

### Videolabs libmicrodns 0.1.0 TXT record RDATA-parsing denial-of-service vulnerability

MARCH 23, 2020

#### CVE NUMBER

CVE-2020-6073

#### Summary

An exploitable denial-of-service vulnerability exists in the TXT record-parsing functionality of Videolabs libmicrodns 0.1.0. When parsing the RDATA section in a TXT record in mDNS messages, multiple integer overflows can be triggered, leading to a denial of service. An attacker can send an mDNS message to trigger this vulnerability.

#### Tested Versions

Videolabs libmicrodns 0.1.0

#### Product URLs

<https://github.com/videolabs/libmicrodns>

#### CVSSv3 Score

7.5 - CVSS:3.1/AV:N/AC:L/PR:N/UI:N/S:U/C:N/I:N/A:H

#### CWE

CWE-190: Integer Overflow or Wraparound

#### Details

The libmicrodns library is an mDNS resolver that aims to be simple and compatible cross-platform.

The function `mdns_rcv` reads and parses an mDNS message:

```
static int
mdns_rcv(const struct mdns_conn* conn, struct mdns_hdr *hdr, struct rr_entry **entries)
{
    uint8_t buf[MDNS_PKT_MAXSZ];
    size_t num_entry, n;
    ssize_t length;
    struct rr_entry *entry;

    *entries = NULL;
    if ((length = rcv(conn->sock, (char *) buf, sizeof(buf), 0)) < 0) // [1]
        return (MDNS_NETERR);

    const uint8_t *ptr = mdns_read_header(buf, length, hdr); // [2]
    n = length;

    num_entry = hdr->num_qn + hdr->num_ans_rr + hdr->num_add_rr;
    for (size_t i = 0; i < num_entry; ++i) {
        entry = calloc(1, sizeof(struct rr_entry));
        if (!entry)
            goto err;
        ptr = rr_read(ptr, &n, buf, entry, i >= hdr->num_qn); // [3]
        if (!ptr) {
            free(entry);
            errno = ENOSPC;
            goto err;
        }
        entry->next = *entries;
        *entries = entry;
    }
    ...
}
```

At [1], a message is read from the network. The 12-bytes mDNS header is then parsed at [2]. Based on the header info, the loop parses each resource record ("RR") using the function `rr_read` [3].

```

const uint8_t *
rr_read(const uint8_t *ptr, size_t *n, const uint8_t *root, struct rr_entry *entry, int8_t ans)
{
    size_t skip;
    const uint8_t *p;

    p = ptr = rr_read_RR(ptr, n, root, entry, ans);          // [4]
    if (ans == 0) return ptr;

    for (size_t i = 0; i < rr_num; ++i) {
        if (rrs[i].type == entry->type) {
            ptr = (*rrs[i].read)(ptr, n, root, entry);      // [5]
            if (!ptr)
                return (NULL);
            break;
        }
    }
    ...
}

```

The function `rr_read_RR` [4] reads the current resource record, except for the RDATA section. This is read by the loop at [5]. For each RR type, a different function is called. When the RR type is 0x10, the function `rr_read_TXT` is called at [5].

```

#define advance(x) ptr += x; *n -= x

static const uint8_t *
rr_read_TXT(const uint8_t *ptr, size_t *n, const uint8_t *root, struct rr_entry *entry)
{
    union rr_data *data = &entry->data;
    uint16_t len = entry->data_len;                          // [8]
    uint8_t l;

    if (*n == 0 || *n < len)
        return (NULL);

    for (; len > 0; len -= l + 1) {                          // [9]
        struct rr_data_txt *text;

        memcpy(&l, ptr, sizeof(l));                          // [6]
        advance(1);
        if (*n < l)
            return (NULL);
        text = malloc(sizeof(struct rr_data_txt));
        if (!text)
            return (NULL);
        text->next = data->TXT;
        data->TXT = text;
        if (l > 0)
            memcpy(text->txt, ptr, l);                        // [7]
        text->txt[l] = '\0';
        advance(l);                                          // [10]
    }
    return (ptr);
}

```

This function expects four parameters:

- `ptr`: the pointer to the start of the label to parse
- `n`: the number of remaining bytes in the message, starting from `ptr`
- `root`: the pointer to the start of the mDNS message
- `entry`: the entry struct, containing the parsed resource record

The function is supposed to extract each variable-length string from the RDATA section. In this case, it extracts a length in position 0 [6], and copies the data found in `text->txt` [7]. During this parsing, `*n` and `len` are decremented accordingly. In this loop, `len` tracks the number of characters left to read in the same RDATA section, as previously declared in the `data_len` field [8].

However, note that both `*n` and `len` are unsigned integers. This means that the loop will only stop when `len` is exactly equal to 0, or when `*n` is less than the length read in RDATA.

Also note that the `advance` macro is moving `ptr` forward and decrements `*n` (the number of bytes left in the packet) accordingly.

So, by making `l` at [6] equal to `*n`, and having at the same time `len` less than or equal to `l`, will cause `*n` to be 0 after [10], since we advance by `l`. Right after this, `len` overflows because `l` is bigger than `len`, making the loop itself cycle indefinitely. Then, a new `l` is read at [6], and `advance` is called again, making `*n` overflow. At this point both `*n` and `len` are overflowed and the program will eventually crash with an out-of-bounds read at [7] or [6].

#### Timeline

2020-01-30 - Vendor Disclosure

2020-03-20 - Vendor Patched

2020-03-23 - Public Release

#### CREDIT

Discovered by Claudio Bozzato of Cisco Talos.

