


- [First message in thread](#)
- [Zhiqiang Liu](#)
- [Coly Li](#)

Patch in this message

- [Get diff 1](#)

From Zhiqiang Liu <>
Subject [PATCH V2] bcache: fix potential deadlock problem in btree_gc_coalesce 
Date Sun, 26 Apr 2020 16:06:27 +0800

From: Zhiqiang Liu <liuzhiqiang26@huawei.com>

coccicheck reports:
drivers/md/bcache/btree.c:1538:1-7: preceding lock on line 1417

btree_gc_coalesce func is designed to coalesce two adjacent nodes in new_nodes[GC_MERGE_NODES] and finally release one node. All nodes'write_lock, new_nodes[i]->write_lock, are holded before coalescing adjacent nodes, and them will be released after coalescing successfully.

However, if the coalescing process fails, such as no enough space of new_nodes[1] to fit all of the remaining keys in new_nodes[0] and realloc keylist failed, we will goto to out_nocoalesce tag directly without releasing new_nodes[i]->write_lock. Then, a deadlock will occur after calling btree_node_free to free new_nodes[i], which also try to acquire new_nodes[i]->write_lock.

Here, we add a new tag 'out_unlock_nocoalesce' before out_nocoalesce tag to release new_nodes[i]->write_lock when coalescing process fails.

--
V1->V2: rewrite commit log (suggested by Coly Li) and rename the patch

Fixes: 2a285686c1 ("bcache: btree locking rework")
Signed-off-by: Zhiqiang Liu <liuzhiqiang26@huawei.com>

drivers/md/bcache/btree.c | 8 +++++--
1 file changed, 6 insertions(+), 2 deletions(-)

```
diff --git a/drivers/md/bcache/btree.c b/drivers/md/bcache/btree.c
index fa872df4e770..cad8b0b97e33 100644
--- a/drivers/md/bcache/btree.c
+++ b/drivers/md/bcache/btree.c
@@ -1447,7 +1447,7 @@ static int btree_gc_coalesce(struct btree *b, struct btree_op *op,
    if (__set_blocks(n1, n1->keys + n2->keys,
                    block_bytes(b->c) >
                    btree_blocks(new_nodes[i]))
        goto out_nocoalesce;
+       goto out_unlock_nocoalesce;
-
    keys = n2->keys;
    /* Take the key of the node we're getting rid of */
@@ -1476,7 +1476,7 @@ static int btree_gc_coalesce(struct btree *b, struct btree_op *op,
    if (__bch_keylist_realloc(&keylist,
                             bkey_u64s(&new_nodes[i]->key)))
        goto out_nocoalesce;
+       goto out_unlock_nocoalesce;
-
    bch_btree_node_write(new_nodes[i], &c1);
    bch_keylist_add(&keylist, &new_nodes[i]->key);
@@ -1522,6 +1522,10 @@ static int btree_gc_coalesce(struct btree *b, struct btree_op *op,
    /* Invalidated our iterator */
    return -EINTR;

+out_unlock_nocoalesce:
+    for(i = 0; i < nodes; i++)
+        mutex_unlock(&new_nodes[i]->write_lock);
+
    out_nocoalesce:
        closure_sync(&c1);

--
2.19.1
```

