

Talos Vulnerability Report

TALOS-2020-1133

Microsoft Azure Sphere Capability access control privilege escalation vulnerability

AUGUST 24, 2020

CVE NUMBER

None

SUMMARY

A privilege escalation vulnerability exists in the Capability access control functionality of Microsoft Azure Sphere 20.06. A set of specially crafted ptrace syscalls can be used to obtain elevated capabilities. An attacker can write a shellcode to trigger this vulnerability.

CONFIRMED VULNERABLE VERSIONS

The versions below were either tested or verified to be vulnerable by Talos or confirmed to be vulnerable by the vendor.

Microsoft Azure Sphere 20.06

PRODUCT URLS

Azure Sphere - <https://azure.microsoft.com/en-us/services/azure-sphere/>

CVSSV3 SCORE

8.1 - CVSS:3.0/AV:L/AC:H/PR:N/UI:N/S:C/C:H/I:H/A:H

CWE

CWE-284 - Improper Access Control

DETAILS

Microsoft's Azure Sphere is a platform for the development of internet-of-things applications. It features a custom SoC that consists of a set of cores that run both high-level and real-time applications, enforces security and manages encryption (among other functions). The high-level applications execute on a custom Linux-based OS, with several modifications to make it smaller and more secure, specifically for IoT applications.

In the Azure Sphere security model an extra layer of protection exists within the traditional Linux capabilities system via custom LSM, a separate Azure Sphere specific credential structure that looks like such:

```
// exposed through /proc/<pid>/attr/exec
struct azure_sphere_task_cred {
    union {
        u8      raw_bytes[16];
        struct azure_sphere_guid guid; // [1]
    } component_id;
    char   daa_tenant_id[64];           // [2]
    bool   is_app_man : 1;              // [3]
    bool   job_control_allowed : 1;
    unsigned int : 0;
    u32 capabilities;                  // [4]
};
```

This `azure_sphere_task_cred` contains the `component_id` of the process that is running [1], and also the `tenant_id` used for attestation [2]. More importantly though are the `is_app_man` field at [3] and `capabilities` field at [4]. Whenever a new application is run, the `application_manager` (Azure Sphere's `init`) will determine the correct UID to assign the process via the `uid_map` in `/mnt/config`. In examining the `application_manager` binary we also see that these Azure Sphere capabilities are also determined on application start:

```
// application_manager.bin
00022e76 05f23511 addw r1, r5, #0x135 {component_uid_???}
00022e7a f3f75bf9 bl #cmp_component_id
00022e7e 0028 cmp r0, #0 // d940e448-10b8-4b85-ac5f-69e6a6e4efc6
00022e80 4ed1 bne #0x22f20

00022e82 baf1090f cmp r10, #9
00022e86 18d1 bne #0x22eba

00022e88 05f24511 addw r1, r5, #0x145 {component_uid_1002}
00022e8c 3046 mov r0, r6 // 641f94d9-7600-4c5b-9955-5163cb7f1d75
00022e8e f3f751f9 bl #cmp_component_id
00022e92 0028 cmp r0, #0
00022e94 46d1 bne #0x22f24

00022e96 05f25511 addw r1, r5, #0x155 {component_uid_1003}
00022e9a 3046 mov r0, r6 // 48a22e96-d078-4e34-9d7a-91b3404031da
00022e9c f3f74af9 bl #cmp_component_id
00022ea0 0028 cmp r0, #0
00022ea2 42d1 bne #0x22f2a

00022ea4 05f26511 addw r1, r5, #0x165 {component_uid_1001}
00022ea8 3046 mov r0, r6 // 7ba05ff7-7835-4b26-9eda-29af0c635280
00022eaa f3f743f9 bl #cmp_component_id
00022eae 0028 cmp r0, #0
00022eb0 14bf ite ne
```

For each given hardcoded component_id in this check there is a corresponding binary in the official firmware that uses it:

```
1001 7ba05ff7-7835-4b26-9eda-29af0c635280 networkd
1002 641f94d9-7600-4c5b-9955-5163cb7f1d75 gatewayd
1003 48a22e96-d078-4e34-9d7a-91b3404031da azured
1004 a65f3686-e50a-4fff-b25d-415c206537af azcore
1005 89ecd022-0bdd-4767-a527-d756dd784a19 rng-tools
```

Continuing on from where the disassembly left off, we can see that the result of the component_id checking results in a value being moved into r0 that corresponds to the previously mentioned azure_sphere_task_cred->capabilities field:

```
00022eb2 4ff49070 mov r0, #0x120 // uid 1001
00022eb6 4ff48070 mov r0, #0x100 // uid 1004, uid 1005
[...]

00022f20 1020 movs r0, #0x10 // SFS tenant
00022f22 cae7 b #0x22eba

00022f24 40f20730 movw r0, #0x307 // uid 1002
00022f28 c7e7 b #0x22eba

00022f2a 40f24330 movw r0, #0x343 // uid 1003
00022f2e c4e7 b #0x22eba
```

If we convert a given UID's AZURE_SPHERE capabilities to readable form, we can see for instance that UID 1003 (azured) is granted the following capabilities:

```
AZURE_SPHERE_CAP_UPDATE_IMAGE
AZURE_SPHERE_CAP_QUERY_IMAGE_INFO
AZURE_SPHERE_CAP_POSTCODE
AZURE_SPHERE_CAP_RECORD_TELEMETRY
AZURE_SPHERE_CAP_MANAGE_AND_LOG_TELEMETRY
```

It's interesting to note that in order to see ones own processes' Azure sphere capabilities, the capget() syscall or /proc/self/status will not suffice, instead you must look at /proc/self/attr/current:

```
> cat /proc/self/attr/current
CID: AAAAAA-BBBB-CCCC-DDDD-123412341234
TID: 111111-1111-1111-4444-555555555555
CAPS: 00000000
```

This kernel file is defined in <kernel_root>/security/azure_sphere/lsm.c in azure_sphere_security_getprocatr. More interesting however is that this file also defines a azure_sphere_security_setprocatr function, one that permanently changes a given processes Azure Sphere capabilities:

```

static int azure_sphere_security_setprocattr(struct task_struct *p, char *name, void *value, size_t size) {
    struct azure_sphere_security_set_process_details *data = value;
    struct cred *cred; // comp_id[16], bool job_control_allowed, azure_sphere_capability_t capabilities
    struct azure_sphere_task_cred *tsec;
    int ret;

    // Can only set in binary format
    if (strcmp(name, "exec") != 0) {
        return -EINVAL;
    }

    if (value == NULL || size < sizeof(*data)) { // user controlled, max 0x1000
        return -EINVAL;
    }

    if (p != current) {
        // You can only change the current process
        return -EPERM;
    }

    cred = prepare_creds();
    if (!cred) {
        return -ENOMEM;
    }
    tsec = cred->security; // [1]

    //if no security entry then fail
    if (!tsec) {
        ret = -ENOENT;
        goto error;
    }

    if (!tsec->is_app_man) { // [2]
        ret = -EPERM;
        goto error;
    }

    memcpy(&tsec->component_id, data->component_id, sizeof(tsec->component_id));
    memset(&tsec->daa_tenant_id, 0, sizeof(tsec->daa_tenant_id));
    memcpy(&tsec->daa_tenant_id, data->daa_tenant_id, strlen(data->daa_tenant_id, sizeof(tsec->daa_tenant_id) - 1));
    tsec->is_app_man = false; // [3]
    tsec->job_control_allowed = data->job_control_allowed;
    tsec->capabilities = data->capabilities; // [4]

    return commit_creds(cred);
}

```

The main thing we care about in the above is that by writing to `/proc/self/attr/exec` we get our process' `azure_sphere_task_cred` struct at [1], verify that the `is_app_man` bool is still true [2] and then switch it to false [3], and finally assign the Azure Sphere specific capabilities at [4]. While simple in implementation, the consequences are quite substantial, as the `is_app_man` bool is never referenced anywhere else in the Azure Sphere kernel source code, leaving us with only one intended way to ever gain Azure Sphere capabilities, these capabilities being the sole way of gaining access to Security Manager and Pluton ioctls.

An interesting situation can occur however by utilizing TALOS-2020-1131, TALOS-2020-1132, and TALOS-2020-1137, in which the Azure Sphere device can be manipulated into running our installed application with a UID normally reserved for one of the system UIDs (e.g. 1003, `azured`). In such a situation, we do not have any of the Azure Sphere capabilities reserved for that UID since our application's `component_id` does not match up to any of the `component_id`'s in the `application_manager` disassembly from above.

As discussed in TALOS-2020-1137, we can actually see the `azured` process running via `procf`s:

```

> id
uid=1003 gid=1003 groups=5(gpio)
> ps aux | grep azure
PID   USER     TIME   COMMAND
 30 1003      0:00 /mnt/sys/azured/bin/azured --update --daemonize

```

Intuitively this makes sense, if we're running as the same UID and GID as another process we should be able to see that other process, but let's examine a bit more in depth why this is exactly:

```
//<kernel_root>/fs/proc/root.c

/* for the /proc/ directory itself, after non-process stuff has been done */
int proc_pid_readdir(struct file *file, struct dir_context *ctx)
{
    struct tgid_iter iter;
    struct pid_namespace *ns = file_inode(file)->i_sb->s_fs_info;
    loff_t pos = ctx->pos;

    // [...]

    iter.tgid = pos - TGID_OFFSET;
    iter.task = NULL;
    for (iter = next_tgid(ns, iter); // [1]
         iter.task;
         iter.tgid += 1, iter = next_tgid(ns, iter)) {
        char name[PROC_NUMBUF];
        int len;

        cond_resched();
        if (!has_pid_permissions(ns, iter.task, 2)) //[2]
            continue;

        len = snprintf(name, sizeof(name), "%d", iter.tgid);
        ctx->pos = iter.tgid + TGID_OFFSET;
        if (!proc_fill_cache(file, ctx, name, len,
                           proc_pid_instantiate, iter.task, NULL)) {
            put_task_struct(iter.task);
            return 0;
        }
    }
    ctx->pos = PID_MAX_LIMIT + TGID_OFFSET;
    return 0;
}
```

The above code is run when we do `ls` on `/proc`, as Linux traditionally lists directories corresponding to the different processes currently running. At [1], we see this iteration occurring over each “Thread Group ID”, and if our current process passes the `has_pid_permission()` check, then we get to see that process’ directory. Continuing into `has_pid_permission()`:

```
/*
 * May current process learn task's sched/cmdline info (for hide_pid_min=1)
 * or euid/egid (for hide_pid_min=2)?
 */
static bool has_pid_permissions(struct pid_namespace *pid,
                               struct task_struct *task,
                               int hide_pid_min)
{
    if (pid->hide_pid < hide_pid_min)
        return true;
    if (in_group_p(pid->pid_gid))
        return true;
    return ptrace_may_access(task, PTTRACE_MODE_READ_FSCREDS); // [1]
}
```

For our purposes we only care about the line at [1], which implies that if we can see the `/proc` directory for a given process then we have a baseline amount of `PTRACE` permissions. Continuing into `__ptrace_may_access()`:

```
/* Returns 0 on success, -errno on denial. */
static int __ptrace_may_access(struct task_struct *task, unsigned int mode)
{
    const struct cred *cred = current_cred(), *tcred;
    struct mm_struct *mm;
    kuid_t caller_uid;
    kgid_t caller_gid;

    [...]
    /* Don't let security modules deny introspection */
    if (same_thread_group(task, current)) // task->signal == current->signal
        return 0;
    rcu_read_lock();
    if (mode & PTTRACE_MODE_FSCREDS) {
        caller_uid = cred->fsuid;
        caller_gid = cred->fsgid;

    [...]
        tcred = __task_cred(task);
        if (uid_eq(caller_uid, tcred->euid) &&
            uid_eq(caller_uid, tcred->suid) &&
            uid_eq(caller_uid, tcred->uid) &&
            gid_eq(caller_gid, tcred->egid) &&
            gid_eq(caller_gid, tcred->sgid) &&
            gid_eq(caller_gid, tcred->gid))
            goto ok;
}
```

We don't really have to go too in depth into `__ptrace_may_access`, the only things it seems to care about are that the `resuid()` and `resgid()` of the accessing process matches up with the process being targeted. Since we can already see the `azured` process directory in `/proc` we know by definition that we pass the `__ptrace_may_access()` check, and this fact might lead one to guess where else `__ptrace_may_access()` is used. In the event that astute readers guessed `ptrace()`, then they would be right, and this brings us to the crux of the vulnerability:

```

// <kernel_root>/security/commoncap.c
/**
 * cap_ptrace_access_check - Determine whether the current process may access
 * another
 * @child: The process to be accessed
 * @mode: The mode of attachment.
 *
 * If we are in the same or an ancestor user_ns and have all the target
 * task's capabilities, then ptrace access is allowed.
 * If we have the ptrace capability to the target user_ns, then ptrace
 * access is allowed.
 * Else denied.
 *
 * Determine whether a process may access another, returning 0 if permission
 * granted, -ve if denied.
 */
int cap_ptrace_access_check(struct task_struct *child, unsigned int mode)
{
    int ret = 0;
    const struct cred *cred, *child_cred;
    const kernel_cap_t *caller_caps;

    rcu_read_lock();
    cred = current_cred();
    child_cred = __task_cred(child);
    if (mode & PTRACE_MODE_FSCREDS)
        caller_caps = &cred->cap_effective;
    else
        caller_caps = &cred->cap_permitted;
    if (cred->user_ns == child_cred->user_ns &&
        cap_issubset(child_cred->cap_permitted, *caller_caps)) // [1]
        goto out;
    if (ns_capable(child_cred->user_ns, CAP_SYS_PTRACE))
        goto out;
    ret = -EPERM;
out:
    rcu_read_unlock();
    return ret;
}

```

With standard Linux capabilities there is an explicit check at [1] to make sure that the attaching process' capabilities are a subset of the target's capabilities, and if this condition is not met, the attaching process must have CAP_SYS_PTRACE in order to attach. On the Azure Sphere device we can never get this CAP_SYS_PTRACE capability, however there is no explicit check that verifies an attaching process' Azure Sphere capability is a subset of the target's Azure Sphere capabilities. If they had been implemented in the same field of the standard Linux credential structure, this would not be an issue, but as pointed out in the beginning of the writeup, the Azure Sphere specific capabilities are located cred->security->capabilities. And this fact allows us to attach to another process with more Azure Sphere capabilities (e.g. any other app running), as long as we share the UID with that process (which can be done by chaining TALOS-2020-1131, TALOS-2020-1132, and TALOS-2020-1137):

```

> /mnt/apps/11223344-1234-1234-1234-aabbccddeeff/bin/busybox id
uid=1001 gid=1001 groups=12(sys-log),500(ipc-appman-reboot),501(ipc-appman-sideload),502(ipc-appman-cloudload),503(ipc-appman-apptart)

> cat /proc/self/attr/current
CID: 48A22E96-D078-4E34-9D7A-91B3404031DA
TID: [...]
CAPS: 00000343

```

Thus in summary: an attacker can use the ptrace() API to gain execution in another Azure Sphere process and use its Azure Sphere capabilities to access an entirely new set of IOCTL requests for /dev/pluton and all the /dev/security-monitor ioctls that likewise could not be utilized before. As a sample of this, in the Security Monitor's SMAPI alone for capability 0x343:

```

// 0x343:
// AZURE_SPHERE_CAP_UPDATE_IMAGE
// AZURE_SPHERE_CAP_QUERY_IMAGE_INFO
// AZURE_SPHERE_CAP_POSTCODE
// AZURE_SPHERE_CAP_RECORD_TELEMETRY
// AZURE_SPHERE_CAP_MANAGE_AND_LOG_TELEMETRY

static azure_sphere_sm_cap_lookup_t azure_sphere_sm_cmd_required_capabilities[] = {

    { .cmd = AZURE_SPHERE_SMAPI_GET_APPLICATION_IMAGE_COUNT, .caps = AZURE_SPHERE_CAP_QUERY_IMAGE_INFO },
    { .cmd = AZURE_SPHERE_SMAPI_LIST_ALL_APPLICATION_IMAGES, .caps = AZURE_SPHERE_CAP_QUERY_IMAGE_INFO },
    { .cmd = AZURE_SPHERE_SMAPI_GET_COMPONENT_IMAGES, .caps = AZURE_SPHERE_CAP_QUERY_IMAGE_INFO },
    { .cmd = AZURE_SPHERE_SMAPI_GET_ABI_TYPE_COUNT, .caps = AZURE_SPHERE_CAP_QUERY_IMAGE_INFO },
    { .cmd = AZURE_SPHERE_SMAPI_GET_ABI_VERSIONS, .caps = AZURE_SPHERE_CAP_QUERY_IMAGE_INFO },
    { .cmd = AZURE_SPHERE_SMAPI_GET_UPDATE_CERT_STORE_IMAGE_INFO, .caps = AZURE_SPHERE_CAP_QUERY_IMAGE_INFO },
    // [...]
    { .cmd = AZURE_SPHERE_SMAPI_SHOULD_IMAGE_BE_UPDATED, .caps = AZURE_SPHERE_CAP_UPDATE_IMAGE },
    { .cmd = AZURE_SPHERE_SMAPI_INVALIDATE_IMAGE, .caps = AZURE_SPHERE_CAP_UPDATE_IMAGE },
    { .cmd = AZURE_SPHERE_SMAPI_OPEN_IMAGE_FOR_STAGING, .caps = AZURE_SPHERE_CAP_UPDATE_IMAGE },
    { .cmd = AZURE_SPHERE_SMAPI_WRITE_BLOCK_TO_STAGE_IMAGE, .caps = AZURE_SPHERE_CAP_UPDATE_IMAGE },
    { .cmd = AZURE_SPHERE_SMAPI_COMMIT_IMAGE_STAGING, .caps = AZURE_SPHERE_CAP_UPDATE_IMAGE },
    { .cmd = AZURE_SPHERE_SMAPI_ABORT_IMAGE_STAGING, .caps = AZURE_SPHERE_CAP_UPDATE_IMAGE },
    { .cmd = AZURE_SPHERE_SMAPI_INSTALL_STAGED_IMAGES, .caps = AZURE_SPHERE_CAP_UPDATE_IMAGE },
    { .cmd = AZURE_SPHERE_SMAPI_GET_COMPONENT_COUNT, .caps = AZURE_SPHERE_CAP_UPDATE_IMAGE },
    { .cmd = AZURE_SPHERE_SMAPI_GET_COMPONENT_SUMMARY, .caps = AZURE_SPHERE_CAP_UPDATE_IMAGE },
    { .cmd = AZURE_SPHERE_SMAPI_STAGE_COMPONENT_MANIFESTS, .caps = AZURE_SPHERE_CAP_UPDATE_IMAGE },
    { .cmd = AZURE_SPHERE_SMAPI_COUNT_OF_MISSING_IMAGES_TO_DOWNLOAD, .caps = AZURE_SPHERE_CAP_UPDATE_IMAGE },
    { .cmd = AZURE_SPHERE_SMAPI_GET_MISSING_IMAGES_TO_DOWNLOAD, .caps = AZURE_SPHERE_CAP_UPDATE_IMAGE },
    { .cmd = AZURE_SPHERE_SMAPI_STAGE_BASE_MANIFESTS, .caps = AZURE_SPHERE_CAP_UPDATE_IMAGE },
    { .cmd = AZURE_SPHERE_SMAPI_COUNT_OF_MISSING_BASE_IMAGES_TO_DOWNLOAD, .caps = AZURE_SPHERE_CAP_UPDATE_IMAGE },
    { .cmd = AZURE_SPHERE_SMAPI_GET_MISSING_BASE_IMAGES_TO_DOWNLOAD, .caps = AZURE_SPHERE_CAP_UPDATE_IMAGE },
    { .cmd = AZURE_SPHERE_SMAPI_GET_SOFTWARE_ROLLBACK_INFO, .caps = AZURE_SPHERE_CAP_UPDATE_IMAGE },
    // [...]
    { .cmd = AZURE_SPHERE_SMAPI_PERIPHERAL_ACQUIRE, .caps = AZURE_SPHERE_CAP_PERIPHERAL_PIN_MAPPING },
    { .cmd = AZURE_SPHERE_SMAPI_PERIPHERAL_RELEASE, .caps = AZURE_SPHERE_CAP_PERIPHERAL_PIN_MAPPING },
    { .cmd = AZURE_SPHERE_SMAPI_PERIPHERAL_GET_AVAILABLE_DOMAINS, .caps = AZURE_SPHERE_CAP_PERIPHERAL_PIN_MAPPING },
    { .cmd = AZURE_SPHERE_SMAPI_PERIPHERAL_LOCK_CONFIG, .caps = AZURE_SPHERE_CAP_PERIPHERAL_PIN_MAPPING },
};

```

TIMELINE

2020-07-31 - Vendor Disclosure

2020-08-24 - Public Release

CREDIT

Discovered by Lilith & Claudio Bozzato and Dave McDaniel of Cisco Talos.

VULNERABILITY REPORTS

PREVIOUS REPORT

NEXT REPORT

TALOS-2020-1128

TALOS-2020-1137
