

CVE-2020-24175: yz1: stack overflow

Feb 22, 2021

Introduction

I recently ran into an old-school buffer overflow while fuzzing the [Yz1](#) archive (de)compression library.

The intention with this write-up is to go from a crash to code execution in one of the archive software that bundles the Yz1 library – namely, [IZArc](#) version 4.4. The target platform is Windows 10 64bit (although both IZArc and Yz1 are 32bit-only).

The analysis is made with [Ghidra](#) coupled with the [PHAROS OoAnalyzer](#) plugin.

The image base for Yz1.dll in our analysis is 0x10000000 and the version is 0.30 (as is shipped with IZArc, but newer versions of Yz1 are also vulnerable).

The Yz1 library

Yz1 is an archaic compression format developed by YAMAZAKI at Binary Technology. It was part of their [DeepFreezer](#) archiver software.

Both of these components are closed-source and proprietary. However, Yz1 is distributed as a shareware binary-only DLL and it's bundled with a few modern file archivers – including [IZArc](#), [ZipGenius](#) and [Explzh](#).

The interface for Yz1 is somewhat interesting. There are a few standalone functions that tries to verify that an archive is valid. There are also functions for retrieving filenames and their metadata in an archive. File (de)compression is performed by a single public function named [Yz1](#):

```
int WINAPI Yz1(const HWND wnd, LPCSTR cmd, LPSTR buf, const DWORD siz);
```

Every argument other than `cmd` can be `NULL` or `0` for window-less use where no feedback is to be received from the module itself.

The `cmd` argument specifies what operation to perform and with what options. Some of these options include:

```
c      - Create archive
x      - Expand Archive
-cN    - Check timestamp according to N
-iN    - Silence status output according to N
```

This makes working with the Yz1 API kind of like working with the command-line interfaces for traditional (un)archivers like `tar` or `zip`.

As mentioned, the focus in this write-up is on [IZArc](#) and its use of the Yz1 library. The functionality in Yz1 that we'll pay attention to is the functionality that's used by IZArc.

The Yz1 header

The (for this write up) first relevant entrypoint, before the `Yz1()` function is reached, is `Yz1CheckArchive()`. IZArc uses this function to validate Yz1 archives before processing them.

The prototype looks like [this](#):

```
BOOL WINAPI Yz1CheckArchive(LPCSTR filename, const int mode);
```

The first argument is the filename of the archive to check. The second argument is the mode to check. There are a number of checking modes defined in [Yz1.h](#):

```
CHECKARCHIVE_BASIC 1
[...]
CHECKARCHIVE_ALL 16
```

`Yz1CheckArchive` returns `true` or `false` for all modes except `CHECKARCHIVE_ALL`. A mode of `CHECKARCHIVE_ALL` introduces other possible return values, despite the function signature.

Our target program, IZArc, seem to always invoke `Yz1CheckArchive` with a mode of `CHECKARCHIVE_BASIC`, so the other modes are ignored.

The `Yz1CheckArchive()` function, as well as the generic `Yz()` function (during decompression), takes us to a class method with the following signature:

```
int __thiscall YzFile_DecodeHeader(yzFileDecode *this, char *x_path);
```

This method is far too complex to distill in its entirety, but it performs a number of noteworthy operations. First off, it starts by reading a 0x14 byte header from the input file:

```
/*
 * 1000e7b7
 */
x_size = _fread(&x_header,1,0x14,x_yz1File->fp);
```

For example, with an archive containing the following three files:

```
>>> from pathlib import Path
>>> for p in Path().glob("*.txt"):
...     print(f"{p}: {p.stat().st_size:#x} bytes")
...
aaaa.txt: 0x25 bytes
bbbbbbbbbbbbbbbb.txt: 0x1d bytes
cccccccccccccccc.txt: 0x21 bytes
```

The header will look something like this:

```
$ hexdump -e '4/1 "%02X" "\n"' demo.yz1
# 0: Archive magic (yz01)
797A3031
# 1: Flags; used to, e.g., indicate whether the archive is password-protected.
30363030
# 2: ???
000000B2
# 3: Number of bytes required to decode the filenames.
#
# > file_count * sizeof(DWORD) * 2 * len(all_filenames_incl_NULL)
```

Both of these functions invoke a method starting at 0x10002de0 that enforce the FNAME_MAX32 (512/0x200) byte limit (sorry for the lack of cleanup!):

```
[...]
/*
 * LAB_10002f17
 */
if (*(uint *) (* (int *) (iVar3 + 0x10) + 0x14 + uVar2 * 0x1c) < 0x200) {
    iVar3 = x getPathInstance((cls 10002bc0 *)
        (*(int *) (this->mbr_34 + 4) + 0xc), this->mbr_48);
    /*
     * NOTE: This is not important right now, but it will matter during
     * exploitation. Filenames shorter than 0x10 bytes are stored
     * inline at iVar3 + 4. Filenames GTE 0x10 are allocated a separate
     * buffer whose address is stored at iVar3 + 4.
     */
    if (*(uint *) (iVar3 + 0x18) < 0x10) {
        x_filenameSrc = (char *) (iVar3 + 4);
    }
    else {
        x_filenameSrc = *(char **) (iVar3 + 4);
    }
    x_filenameDst = x_dst->szFileName;
    do {
        x_chr = *x_filenameSrc;
        *x_filenameDst = x_chr;
        x_filenameSrc = x_filenameSrc + 1;
        x_filenameDst = x_filenameDst + 1;
    } while (x_chr != '\0');
}
else {
    /*
     * x_dst->szFileName = "too_long_file_name\0"
     */
    *(undefined4 *) x_dst->szFileName = 0x5f6f6f74; /* _oot */
    *(undefined4 *) (x_dst->szFileName + 4) = 0x676e6f6c; /* _gnol */
    *(undefined4 *) (x_dst->szFileName + 8) = 0x6c69666f; /* _iff */
    *(undefined4 *) (x_dst->szFileName + 0xc) = 0x616e5f65; /* _an_e */
    *(undefined2 *) (x_dst->szFileName + 0x10) = 0x656d; /* _em */
    x_dst->szFileName[0x12] = '\0';
}
[...]
```

A stack-based buffer overflow

Not all code paths pay attention to the recorded lengths of the filenames. The one my fuzzer ran into is a function that starts at 0x10005080. It `printf(..., "expanding %s", ...)` with the file currently being extracted for a logging message.

It's kind of interesting too, because – similar to the snippet above – the call to `printf()` also checks whether the filename is inline (that is, if its length is below 0x10). But it doesn't check that the filename is below `FNAME_MAX32`.

```
/*
 * 100055b5
 */
if (this_00->mbr_18 < 0x10) {
    pDVar6 = &this_00->mbr_4;
}
else {
    pDVar6 = (DWORD *) this_00->mbr_4;
}
_printf(&local_264, "expanding %s", pDVar6)
```

Mo' bugs mo' problems

In working to exploit the fuzzed bug in the last section, I ran into a situation where we had written `N` bytes on the stack before the first `[RJC]OP` gadget. However, after the first gadget we could only write a handful of subsequent gadgets. Otherwise, we'd run into another bug earlier in the extraction process.

Similar to the previous flaw, this flaw is caused by a stack overflow. It happens in `yzFileDecode::DecodeFile`. Ghidra produces a somewhat wonky decompilation of this method, so the following C-ish code has been rewritten for clarity (at the expense of not being an accurate representation of its disassembly – although the important locations are commented):

```
/*
 * 1000eec0
 */
int yzFileDecode::DecodeFile(char *param_1, int *param_2)
{
    int rc;
    int duplicateCount = 0;
    unsigned int i = 0;
    char buf[XXX];

    [...];

    /*
     * LAB_1000efa0
     */
    do {
        if (this->x_yzDecHead->filenames == NULL) {
            [...];
        }

        /*
         * 1000f170
         */
        sprintf(buf, "%s%s", this->x_dirname,
            this->x_yzDecHead->x_filename[4 + i * 0x1c]);

        rc = x_hasFile(buf);
        if (rc) {
            duplicateCount += 1;
        }
    } while (i < this->x_yzDecHead->x_fileCount); /* 1000f02e */

    [...];

    /*
     * 1000f036
     */
    if (duplicateCount > 0) {
        x_overwriteWarning();
    }

    [...];
}
```

In the snippet above, each filename in the archive is checked for existence on disk. If it already exist, a warning message *may be* presented to the user (`yz1` only shows GUI messages if it's been given a `HWND`).

As with the previous bug, the call to `printf()` is unchecked. If a path in 1000f170 is large enough, we'll overflow the stack.

However, one major issue with this flaw is that the call to `printf()` at 1000f170 prepends the extraction directory to the filename. This complicates exploitation. It also makes it difficult to exploit the first bug mentioned in this writeup because this bug could be triggered earlier in the execution if the user chooses a long extraction directory.

With that in mind, one positive aspect of this bug is that we can overwrite `this->x_yzDecHead` and cause an invalid memory access in the `do-while()` conditional. This leads to quick control of execution if we overwrite a SEH. So this is the bug that's exploited in the PoC.

Splittin' like its the 00s

There are two important aspects of the decoding process of the archive header and its filenames:

1. As mentioned above, filenames are separated by their terminating NUL-byte in the initial processing.
2. The chunk referenced as `x_filenames` above will contain as much decoded data as is specified by the third DWORD in the archive header (excl. leading metadata).

As will be seen in the PoC, I haven't bothered reverse engineering and reimplementing the (de)en)coding algorithm (presumably based on Huffman). However, it seems that the archive filenames and their content are adjacent each other such that:

```
- filename 0
- filename 1
- filename_2
- ...
- content_of_filename_0
- content_of_filename_1
- content_of_filename_2
- ...
```

If we'd modify the third DWORD in the header (0x4F in the demonstrative archive above) to a larger value, the buffer referenced as `x_filenames` would not only contain the decoded filenames, but also (part of, depending on the value) their decoded contents.

This means that we can use the `Yz1` library itself to write our exploit for the unchecked calls to `sprintf()`. The general approach looks like:

1. Create an archive with N files.
2. Set a breakpoint *before* the filenames are encoded, but *after* the metadata has been constructed.
3. Remove the terminating NUL-byte for one of the filenames (effectively concatenating them).
4. Let the process finish.
5. Increase the third DWORD in the header to a size that includes the length of all file content.

The result is that the decoding process will interpret file contents as filenames. This gives us ample opportunity to create a source buffer large enough to overflow the stack in the call to `sprintf()`.

The PoC creation is accomplished with `pykd` – which is not only a plugin for `WinDbg` but also very usable as a standalone Python module for automated debugging.

As for exploit mitigations, the changelog for `IZArc` mentions that ASLR and DEP was introduced in `IZArc` version 4.3. However, that only applies to the main executable and *some* plugins (presumably the plugins for which the author has access to the source code).

With that said, only two of the shipped modules are non-rebased: `Tar32.dll` and `cabinet5.dll`.

Anyway, after having removed the `NUL` between two filenames in the archive, the file contents that will later be interpreted as a filename will contain the following:

1. Enough data to overflow the stack (incl. SEH).
2. A SEH gadget that adjusts `esp` and returns into our ROP sled.
3. Gadgets that prepares the stack with appropriate arguments for `VirtualAlloc()`
4. Gadget to invoke `VirtualAlloc()` by using its IAT slot in `Tar32.dll`.
5. Our shellcode.

Unfortunately, it's difficult to write a reliable exploit due to the extraction directory being prepended to our overflowing "filename". The approach taken in the PoC is to spray the SEH overwrite after adjusting the initial bogus data in an attempt for the overwrite to land on an appropriate DWORD boundary. The alignment is done in the interval `[0,4)` – i.e. `len(path) % 4` (where `path` includes the trailing `\`). So, there's a 1 in 4 shot for success if the extraction path is unpredictable.


Demonstration

Because the PoC uses `Yz1.dll` and `pykd` to create the payload, and because `Yz1.dll` is a 32-bit Windows-only module, the payload has to be created on a Windows system with a 32-bit Python >=3.6.

Example:

```
> "C:\Program Files (x86)\Python38\python.exe" exploit.py \
--dll "C:\Program Files (x86)\IZArc\Yz1.dll" \
--output C:\Users\user\Downloads\archive.yz1 \
--align C:\Users\user\Downloads\archive
=> created: C:\Users\user\Downloads\archive.yz1
=> extraction path alignment: 0
```

Note that `--align` can also be an integer `[0, 4)` or left out completely (in which case it's derived from the `--output path`).



See [source](#).

Solution

These flaws were assigned CVE-2020-24175 and no solution exist for either [Yz1](#) or [lZArc](#) at the time of writing.