

New issue

[Jump to bottom](#)

Memory leak in ndpi_free_flow() #994

Closed

FitzBC opened this issue on Aug 26, 2020 · 5 comments

Assignees



FitzBC commented on Aug 26, 2020 • edited

I am using `nfstream` to detect traffic, and the `dissect` option is turned on, which means that `nfstream` will use `nDPI` to complete this task. But I encountered a **memory leak** problem during use. `Nfstream` is developed based on the python language, memory leaks are rarely encountered, so I started to locate the cause.

In `nfstream`, I located the reason for this line of code: `plugin.py`. In fact, the code calls the `ndpi_flow_free()` function in `nDPI`: `ndpi_main.c`. By comparing the memory release operation in `ndpi_flow_free()` and the [prototype of the ndpi_flow_struct structure](#), I think that `ndpi_flow_free()` has not completely released all the variables in `ndpi_flow_struct`.

I think this may be the root cause of the memory leak.

aouinized commented on Aug 26, 2020

Collaborator

@FitzBC Hi, can you please provide the version of `nfstream` you use. Which interpreter (PyPy or CPython) and its version. How did you confirm that it's a memory leak? It was on live or offline mode?

BR,
Zied

lucaderi assigned aouinized on Aug 26, 2020

FitzBC commented on Aug 26, 2020 • edited

Author

@aouinized Thanks for your reply, I am using offline mode, and my relevant version is as follows:

- `nfstream`: 5.2.0
- CPython: 3.7.0

For the confirm of a memory leak, first, during the running of my program, the memory usage rate keeps rising and eventually is forced to end. You can use the following simple test code to verify (Because it is a single process, the memory usage rate increases slowly, about 2 MB/s. You can use the process pool to speed up this process):

```
from nfstream import NFStreamer
import os

def nfs_judge_process(pcap_name, return_dict=None):

    my_awesome_streamer = NFStreamer(source=pcap_name, # or network interface
                                     snaplen=1,
                                     idle_timeout=0,
                                     active_timeout=0,
                                     # plugins=(),
                                     dissect=True,
                                     max_tcp_dissections=1,
                                     max_udp_dissections=1,
                                     statistics=False,
                                     enable_guess=True,
                                     decode_tunnels=True,
                                     bpf_filter=None,
                                     promisc=False
                                    )

    data_pandas = my_awesome_streamer.to_pandas(ip_anonymization=False)
    application_name = data_pandas[0:1]['application_name'][0]

    return application_name

def main():

    pcap_dir='your_pcap_dir'
    while(1):
        for path, dir_list, file_list in os.walk(pcap_dir):
            file_list.sort()
            for pcap_name in file_list:

                nfs_judge_process(os.path.join(pcap_dir, pcap_name))

if __name__ == "__main__":
    main()
```

Second, I used `muppy` to help locate. My code is as follows:

```
from pymler import muppy, summary

all_objects_1 = muppy.get_objects()
sum1 = summary.summarize(all_objects_1)

nfs_judge_process(os.path.join(pcap_dir, pcap_name))

sum2 = summary.summarize(muppy.get_objects())
```

```
diff = summary.get_diff(sum1, sum2)
summary.print_(diff)
```

The output of each loop is similar to the following output:

	types	# objects	total size
=====	=====	=====	=====
	list	7757	1.60 MB
	dict	1431	793.80 KB
	str	10037	760.77 KB
	tuple	1141	94.01 KB
	code	545	76.77 KB
	type	77	76.48 KB
	int	2478	67.79 KB
	set	37	21.09 KB
	re.Pattern	25	20.81 KB
	pycparser.ply.yacc.MinProduction	310	16.95 KB
	_cffi_backend.CType	127	16.16 KB
	weakref	165	12.89 KB
	method	206	12.88 KB
	weakref.KeyedRef	132	11.34 KB
	member_descriptor	142	9.98 KB

aouinized commented on Aug 26, 2020 • edited

Collaborator

@FitzBC Thanks for the information.

- Do you confirm the "leak" disappears when dissect is set to False?
- With your current parameters (idle=0 and active=0), each packet will be treated as a single flow. which explain why memory grows is a fast fashion. Is that to answer a specific need?

FitzBC commented on Aug 26, 2020 • edited

Author

@aouinized

- When dissect is set to False, the leak disappeared, so I think the cause may be on the nDPI side.
- In my needs, I need to judge each packet. These packets may have different sources, so I used idle=0 and active=0.

In addition, I think that even if each packet has a single flow, if each flow will release the memory normally after it is used, it will not cause the memory to rise rapidly.

Finally, I found an interesting thing. In the `nfs_judge_process()` function of the sample code I provided, if you delete the following two lines and replace them with `del`, then the memory rise will become very rapid, I think this information may be helpful to you.

```
# data_pandas = my_awesome_streamer.to_pandas(ip_anonymization=False)
# application_name = data_pandas[0:1]['application_name'][0]
del my_awesome_streamer
```

aouinized commented on Aug 27, 2020


Collaborator

@FitzBC Thank you for the provided details.

After investigation, this issue has nothing to deal with nDPI.

- Memory was not leaking. What you observed was specific to your extreme settings (idle=0, active=0) and how backpressure was handled in `nfstream` in such a scenario. I fixed it and the fix will be integrated into the next release of `nfstream`.
- Your second observation is due to another reason that is independent of the current issue. In your loop, you instantiate an `NFStreamer` object and you release it with the `del` without any processing. The issue is that some allocated modules are not correctly freed when no processing is done. This is also fixed and will be integrated into the next release.

BR,
Zied

 aouinized closed this as completed on Aug 27, 2020

Assignees

 aouinized

Labels

None yet

Projects

None yet

Milestone

No milestone

Development

No branches or pull requests

2 participants