

Talos Vulnerability Report

TALOS-2020-1224

Openscad import_stl.cc:import_stl() out-of-bounds stack write vulnerability

FEBRUARY 23, 2021

CVE NUMBER

CVE-2020-28600

Summary

An out-of-bounds write vulnerability exists in the `import_stl.cc:import_stl()` functionality of Openscad openscad-2020.12-RC2. A specially crafted STL file can lead to code execution. An attacker can provide a malicious file to trigger this vulnerability.

Tested Versions

Openscad openscad-2020.12-RC2

Product URLs

<https://github.com/openscad/openscad>

CVSSv3 Score

8.8 - CVSS:3.0/AV:N/AC:L/PR:N/UI:R/S:U/C:H/I:H/A:H

CWE

CWE-119 - Improper Restriction of Operations within the Bounds of a Memory Buffer

Details

Openscad is an open-source program for creating 3-D CAD models, available for all platforms. Aside from describing and creating objects from scripts, it's also possible to import existing .stl, .amf, .3mf, .svg, and .dxf files into a scene for rendering.

When importing a given .stl file into a scene via the `import("file.stl");` command, the first stl-specific function we hit is `PolySet *import_stl(const std::string &filename, const Location &loc):`

```
PolySet *import_stl(const std::string &filename, const Location &loc)
{
    PolySet *p = new PolySet(3);

    // Open file and position at the end
    std::ifstream f(filename.c_str(), std::ios::in | std::ios::binary | std::ios::ate); // [1]
    if (!f.good()) {
        LOG(message_group::Warning, Location::NONE, "", "Can't open import file '%1$s', import() at line %2$d", filename, loc.firstLine());
        return p;
    }

    boost::regex ex_sfe("solid|facet|endloop"); // [2]
    boost::regex ex_outer("outer loop");
    boost::regex ex_vertex("vertex");
    boost::regex ex_vertices("\\s*vertex\\s+([^\\s]+)\\s+([^\\s]+)\\s+([^\\s]+)"); // [3]

    bool binary = false;
    std::streampos file_size = f.tellg();
    f.seekg(80);
    if (f.good() && !f.eof()) { // [4]
        uint32_t facenum = 0;
        f.read((char *)&facenum, sizeof(uint32_t));
    #if BOOST_ENDIAN_BIG_BYTE
        uint32_byte_swap( facenum );
    #endif
        if (file_size == static_cast<std::streamoff>(80 + 4 + 50*facenum)) {
            binary = true;
        }
    }
}
```

At [1], our input file is opened, and at [2] through [3] we notice some important regexes that will be used further on. Assuming we pass the check at [4], which makes sure our file is at least 80 bytes, then we move on to the following code:

```

PolySet *import_stl(const std::string &filename, const Location &loc)
{
    // [...]
    char data[5];
    f.read(data, 5);
    if (!binary && !f.eof() && f.good() && !memcmp(data, "solid", 5)) {
        int i = 0;
        double vdata[3][3]; // [1]
        std::string line;
        std::getline(f, line);
        while (!f.eof()) { // [2]
            std::getline(f, line);
            boost::trim(line);
            if (boost::regex_search(line, ex_sfe)) { // "solid|facet|endloop" // [3]
                continue;
            }
            if (boost::regex_search(line, ex_outer)) { // "outer loop" // [4]
                i = 0;
                continue;
            }
            boost::smatch results;
            if (boost::regex_search(line, results, ex_vertices)) { // "\\s*vertex\\s+([^\s]+)\\s+([^\s]+)\\s+([^\s]+)" // [5]
                // [...]
            }
        }
    }
}

```

At [2] we hit our parsing loop, iterating over each line of the input .stl file, looking for different regexes as we go along. Lines matching the regex at [3], "solid|facet|endloop", are completely ignored, lines matching at [4], "outer loop", reset the i variable, but that's about it. The only regex that is actually read in is at [5], "\\s*vertex\\s+([^\s]+)\\s+([^\s]+)\\s+([^\s]+)". To give an example:

```

facet normal 1.000000e+00 0.000000e+00 -0.000000e+00
outer loop
vertex 2.000000e+01 2.000000e+01 0.000000e+00
vertex 2.000000e+01 2.000000e+01 2.000000e+01
vertex 2.000000e+01 0.000000e+00 2.000000e+01
endloop
endfacet

```

To proceed, let us now examine the code hit when the ex_vertices regex is hit:

```

if (boost::regex_search(line, results, ex_vertices)) { // "\\s*vertex\\s+([^\s]+)\\s+([^\s]+)\\s+([^\s]+)"
    try {
        for (int v=0; v<3; ++v) {
            vdata[i][v] = boost::lexical_cast<double>(results[v+1]); // [1]
        }
    }
    catch (const boost::bad_lexical_cast &bblc) { // [2]
        LOG(message_group::Warning, Location::NONE, "", "Can't parse vertex line '%1$s', import() at line %2$d", line, loc.firstLine());
        i = 10; // [3]
        continue;
    }
    if (++i == 3) { // [4]
        p->append_poly();
        p->append_vertex(vdata[0][0], vdata[0][1], vdata[0][2]);
        p->append_vertex(vdata[1][0], vdata[1][1], vdata[1][2]);
        p->append_vertex(vdata[2][0], vdata[2][1], vdata[2][2]);
    }
}

```

Each of the vertex numbers are populated into the vdata variable at [1], and if we have three vertexes read in (forming a triangle) at [4], we append these vertexes into the PolySet *p object.

A curious thing happens though if we have a given vertex co-ordinate that raises an error from boost::lexical_cast<double> [2], the i variable is set to 10 at [3]. Thus, once we hit the next line in the file that hits the ex_vertices regex, the line at [1] will write a user controlled value to vdata[10][v], the consequences thereof depending on the compiler.

At least for our testing build, this overwrote a pointer inside a boost::shared_ptr, allowing us to control exactly what happened during the boost::shared_ptr's destructor.

Crash Information

```
*****
rax      : 0x4034000000000034 | r13[S]   : 0x7ffd8634d370
rbx[S]   : 0x7ffd8634d600    | r14[S]   : 0x7ffd8634d310
rcx      : 0x4034000000000008 | r15[S]   : 0x7ffd8634d330
rdx      : 0x1               | rip[L]   : 0x7f6aaebad51f
rsi      : 0x8068000000000001 | eflags   : 0x10202
rdi      : 0x4034000000000008 | cs       : 0x33
rbp[S]   : 0x7ffd8634ceb0    | ss       : 0x2b
rsp[S]   : 0x7ffd8634ce90    | ds       : 0x0
r8       : 0x5c4000          | es       : 0x0
r9       : 0x21             | fs       : 0x0
r10      : 0x7ffd85b50000    | gs       : 0x0
r11      : 0xffffffff01      | fs_base  : 0x7f6aa7583e00
r12[S]   : 0x7ffd8634d350    | gs_base  : 0x0
*****
0x7f6aaebad50e : mov     DWORD PTR [rbp-0xc],0x1
0x7f6aaebad515 : mov     edx,DWORD PTR [rbp-0xc]
0x7f6aaebad518 : mov     rsi,rcx
0x7f6aaebad51b : shr     rsi,0x3
=>0x7f6aaebad51f : mov     al,BYTE PTR [rsi+0x7fff8000]
0x7f6aaebad525 : cmp     al,0x0
0x7f6aaebad527 : mov     QWORD PTR [rbp-0x18],rcx
0x7f6aaebad52b : mov     DWORD PTR [rbp-0x1c],edx
0x7f6aaebad52e : mov     BYTE PTR [rbp-0x1d],al
0x7f6aaebad531 : je      0x7f6aaebad560 <boost::detail::atomic_decrement(int _Atomic*)+112>
*****
#0  0x00007f6aaebad51f in boost::detail::atomic_decrement(int _Atomic*) (pw=0x4034000000000008) at/sp_counted_base_clang.hpp:38
#1  0x00007f6aaebad426 in boost::detail::sp_counted_base::release (this=0x4034000000000000) at/sp_counted_base_clang.hpp:118
#2  0x00007f6aaebacc7a in boost::detail::shared_count::~~shared_count (this=0x7ffd8634d4f0) at/shared_count.hpp:427
#3  0x00007f6aaee27a8f in boost::shared_ptr<boost::re_detail_107100::named_subexpressions>::~~shared_ptr (this=0x7ffd8634d4e8)
    at/shared_ptr.hpp:341
#4  0x00007f6aaed88db6 in boost::match_results<__gnu_cxx::__normal_iterator<char const*, std::__cxx11::basic_string<char,
    std::char_traits<char>, std::allocator<char> > >, std::allocator<boost::sub_match<__gnu_cxx::__normal_iterator<char const*,
    std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char> > > >::match_results (this=0x7ffd8634d4b0)
    at/match_results.hpp:119
#5  0x00007f6aaf1c7aa2 in import_stl (filename= at/import_stl.cc:129
#6  0x000000000055bb6d in LLVMFuzzerTestOneInput (Data=0x61f000000e0 "solid STL generated by MeshLab\n facet normal 0.000000e+00
-0.000000e+00 1.000000e+00\n outer loop\n vertex 2.000000e+01 2.000000e+01 0.000000e+00\n vertex 2.000000e+01
0.000000e+00"... Size=3342) at/fuzz_stl_harness.cpp:71
#7  0x0000000000461ae2 in fuzzer::Fuzzer::ExecuteCallback(unsigned char const*, unsigned long) () at/optional.hpp:99
#8  0x000000000044d253 in fuzzer::RunOneTest(fuzzer::Fuzzer*, char const*, unsigned long) () at/optional.hpp:99
#9  0x0000000000452d07 in fuzzer::FuzzerDriver(int*, char***, int (*)(unsigned char const*, unsigned long)) () at/optional.hpp:99
#10 0x000000000047b9c3 in main () at/optional.hpp:99
*****
[ ^_^ ] 2020_12_16_12_42_26_078463 - Got a crash! SIGSEGV, 0x7f6aaebad51f
```

Timeline

2021-01-11 - Vendor Disclosure

2021-02-23 - Public Release

CREDIT

Discovered by Liliith >_> of Cisco Talos.

VULNERABILITY REPORTS

PREVIOUS REPORT

NEXT REPORT

TALOS-2020-1223

TALOS-2020-1225

