



index : kernel/git/torvalds/linux.git

Linux kernel source tree

master switch

Linus Torvalds

about summary refs log tree commit diff stats

log msg search

author Jann Horn <jannh@google.com> 2020-12-03 02:25:05 +0100
committer Greg Kroah-Hartman <gregkh@linuxfoundation.org> 2020-12-04 17:39:58 +0100
commit c8bcd9c5be24fb9e6132e97da5a35e5a83e36b9 (patch)
tree 1b7e3191d3fd63c02d3029b19e45d88d322397df
parent 54ffccbf053b5b6ca4f6e45094b942fab92a25fc (diff)
download linux-c8bcd9c5be24fb9e6132e97da5a35e5a83e36b9.tar.gz

diff options

context: 3
space: include
mode: unified

tty: Fix ->session locking

Currently, locking of ->session is very inconsistent; most places protect it using the legacy tty mutex, but disassociate_ctty(), __do_SAK(), tiocspgrp() and tiocgsid() don't.
Two of the writers hold the ctrl_lock (because they already need it for ->pgrp), but __proc_set_tty() doesn't do that yet.

On a PREEMPT=y system, an unprivileged user can theoretically abuse this broken locking to read 4 bytes of freed memory via TIOCGSID if tiocgsid() is preempted long enough at the right point. (Other things might also go wrong, especially if root-only ioctls are involved; I'm not sure about that.)

Change the locking on ->session such that:

- tty_lock() is held by all writers: By making disassociate_ctty() hold it. This should be fine because the same lock can already be taken through the call to tty_vhangup_session().
The tricky part is that we need to shorten the area covered by siglock to be able to take tty_lock() without ugly retry logic; as far as I can tell, this should be fine, since nothing in the signal_struct is touched in the 'if (tty)' branch.
- ctrl_lock is held by all writers: By changing __proc_set_tty() to hold the lock a little longer.
- All readers that aren't holding tty_lock() hold ctrl_lock: By adding locking to tiocgsid() and __do_SAK(), and expanding the area covered by ctrl_lock in tiocspgrp().

Cc: stable@kernel.org

Signed-off-by: Jann Horn <jannh@google.com>

Reviewed-by: Jiri Slaby <jirislaby@kernel.org>

Signed-off-by: Greg Kroah-Hartman <gregkh@linuxfoundation.org>

Diffstat

```
-rw-r--r-- drivers/tty/tty_io.c 7  
-rw-r--r-- drivers/tty/tty_jobctrl.c 44  
-rw-r--r-- include/linux/tty.h 4
```

3 files changed, 41 insertions, 14 deletions

diff --git a/drivers/tty/tty_io.c b/drivers/tty/tty_io.c
index 9f8b9a567b359..56ade99ef99f4 100644

```
--- a/drivers/tty/tty_io.c  
+++ b/drivers/tty/tty_io.c  
@@ -2897,10 +2897,14 @@ void __do_SAK(struct tty_struct *tty)  
     struct task_struct *g, *p;  
     struct pid *session;  
     int i;  
+     unsigned long flags;  
  
     if (!tty)  
         return;  
-     session = tty->session;  
+  
+     spin_lock_irqsave(&tty->ctrl_lock, flags);  
+     session = get_pid(tty->session);  
+     spin_unlock_irqrestore(&tty->ctrl_lock, flags);  
  
     tty_ldisc_flush(tty);  
  
@@ -2932,6 +2936,7 @@ void __do_SAK(struct tty_struct *tty)  
     task_unlock(p);  
     } while_each_thread(g, p);  
     read_unlock(&tasklist_lock);  
+     put_pid(session);  
+ #endif  
 }
```

diff --git a/drivers/tty/tty_jobctrl.c b/drivers/tty/tty_jobctrl.c
index baadeea4a289b..aa6d0537b379e 100644

```
--- a/drivers/tty/tty_jobctrl.c  
+++ b/drivers/tty/tty_jobctrl.c  
@@ -103,8 +103,8 @@ static void __proc_set_tty(struct tty_struct *tty)  
     put_pid(tty->session);  
     put_pid(tty->pgrp);  
     tty->pgrp = get_pid(task_pgrp(current));  
-     spin_unlock_irqrestore(&tty->ctrl_lock, flags);  
+     spin_unlock_irqrestore(&tty->ctrl_lock, flags);  
     tty->session = get_pid(task_session(current));  
+     spin_unlock_irqrestore(&tty->ctrl_lock, flags);  
     if (current->signal->tty) {  
         tty_debug(tty, "current tty %s not NULL!!\n",  
                 current->signal->tty->name);  
@@ -293,20 +293,23 @@ void disassociate_ctty(int on_exit)  
     spin_lock_irq(&current->sigband->siglock);  
     put_pid(current->signal->tty_old_pgrp);  
     current->signal->tty_old_pgrp = NULL;  
  
-  
+     tty = tty_kref_get(current->signal->tty);  
+     spin_unlock_irq(&current->sigband->siglock);  
+  
     if (tty) {  
         unsigned long flags;  
  
+         tty_lock(tty);  
+         spin_lock_irqsave(&tty->ctrl_lock, flags);  
+         put_pid(tty->session);  
+         put_pid(tty->pgrp);  
+         tty->session = NULL;  
+         tty->pgrp = NULL;  
+         spin_unlock_irqrestore(&tty->ctrl_lock, flags);  
+         tty_unlock(tty);  
+         tty_kref_put(tty);  
     }
```

```

    }

-   spin_unlock_irq(&current->sigband->siglock);
    /* Now clear signal->tty under the lock */
    read_lock(&tasklist_lock);
    session_clear_tty(task_session(current));
@@ -477,14 +480,19 @@ static int tiocspgrp(struct tty_struct *tty, struct tty_struct *real_tty, pid_t
    return -ENOTTY;
    if (retval)
        return retval;
-   if (!current->signal->tty ||
-       (current->signal->tty != real_tty) ||
-       (real_tty->session != task_session(current)))
-       return -ENOTTY;
+
    if (get_user(pgrp_nr, p))
        return -EFAULT;
    if (pgrp_nr < 0)
        return -EINVAL;

+   spin_lock_irq(&real_tty->ctrl_lock);
+   if (!current->signal->tty ||
+       (current->signal->tty != real_tty) ||
+       (real_tty->session != task_session(current))) {
+       retval = -ENOTTY;
+       goto out_unlock_ctrl;
+   }
    rcu_read_lock();
    pgrp = find_vpid(pgrp_nr);
    retval = -ESRCH;
@@ -494,12 +502,12 @@ static int tiocspgrp(struct tty_struct *tty, struct tty_struct *real_tty, pid_t
    if (session_of_pgrp(pgrp) != task_session(current))
        goto out_unlock;
    retval = 0;
-   spin_lock_irq(&real_tty->ctrl_lock);
    put_pid(real_tty->pgrp);
    real_tty->pgrp = get_pid(pgrp);
-   spin_unlock_irq(&real_tty->ctrl_lock);
out_unlock:
    rcu_read_unlock();
+out_unlock_ctrl:
+   spin_unlock_irq(&real_tty->ctrl_lock);
    return retval;
}

@@ -511,20 +519,30 @@ out_unlock:
 *
 * Obtain the session id of the tty. If there is no session
 * return an error.
- *
- * Locking: none. Reference to current->signal->tty is safe.
 */
static int tiocgsid(struct tty_struct *tty, struct tty_struct *real_tty, pid_t __user *p)
{
+   unsigned long flags;
+   pid_t sid;
+
    /*
     * (tty == real_tty) is a cheap way of
     * testing if the tty is NOT a master pty.
     */
    if (tty == real_tty && current->signal->tty != real_tty)
        return -ENOTTY;

+   spin_lock_irqsave(&real_tty->ctrl_lock, flags);
    if (!real_tty->session)
        return -ENOTTY;
-   return put_user(pid_vnr(real_tty->session), p);
+   goto err;
+   sid = pid_vnr(real_tty->session);
+   spin_unlock_irqrestore(&real_tty->ctrl_lock, flags);
+
    return put_user(sid, p);
+err:
+   spin_unlock_irqrestore(&real_tty->ctrl_lock, flags);
+   return -ENOTTY;
}

/*
diff --git a/include/linux/tty.h b/include/linux/tty.h
index a99e9b8e4e316..eb33d948788cc 100644
--- a/include/linux/tty.h
+++ b/include/linux/tty.h
@@ -306,6 +306,10 @@ struct tty_struct {
    struct termiox *termiox;          /* May be NULL for unsupported */
    char name[64];
    struct pid *pgrp;                 /* Protected by ctrl lock */
+   /*
+    * Writes protected by both ctrl lock and legacy mutex, readers must use
+    * at least one of them.
+    */
    struct pid *session;
    unsigned long flags;
    int count;

```