

[Products](#)[Services](#)[Publications](#)[Resources](#)[What's new](#)

Follow @Openwall on Twitter for new release announcements and other news

[<prev](#)] [next>\]](#) [\[thread-next>\]](#) [\[day\]](#) [\[month\]](#) [\[year\]](#) [\[list\]](#)

Date: Tue, 25 May 2021 15:18:22 +0800
From: Martillin <mmmartillin@openwall.com>
To: oss-security@openwall.com
Subject: CVE-2021-3564 Linux Bluetooth device initialization implementation bug

Hello there,

Our team (BlockSec) found an implementation bug that resides in the kernel Bluetooth subsystem when the HCI device initialization fails. It can lead to unexpected results, like double-free memory corruption vulnerability.

===== BUG DETAILS =====

This implementation bug is inside hci_dev_do_open() function.

```
static int hci_dev_do_open(struct hci_dev *hdev)
{
    ...
} else {
    /* Init failed, cleanup */
    flush_work(&hdev->tx_work);
    flush_work(&hdev->cmd_work); // (1)
    flush_work(&hdev->rx_work); // (2)

    skb_queue_purge(&hdev->cmd_q);
    skb_queue_purge(&hdev->rx_q);

    if (hdev->flush)
        hdev->flush(hdev);

    if (hdev->sent_cmd) {
        kfree_skb(hdev->sent_cmd);
        hdev->sent_cmd = NULL;
    }
}
...
}
```

The purpose of flush_work(struct work_struct *work) is to wait for the accomplishment of the work_struct. Hence, the accomplishment of the code flush_work(&hdev->cmd_work) (1) means the cmd_work is finished. However, we discover an implementation bug that can result in activating hci_cmd_work() even the hdev->cmd_work has already been flushed (2).

The process is as follows:

```
hci_rx_work() -> hci_event_packet() -> hci_event_packet() ->
hci_cmd_complete_evt() -> queue_work(hdev->workqueue, &hdev->cmd_work)
```

We found this implementation bug can lead to double-free memory corruption, which resulted from a data race of the hdev->sent_cmd. Here is the code snippet for this race.

```
static void hci_cmd_work(struct work_struct *work)
{
    ...
    if (atomic_read(&hdev->cmd_cnt)) {
        skb = skb_dequeue(&hdev->cmd_q);
        if (!skb)
            return;

        kfree_skb(hdev->sent_cmd);

        hdev->sent_cmd = skb_clone(skb, GFP_KERNEL);
    }
    ...
}
```

We use thread-A to represent hci_dev_do_open() function and the thread-B for hci_cmd_work(). The normal sequence should be like this:

thread-A	thread-B
	kfree_skb(hdev->sent_cmd); (FREE)
	hdev->sent_cmd = skb_clone(skb,
GFP_KERNEL); (WRITE)	
if (hdev->sent_cmd) { (READ)	
kfree_skb(hdev->sent_cmd); (FREE)	
hdev->sent_cmd = NULL; (WRITE)	

However, if the sequence is like this:

thread-A	thread-B
	kfree_skb(hdev->sent_cmd); (FREE)
if (hdev->sent_cmd) { (READ)	
kfree_skb(hdev->sent_cmd); (FREE)	hdev->sent_cmd = skb_clone(skb,
GFP_KERNEL); (WRITE)	
hdev->sent_cmd = NULL; (WRITE)	

If the FREE operation in thread-A is before WRITE operation in thread-B, it can lead to double-free memory corruption in the kernel.

===== BUG EFFECTS =====

For now, we can successfully trigger the vulnerability to corrupt the kernel memory and thus crash the kernel. Although this bug is related to Bluetooth device initialization, the attacker can trigger it without extra privileges.

That is because the Linux kernel does not ask for the privilege when attaching the HCI device as the attached device is default set to HCI_AUTO_OFF state. This bug is inside in the very first attaching procedure and requires no syscalls.

The crash log is presented below.

```
=====
[ 500.906562] hci0 type 1 len 3
[ 500.904986] BUG: KASAN: use-after-free in kfree_skb+0x33/0x1c0
[ 500.904986] Read of size 4 at addr ffff888009d3599c by task
kworker/u5:0/54
[ 500.904986]
[ 500.909997] CPU: 0 PID: 54 Comm: kworker/u5:0 Not tainted 5.11.11+ #16
[ 500.909997] Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS
1.13.0-lubuntu1.1 04/01/2014
[ 500.909997] Workqueue: hci0 hci_power_on
[ 500.909997] Call Trace:
[ 500.909997] dump_stack+0x16c/0x1be
```

```
[ 500.924511] print_address_description+0x7b/0x3a0
[ 500.924511]   kasan_report+0x14e/0x200
[ 500.924511]   ? kfree_skb+0x33/0x1c0
[ 500.924511]   ? skb_queue_purge+0x193/0x1c0
[ 500.924511] kasan_report+0x47/0x60
[ 500.924511]   ? skb_queue_purge+0x193/0x1c0
[ 500.924511] check_memory_region+0x2e2/0x330
[ 500.924511] kfree_skb+0x33/0x1c0
[ 500.924511] hci_dev_do_open+0x1008/0x1570
[ 500.924511]   ? printk+0x62/0x83
[ 500.924511] hci_power_on+0x183/0x580
[ 500.924511]   ? strscpy+0x7f/0x240
[ 500.924511] process_one_work+0x722/0x1150
[ 500.924511] worker_thread+0xb5c/0x17d0
[ 500.924511]   ? process_one_work+0x1150/0x1150
[ 500.924511] kthread+0x2fc/0x320
[ 500.924511]   ? process_one_work+0x1150/0x1150
[ 500.924511]   ? kthread_unuse_mm+0x1d0/0x1d0
[ 500.924511] ret_from_fork+0x22/0x30
[ 500.924511]
[ 500.924511] Allocated by task 273:
[ 500.924511]   kasan_kmalloc+0xc6/0x100
[ 500.924511] kmem_cache_alloc+0xfe/0x1f0
[ 500.924511] skb_clone+0x1b5/0x360
[ 500.924511] hci_cmd_work+0x15d/0x350
[ 500.924511] process_one_work+0x722/0x1150
[ 500.924511] worker_thread+0xb5c/0x17d0
[ 500.924511] kthread+0x2fc/0x320
[ 500.924511] ret_from_fork+0x22/0x30
[ 500.924511]
[ 500.924511] Freed by task 273:
[ 500.924511] kasan_set_track+0x3d/0x70
[ 500.924511] kasan_set_free_info+0x1f/0x40
[ 500.924511]   kasan_slab_free+0x10e/0x140
[ 500.924511] kmem_cache_free+0xca/0x210
[ 500.924511] hci_cmd_work+0x150/0x350
[ 500.924511] process_one_work+0x722/0x1150
[ 500.924511] worker_thread+0xb5c/0x17d0
[ 500.924511] kthread+0x2fc/0x320
[ 500.924511] ret_from_fork+0x22/0x30
[ 500.924511]
[ 500.924511] The buggy address belongs to the object at ffff888009d358c0
[ 500.924511]   which belongs to the cache skbuff_head_cache of size 232
[ 500.924511] The buggy address is located 220 bytes inside of
[ 500.924511]   232-byte region [ffff888009d358c0, ffff888009d359a8)
[ 500.924511] The buggy address belongs to the page:
[ 500.924511] page:00000000b691648a refcount:1 mapcount:0
mapping:0000000000000000 index:0x0 pfn:0x9d35
[ 500.924511] flags: 0x100000000000200 dead000000000100 (slab)
[ 500.924511] raw: 0100000000000200 dead000000000100 dead000000000122
ffff888006d64640
[ 500.924511] raw: 0000000000000000 00000000000c000c 00000001ffffffff
0000000000000000
[ 500.924511] page dumped because: kasan: bad access detected
[ 500.924511]
[ 500.924511] Memory state around the buggy address:
[ 500.924511] ffff888009d35880: fc fc fc fc fc fc fc fa fb fb fb fb
fb fb
[ 500.924511] ffff888009d35900: fb fb fb fb fb fb fb fb fb fb fb fb
fb fb
[ 500.924511] >ffff888009d35980: fb fb fb fb fb fc fc fc fc fc fc fc
fc fc
[ 500.924511]                                     ^
[ 500.924511] ffff888009d35a00: fa fb fb fb fb fb fb fb fb fb fb fb
fb fb
[ 500.924511] ffff888009d35a80: fb fb fb fb fb fb fb fb fb fb fb fc
fc fc
[ 500.924511]
=====
[ 500.924511] Disabling lock debugging due to kernel taint
[ 501.014277]
=====
[ 501.014929] BUG: KASAN: double-free or invalid-free in
hci_dev_do_open+0x1008/0x1570
[ 501.014929]
[ 501.014929] CPU: 0 PID: 54 Comm: kworker/u5:0 Tainted: G   B
5.11.11+ #16
[ 501.014929] Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS
1.13.0-lubuntu1.1 04/01/2014
[ 501.014929] Workqueue: hci0 hci_power_on
[ 501.014929] Call Trace:
[ 501.014929]   dump_stack+0x16c/0x1be
[ 501.014929]   ? hci_dev_do_open+0x1008/0x1570
[ 501.014929]   ? hci_dev_do_open+0x1008/0x1570
[ 501.014929]   ? hci_dev_do_open+0x1008/0x1570
[ 501.014929]   ? hci_dev_do_open+0x1008/0x1570
[ 501.014929]   ? hci_dev_do_open+0x1008/0x1570
[ 501.014929]   kasan_report_invalid_free+0x54/0xd0
[ 501.014929]   kasan_slab_free+0xe7/0x140
[ 501.014929]   kmem_cache_free+0xca/0x210
[ 501.014929]   ? hci_dev_do_open+0x1008/0x1570
[ 501.014929]   hci_dev_do_open+0x1008/0x1570
[ 501.014929]   ? printk+0x62/0x83
[ 501.014929]   hci_power_on+0x183/0x580
[ 501.014929]   ? strscpy+0x7f/0x240
[ 501.014929]   process_one_work+0x722/0x1150
[ 501.014929]   worker_thread+0xb5c/0x17d0
[ 501.014929]   ? process_one_work+0x1150/0x1150
[ 501.014929]   kthread+0x2fc/0x320
[ 501.014929]   ? process_one_work+0x1150/0x1150
[ 501.014929]   ? kthread_unuse_mm+0x1d0/0x1d0
[ 501.014929]   ret_from_fork+0x22/0x30
[ 501.014929]
[ 501.014929] Allocated by task 273:
[ 501.014929]   kasan_kmalloc+0xc6/0x100
[ 501.014929] kmem_cache_alloc+0xfe/0x1f0
[ 501.014929] skb_clone+0x1b5/0x360
[ 501.014929] hci_cmd_work+0x15d/0x350
[ 501.014929] process_one_work+0x722/0x1150
[ 501.014929] worker_thread+0xb5c/0x17d0
[ 501.014929] kthread+0x2fc/0x320
[ 501.014929] ret_from_fork+0x22/0x30
[ 501.014929]
[ 501.014929] Freed by task 273:
[ 501.014929] kasan_set_track+0x3d/0x70
[ 501.066803] kasan_set_free_info+0x1f/0x40
[ 501.066803]   kasan_slab_free+0x10e/0x140
[ 501.066803] kmem_cache_free+0xca/0x210
[ 501.066803] hci_cmd_work+0x150/0x350
[ 501.066803] process_one_work+0x722/0x1150
[ 501.066803] worker_thread+0xb5c/0x17d0
[ 501.066803] kthread+0x2fc/0x320
[ 501.066803] ret_from_fork+0x22/0x30
[ 501.066803]
[ 501.066803] The buggy address belongs to the object at ffff888009d358c0
[ 501.066803]   which belongs to the cache skbuff_head_cache of size 232
[ 501.066803] The buggy address is located 0 bytes inside of
[ 501.066803]   232-byte region [ffff888009d358c0, ffff888009d359a8)
[ 501.066803] The buggy address belongs to the page:
[ 501.066803] page:00000000b691648a refcount:1 mapcount:0
mapping:0000000000000000 index:0x0 pfn:0x9d35
[ 501.066803] flags: 0x100000000000200 dead000000000100 dead000000000122
ffff888006d64640
[ 501.066803] raw: 0000000000000000 00000000000c000c 00000001ffffffff
0000000000000000
[ 501.066803] page dumped because: kasan: bad access detected
[ 501.066803]
[ 501.066803] Memory state around the buggy address:
[ 501.066803] ffff888009d35780: fa fb fb fb fb fb fb fb fb fb fb fb
fb fb
```

```
[ fb  
[ 501.066803] ffff888009d35800: fb fb fb fb fb fb fb fb fb fb fc  
fc  
[ 501.066803] >ffff888009d35880: fc fc fc fc fc fc fa fb fb fb fb  
fb  
fb  
[ 501.066803] ^  
[ 501.066803] ffff888009d35900: fb fb fb fb fb fb fb fb fb fb  
fb  
fb  
[ 501.066803] ffff888009d35980: fb fb fb fb fb fc fc fc fc fc fc fc  
fc  
fc  
[ 501.066803]
```

***** Timeline *****

2021-05-17: Bug reported to security () kernel.org and linux-distros
() vs openwall.org

2021-05-25: CVE-2021-3564 assigned

We informed security@...nel.org on May 17, 2021. Now the 7-day embargo period is over, we are being asked to bring the issue to public.

Since our patch has not been applied to upstream yet, we will release the POC later.

***** Credit *****

HaoXiong@...ckSec Team

LinMa@...cksec Team

syzkaller

Best regards.

Mart111n

Powered by blists - more mailing lists

Please check out the [Open Source Software Security Wiki](#), which is counterpart to this [mailing list](#).

Confused about [mailing lists](#) and their use? [Read about mailing lists on Wikipedia](#) and check out these [guidelines on proper formatting of your messages](#).

