

[Products](#)[Services](#)[Publications](#)[Resources](#)[What's new](#)

Hash Suite - Windows password security audit tool. GUI, reports in PDF.

[\[<prev\]](#) [\[next>\]](#) [\[<thread-prev\]](#) [\[day\]](#) [\[month\]](#) [\[year\]](#) [\[list\]](#)

Date: Tue, 1 Jun 2021 09:17:12 +0800
From: Martillin <mmmartillin@il.com>
To: oss-security@openwall.com
Subject: Re: CVE-2021-3564 Linux Bluetooth device initialization implementation bug

Because it's been a week since I reported this issue, it's time to post the POC to the public and the POC is in attachment.

Thanks,

Martillin

On Tue, 25 May 2021 at 15:18, Martillin <mmmartillin@il.com> wrote:

```
> Hello there,
>
> Our team (BlockSec) found an implementation bug that resides in the kernel
> Bluetooth subsystem when the HCI device initialization fails. It can lead
> to unexpected results, like double-free memory corruption vulnerability.
>
> ===== BUG DETAILS =====
> This implementation bug is inside hci_dev_do_open() function.
>
> static int hci_dev_do_open(struct hci_dev *hdev)
> {
> ...
> } else {
>     /* Init failed, cleanup */
>     flush_work(&hdev->tx_work);
>     flush_work(&hdev->cmd_work); // {1}
>     flush_work(&hdev->rx_work); // {2}
>
>     skb_queue_purge(&hdev->cmd_q);
>     skb_queue_purge(&hdev->rx_q);
>
>     if (hdev->flush)
>         hdev->flush(hdev);
>
>     if (hdev->sent_cmd) {
>         kfree_skb(hdev->sent_cmd);
>         hdev->sent_cmd = NULL;
>     }
> ...
> }
```

> The purpose of flush_work(struct work_struct *work) is to wait for the accomplishment of the work_struct. Hence, the accomplishment of the code > flush_work(&hdev->cmd_work) {1} means the cmd work is finished. However, we > discover an implementation bug that can result in activating hci_cmd_work() > even the hdev->cmd_work has already been flushed {2}.

> The process is as follows:
> hci_rx_work() -> hci_event_packet() -> hci_event_packet() ->
> hci_cmd_complete_evt() -> queue_work(hdev->workqueue, &hdev->cmd_work)

> We found this implementation bug can lead to double-free memory corruption, which resulted from a data race of the hdev->sent_cmd. Here is the code snippet for this race.

```
> static void hci_cmd_work(struct work_struct *work)
> {
> ...
>     if (atomic_read(&hdev->cmd_cnt)) {
>         skb = skb_dequeue(&hdev->cmd_q);
>         if (!skb)
>             return;
>
>         kfree_skb(hdev->sent_cmd);
>
>         hdev->sent_cmd = skb_clone(skb, GFP_KERNEL);
> ...
> }
```

> We use thread-A to represent hci_dev_do_open() function and the thread-B > for hci_cmd_work().
> The normal sequence should be like this:

```
> -----
> thread-A | thread-B
> | kfree_skb(hdev->sent_cmd); (FREE)
> | hdev->sent_cmd = skb_clone(skb,
> GFP_KERNEL); (WRITE)
> if (hdev->sent_cmd) { (READ) |
> |
> kfree_skb(hdev->sent_cmd); (FREE) |
> |
> hdev->sent_cmd = NULL; (WRITE) |
> |
> -----
```

> However, if the sequence is like this:

```
> -----
> thread-A | thread-B
> | kfree_skb(hdev->sent_cmd); (FREE)
> if (hdev->sent_cmd) { (READ) |
> |
> kfree_skb(hdev->sent_cmd); (FREE) |
> | hdev->sent_cmd = skb_clone(skb,
> GFP_KERNEL); (WRITE) |
> |
> hdev->sent_cmd = NULL; (WRITE) |
> |
> -----
```

> If the FREE operation in thread-A is before WRITE operation in thread-B, > it can lead to double-free memory corruption in the kernel.

> ===== BUG EFFECTS =====

> For now, we can successfully trigger the vulnerability to corrupt the > kernel memory and thus crash the kernel. Although this bug is related to > Bluetooth device initialization, the attacker can trigger it without extra > privileges.

> That is because the Linux kernel does not ask for the privilege when > attaching the HCI device as the attached device is default set to > HCI_AUTO_OFF state. This bug is inside in the very first attaching > procedure and requires no syscalls.

```

> The crash log is presented below.
>
> =====
> [ 500.906562] hci0 type 1 len 3
> [ 500.904986] BUG: KASAN: use-after-free in kfree_skb+0x33/0x1c0
> [ 500.904986] Read of size 4 at addr ffff888009d3599c by task
> kworker/u5:0/54
> [ 500.904986] CPU: 0 PID: 54 Comm: kworker/u5:0 Not tainted 5.11.11+ #16
> [ 500.909997] Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS
> 1.13.0-lubuntul.1 04/01/2014
> [ 500.909997] Workqueue: hci0 hci_power_on
> [ 500.909997] Call Trace:
> [ 500.909997] dump_stack+0x16c/0x1be
> [ 500.924511] print_address_description+0x7b/0x3a0
> [ 500.924511] kasan_report+0x14e/0x200
> [ 500.924511] ? kfree_skb+0x33/0x1c0
> [ 500.924511] ? skb_queue_purge+0x193/0x1c0
> [ 500.924511] kasan_report+0x47/0x60
> [ 500.924511] ? skb_queue_purge+0x193/0x1c0
> [ 500.924511] check_memory_region+0x2e2/0x330
> [ 500.924511] kfree_skb+0x33/0x1c0
> [ 500.924511] hci_dev_do_open+0x1008/0x1570
> [ 500.924511] ? printk+0x62/0x83
> [ 500.924511] hci_power_on+0x183/0x580
> [ 500.924511] ? strcpy+0x7f/0x240
> [ 500.924511] process_one_work+0x722/0x1150
> [ 500.924511] worker_thread+0xb5c/0x17d0
> [ 500.924511] ? process_one_work+0x1150/0x1150
> [ 500.924511] kthread+0x2fc/0x320
> [ 500.924511] ? process_one_work+0x1150/0x1150
> [ 500.924511] ? kthread_unuse_mm+0x1d0/0x1d0
> [ 500.924511] ret_from_fork+0x22/0x30
> [ 500.924511]
> [ 500.924511] Allocated by task 273:
> [ 500.924511] kasan_kmalloc+0xc6/0x100
> [ 500.924511] kmem_cache_alloc+0xfe/0x1f0
> [ 500.924511] skb_clone+0x1b5/0x360
> [ 500.924511] hci_cmd_work+0x15d/0x350
> [ 500.924511] process_one_work+0x722/0x1150
> [ 500.924511] worker_thread+0xb5c/0x17d0
> [ 500.924511] kthread+0x2fc/0x320
> [ 500.924511] ret_from_fork+0x22/0x30
> [ 500.924511]
> [ 500.924511] Freed by task 273:
> [ 500.924511] kasan_set_track+0x3d/0x70
> [ 500.924511] kasan_set_free_info+0x1f/0x40
> [ 500.924511] kasan_slab_free+0x10e/0x140
> [ 500.924511] kmem_cache_free+0xca/0x210
> [ 500.924511] hci_cmd_work+0x150/0x350
> [ 500.924511] process_one_work+0x722/0x1150
> [ 500.924511] worker_thread+0xb5c/0x17d0
> [ 500.924511] kthread+0x2fc/0x320
> [ 500.924511] ret_from_fork+0x22/0x30
> [ 500.924511]
> [ 500.924511] The buggy address belongs to the object at ffff888009d358c0
> [ 500.924511] which belongs to the cache skbuff_head_cache of size 232
> [ 500.924511] The buggy address is located 220 bytes inside of
> [ 500.924511] 232-byte region [ffff888009d358c0, ffff888009d359a8)
> [ 500.924511] The buggy address belongs to the page:
> [ 500.924511] page:00000000b691648a refcount:1 mapcount:0
> mapping:0000000000000000 index:0x0 pfn:0x9d35
> [ 500.924511] flags: 0x1000000000000200 (slab)
> [ 500.924511] raw: 01000000000000200 dead000000000100 dead000000000122
> ffff888006a64040
> [ 500.924511] raw: 0000000000000000 00000000000c000c 00000001ffffffffff
> 0000000000000000
> [ 500.924511] page dumped because: kasan: bad access detected
> [ 500.924511]
> [ 500.924511] Memory state around the buggy address:
> [ 500.924511] ffff888009d35880: fc fc fc fc fc fc fc fa fb fb fb
> fb fb fb
> [ 500.924511] ffff888009d35900: fb fb fb fb fb fb fb fb fb fb fb
> fb fb fb
> [ 500.924511] >ffff888009d35980: fb fb fb fb fb fc fc fc fc fc fc fc
> fc fc fc
> [ 500.924511] ^
> [ 500.924511] ffff888009d35a00: fa fb fb fb fb fb fb fb fb fb fb fb
> fb fb fb
> [ 500.924511] ffff888009d35a80: fb fb fb fb fb fb fb fb fb fb fb fb
> fc fc fc
> [ 500.924511]
> =====
> [ 500.924511] Disabling lock debugging due to kernel taint
> [ 501.014277]
> =====
> [ 501.014929] BUG: KASAN: double-free or invalid-free in
> hci_dev_do_open+0x1008/0x1570
> [ 501.014929]
> [ 501.014929] CPU: 0 PID: 54 Comm: kworker/u5:0 Tainted: G B
> 5.11.11+ #16
> [ 501.014929] Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS
> 1.13.0-lubuntul.1 04/01/2014
> [ 501.014929] Workqueue: hci0 hci_power_on
> [ 501.014929] Call Trace:
> [ 501.014929] dump_stack+0x16c/0x1be
> [ 501.014929] ? hci_dev_do_open+0x1008/0x1570
> [ 501.014929] ? hci_dev_do_open+0x1008/0x1570
> [ 501.014929] print_address_description+0x7b/0x3a0
> [ 501.014929] ? hci_dev_do_open+0x1008/0x1570
> [ 501.014929] ? hci_dev_do_open+0x1008/0x1570
> [ 501.014929] kasan_report_invalid_free+0x54/0xd0
> [ 501.014929] kasan_slab_free+0xe7/0x140
> [ 501.014929] kmem_cache_free+0xca/0x210
> [ 501.014929] ? hci_dev_do_open+0x1008/0x1570
> [ 501.014929] hci_dev_do_open+0x1008/0x1570
> [ 501.014929] ? printk+0x62/0x83
> [ 501.014929] hci_power_on+0x183/0x580
> [ 501.014929] ? strcpy+0x7f/0x240
> [ 501.014929] process_one_work+0x722/0x1150
> [ 501.014929] worker_thread+0xb5c/0x17d0
> [ 501.014929] ? process_one_work+0x1150/0x1150
> [ 501.014929] kthread+0x2fc/0x320
> [ 501.014929] ? process_one_work+0x1150/0x1150
> [ 501.014929] ? kthread_unuse_mm+0x1d0/0x1d0
> [ 501.014929] ret_from_fork+0x22/0x30
> [ 501.014929]
> [ 501.014929] Allocated by task 273:
> [ 501.014929] kasan_kmalloc+0xc6/0x100
> [ 501.014929] kmem_cache_alloc+0xfe/0x1f0
> [ 501.014929] skb_clone+0x1b5/0x360
> [ 501.014929] hci_cmd_work+0x15d/0x350
> [ 501.014929] process_one_work+0x722/0x1150
> [ 501.014929] worker_thread+0xb5c/0x17d0
> [ 501.014929] kthread+0x2fc/0x320
> [ 501.014929] ret_from_fork+0x22/0x30
> [ 501.014929]
> [ 501.014929] Freed by task 273:
> [ 501.014929] kasan_set_track+0x3d/0x70
> [ 501.014929] kasan_set_free_info+0x1f/0x40
> [ 501.066803] kasan_slab_free+0x10e/0x140
> [ 501.066803] kmem_cache_free+0xca/0x210
> [ 501.066803] hci_cmd_work+0x150/0x350
> [ 501.066803] process_one_work+0x722/0x1150
> [ 501.066803] worker_thread+0xb5c/0x17d0
> [ 501.066803] kthread+0x2fc/0x320
> [ 501.066803] ret_from_fork+0x22/0x30
> [ 501.066803]
> [ 501.066803] The buggy address belongs to the object at ffff888009d358c0
> [ 501.066803] which belongs to the cache skbuff_head_cache of size 232

```

```

> [ 501.066803] The buggy address is located 0 bytes inside of
> [ 501.066803] 232-byte region [ffff888009d358c0, ffff888009d359a8)
> [ 501.066803] The buggy address belongs to the page:
> [ 501.066803] page:00000000b691648a refcount:1 mapcount:0
> mapping:0000000000000000 index:0x0 pfn:0x9d35
> [ 501.066803] flags: 0x1000000000000200 (slab)
> [ 501.066803] raw: 0100000000000200 dead000000000100 dead000000000122
> ffff888006d64640
> [ 501.066803] raw: 0000000000000000 00000000000c000c 00000001ffffffff
> 0000000000000000
> [ 501.066803] page dumped because: kasan: bad access detected
> [ 501.066803]
> [ 501.066803] Memory state around the buggy address:
> [ 501.066803] ffff888009d35780: fa fb fb fb fb fb fb fb fb fb
> fb fb fb
> [ 501.066803] ffff888009d35800: fb fb fb fb fb fb fb fb fb fb fb
> fc fc fc
> [ 501.066803] >ffff888009d35880: fc fc fc fc fc fc fc fa fb fb fb
> fb fb fb
> [ 501.066803]
> [ 501.066803] ffff888009d35900: fb fb fb fb fb fb fb fb fb fb fb
> fb fb fb
> [ 501.066803] ffff888009d35980: fb fb fb fb fb fc fc fc fc fc fc
> fc fc fc
> [ 501.066803]
> =====
>
> =====*====* Timeline *====*=====
>
> 2021-05-17: Bug reported to security () kernel.org and linux-distros () vs openwall.org
>
> 2021-05-25: CVE-2021-3564 assigned
>
> We informed security@...nel.org on May 17, 2021. Now the 7-day embargo period is over, we are being asked to bring the issue to public.
>
> Since our patch has not been applied to upstream yet, we will release the POC later.
>
> =====*====* Credit *====*=====
>
> HaoXiong@...ckSec Team
>
> LinMa@...cksec Team
>
> syzkaller
>
>
> Best regards.
>
> Mart111n
>

```

Content of type "text/html" skipped

Download attachment "[reproduce.zip](#)" of type "application/zip" (8680 bytes)

Powered by [blists](#) - [more mailing lists](#)

Please check out the [Open Source Software Security Wiki](#), which is counterpart to this [mailing list](#).

Confused about [mailing lists](#) and their use? [Read about mailing lists on Wikipedia](#) and check out these [guidelines on proper formatting of your messages](#).

