

[Products](#)[Services](#)[Publications](#)[Resources](#)[What's new](#)

[Hash Suite - Windows password security audit tool. GUI, reports in PDF.](#)

[\[<prev\]](#) [\[next>\]](#) [\[day\]](#) [\[month\]](#) [\[year\]](#) [\[list\]](#)

Date: Thu, 22 Jul 2021 14:28:05 +0800
From: Lin Horse <kylin.formalin@...il.com>
To: oss-security@...ts.openwall.com
Subject: CVE-2021-3640: Linux kernel: UAF in sco_send_frame function

Hello there,

Just like the previous, tedious race condition vulnerability caused by the unexpected locking behavior (CVE-2021-3573), a similar one is found this time.

===== BUG DETAILS =====

We can find another place that uses bh_lock_sock() in the Linux Bluetooth stacks.

```
static void sco_conn_del(struct hci_conn *hcon, int err)
{
    ...
    if (sk) {
        sock_hold(sk);
        bh_lock_sock(sk); // {1} LOCK
        sco_sock_clear_timer(sk);
        sco_chan_del(sk, err);
        bh_unlock_sock(sk); // {2} UNLOCK
        sco_sock_kill(sk);
        sock_put(sk);
    }
    ...
    hcon->sco_data = NULL;
    kfree(conn);
}
```

Between these lock pairs, sco_chan_del() is called, which will delete the channel associated with this sk. At the end of this function, the conn will be released by kfree().

Similar to the CVE-2021-3573, there is another thread that can be controlled by the attacker. It will wait for the kfree() and thereafter, race to cause UAF.

For example, the sco_sock_sendmsg() function.

```
static int sco_sock_sendmsg(struct socket *sock, struct msghdr *msg,
                           size_t len)
{
    ...
    lock_sock(sk);

    if (sk->sk_state == BT_CONNECTED)
        err = sco_send_frame(sk, msg, len);
    else
        err = -ENOTCONN;

    release_sock(sk);
    return err;
}

static int sco_send_frame(struct sock *sk, struct msghdr *msg, int len)
{
    ...

    skb = bt_skb_send_alloc(sk, len, msg->msg_flags & MSG_DONTWAIT, &err);
    if (!skb)
        return err;

    if (memcpy_from_msg(skb_put(skb, len), msg, len)) { // {3}
        kfree_skb(skb);
        return -EFAULT;
    }

    hci_send_sco(conn->hcon, skb);

    ...
}
```

As you can see, the attacker can adopt userfaultfd technique to stop the thread at {3} point.

Because the sco_send_frame() is protected by the lock_sock() and release_sock(), which will not block the sco_conn_del() to release the conn.

One vulnerable race window is shown below:

sco_sock_sendmsg thread	sco_conn_del thread
lock_sock(sk);	
...	
	bh_lock_sock(sk);
	...
	bh_unlock_sock(sk);
	...
	kfree(conn);
// UAF	
hci_send_sco(conn->hcon, skb);	

===== BUG EFFECTS =====

Similar to CVE-2021-3573, the attacker may stably cause the UAF and do further exploitation.

As the sco_conn struct is pretty juicy (two previous data pointers inside)

```
struct sco_conn {
    struct hci_conn *hcon;

    spinlock_t lock;
    struct sock *sk;

    unsigned int mtu;
};
```

The attacker can easily spray these kcalloc-32 objects with the malicious payload, with CAP_NET_ADMIN privilege.

The provided POC code can cause the crash report below:

```
[ 62.856933]
=====
[ 62.857336] BUG: KASAN: use-after-free in sco_sock_sendmsg+0x1d6/0x2c0
[ 62.858202] Read of size 8 at addr ffff888002478540 by task
poc.sco.new/120
[ 62.858663]
[ 62.859014] CPU: 0 PID: 120 Comm: poc.sco.new Not tainted 5.13.0+ #1
[ 62.859405] Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS
```

```
1.10.2-lubuntu1 04/01/2014
[ 62.859884] Call Trace:
[ 62.860168] dump_stack_lvl+0x73/0x9e
[ 62.860525] print_address_description+0x82/0x3a0
[ 62.860879] kasan_report+0x154/0x240
[ 62.861115] ? lock_sock_nested+0x100/0x140
[ 62.861446] ? sco_sock_sendmsg+0x1d6/0x2c0
[ 62.861811] kasan_report+0x45/0x60
[ 62.862133] sco_sock_sendmsg+0x1d6/0x2c0
[ 62.862461] ? sco_sock_getsockopt+0x410/0x410
[ 62.862748] ? inet_send_prepare+0x190/0x190
[ 62.863000] sock_write_iter+0x21b/0x230
[ 62.863232] vfs_write+0x53a/0x5c0
[ 62.863479] ksys_write+0x8b/0x100
[ 62.863723] ? __fpregs_load_activate+0xc2/0x150
[ 62.864017] do_syscall_64+0x43/0x90
[ 62.864287] entry_SYSCALL_64_after_hwframe+0x44/0xae
[ 62.864615] RIP: 0033:0x7f9b6c8d4abf
[ 62.865073] Code: 89 54 24 18 48 89 74 24 10 89 7c 24 08 e8 69 fd ff ff
48 8b 54 24 18 48 8b 74 24 10 41 89 c0 8b 7c 24 08 b8 01 00 00 0f 05
<48> 3d 00 f0 ff ff 77 2d 44 89 c7 48 89 44 24 08 e8 9c fd ff ff 48
[ 62.865843] RSP: 002b:00007ffdb6b0133a0 EFLAGS: 00000293 ORIG_RAX:
0000000000000001
[ 62.866304] RAX: ffffffff8800000000 RBX: 000055be494024e0 RCX:
00007f9b6c8d4abf
[ 62.866660] RDX: 0000000000000010 RSI: 00007f9b6c90e000 RDI:
0000000000000005
[ 62.866992] RBP: 00007ffdb6b013480 R08: 0000000000000000 R09:
00007f9b6c703700
[ 62.867293] R10: 00007f9b6c7039d0 R11: 0000000000000293 R12:
000055be49400d10
[ 62.867576] R13: 00007ffdb6b013570 R14: 0000000000000000 R15:
0000000000000000
[ 62.868106]
[ 62.868302] Allocated by task 120:
[ 62.868586] kasan_kmalloc+0xb5/0xe0
[ 62.868999] kmem_cache_alloc_trace+0x12d/0x210
[ 62.869349] sco_sock_connect+0x1f7/0x4a0
[ 62.869647] __sys_connect+0x16f/0x1a0
[ 62.869944] __x64_sys_connect+0x38/0x40
[ 62.870243] do_syscall_64+0x43/0x90
[ 62.870556] entry_SYSCALL_64_after_hwframe+0x44/0xae
[ 62.870883]
[ 62.871020] Freed by task 125:
[ 62.871192] kasan_set_track+0x3d/0x70
[ 62.871432] kasan_set_free_info+0x1f/0x40
[ 62.871708] kasan_slab_free+0x111/0x150
[ 62.871956] kfree+0xf3/0x2d0
[ 62.872208] hci_conn_hash_flush+0xbf/0x120
[ 62.872529] hci_dev_do_close+0x51a/0x870
[ 62.872789] hci_unregister_dev+0x23a/0xb70
[ 62.873054] vhci_release+0x3f/0x70
[ 62.873334] fput+0x197/0x360
[ 62.873598] task_work_run+0xc0/0xe0
[ 62.873919] exit_to_user_mode_prepare+0xf0/0x130
[ 62.874253] syscall_exit_to_user_mode+0x20/0x40
[ 62.874511] do_syscall_64+0x52/0x90
[ 62.874768] entry_SYSCALL_64_after_hwframe+0x44/0xae
[ 62.875160]
[ 62.875352] The buggy address belongs to the object at ffff888002478540
[ 62.875352] which belongs to the cache kmalloc-32 of size 32
[ 62.875900] The buggy address is located 0 bytes inside of
[ 62.875900] 32-byte region [ffff888002478540, ffff888002478560)
[ 62.876472] The buggy address belongs to the page:
[ 62.876885] page:00000000b13206d refcount:1 mapcount:0
mapping:0000000000000000 index:0x0 pfn:0x2478
[ 62.877481] flags: 0x100000000000200 (slab,node=0|zone=1)
[ 62.878361] raw: 01000000000000200 ffffea0000078d00 0000000e0000000e
ffff888001041500
[ 62.878901] raw: 0000000000000000 0000000080400040 00000001fffffff
0000000000000000
[ 62.879313] page dumped because: kasan: bad access detected
[ 62.879588]
[ 62.879704] Memory state around the buggy address:
[ 62.880003] ffff888002478400: fb fb fb fb fc fc fb fb fb fb fc fc
fc fc
[ 62.880286] ffff888002478480: fb fb fb fb fc fc fc fb fb fb fb fc fc
fc fc
[ 62.880532] >ffff888002478500: fa fb fb fb fc fc fc fa fb fb fb fc fc
fc fc
[ 62.880785] ^
[ 62.881199] ffff888002478580: 00 00 00 00 fc fc fc 00 00 00 fc fc fc
fc fc
[ 62.881457] ffff888002478600: 00 00 00 fc fc fc fc 00 00 00 fc fc fc
fc fc
[ 62.881716]
=====
[ 62.881991] Disabling lock debugging due to kernel taint
[ 62.883072] BUG: unable to handle page fault for address:
ffffbfef22fa79f
[ 62.883427] #PF: supervisor read access in kernel mode
[ 62.883774] #PF: error code(0x0000) - not-present page
[ 62.884165] PGD 36fd0067 P4D 36fd0067 PUD 36fd4067 PMD 0
[ 62.884827] Oops: 0000 [#1] SMP KASAN NOPTI
[ 62.885132] CPU: 0 PID: 120 Comm: poc.sco.new Tainted: G B
5.13.0+ #1
[ 62.885528] Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS
1.10.2-lubuntu1 04/01/2014
[ 62.885901] RIP: 0010: kasan_store8+0x6c/0xb0
[ 62.886184] Code: be 00 00 00 00 00 ff df 0f be 14 32 85 d2 74 07 83
e0 07 39 d0 7d 29 c3 48 89 fe 48 c1 ee 03 48 ba 00 00 00 00 ff df
<80> 3c 16 00 75 11 48 89 c6 48 c1 ee 03 0f be 14 16 85 d2 75 d2 eb
[ 62.886853] RSP: 0018:ffff8880030fbfb8 EFLAGS: 00000006
[ 62.887244] RAX: ffffffff917d3d02 RBX: 0000000000040000 RCX:
ffff88800337d86
[ 62.887524] RDX: dffffc0000000000 RSI: 1fffffff22fa79f RDI:
ffff88800317d3c3b
[ 62.887855] RBP: 0000000000000030 R08: dffffc0000000000 R09:
0000000000000007
[ 62.888162] R10: ffffd100035159c R11: 00000000000000fb R12:
ffff88800317a1b4b
[ 62.888476] R13: ffffffff8800000000 R14: ffff888001a8acdc R15:
ffff888036432188
[ 62.888838] FS: 00007f9b6c704740 (0000) GS:ffff888036400000 (0000)
knlGS:0000000000000000
[ 62.889331] CS: 0010 DS: 0000 ES: 0000 CR0: 000000080050033
[ 62.889908] CR2: fffffbfff22fa79f CR3: 00000000011c0000 CR4:
00000000003006f0
[ 62.890341] Call Trace:
[ 62.890617] queued_spin_lock_slowpath+0x286/0x410
[ 62.890915] raw_spin_lock_irqsave+0x9f/0xb0
[ 62.891201] skb_queue_tail+0x1c/0x90
[ 62.891548] hci_send_sco+0xd6/0x110
[ 62.891871] sco_sock_sendmsg+0x1e1/0x2c0
[ 62.892170] ? sco_sock_getsockopt+0x410/0x410
[ 62.892511] ? inet_send_prepare+0x190/0x190
[ 62.892796] sock_write_iter+0x21b/0x230
[ 62.893156] vfs_write+0x53a/0x5c0
[ 62.893533] ksys_write+0x8b/0x100
[ 62.893870] ? __fpregs_load_activate+0xc2/0x150
[ 62.894258] do_syscall_64+0x43/0x90
[ 62.894523] entry_SYSCALL_64_after_hwframe+0x44/0xae
[ 62.894929] RIP: 0033:0x7f9b6c8d4abf
[ 62.895178] Code: 89 54 24 18 48 89 74 24 10 89 7c 24 08 e8 69 fd ff ff
48 8b 54 24 18 48 8b 74 24 10 41 89 c0 8b 7c 24 08 b8 01 00 00 0f 05
<48> 3d 00 f0 ff ff 77 2d 44 89 c7 48 89 44 24 08 e8 9c fd ff ff 48
[ 62.895930] RSP: 002b:00007ffdb6b0133a0 EFLAGS: 00000293 ORIG_RAX:
0000000000000001
[ 62.896396] RAX: ffffffff8800000000 RBX: 000055be494024e0 RCX:
00007f9b6c8d4abf
```

```

==*==*==*==*==*==*== BUG REPRODUCE ==*==*==*==*==*==*==

```

The attacker has to fake an SCO connection and then calls `sco_sock_sendmsg()` with the expected controllable faulting page. After that, the attacker just needs to detach the controller to call `sco_conn_del()`.

```
hci_unregister_dev() -> hci_dev_do_close() -> hci_conn_hash_flush() ->
hci_disconn_cfm() ->
sco_disconn_cfm() -> sco_conn_del().
```

You can refer to the provided POC code for the details.

***** Timeline *****

```
2021-07-08: Bug reported to security@...nel.org and
linux-distros@...openwall.org
2021-07-09: CVE-2021-3640 is assigned
2021-07-22: 14 days of the embargo is over
```

One sad thing is that the bluez team is currently focused on fixing up the CVE-2021-3573, which I failed to properly patched, and the patch for this new is not yet fully discussed. I hope the patch will be settled down and merged to the mainline in the near future.

```

===== Credit =====
LinMa@...ckSec Team

```

Best Regards

Content of type "text/html" skipped

Download attachment "[reproduce.tar](#)" of type "application/x-tar" (51200 bytes)

Powered by blists - more mailing lists

Please check out the [Open Source Software Security Wiki](#), which is counterpart to this [mailing list](#).

Confused about [mailing lists](#) and their use? [Read about mailing lists on Wikipedia](#) and check out these [guidelines on proper formatting of your messages](#).

