

[Products](#)[Services](#)[Publications](#)[Resources](#)[What's new](#)

Follow @Openwall on Twitter for new release announcements and other news

[\[<prev\]](#) [\[next>\]](#) [\[day\]](#) [\[month\]](#) [\[year\]](#) [\[list\]](#)

Date: Mon, 24 Jan 2022 14:05:01 +0000
From: Qualys Security Advisory <gsa@...lys.com>
To: "oss-security@...ts.openwall.com" <oss-security@...ts.openwall.com>
Subject: CVE-2021-3998 and CVE-2021-3999 in glibc's realpath() and getcwd()

Hi all,

We discovered two vulnerabilities in the glibc, CVE-2021-3998 in realpath() and CVE-2021-3999 in getcwd(). Patches are now available at (many thanks to Siddhesh Foyarekar and Red Hat Product Security):

<https://sourceware.org/git/gitweb.cgi?p=glibc.git;h=ee8d5e33adb284601c00c94687bc907e10aec9bb>
<https://sourceware.org/git/gitweb.cgi?p=glibc.git;h=f7a79879c0b2bef0dadd6caaaeb0d26423e04e5>

<https://sourceware.org/git/gitweb.cgi?p=glibc.git;h=23e0e8f5f1fb5ed150253d986eccdc90c2dcd5e>
<https://sourceware.org/git/gitweb.cgi?p=glibc.git;h=472e799a5f2102bc0c3206dbd5a801765fceb39c>

Below is a short write-up (which is part of a longer advisory that is mostly unrelated to the glibc and that we will publish at a later date):

=====

CVE-2021-3998: Unexpected return value from glibc's realpath()

=====

While auditing umount and fusermount, we also discovered a vulnerability in the glibc's realpath() function, which is used internally by various programs. Normally, when the output buffer "resolved" that is passed to realpath() is not NULL, then realpath() either returns NULL on failure, or it returns the output buffer "resolved" on success. Unfortunately, since commit c6e0b0b ("stdlib: Sync canonicalize with gnu lib") from January 2021, realpath() can mistakenly return a malloc()ated buffer that is neither NULL nor the output buffer "resolved":

```
-----
430 char *
431 __realpath (const char *name, char *resolved)
432 {
433     ...
437     struct scratch_buffer rname_buffer;
438     return realpath_stk (name, resolved, &rname_buffer);
439 }
-----
197 static char *
198 realpath_stk (const char *name, char *resolved,
199               struct scratch_buffer *rname_buf)
200 {
201     ...
399     failed = false;
202     ...
403     if (resolved != NULL && dest - rname <= get_path_max ())
404         rname = strcpy (resolved, rname);
205     ...
410     if (failed || rname == resolved)
411     {
412         scratch_buffer_free (rname_buf);
413         return failed ? NULL : resolved;
414     }
415
416     return scratch_buffer_dupfree (rname_buf, dest - rname);
417 }
-----
```

For example, if the input path "name" is "." and if the current working directory is longer than PATH_MAX, then:

- at line 399, "failed" is set to false;
- at lines 403-404, "rname" is NOT set to "resolved" and "resolved" is left untouched and uninitialized (because "dest - rname" is longer than PATH_MAX);
- the code block at lines 410-414 is skipped (because "failed" is false and "rname" is not "resolved");
- at line 416, scratch_buffer_dupfree() returns a malloc()ated buffer that is NOT the output buffer "resolved".

The consequences of this vulnerability depend on the affected programs; for example, fusermount (a SUID-root program) can disclose sensitive information (pointers) when displaying the contents of a stack-based buffer that is mistakenly left uninitialized by realpath() (we tested this proof of concept on Ubuntu 21.04):

```
-----
$ gcc -o CVE-2021-3998-fusermount CVE-2021-3998-fusermount.c
$ ./CVE-2021-3998-fusermount > CVE-2021-3998-fusermount.output
...

$ hexdump -C CVE-2021-3998-fusermount.output
00000000 2f 75 73 72 2f 62 69 6e 2f 66 75 73 65 72 6d 6f |/usr/bin/fusermo|
00000010 75 6e 74 3a 20 65 6e 74 72 79 20 66 6f 72 20 f0 |unt: entry for .|
00000020 83 9b 99 ff 7f 20 6e 6f 74 20 66 6f 75 6e 64 20 |..... not found|
00000030 69 6e 20 2f 65 74 63 2f 6d 74 61 62 0a 0a 2f 75 |in /etc/mtab../u|
00000040 73 72 2f 62 69 6e 2f 66 75 73 65 72 6d 6f 75 6e |sr/bin/fusermount|
00000050 74 3a 20 65 6e 74 72 79 20 66 6f 72 20 39 ac b7 |t: entry for 9..|
00000060 a5 a2 7f 20 6e 6f 74 20 66 6f 75 6e 64 20 69 6e |... not found in|
00000070 20 2f 65 74 63 2f 6d 74 61 62 0a 0a | /etc/mtab..|
-----
```

=====

CVE-2021-3999: Off-by-one buffer overflow/underflow in glibc's getcwd()

=====

While studying the vulnerability in realpath(), we also discovered a vulnerability in the glibc's getcwd() function (which is used internally by realpath() to resolve relative pathnames) -- an off-by-one buffer overflow and underflow, but if and only if the "size" of "buf" is exactly 1:

```
-----
48 __getcwd (char *buf, size_t size)
49 {
50     ...
54     size_t alloc_size = size;
55     ...
76     path = buf;
77     ...
80     retval = INLINE_SYSCALL (getcwd, 2, path, alloc_size);
81     ...
100     if (retval >= 0 || errno == ENAMETOOLONG)
101     {
102         ...
110         result = __getcwd_generic (path, size);
111     }
112     ...
158 __getcwd_generic (char *buf, size_t size)
159 {
160     ...
-----
```

```

187  size_t allocated = size;
...
247  dir = buf;
248
249  dirp = dir + allocated;
250  *--dirp = '\0';
...
262  while (!(thisdev == rootdev && thisino == rootino))
263  {
...
441  }
...
449  if (dirp == &dir[allocated - 1])
450  *--dirp = '/';
...
457  used = dir + allocated - dirp;
458  memmove (dir, dirp, used);

```

If, at line 48, the "size" of "buf" is exactly 1:

- and if, at line 80, the kernel's getcwd() syscall fails with the error ENAMETOOLONG (because the current working directory is longer than PATH_MAX),
- then, at line 110, a generic implementation of getcwd() is called;
- at line 250, a null byte is written to "dirp", which points exactly to "buf" (because "size", and hence "allocated", are exactly 1);
- if the code block at lines 262-441 is skipped entirely (if the current working directory corresponds to the "/" directory),
- then, at lines 449-450, a slash is written to "buf+1" (an off-by-one buffer underflow, because at line 449 "dirp" was still pointing exactly to "buf"),
- and, at lines 457-458, a null byte is written to "buf+1" (an off-by-one buffer overflow, because at line 457 "used" is exactly 2).

It may seem impossible to satisfy the condition at line 100 (the current working directory is longer than PATH_MAX) and the condition at line 262 (the current working directory corresponds to the "/" directory), but in reality we can:

- in a child process:
 - create an unprivileged mount namespace;
 - create a directory longer than PATH_MAX;
 - bind-mount "/" onto this directory;
 - open() this directory and send its file descriptor to the parent process (outside the unprivileged mount namespace);
- in the parent process:
 - receive the file descriptor of this directory (which corresponds to "/" and is longer than PATH_MAX) and fchdir() to it;
 - execute a SUID program that calls getcwd() with a buffer of size 1, which triggers the off-by-one buffer overflow and underflow.

Apparently, this vulnerability was introduced in February 1995 by the very first commit in the glibc's git history (28f540f, "initial import") and could be triggered without an unprivileged mount namespace, by simply chdir()ing to the "/" directory:

```

190 getcwd (buf, size)
...
218  path = buf;
...
226  pathp = path + size;
227  *--pathp = '\0';
...
242  while (!(thisdev == rootdev && thisino == rootino))
243  {
...
351  }
352
353  if (pathp == &path[size - 1])
354  *--pathp = '/';
...
359  memmove (path, pathp, path + size - pathp);

```

Although "the size of buf is exactly 1" is a strong requirement, vulnerable code like the following may exist in the wild:

```

#include <unistd.h>
#include <stdio.h>

int main(int argc, char * argv[]) {
    char buf[4096];
    int len = snprintf(buf, sizeof(buf), "%s: cwd is ", argv[0]);
    if (len <= 0 || (unsigned)len >= sizeof(buf)) return _LINE__;
    if (!getcwd(buf + len, sizeof(buf) - len)) return _LINE__;
    puts(buf);
    return 0;
}

```

Thank you very much! We are at your disposal for questions, comments, and further discussions.

With best regards,

--
the Qualys Security Advisory team

[Powered by blists](#) - [more mailing lists](#)

Please check out the [Open Source Software Security Wiki](#), which is counterpart to this [mailing list](#).

Confused about [mailing lists](#) and their use? Read about [mailing lists on Wikipedia](#) and check out these [guidelines on proper formatting of your messages](#).

