

## ✓ s3v4: read and verify S3 signature v4 chunks separately (#11801)

[Browse files](#)

This commit fixes a security issue in the signature v4 chunked reader. Before, the reader returned unverified data to the caller and would only verify the chunk signature once it has encountered the end of the chunk payload.

Now, the chunk reader reads the entire chunk into an in-memory buffer, verifies the signature and then returns data to the caller.

In general, this is a common security problem. We verifying data streams, the verifier MUST NOT return data to the upper layers / its callers as long as it has not verified the current data chunk / data segment:

```
...
func (r *Reader) Read(buffer []byte) {
    if err := r.readNext(r.internalBuffer); err != nil {
        return err
    }
    if err := r.verify(r.internalBuffer); err != nil {
        return err
    }
    copy(buffer, r.internalBuffer)
}
...
```

👤 master (#11801)

🕒 RELEASE.2022-12-12T19:27-27Z ... RELEASE.2021-03-17T02-33-02Z

👤 aead committed on Mar 16, 2021 parent 980311f commit e197800f9055489415b53cf137e31e194aaf7ba0

Showing 1 changed file with 165 additions and 132 deletions.

[Split](#) [Unified](#)

```
cmd/streaming-signature-v4.go
166 166         seedDate:      seedDate,
167 167         region:        region,
168 168         chunkSHA256Writer: sha256.New(),
169 -         state:         readChunkHeader,
169 +         buffer:        make([]byte, 64*1024),
170 170     }, ErrNone
171 171 }
172 172
173 173 // Represents the overall state that is required for decoding a
174 174 // AWS Signature V4 chunked reader.
175 175 type s3ChunkedReader struct {
176 -     reader      *bufio.Reader
177 -     cred         auth.Credentials
178 -     seedSignature string
179 -     seedDate     time.Time
180 -     region       string
181 -     state        chunkState
182 -     lastChunk    bool
183 -     chunkSignature string
176 +     reader      *bufio.Reader
177 +     cred         auth.Credentials
178 +     seedSignature string
179 +     seedDate     time.Time
180 +     region       string
181 +
182 +     chunkSHA256Writer hash.Hash // Calculates sha256 of chunk data.
183 +     n                  uint64 // Unread bytes in chunk
184 +     buffer             []byte
185 +     offset             int
186 +     err                error
187 }
188
189 // Read chunk reads the chunk token signature portion.
190 func (cr *s3ChunkedReader) readS3ChunkHeader() {
191     // Read the first chunk line until CRLF.
192     var hexChunkSize, hexChunkSignature []byte
193     hexChunkSize, hexChunkSignature, cr.err = readChunkLine(cr.reader)
194     if cr.err != nil {
195         return
196     }
197     // <hex>;token=value - converts the hex into its uint64 form.
198     cr.n, cr.err = parseHexUint(hexChunkSize)
199     if cr.err != nil {
200         return
201     }
202     if cr.n == 0 {
203         cr.err = io.EOF
204     }
205     // Save the incoming chunk signature.
206     cr.chunkSignature = string(hexChunkSignature)
207 }
208
209 type chunkState int
210
211 const (
212     readChunkHeader chunkState = iota
```

```

213     readChunkTrailer
214     readChunk
215     verifyChunk
216     eofChunk
217 }
218
219 func (cs chunkState) String() string {
220     stateString := ""
221     switch cs {
222     case readChunkHeader:
223         stateString = "readChunkHeader"
224     case readChunkTrailer:
225         stateString = "readChunkTrailer"
226     case readChunk:
227         stateString = "readChunk"
228     case verifyChunk:
229         stateString = "verifyChunk"
230     case eofChunk:
231         stateString = "eofChunk"
232     }
233     return stateString
234 }
235
236
237 func (cr *s3ChunkedReader) Close() (err error) {
238     return nil
239 }
240
241 // Read implements 'io.Reader', which transparently decodes
242 // the incoming AWS Signature V4 streaming signature.
243 func (cr *s3ChunkedReader) Read(buf []byte) (n int, err error) {
244     // First, if there is any unread data, copy it to the client
245     // provided buffer.
246     if cr.offset > 0 {
247         n = copy(buf, cr.buffer[cr.offset:])
248         if n == len(buf) {
249             cr.offset += n
250             return n, nil
251         }
252     }
253     cr.offset = 0
254     buf = buf[n:]
255 }
256
257 // Now, we read one chunk from the underlying reader.
258 // A chunk has the following format:
259 // <chunk-size-as-hex> + ";chunk-signature=" + <signature-as-hex> + "\r\n" + <payload> + "\r\n"
260 //
261 // First, we read the chunk size but fail if it is larger
262 // than 1 MB. We must not accept arbitrary large chunks.
263 // One 1 MB is a reasonable max limit.
264 //
265 // Then we read the signature and payload data. We compute the SHA256 checksum
266 // of the payload and verify that it matches the expected signature value.
267 //
268 // The last chunk is *always* 0-sized. So, we must only return io.EOF if we have encountered
269 // a chunk with a chunk size = 0. However, this chunk still has a signature and we must
270 // verify it.
271 const MaxSize = 1 << 20 // 1 MB
272 var size int
273 for {
274     switch cr.state {
275     case readChunkHeader:
276         cr.readS3ChunkHeader()
277         // If we're at the end of a chunk.
278         if cr.n == 0 && cr.err == io.EOF {
279             cr.state = readChunkTrailer
280             cr.lastChunk = true
281             continue
282         }
283         if cr.err != nil {
284             return 0, cr.err
285         }
286         cr.state = readChunk
287     case readChunkTrailer:
288         cr.err = readCRLF(cr.reader)
289         if cr.err != nil {
290             return 0, errMalformedEncoding
291         }
292         cr.state = verifyChunk
293     case readChunk:
294         // There is no more space left in the request buffer.
295         if len(buf) == 0 {
296             return n, nil
297         }
298         rbuf := buf
299         // The request buffer is larger than the current chunk size.
300         // Read only the current chunk from the underlying reader.
301         if uint64(len(rbuf)) > cr.n {
302             rbuf = rbuf[:cr.n]
303         }
304         var n0 int
305         n0, cr.err = cr.reader.Read(rbuf)
306         if cr.err != nil {
307             // We have lesser than chunk size advertised in chunkHeader, this is 'unexpected'.
308             if cr.err == io.EOF {
309                 cr.err = io.ErrUnexpectedEOF
310             }
311             return 0, cr.err
312         }

```

```

283 -         }
284 -
285 -         // Calculate sha256.
286 -         cr.chunkSHA256Writer.Write(rbuf[:n0])
287 -         // Update the bytes read into request buffer so far.
288 -         n += n0
289 -         buf = buf[n0:]
290 -         // Update bytes to be read of the current chunk before verifying chunk's signature.
291 -         cr.n -= uint64(n0)
292 -
293 -         // If we're at the end of a chunk.
294 -         if cr.n == 0 {
295 -             cr.state = readChunkTrailer
296 -             continue
297 -         }
298 -         case verifyChunk:
299 -             // Calculate the hashed chunk.
300 -             hashedChunk := hex.EncodeToString(cr.chunkSHA256Writer.Sum(nil))
301 -             // Calculate the chunk signature.
302 -             newSignature := getChunkSignature(cr.cred, cr.seedSignature, cr.region, cr.seedDate, hashedChunk)
303 -             if !compareSignatureV4(cr.chunkSignature, newSignature) {
304 -                 // Chunk signature doesn't match we return signature does not match.
305 -                 cr.err = errSignatureMismatch
306 -                 return 0, cr.err
307 -             }
308 -             // Newly calculated signature becomes the seed for the next chunk
309 -             // this follows the chaining.
310 -             cr.seedSignature = newSignature
311 -             cr.chunkSHA256Writer.Reset()
312 -             if cr.lastChunk {
313 -                 cr.state = eofChunk
314 -             } else {
315 -                 cr.state = readChunkHeader
316 -             }
317 -         case eofChunk:
318 -             return n, io.EOF
224 + b, err := cr.reader.ReadByte()
225 + if err == io.EOF {
226 +     err = io.ErrUnexpectedEOF
227 + }
319 + if err != nil {
228 +     cr.err = err
229 +     return n, cr.err
230 + }
231 + }
232 + if b == ';' { // separating character
233 +     break
234 + }
235 +
236 + // Manually deserialize the size since AWS specified
237 + // the chunk size to be of variable width. In particular,
238 + // a size of 16 is encoded as '10' while a size of 64 KB
239 + // is '10000'.
240 + switch {
241 + case b >= '0' && b <= '9':
242 +     size = size<<4 | int(b-'0')
243 + case b >= 'a' && b <= 'f':
244 +     size = size<<4 | int(b-('a'-10))
245 + case b >= 'A' && b <= 'F':
246 +     size = size<<4 | int(b-('A'-10))
247 + default:
248 +     cr.err = errMalformedEncoding
249 +     return n, cr.err
250 + }
251 + if size > MaxSize {
252 +     cr.err = errMalformedEncoding
253 +     return n, cr.err
254 + }
255 + }
256 +
257 + // Now, we read the signature of the following payload and expect:
258 + // chunk-signature=" + <signature-as-hex> + "\r\n"
259 + //
260 + // The signature is 64 bytes long (hex-encoded SHA256 hash) and
261 + // starts with a 16 byte header: len("chunk-signature=") + 64 == 80.
262 + var signature [80]byte
263 + _, err = io.ReadFull(cr.reader, signature[:])
264 + if err == io.EOF {
265 +     err = io.ErrUnexpectedEOF
320 + }
267 + if err != nil {
268 +     cr.err = err
269 +     return n, cr.err
270 + }
271 + if !bytes.HasPrefix(signature[:], []byte("chunk-signature=")) {
272 +     cr.err = errMalformedEncoding
273 +     return n, cr.err
274 + }
275 + b, err := cr.reader.ReadByte()
276 + if err == io.EOF {
277 +     err = io.ErrUnexpectedEOF
278 + }
279 + if err != nil {
280 +     cr.err = err
281 +     return n, cr.err
282 + }
283 + if b != '\r' {
284 +     cr.err = errMalformedEncoding
285 +     return n, cr.err

```

```

286 +     }
287 +     b, err = cr.reader.ReadByte()
288 +     if err == io.EOF {
289 +         err = io.ErrUnexpectedEOF
290 +     }
291 +     if err != nil {
292 +         cr.err = err
293 +         return n, cr.err
294 +     }
295 +     if b != '\n' {
296 +         cr.err = errMalformedEncoding
297 +         return n, cr.err
298 +     }
299 +
300 +     if cap(cr.buffer) < size {
301 +         cr.buffer = make([]byte, size)
302 +     } else {
303 +         cr.buffer = cr.buffer[:size]
304 +     }
305 +
306 +     // Now, we read the payload and compute its SHA-256 hash.
307 +     _, err = io.ReadFull(cr.reader, cr.buffer)
308 +     if err == io.EOF && size != 0 {
309 +         err = io.ErrUnexpectedEOF
310 +     }
311 +     if err != nil && err != io.EOF {
312 +         cr.err = err
313 +         return n, cr.err
314 +     }
315 +     b, err = cr.reader.ReadByte()
316 +     if b != '\r' {
317 +         cr.err = errMalformedEncoding
318 +         return n, cr.err
319 +     }
320 +     b, err = cr.reader.ReadByte()
321 +     if err == io.EOF {
322 +         err = io.ErrUnexpectedEOF
323 +     }
324 +     if err != nil {
325 +         cr.err = err
326 +         return n, cr.err
327 +     }
328 +     if b != '\n' {
329 +         cr.err = errMalformedEncoding
330 +         return n, cr.err
331 +     }
332 +
333 +     // Once we have read the entire chunk successfully, we verify
334 +     // that the received signature matches our computed signature.
335 +     cr.chunkSHA256Writer.Write(cr.buffer)
336 +     newSignature := getChunkSignature(cr.cred, cr.seedSignature, cr.region, cr.seedDate, hex.EncodeToString(cr.chunkSHA256Writer.Sum(nil)))
337 +     if !compareSignatureV4(string(signature[16:]), newSignature) {
338 +         cr.err = errSignatureMismatch
339 +         return n, cr.err
340 +     }
341 +     cr.seedSignature = newSignature
342 +     cr.chunkSHA256Writer.Reset()
343 +
344 +     // If the chunk size is zero we return io.EOF. As specified by AWS,
345 +     // only the last chunk is zero-sized.
346 +     if size == 0 {
347 +         cr.err = io.EOF
348 +         return n, cr.err
349 +     }
350 +
351 +     cr.offset = copy(buf, cr.buffer)
352 +     n += cr.offset
353 +     return n, err
354 + }
355 +
356 + // readCRLF - check if reader only has '\r\n' CRLF character.

```

0 comments on commit [e197800](#)

Please [sign in](#) to comment.