

Talos Vulnerability Report

TALOS-2021-1243

Linux Kernel Arm SIGPAGE information disclosure vulnerability

MAY 28, 2021

CVE NUMBER

CVE-2021-21781

SUMMARY

An information disclosure vulnerability exists in the ARM SIGPAGE functionality of Linux Kernel v5.4.66 and v5.4.54. The latest version (5.11-rc4) seems to still be vulnerable. A userland application can read the contents of the sigpage, which can leak kernel memory contents. An attacker can read a process's memory at a specific offset to trigger this vulnerability.

CONFIRMED VULNERABLE VERSIONS

The versions below were either tested or verified to be vulnerable by Talos or confirmed to be vulnerable by the vendor.

Linux Kernel v5.4.54

Linux Kernel v5.4.66

PRODUCT URLS

Kernel - <https://github.com/torvalds/linux>

CVSSV3 SCORE

4.0 - CVSS:3.0/AV:L/AC:L/PR:N/UI:N/S:U/C:L/I:N/A:N

CWE

CWE-908 - Use of Uninitialized Resource

DETAILS

The Linux Kernel is the free and open-source core of Unix-like operating systems.

When examining a given Linux process' virtual memory space on an ARMv7 processor, the quickest ways to look are `info proc map` in gdb, and also by reading a process' `/proc/self/maps` file. The output thereof might look approximately like so:

```
[^.^]> info proc map
process 37
Mapped address spaces:

   Start Addr   End Addr       Size     Offset objfile
   0xbbee54000 0xbbee7c000     0x28000      0x0 /lib/libgcc_s.so.1
   0xbbee7c000 0xbbee7d000     0x1000      0x18000 /lib/libgcc_s.so.1
   0xbbee7d000 0xbbeedb000     0x5e000      0x0 /usr/lib/libc.so
   0xbbeea0000 0xbbeec0000     0x2000      0x5d000 /usr/lib/libc.so
   0xbbeec0000 0xbbeed0000     0x1000      0x0
   0xbbeed0000 0xbbeee0000     0x1000      0x0 /mnt/apps/[...]/app
   0xbbeefd000 0xbbeefe000     0x1000      0x0 /mnt/apps/[...]/app
   0xbbeefe000 0xbbeeff000     0x1000      0x1000 /mnt/apps/[...]/app
   0xbbe2000 0xbbe23000     0x1000      0x0 [sigpage] // [1]
   0xbbe25000 0xbbe26000     0x1000      0x0 [stack]
   0xbbe26000 0xbbe27000     0x1000      0x0 [vectors]
```

Nothing really out of place, but let us examine the `[sigpage]` segment of memory at `[1]`, which is used for storing signal handler information in userland:

```
[o.o]> x/200wx 0xbbe2000
0xbbe2000: 0x00000005 0x00000000 0x00000000 0xa649a76f
0xbbe2010: 0x00000000 0x00000000 0x00000000 0x00000000
0xbbe2020: 0x00000000 0xa64868c9 0xa649a7c5 0xa64c48f7
0xbbe2030: 0x00000000 0x00000000 0x00000000 0x00000002
0xbbe2040: 0x00000000 0x00000000 0x00000000 0x00000000
0xbbe2050: 0xffffffff 0x00000000 0x00000000 0x00000000
0xbbe2060: 0x00000000 0x00000000 0x00000000 0x00000000
0xbbe2070: 0x00000000 0x00000000 0x00000000 0x00000000
0xbbe2080: 0x00000000 0x00000000 0x00000000 0x00000000
0xbbe2090: 0x00000009 0x00000000 0x00000000 0xa649a76f
0xbbe20a0: 0x00000000 0x00000000 0x00000000 0x00000000
0xbbe20b0: 0xa6486819 0x00000000 0xa649a7c5 0xa64c48ff
0xbbe20c0: 0x00000100 0x00000000 0x00000000 0x00000000
0xbbe20d0: 0x00000000 0x00000000 0x00000000 0x00000000
0xbbe20e0: 0x00000000 0x00000000 0x00000000 0x00000000
0xbbe20f0: 0x00000000 0x00000000 0x00000000 0x00000000
0xbbe2100: 0x00000000 0x00000000 0x00000000 0x00000000
0xbbe2110: 0x00000000 0x00000000 0x00000000 0x00000000
0xbbe2120: 0x00000005 0x00000000 0x00000000 0xa649a76f
0xbbe2130: 0xa64c4b07 0xa64c4a07 0x00000000 0xa64c4a07
0xbbe2140: 0x00000000 0xa64868c9 0xa649a7c5 0xa64c4a07
// [...]
```

While this above output does not have any clear structure or meaning to it, we can examine the kernel code to explain exactly what we are seeing. The initialization of this page occurs within `arch_setup_additional_pages` of `arch/arm/kernel/process.c`:

```
static struct page *signal_page;
extern struct page *get_signal_page(void);

int arch_setup_additional_pages(struct linux_binprm *bprm, int uses_interp)
{
    struct mm_struct *mm = current->mm;
    struct vm_area_struct *vma;
    unsigned long npages;
    unsigned long addr;
    unsigned long hint;
    int ret = 0;

    if (!signal_page)
        signal_page = get_signal_page(); // [1]
    if (!signal_page)
        return -ENOMEM;

    // [...]
```

Since `static struct page *signal_page` is static, the page can only be initialized once, which occurs at [1] with `get_signal_page()`: The pointer to `sigpage` is assigned to static `struct page *signal_page` at [1], via function `get_signal_page()`:

```
struct page *get_signal_page(void)
{
    unsigned long ptr;
    unsigned offset;
    struct page *page;
    void *addr;

    page = alloc_pages(GFP_KERNEL, 0); // [1]

    if (!page)
        return NULL;

    addr = page_address(page);

    /* Give the signal return code some randomness */
    offset = 0x200 + (get_random_int() & 0x7fc);
    signal_return_offset = offset;

    /*
     * Copy signal return handlers into the vector page, and
     * set sigreturn to be a pointer to these.
     */
    memcpy(addr + offset, sigreturn_codes, sizeof(sigreturn_codes)); // [2]

    ptr = (unsigned long)addr + offset;
    flush_icache_range(ptr, ptr + sizeof(sigreturn_codes));

    return page;
}
```

At [1], the buddy allocator grabs a single page of memory, and at [2], a set of instructions are copied from `extern const unsigned long sigreturn_codes[17]`; into a random spot inside of our `sigpage`. After this, nothing else of import materially occurs on the page, and it's returned back up to `arch_setup_additional_pages`. Continuing therein:

```
if (!signal_page)
    signal_page = get_signal_page(); // [1]
if (!signal_page)
    return -ENOMEM;

npages = 1; /* for sigpage */
npages += vdso_total_pages;

if (down_write_killable(&mm->mmap_sem))
    return -EINTR;
hint = sigpage_addr(mm, npages);
addr = get_unmapped_area(NULL, hint, npages << PAGE_SHIFT, 0, 0);
if (IS_ERR_VALUE(addr)) {
    ret = addr;
    goto up_fail;
}

vma = _install_special_mapping(mm, addr, PAGE_SIZE, // [2]
    VM_READ | VM_EXEC | VM_MAYREAD | VM_MAYWRITE | VM_MAYEXEC,
    &sigpage_mapping);
```

We grab our `sigpage` at [1] (assuming it didn't already exist), and insert it into an appropriate spot of the current memory map at [2]. To reiterate, while the signal page is only initialized once (when the `init` binary is run), it is mapped into every process that is created. This brings us back to the initial question of what exactly is within the `[sigpage]` mapping of our userland process:

```

[o.o]> x/200wx 0xbfe2000
0xbfe2000: 0x00000005 0x00000000 0x00000000 0xa649a76f
0xbfe2010: 0x00000000 0x00000000 0x00000000 0x00000000
0xbfe2020: 0x00000000 0xa64868c9 0xa649a7c5 0xa64c48f7
// [...]
0xbfe25b0: 0x00000000 0x00000000 0x00000000 0x00000000
0xbfe25c0: 0x00000000 0x00000000 0xbeaa3008 0xbeaa3008
0xbfe25d0: 0xe3a07077 0xf900077 0xdf002777 0xe3a070ad // [1]
0xbfe25e0: 0xf9000ad 0xdf0027ad 0xe59d32f4 0xe8930208
0xbfe25f0: 0xe12fff13 0xcb0c9bdd 0x47104699 0xe59d3374
0xbfe2600: 0xe8930208 0xe12fff13 0xcb0c9bdd 0x47104699
0xbfe2610: 0x00000000 0x00000000 0x00000000 0x00000000
0xbfe2620: 0x00000000 0x00000000 0x00000000 0x00000000
//[...]

```

At [1], we do actually see the signal handler instructions that got copied in initially, disassembled they look like so:

```

[-.-]> x/40i 0xc04004ac
0xc04004ac <sigreturn_codes>: mov r7, #119 ; 0x77
0xc04004b0 <sigreturn_codes+4>: svc 0x00900077
0xc04004b4 <sigreturn_codes+8>: movs r7, #119 ; 0x77
0xc04004b6 <sigreturn_codes+10>: svc 0
0xc04004b8 <sigreturn_codes+12>: mov r7, #173 ; 0xad
0xc04004bc <sigreturn_codes+16>: svc 0x009000ad
0xc04004c0 <sigreturn_codes+20>: movs r7, #173 ; 0xad
0xc04004c2 <sigreturn_codes+22>: svc 0
0xc04004c4 <sigreturn_codes+24>: ldr r3, [sp, #756] ; 0x2f4
0xc04004c8 <sigreturn_codes+28>: ldm r3, {r3, r9}
0xc04004cc <sigreturn_codes+32>: bx r3
0xc04004d0 <sigreturn_codes+36>: ldr r3, [sp, #756] ; 0x2f4
0xc04004d2 <sigreturn_codes+38>: ldmbia r3, {r2, r3}
0xc04004d4 <sigreturn_codes+40>: mov r9, r3
0xc04004d6 <sigreturn_codes+42>: bx r2
0xc04004d8 <sigreturn_codes+44>: ldr r3, [sp, #884] ; 0x374
0xc04004dc <sigreturn_codes+48>: ldm r3, {r3, r9}
0xc04004e0 <sigreturn_codes+52>: bx r3
0xc04004e4 <sigreturn_codes+56>: ldr r3, [sp, #884] ; 0x374
0xc04004e6 <sigreturn_codes+58>: ldmbia r3, {r2, r3}
0xc04004e8 <sigreturn_codes+60>: mov r9, r3
0xc04004ea <sigreturn_codes+62>: bx r2
0xc04004ec <sigreturn_codes+64>: andeq r0, r0, r0

```

However as might be obvious, the rest of the memory is uninitialized, having not been zeroed after being grabbed by `page = alloc_pages(GFP_KERNEL, 0);` inside `get_signal_page()`.

Thus, any userland process can read the `[sigpage]` mapping within their own virtual memory space to leak kernel data (that does not change until the device reboots). It's worth noting that this page's contents depend entirely on the device itself, and potentially might contain data from a previous boot if the device is not shut down for too long. To determine the freshness of the memory, one can count how many instances of the `sigreturn_codes` have been copied in.

TIMELINE

2021-01-28 - Vendor Disclosure

2021-02-05 - Vendor Patched

2021-06-25 - Public Release

CREDIT

Discovered by Lilith >_> of Cisco Talos.

