

## Talos Vulnerability Report

TALOS-2021-1279

### AT&T Labs Xmill XML parsing CreateLabelOrAttrib memory corruption vulnerability

AUGUST 11, 2021

CVE NUMBER

CVE-2021-21811

#### Summary

A memory corruption vulnerability exists in the XML-parsing CreateLabelOrAttrib functionality of AT&T Labs' Xmill 0.7. A specially crafted XML file can lead to a heap buffer overflow. An attacker can provide a malicious file to trigger this vulnerability.

#### Tested Versions

AT&T Labs Xmill 0.7

#### Product URLs

None

#### CVSSv3 Score

8.1 - CVSS:3.0/AV:N/AC:H/PR:N/UI:N/S:U/C:H/I:H/A:H

#### CWE

CWE-191 - Integer Underflow (Wrap or Wraparound)

#### Details

Xmill and Xdemill are utilities that are purpose-built for XML compression and decompression, respectively. These utilities claim to be roughly two times more efficient at compressing XML than other compression methods.

While this software is old, released in 1999, it can be found in modern software suites, such as Schneider Electric's EcoStruxure Control Expert.

Xmill utilizes custom heap management code within the software to distribute pointers within the normal heap. Heap chunks are allocated in set block sizes as follows:

```
unsigned long blocksizes[BLOCKSIZE_NUM]= {256, 1024, 8192, 65536};
```

These blocks are then used for various data within the Xmill process for XML compression.

During the process of parsing the provided XML file, tags are stored within the custom heap allocations in order to best compress them. LabelDict::CreateLabelOrAttrib is used within XMLParse::DoParsing in order to allocate and write these labels and attributes into the custom heap chunks. The custom heap chunk management has a bug which allows for an arbitrary heap write.

This bug starts in LabelDict::CreateLabelOrAttrib when allocating space for a label or attribute of size 0xFFC9

```
TLabelID CreateLabelOrAttrib(char *label,unsigned len,unsigned char isattrib)
// Creates a new label in the hash table
{
    // Let's get some memory first
    mainmem.WordAlign();
    CompressLabelDictItem *item=(CompressLabelDictItem *)mainmem.GetByteBlock(sizeof(CompressLabelDictItem)+len);
    mainmem.WordAlign();
    ...
}
```

As seen above, the call to GetByteBlock takes into account the input length and the size of CompressLabelDictItem which is of size 0x18, adding these sizes together we get a total of 0xFFE1 bytes.

```

char *GetByteBlock(unsigned len)
// The main function for allocated more memory of size 'len'.
// The function checks the current block and if there is not enough space,
// the function 'AllocateNewBlock' is called.
{
    ...

    else // We add a new block at the end
    {
        do
        {
            curblock->next=AllocateNewBlock();
            curblock=curblock->next;
        }
        while(curblock->blocksize<len); // If we don't have enough space,
                                        // we simply create a bigger one!

        curblock->cursize=len;
        overallsize+=len;

        return curblock->data;
    }
}

```

Within GetByteBlock a new heap chunk will be allocated of size 0x10000 (65536) to hold this data.

```

MemStreamBlock *AllocateNewBlock()
// Allocates a new block
// The size is determined by the next index in blocksizeidxs
{
    MemStreamBlock *newblock;

    // Let's get the new block
    newblock=(MemStreamBlock *)AllocateBlock(blocksizeidxs[curblocksizeidx]);

    // The usable data size is the block size - the size of the header
    newblock->blocksize=GetBlockSize(blocksizeidxs[curblocksizeidx])-(sizeof(MemStreamBlock)-1);
    ...
}

```

While the heap chunk size is technically 0x10000, within AllocateNewBlock the size is actually the total size (0x10000) minus the total size of the MemStreamBlock minus 1, which makes the total storable data equal to 0xFFE1 in a 0x10000 byte chunk. This means that the data provided in CreateLabelOrAttrib completely fills the custom chunk. Once returned from GetByteBlock the memory is aligned using WordAlign

```

void WordAlign()
// Allocated 1 to 3 bytes to align the current memory pointer
// to an address divisible by 4.
{
    if(curblock==NULL)
        return;

    int addsize=3-((curblock->cursize+3)&3);
    if(addsize>0)
    {
        curblock->cursize+=addsize;
        overallsize+=addsize;
    }
}

```

This sets the curblock->cursize to 0xFFE4. The rest of CreateLabelOrAttrib is executed without issue, copying the len bytes into the new buffer. When the next call to CreateLabelOrAttrib is made, to create another label is when the bug arises.

When calling into GetByteBlock from CreateLabelOrAttrib we first check to see if the data can fit in the current block, due to the WordAlign moving curblock->cursize to an aligned value, we now have a curblock->cursize of 0xFFE4 and a curblock->blocksize of 0xFFE1 which causes an integer underflow with an unsigned comparison.

```

char *GetByteBlock(unsigned len)
// The main function for allocated more memory of size 'len'.
// The function checks the current block and if there is not enough space,
// the function 'AllocateNewBlock' is called.
{
    ...

    if(curblock->blocksize-curblock->cursize>=len)
        // Enough space in current block?
        {
            char *ptr=curblock->data+curblock->cursize;
            curblock->cursize+=len;
            overallsize+=len;
            return ptr;
        }
    ...
}

```

With the integer underflow, this block is considered large enough to hold any amount of data, a pointer is returned, aligned, and then data is copied when finishing CreateLabelOrAttrib

```

TLabelID CreateLabelOrAttrib(char *label,unsigned len,unsigned char isattrib)
// Creates a new label in the hash table
{
    // Let's get some memory first
    mainmem.WordAlign();
    CompressLabelDictItem *item=(CompressLabelDictItem *)mainmem.GetByteBlock(sizeof(CompressLabelDictItem)*len);
    mainmem.WordAlign();

    // We copy the label name
    item->len=(unsigned short)len;
    memcpy(item->GetLabelPtr(),label,len);
    ...
}

```

This integer underflow results in an exploitable heap corruption.

#### Crash Information

```

=====
==2676449==ERROR: AddressSanitizer: heap-buffer-overflow on address 0x631000074814 at pc 0x000000495dca bp 0x7fffffffd330 sp 0x7ffffffeca8b
WRITE of size 6 at 0x631000074814 thread T0
#0 0x495dc9 in __asan_memcpy (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x495dc9)
#1 0x51737d in LabelDict::CreateLabelOrAttrib(char*, unsigned int, unsigned char) (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x51737d)
#2 0x5162d2 in SAXClient::HandleStartLabel(char*, int, char) (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x5162d2)
#3 0x527175 in XMLParser::ParseLabel() (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x527175)
#4 0x52364c in XMLParser::DoParsing(SAXClient*) (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x52364c)
#5 0x52161e in Compress(char*, char*) (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x52161e)
#6 0x5211a6 in HandleSingleFile(char*) (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x5211a6)
#7 0x521b91 in HandleFileArg(char*) (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x521b91)
#8 0x522059 in main (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x522059)
#9 0x7ffff7a1cb1 in __libc_start_main csu/../csu/libc-start.c:314:16
#10 0x41c84d in _start (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x41c84d)

0x631000074814 is located 20 bytes to the right of 65536-byte region [0x631000064800,0x631000074800)
allocated by thread T0 here:
#0 0x4968bd in malloc (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x4968bd)
#1 0x52ec6a in AllocateBlockRecur(unsigned char) (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x52ec6a)
#2 0x50d4d5 in AllocateBlock(unsigned char) (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x50d4d5)
#3 0x50d0c8 in MemStreamer::AllocateNewBlock() (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x50d0c8)
#4 0x50de9c in MemStreamer::GetByteBlock(unsigned int) (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x50de9c)
#5 0x5172fb in LabelDict::CreateLabelOrAttrib(char*, unsigned int, unsigned char) (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x5172fb)
#6 0x5162d2 in SAXClient::HandleStartLabel(char*, int, char) (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x5162d2)
#7 0x526f1c in XMLParser::ParseLabel() (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x526f1c)
#8 0x52364c in XMLParser::DoParsing(SAXClient*) (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x52364c)
#9 0x52161e in Compress(char*, char*) (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x52161e)
#10 0x5211a6 in HandleSingleFile(char*) (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x5211a6)
#11 0x521b91 in HandleFileArg(char*) (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x521b91)
#12 0x522059 in main (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x522059)
#13 0x7ffff7a1cb1 in __libc_start_main csu/../csu/libc-start.c:314:16

SUMMARY: AddressSanitizer: heap-buffer-overflow (/home/fuzz/Desktop/xmill/unix/xmill-asan+0x495dc9) in __asan_memcpy
Shadow bytes around the buggy address:
 0x0c62800068b0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
 0x0c62800068c0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
 0x0c62800068d0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
 0x0c62800068e0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
 0x0c62800068f0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
=>0x0c6280006900: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
 0x0c6280006910: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
 0x0c6280006920: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
 0x0c6280006930: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
 0x0c6280006940: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
 0x0c6280006950: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
Shadow byte legend (one shadow byte represents 8 application bytes):
Addressable: 00
Partially addressable: 01 02 03 04 05 06 07
Heap left redzone: fa
Freed heap region: fd
Stack left redzone: f1
Stack mid redzone: f2
Stack right redzone: f3
Stack after return: f5
Stack use after scope: f8
Global redzone: f9
Global init order: f6
Poisoned by user: f7
Container overflow: fc
Array cookie: ac
Intra object redzone: bb
ASan internal: fe
Left alloca redzone: ca
Right alloca redzone: cb
Shadow gap: cc
==2676449==ABORTING

```

#### Timeline

2021-04-30 - Vendor Disclosure  
2021-08-10 - Public Release

#### CREDIT

Discovered by Carl Hurd of Cisco Talos.

