



[Reporting Issues](#)

## Bug 1816 (CVE-2021-28211) - Possible heap corruption with LzmaUefiDecompressGetInfo

**Status:** RESOLVED FIXED

**Alias:** CVE-2021-28211

**Product:** Tianocore Security Issues

**Component:** Security Issue ([show other bugs](#))

**Version:** Current

**Hardware:** All All

**Importance:** Normal normal

**Assignee:** Laszlo Ersek

**URL:**

**Keywords:**

**Depends on:**

**Blocks:**

**Reported:** 2019-05-15 13:55 UTC by Satoshi Tanda

**Modified:** 2021-09-22 15:47 UTC ([History](#))

**CC List:** 16 users ([show](#))

**See Also:**

**Release(s) the issue is observed:** edk2-stable202008

**The OS the target platform is running:** ---

**Package:** IntelFrameworkModulePkg, MdeModulePkg

**Release(s) the issues must be fixed:** edk2-stable202011

### Attachments

<a href="#">[PATCH] MdeModulePkg/LzmaCustomDecompressLib: catch 4GB+ uncompressed buffer sizes</a> (3.85 KB, patch)	<a href="#">Details</a>   <a href="#">Diff</a>
<a href="#">2020-09-28 12:22 UTC</a> , Laszlo Ersek	
<a href="#">CVE json file 2</a> (907 bytes, application/json)	<a href="#">Details</a>
<a href="#">2021-03-04 14:29 UTC</a> , KevinJ	
<a href="#">Add an attachment</a> (proposed patch, testcase, etc.)	<a href="#">Show Obsolete</a> (1)

Note  
You need to [log in](#) before you can comment on or make changes to this bug.

Satoshi Tanda 2019-05-15 13:55:17 UTC

[Description](#)

The LzmaUefiDecompressGetInfo function fetches 8 bytes (64bit-wide) of the uncompressed data size from the LZMA header then silently drops the upper 32bit by casting it to UINT32.

<https://github.com/tianocore/edk2/blob/master/MdeModulePkg/Library/LzmaCustomDecompressLib/LzmaDecompress.c#L146>

Since the caller of the function typically allocates heap based on the reported size by this function to subsequently decompress data, the caller may allocate smaller buffer than necessary when the input LZMA headers is greater than 32bit-wide. For example, when the uncompressed size claimed in the header is 0x1\_00000100, the LzmaUefiDecompressGetInfo function reports 0xi00 to the caller. The caller would allocate that heap and call the LzmaUefiDecompress function. Because the LzmaUefiDecompress function fetches an uncompressed size from the header again which will be 0x1\_00000100 and decompresses data accordingly (without taking the size of allocated buffer), it is possible to overflow the allocated heap during the decompression process.

It is probably challenging to exploit this in code compiled into firmware as you need to control input in the first place, but the LzmaUefiDecompressGetInfo function may be used to 3rd party software that parses UEFI firmware image files. For example, UEFITool has its own copy of this function (and fixed it by themselves <https://github.com/LongSoft/UEFITool/commit/f507d71ead4dccc516d6be1b5b3cd61f900abb4a4#diff-fa894f45e026a9e683e0f552b418626>).

I do not have a PoC firmware image and a standalone parser to demonstrate the issue handy, but hopefully, the issue is clear.

Vincent Zimmer 2019-05-29 12:46:37 UTC

[Comment 1](#)

5/29/19 tianocore infosec bugscrub - brent h. to dig in further.  
Kinney - looks like > 4gb input data would cause issue  
Brett - ok if 32-bit everywhere. mixing would be an issue

Bret Barkelew 2019-05-29 12:47:14 UTC

[Comment 2](#)

```
Bret*  
;)
```

brent.holtscaw 2019-05-29 17:41:05 UTC

[Comment 3](#)

It appears to be a security issue if using compilers size\_t on a 64bit system. Suggest either casting to SizeT within LzmaUefiDecompressGetInfo or casting to UINT32 within LzmaUefiDecompress.

```
LzmaUefiDecompressGetInfo function casts DestinationSize to UINT32  
..  
DecodedSize = GetDecodedSizeOfBuf((UINT8*)Source);  
*DestinationSize = (UINT32)DecodedSize;
```

```
LzmaUefiDecompress function casts DecodedBufSize to SizeT which appears to also be  
defined as UINTN  
DecodedBufSize = (SizeT)GetDecodedSizeOfBuf((UINT8*)Source);  
It then uses the DecodedBufSize in call to LzmaDecode  
LzmaResult = LzmaDecode(  
    Destination,  
    &DecodedBufSize,  
    (Byte*)((UINT8*)Source + LZMA_HEADER_SIZE),  
    &EncodedDataSize,  
    Source,  
    LZMA_PROPS_SIZE,  
    LZMA_FINISH_END,  
    &Status,  
    &(AllocFuncs.Functions)  
)
```

```
Function LzmaDecode uses DecodedBufSize as the size of the destination buffer and  
also treats it as SizeT  
SRes LzmaDecode(Byte *dest, SizeT *destLen, const Byte *src, SizeT *srcLen,  
    const Byte *propData, unsigned propSize, ELzmaFinishMode finishMode,  
    ELzmaStatus *status, ISzAlloc *alloc)
```

```
SizeT outSize = *destLen  
p.dicBufSize = outSize;  
res = LzmaDec_DecodeToDic(&p, outSize, src, srcLen, finishMode, status);
```

Bret Barkelew 2019-06-25 16:48:12 UTC

[Comment 4](#)

Agree with Brent about explicitly casting consistently in this library.

Laszlo Ersek 2019-08-01 19:19:02 UTC

[Comment 5](#)

(comment made against edk2 @ 3d34b5f32692)

I think I disagree about both impact ([comment 0](#)) and mitigation (comments 3-4).

Regarding impact: 3rd party UEFI binaries may be LZMA compressed. They are decompressed before e.g. secure boot image verification. This means that a crafted EFI binary can look small but cause buffer overflow in section extraction, before validation. A similar issue was [bug-666](#), which merited multiple CVE numbers.

(In particular, see:

- [https://bugzilla.kernel.org/show\\_bug.cgi?id=666111](https://bugzilla.kernel.org/show_bug.cgi?id=666111)  
- <https://lists.01.org/pipermail/edk2-devel/2018-October/031095.html>  
- <https://lists.01.org/pipermail/edk2-devel/2018-October/031138.html>  
)

Regarding mitigation:

- I don't think we should ever truncate to UINT32 if the input is provided as UINT64. (As a special case, I don't think we should truncate to UINT32 within LzmaUefiDecompress().)

- Outputting UINT64 from LzmaUefiDecompressGetInfo() is theoretically safe, but messy in practice. While LzmaUefiDecompressGetInfo() is indeed an internal function, declared in:

```
MdeModulePkg/Library/LzmaCustomDecompressLib/LzmaDecompressLibInternal.h
```

its output parameter called "DestinationSize" is directly propagated out of two public functions. Namely (all pathnames relative to "MdeModulePkg/Library/LzmaCustomDecompressLib"):

```
- LzmaGuidedSectionGetInfo() [GuidedSectionExtraction.c]
- LzmaArchGuidedSectionGetInfo [F86GuidedSectionExtraction.c]
```

Both of these functions are registered with the ExtractGuidedSectionRegisterHandlers() API of the ExtractGuidedSectionLib class.

The registered functions need to conform to the EXTRACT\_GUIDED\_SECTION\_GET\_INFO\_HANDLER type.

The registered functions are searched for, and called, through the public ExtractGuidedSectionGetInfo() API, from:

```
MdePkg/Include/Library/ExtractGuidedSectionLib.h
```

The ExtractGuidedSectionGetInfo() function is called from several places in edk2. I guess the most important could be that the DXE Core uses ExtractGuidedSectionGetInfo() for implementing EFI GUIDED\_SECTION\_EXTRACTION\_PROTOCOL -- see "MdeModulePkg/Core/Dxe/SectionExtraction/CoreSectionExtraction.c".

This protocol comes from the PI spec, and the "OutputSize" output parameter has type UINTN. Therefore, even if all the internal propagation of the original "DecodedSize" were widened correctly to 64-bit, the PI protocol would still be a bottleneck on 32-bit architectures.

Therefore, I would catch >= 4GB DestinationSize values right in LzmaUefiDecompressGetInfo(), and return an error. The function's return type is already RETURN\_STATUS, and all of the call sites correctly propagate errors. This wouldn't enable edk2 to extract GUIDed (LZMA) sections larger than 4GB, but we'd replace the potential buffer overflow with clean errors, at minimal cost to dependent code. (4GB+ GUIDed sections are clearly not needed in practice.)

Laszlo Ersek 2019-08-09 09:42:31 UTC

[Comment 6](#)

The CustomGuidedSectionExtract() function in "MdeModulePkg/Core/Dxe/SectionExtraction/CoreSectionExtraction.c" at commit 4b1b7c191309 is vulnerable to the exact issue described in [comment 0](#); therefore I'm marking this as CONFIRMED.

```
UINT32      OutputBufferSize;
...
Status = ExtractGuidedSectionGetInfo (
    InputSection,
    &OutputBufferSize,
    &ScratchBufferSize,
    &SectionAttribute
);
...
if (OutputBufferSize > 0) {
    //
    // Allocate output buffer
    //
    AllocatedOutputBuffer = AllocatePool (OutputBufferSize);
...
*OutputBuffer = AllocatedOutputBuffer;
}
...
Status = ExtractGuidedSectionDecode (
    InputSection,
    OutputBuffer,
    ScratchBuffer,
    AuthenticationStatus
);
```

and see [comment 3](#) for where the last function call quoted above ends.

Alex Ionescu 2020-07-01 16:29:05 UTC

[Comment 7](#)

Any plan on actually fixing this issue? It seems to be bouncing around. Since it seems to have real security merit, perhaps we should plan a disclosure timeline since many 3rd party libraries are affected.

Laszlo Ersek 2020-09-25 11:30:39 UTC

[Comment 8](#)

I'll try to propose a patch (as an attachment) next week.

Laszlo Ersek 2020-09-28 12:22:25 UTC

[Comment 9](#)

Created [attachment 564 \[details\]](#)

[PATCH] MdeModulePkg/LzmaCustomDecompressLib: catch 4GB+ uncompressed buffer sizes

Proposed patch, applies on top of commit 1d058c3e86b0 ("IntelFsp2Pkg GenCfgOpt.py: Initialize Inclines as empty list", 2020-09-25).

Laszlo Ersek 2020-09-28 12:26:06 UTC

[Comment 10](#)

(In reply to Laszlo Ersek from [comment #9](#))

> Created [attachment 564 \[details\]](#)

> [PATCH] MdeModulePkg/LzmaCustomDecompressLib: catch 4GB+ uncompressed buffer sizes

> Proposed patch, applies on top of commit 1d058c3e86b0 ("IntelFsp2Pkg GenCfgOpt.py: Initialize Inclines as empty list", 2020-09-25).

Cc: Dandan Bi <[dandan.bi@intel.com](mailto:dandan.bi@intel.com)>

Cc: Hao A Wu <[hao.a.wu@intel.com](mailto:hao.a.wu@intel.com)>  
Cc: Jian J Wang <[jian.j.wang@intel.com](mailto:jian.j.wang@intel.com)>  
Cc: Liming Gao <[gaoliming@byosoft.com.cn](mailto:gaoliming@byosoft.com.cn)>  
Cc: Philippe Mathieu-Daudé <[philmd@redhat.com](mailto:philmd@redhat.com)>

Please review.

Note: I have only regression-tested this patch, using OVMF's automatic LZMA decompression (covering FE1FV and DXEFV) during boot, and S3 resume. (See the message on commit b24fca05751f for a few more details on FE1FV and DXEFV.)

Alex Ionescu 2020-09-28 14:46:34 UTC

[Comment 11](#)

Thanks for the patch.

Team/Vincent -- would it help to get a CVE assigned? I can work with MITRE on that...

gaoliming 2020-09-28 21:29:37 UTC

[Comment 12](#)

I agree this fix. Reviewed-by: Liming Gao <[gaoliming@byosoft.com.cn](mailto:gaoliming@byosoft.com.cn)>

Laszlo Ersek 2020-09-29 14:56:01 UTC

[Comment 13](#)

Hi Alex,

(In reply to Alex Ionescu from [comment #11](#))

> Thanks for the patch.

>

> Team/Vincent -- would it help to get a CVE assigned? I can work with MITRE on that...

I agree that a CVE should be assigned (see my reference to [bug#686](#) in [comment#5](#) above).

Luckily, Tianocore has recently become a CNA <[https://cve.mitre.org/cve/request\\_id.html#t](https://cve.mitre.org/cve/request_id.html#t)> thanks to the work of the infosec group. From Vincent's announcement on the edk2-infosec mailing list, "Points of contact for infosec and CVE's will include Kevin Davis from Insyde, Eric Johnson from AMI, Dick Wilkins from Phoenix, and Vincent Zimmer from Intel". So (I think?) we need not contact an external org for assigning a CVE.

Thanks.

Laszlo Ersek 2020-09-29 14:58:27 UTC

[Comment 14](#)

(In reply to gaoliming from [comment #12](#))

> I agree this fix. Reviewed-by: Liming Gao <[gaoliming@byosoft.com.cn](mailto:gaoliming@byosoft.com.cn)>

Thank you, Liming!

So if there are no more comments on the patch, I think we should set up an embargo end date, so that people can start applying the patch.

Once the embargo elapses, I'll post the patch to edk2-devel for public review, with your R-b from [comment#12](#) included. Thanks!

Alex Ionescu 2020-09-29 17:23:31 UTC

[Comment 15](#)

(In reply to Laszlo Ersek from [comment #13](#))

> Hi Alex,

>

> (In reply to Alex Ionescu from [comment #11](#))

> > Thanks for the patch.

>

> > Team/Vincent -- would it help to get a CVE assigned? I can work with MITRE

> > on that...

>

> > I agree that a CVE should be assigned (see my reference to [bug#686](#) in [comment#5](#) above).

>

> > Luckily, Tianocore has recently become a CNA

> > <[https://cve.mitre.org/cve/request\\_id.html#t](https://cve.mitre.org/cve/request_id.html#t)> thanks to the work of the infosec group. From Vincent's announcement on the edk2-infosec mailing list, "Points of contact for infosec and CVE's will include Kevin Davis from Insyde, Eric Johnson from AMI, Dick Wilkins from Phoenix, and Vincent Zimmer from Intel". So (I think?) we need not contact an external org for assigning a CVE.

>

> > Thanks.

That's awesome. Let me know if this is something I need to ping Vincent on or if it will be handled automatically as part of the bug/triage process. I want to make sure the bug will be correctly credited to Satoshi Tanda -- CrowdStrike and to be aware of disclosure timeline so that we can also absorb the patch in our own product -- for now we used a workaround to avoid showcasing the issue (it's a closed-source component).

Laszlo Ersek 2020-10-01 03:33:42 UTC

[Comment 16](#)

According to the flowchart at

<https://github.com/tianocore/tianocore.github.io/wiki/Reporting-Security-Issues>

(direct link:

<https://raw.githubusercontent.com/jwang36/tianocore.github.io/master/security/flowchart.svg>)

from phase 2, we're still missing the CVSS evaluation step on this BZ. Once that is done (or may in parallel to it), we should indeed figure out the embargo length (phase 3), and get a CVE number (for phase 4).

I'm going to send an email to the edk2-infosec mailing list now, CC'ing the CVE contacts, about this BZ. Thanks.

Philippe Mathieu-Daudé 2020-10-07 07:17:09 UTC

[Comment 17](#)

(In reply to Laszlo Ersek from [comment #9](#))

> Created [attachment 564 \[details\]](#)

> [PATCH] MdeModulePkg/LzmaCustomDecompressLib: catch 4GB+ uncompressed buffer sizes

>

> Proposed patch, applies on top of commit 1d058c3e86b0 ("IntelFsp2Pkg GenCfgOpt.py: Initialize Inclines as empty list", 2020-09-25).

Reviewed-by: Philippe Mathieu-Daudé <[philmd@redhat.com](mailto:philmd@redhat.com)>

Laszlo Ersek 2020-10-07 11:52:56 UTC

[Comment 18](#)

Thank you, Phil!

Laszlo Ersek 2020-11-12 12:16:52 UTC

[Comment 19](#)

PROPOSED PUBLIC DATE (opening up the BZ and posting the patch to edk2-devel):

Thursday 2020-Nov-19 07:00 UTC

in order to get the fix into edk2-stable202011.

Riccardo Schirone 2020-11-13 06:14:53 UTC

[Comment 20](#)

Is this going to have a CVE before it goes public?

Laszlo Ersek 2020-11-13 15:21:13 UTC

[Comment 21](#)

(In reply to Riccardo Schirone from [comment #20](#))  
> Is this going to have a CVE before it goes public?

CC'ing Eric.

Laszlo Ersek 2020-11-19 06:53:01 UTC

[Comment 22](#)

(In reply to Laszlo Ersek from [comment #3](#))  
> Created [attachment 564](#) [\[details\]](#)  
> [PATCH] MdeModulePkg/LzmaCustomDecompressLib: catch 4GB+ uncompressed buffer  
> sizes  
>  
> Proposed patch, applies on top of commit 1d058c3e86b0 ("IntelFsp2Pkg  
> GenCfgOpt.py: Initialize InLines as empty list", 2020-09-25).

Public posting:

\* [edk2-devel] [PATCH RESEND 0/1]  
security fix: possible heap corruption with LzmaUefiDecompressGetInfo

msgid <20201119115034.12897-1-lerssek@redhat.com>  
<https://edk2.groups.io/g/devlist/message/67708>  
<https://www.redhat.com/archives/edk2-devel-archive/2020-November/msg00866.html>

Alex Ionescu 2020-11-19 11:25:12 UTC

[Comment 23](#)

@Ericj -- what's the CVE/Ack for this? I see this is now public.

Laszlo Ersek 2020-11-20 21:06:10 UTC

[Comment 24](#)

Merged as commit e7bd0dd26db7, via <https://github.com/tianocore/edk2/pull/1138>.

kevinj 2021-03-03 11:46:54 UTC

[Comment 25](#)

Created [attachment 665](#) [\[details\]](#)  
CVE .json file

I have attached the .json file for CVE classification. Please review and provide feedback, especially for the version.

Laszlo Ersek 2021-03-03 12:02:57 UTC

[Comment 26](#)

Hello Kevin,

last vulnerable release: edk2-stable202008  
first fixed release: edk2-stable202011

\$ git tag --contains e7bd0dd26db7  
edk2-stable202011

Laszlo Ersek 2021-03-03 12:05:23 UTC

[Comment 27](#)

Upon reviewing commit e7bd0dd26db7, I think CWE-122 is a good match.

kevinj 2021-03-04 14:29:59 UTC

[Comment 28](#)

Created [attachment 666](#) [\[details\]](#)  
CVE .json file\_2

Thank you Laszlo for your feedback. I have updated the version in the .json file and re-uploaded it.

Alex Ionescu 2021-03-10 09:26:36 UTC

[Comment 29](#)

Hi,

Still wondering when it's possible to have the CVE assigned? This has almost taken two years now since reporting.

kevinj 2021-03-11 16:59:03 UTC

[Comment 30](#)

(In reply to Alex Ionescu from [comment #29](#))  
> Hi,  
>  
> Still wondering when it's possible to have the CVE assigned? This has almost  
> taken two years now since reporting.

We have requested CVE-IDs from MITRE. We are still waiting to receive said ID number from them.

Laszlo Ersek 2021-03-12 16:03:49 UTC

[Comment 31](#)

Thanks for the CVE number, Kevin!

kevinj 2021-03-12 16:05:15 UTC

[Comment 32](#)

Laszlo,

Please review the .json file again, especially the version this bug is observed in and inform me when you plan to publicly disclose this bug, so we know when to submit this CVE back to MITRE. Thank you!

Laszlo Ersek 2021-03-12 16:56:54 UTC

[Comment 33](#)

The release info (edk2-stable202008) in the JSON from [comment 28](#) seems OK, matching my [comment 26](#).

Regarding public disclosure, we're past that -- please see [comment 19](#) and [comment 22](#). The issue was disclosed on November 19, 2020.

Alex Ionescu 2021-03-12 17:09:26 UTC

[Comment 34](#)

Since Tianocore is a CNA, I'm unclear why you're contacting MITRE to get a CVE -- can't Tianocore issue one directly?

Additionally, when will the information be published here:

<https://edk2-docs.gitbook.io/security-advisory/>

Satoshi Tanda 2021-04-18 14:27:24 UTC

[Comment 35](#)

Just noticed CVE-2021-28211 was assigned. It is still reserved in MITRE, but some vendors have public info  
<https://access.redhat.com/security/cve/CVE-2021-28211>

Kim Olsun    2021-09-22 15:47:43 UTC [Comment 36](#)

<http://www.compilatori.com/tech/nvidia-and-samsung/>  
<http://www.wearielondonmade.com/tech/nvidia-and-samsung/>  
<http://www.lopspeach.com/tech/nvidia-and-samsung/>  
<http://icord.li/tech/nvidia-and-samsung/>    <http://connstr.net/tech/nvidia-and-samsung/>  
<http://embermanchester.uk/tech/nvidia-and-samsung/>  
<http://www.slipstone.co.uk/tech/nvidia-and-samsung/>  
<http://www.iqoarts.co.uk/tech/nvidia-and-samsung/>  
<http://www.acpizateradio.co.uk/health/covid-and-tech/>  
<https://waytowhatsnext.com/category/services/>    <https://www.wehb-dev.co.uk/category/shopping/>  
<http://www.lu-bloomington.com/category/services/>  
<http://www-look-4.com/tech/nvidia-and-samsung/>