

15e3b9dbcc

...

envoy / api / envoy / api / v2 / cluster.proto



mattklein123 http: remove getAll() header map API and switch all usages to get() (#... .. ✓

History

10 contributors



867 lines (748 sloc) 41.5 KB

...

```
1 syntax = "proto3";
2
3 package envoy.api.v2;
4
5 import "envoy/api/v2/auth/tls.proto";
6 import "envoy/api/v2/cluster/circuit_breaker.proto";
7 import "envoy/api/v2/cluster/filter.proto";
8 import "envoy/api/v2/cluster/outlier_detection.proto";
9 import "envoy/api/v2/core/address.proto";
10 import "envoy/api/v2/core/base.proto";
11 import "envoy/api/v2/core/config_source.proto";
12 import "envoy/api/v2/core/health_check.proto";
13 import "envoy/api/v2/core/protocol.proto";
14 import "envoy/api/v2/endpoint.proto";
15 import "envoy/type/percent.proto";
16
17 import "google/protobuf/any.proto";
18 import "google/protobuf/duration.proto";
19 import "google/protobuf/struct.proto";
20 import "google/protobuf/wrappers.proto";
21
22 import "envoy/annotations/deprecation.proto";
23 import "udpa/annotations/migrate.proto";
24 import "udpa/annotations/status.proto";
25 import "validate/validate.proto";
26
27 option java_package = "io.envoyproxy.envoy.api.v2";
28 option java_outer_classname = "ClusterProto";
29 option java_multiple_files = true;
30 option (udpa.annotations.file_migrate).move_to_package = "envoy.config.cluster.v3";
31 option (udpa.annotations.file_status).package_version_status = FROZEN;
32
33 // [#protodoc-title: Cluster configuration]
34
35 // Configuration for a single upstream cluster.
36 // [#next-free-field: 48]
37 message Cluster {
38   // Refer to :ref:`service discovery type <arch_overview_service_discovery_types>`
39   // for an explanation on each type.
40   enum DiscoveryType {
41     // Refer to the :ref:`static discovery type<arch_overview_service_discovery_types_static>`
42     // for an explanation.
43     STATIC = 0;
44
45     // Refer to the :ref:`strict DNS discovery
46     // type<arch_overview_service_discovery_types_strict_dns>`
47     // for an explanation.
48     STRICT_DNS = 1;
49
50     // Refer to the :ref:`logical DNS discovery
51     // type<arch_overview_service_discovery_types_logical_dns>`
52     // for an explanation.
53     LOGICAL_DNS = 2;
54
55     // Refer to the :ref:`service discovery type<arch_overview_service_discovery_types_eds>`
56     // for an explanation.
57     EDS = 3;
58
59     // Refer to the :ref:`original destination discovery
60     // type<arch_overview_service_discovery_types_original_destinations>`
61     // for an explanation.
62     ORIGINAL_DST = 4;
63   }
64
65   // Refer to :ref:`load balancer type <arch_overview_load_balancing_types>` architecture
66   // overview section for information on each type.
67   enum LbPolicy {
68     // Refer to the :ref:`round robin load balancing
69     // policy<arch_overview_load_balancing_types_round_robin>`
70     // for an explanation.
71     ROUND_ROBIN = 0;
72
73     // Refer to the :ref:`least request load balancing
74     // policy<arch_overview_load_balancing_types_least_request>`
75     // for an explanation.
76     LEAST_REQUEST = 1;
77
78     // Refer to the :ref:`ring hash load balancing
```

```

79 // policy<arch_overview_load_balancing_types_ring_hash>`
80 // for an explanation.
81 RING_HASH = 2;
82
83 // Refer to the :ref:`random load balancing
84 // policy<arch_overview_load_balancing_types_random>`
85 // for an explanation.
86 RANDOM = 3;
87
88 // Refer to the :ref:`original destination load balancing
89 // policy<arch_overview_load_balancing_types_original_destination>`
90 // for an explanation.
91 //
92 // .. attention::
93 //
94 // **This load balancing policy is deprecated**. Use CLUSTER_PROVIDED instead.
95 //
96 ORIGINAL_DST_LB = 4 [deprecated = true, (envoy.annotations.disallowed_by_default_enum) = true];
97
98 // Refer to the :ref:`Maglev load balancing policy<arch_overview_load_balancing_types_maglev>`
99 // for an explanation.
100 MAGLEV = 5;
101
102 // This load balancer type must be specified if the configured cluster provides a cluster
103 // specific load balancer. Consult the configured cluster's documentation for whether to set
104 // this option or not.
105 CLUSTER_PROVIDED = 6;
106
107 // [#not-implemented-hide:] Use the new :ref:`load_balancing_policy
108 // <envoy_api_field_Cluster.load_balancing_policy>` field to determine the LB policy.
109 // [#next-major-version: In the v3 API, we should consider deprecating the lb_policy field
110 // and instead using the new load_balancing_policy field as the one and only mechanism for
111 // configuring this.]
112 LOAD_BALANCING_POLICY_CONFIG = 7;
113 }
114
115 // When V4_ONLY is selected, the DNS resolver will only perform a lookup for
116 // addresses in the IPv4 family. If V6_ONLY is selected, the DNS resolver will
117 // only perform a lookup for addresses in the IPv6 family. If AUTO is
118 // specified, the DNS resolver will first perform a lookup for addresses in
119 // the IPv6 family and fallback to a lookup for addresses in the IPv4 family.
120 // For cluster types other than
121 // :ref:`STRICT_DNS<envoy_api_enum_value_Cluster.DiscoveryType.STRICT_DNS>` and
122 // :ref:`LOGICAL_DNS<envoy_api_enum_value_Cluster.DiscoveryType.LOGICAL_DNS>`,
123 // this setting is
124 // ignored.
125 enum DnsLookupFamily {
126     AUTO = 0;
127     V4_ONLY = 1;
128     V6_ONLY = 2;
129 }
130
131 enum ClusterProtocolSelection {
132     // Cluster can only operate on one of the possible upstream protocols (HTTP1.1, HTTP2).
133     // If :ref:`http2_protocol_options <envoy_api_field_Cluster.http2_protocol_options>` are
134     // present, HTTP2 will be used, otherwise HTTP1.1 will be used.
135     USE_CONFIGURED_PROTOCOL = 0;
136
137     // Use HTTP1.1 or HTTP2, depending on which one is used on the downstream connection.
138     USE_DOWNSTREAM_PROTOCOL = 1;
139 }
140
141 // TransportSocketMatch specifies what transport socket config will be used
142 // when the match conditions are satisfied.
143 message TransportSocketMatch {
144     // The name of the match, used in stats generation.
145     string name = 1 [(validate.rules).string = {min_len: 1}];
146
147     // Optional endpoint metadata match criteria.
148     // The connection to the endpoint with metadata matching what is set in this field
149     // will use the transport socket configuration specified here.
150     // The endpoint's metadata entry in *envoy.transport_socket_match* is used to match
151     // against the values specified in this field.
152     google.protobuf.Struct match = 2;
153
154     // The configuration of the transport socket.
155     core.TransportSocket transport_socket = 3;
156 }
157
158 // Extended cluster type.
159 message CustomClusterType {
160     // The type of the cluster to instantiate. The name must match a supported cluster type.
161     string name = 1 [(validate.rules).string = {min_bytes: 1}];
162
163     // Cluster specific configuration which depends on the cluster being instantiated.
164     // See the supported cluster for further documentation.
165     google.protobuf.Any typed_config = 2;
166 }
167
168 // Only valid when discovery type is EDS.
169 message EdsClusterConfig {
170     // Configuration for the source of EDS updates for this Cluster.
171     core.ConfigSource eds_config = 1;
172
173     // Optional alternative to cluster name to present to EDS. This does not
174     // have the same restrictions as cluster name, i.e. it may be arbitrary
175     // length.
176     string service_name = 2;

```

```

177 }
178
179 // Optionally divide the endpoints in this cluster into subsets defined by
180 // endpoint metadata and selected by route and weighted cluster metadata.
181 // [#next-free-field: 8]
182 message LbSubsetConfig {
183     // If NO_FALLBACK is selected, a result
184     // equivalent to no healthy hosts is reported. If ANY_ENDPOINT is selected,
185     // any cluster endpoint may be returned (subject to policy, health checks,
186     // etc). If DEFAULT_SUBSET is selected, load balancing is performed over the
187     // endpoints matching the values from the default_subset field.
188     enum LbSubsetFallbackPolicy {
189         NO_FALLBACK = 0;
190         ANY_ENDPOINT = 1;
191         DEFAULT_SUBSET = 2;
192     }
193
194     // Specifications for subsets.
195     message LbSubsetSelector {
196         // Allows to override top level fallback policy per selector.
197         enum LbSubsetSelectorFallbackPolicy {
198             // If NOT_DEFINED top level config fallback policy is used instead.
199             NOT_DEFINED = 0;
200
201             // If NO_FALLBACK is selected, a result equivalent to no healthy hosts is reported.
202             NO_FALLBACK = 1;
203
204             // If ANY_ENDPOINT is selected, any cluster endpoint may be returned
205             // (subject to policy, health checks, etc).
206             ANY_ENDPOINT = 2;
207
208             // If DEFAULT_SUBSET is selected, load balancing is performed over the
209             // endpoints matching the values from the default_subset field.
210             DEFAULT_SUBSET = 3;
211
212             // If KEYS_SUBSET is selected, subset selector matching is performed again with metadata
213             // keys reduced to
214             // :ref:`fallback_keys_subset`envoy_api_field_Cluster.LbSubsetConfig.LbSubsetSelector.fallback_keys_subset`.
215             // It allows for a fallback to a different, less specific selector if some of the keys of
216             // the selector are considered optional.
217             KEYS_SUBSET = 4;
218         }
219
220         // List of keys to match with the weighted cluster metadata.
221         repeated string keys = 1;
222
223         // The behavior used when no endpoint subset matches the selected route's
224         // metadata.
225         LbSubsetSelectorFallbackPolicy fallback_policy = 2
226             [(validate.rules).enum = {defined_only: true}];
227
228         // Subset of
229         // :ref:`keys`envoy_api_field_Cluster.LbSubsetConfig.LbSubsetSelector.keys` used by
230         // :ref:`KEYS_SUBSET`envoy_api_enum_value_Cluster.LbSubsetConfig.LbSubsetSelector.LbSubsetSelectorFallbackPolicy.KEYS_SUBSET`.
231         // fallback policy.
232         // It has to be a non empty list if KEYS_SUBSET fallback policy is selected.
233         // For any other fallback policy the parameter is not used and should not be set.
234         // Only values also present in
235         // :ref:`keys`envoy_api_field_Cluster.LbSubsetConfig.LbSubsetSelector.keys` are allowed, but
236         // `fallback_keys_subset` cannot be equal to `keys`.
237         repeated string fallback_keys_subset = 3;
238     }
239
240     // The behavior used when no endpoint subset matches the selected route's
241     // metadata. The value defaults to
242     // :ref:`NO_FALLBACK`envoy_api_enum_value_Cluster.LbSubsetConfig.LbSubsetFallbackPolicy.NO_FALLBACK`.
243     LbSubsetFallbackPolicy fallback_policy = 1 [(validate.rules).enum = {defined_only: true}];
244
245     // Specifies the default subset of endpoints used during fallback if
246     // fallback_policy is
247     // :ref:`DEFAULT_SUBSET`envoy_api_enum_value_Cluster.LbSubsetConfig.LbSubsetFallbackPolicy.DEFAULT_SUBSET`.
248     // Each field in default_subset is
249     // compared to the matching LbEndpoint.Metadata under the *envoy.lb*
250     // namespace. It is valid for no hosts to match, in which case the behavior
251     // is the same as a fallback_policy of
252     // :ref:`NO_FALLBACK`envoy_api_enum_value_Cluster.LbSubsetConfig.LbSubsetFallbackPolicy.NO_FALLBACK`.
253     google.protobuf.Struct default_subset = 2;
254
255     // For each entry, LbEndpoint.Metadata's
256     // *envoy.lb* namespace is traversed and a subset is created for each unique
257     // combination of key and value. For example:
258     //
259     // .. code-block:: json
260     //
261     // { "subset_selectors": [
262     //   { "keys": [ "version" ] },
263     //   { "keys": [ "stage", "hardware_type" ] }
264     // ]}
265     //
266     // A subset is matched when the metadata from the selected route and
267     // weighted cluster contains the same keys and values as the subset's
268     // metadata. The same host may appear in multiple subsets.
269     repeated LbSubsetSelector subset_selectors = 3;
270
271     // If true, routing to subsets will take into account the localities and locality weights of the
272     // endpoints when making the routing decision.
273     //
274     // There are some potential pitfalls associated with enabling this feature, as the resulting

```

```

275 // traffic split after applying both a subset match and locality weights might be undesirable.
276 //
277 // Consider for example a situation in which you have 50/50 split across two localities X/Y
278 // which have 100 hosts each without subsetting. If the subset LB results in X having only 1
279 // host selected but Y having 100, then a lot more load is being dumped on the single host in X
280 // than originally anticipated in the load balancing assignment delivered via EDS.
281 bool locality_weight_aware = 4;
282
283 // When used with locality_weight_aware, scales the weight of each locality by the ratio
284 // of hosts in the subset vs hosts in the original subset. This aims to even out the load
285 // going to an individual locality if said locality is disproportionately affected by the
286 // subset predicate.
287 bool scale_locality_weight = 5;
288
289 // If true, when a fallback policy is configured and its corresponding subset fails to find
290 // a host this will cause any host to be selected instead.
291 //
292 // This is useful when using the default subset as the fallback policy, given the default
293 // subset might become empty. With this option enabled, if that happens the LB will attempt
294 // to select a host from the entire cluster.
295 bool panic_mode_any = 6;
296
297 // If true, metadata specified for a metadata key will be matched against the corresponding
298 // endpoint metadata if the endpoint metadata matches the value exactly OR it is a list value
299 // and any of the elements in the list matches the criteria.
300 bool list_as_any = 7;
301 }
302
303 // Specific configuration for the LeastRequest load balancing policy.
304 message LeastRequestLbConfig {
305     // The number of random healthy hosts from which the host with the fewest active requests will
306     // be chosen. Defaults to 2 so that we perform two-choice selection if the field is not set.
307     google.protobuf.UInt32Value choice_count = 1 [(validate.rules).uint32 = {gte: 2}];
308 }
309
310 // Specific configuration for the :ref:`RingHash<arch_overview_load_balancing_types_ring_hash>`
311 // load balancing policy.
312 message RingHashLbConfig {
313     // The hash function used to hash hosts onto the ketama ring.
314     enum HashFunction {
315         // Use `xxHash` <https://github.com/Cyan4973/xxHash>`, this is the default hash function.
316         XX_HASH = 0;
317
318         // Use `MurmurHash2` <https://sites.google.com/site/murmurhash/>`, this is compatible with
319         // std::hash<string> in GNU libstdc++ 3.4.20 or above. This is typically the case when compiled
320         // on Linux and not macOS.
321         MURMUR_HASH_2 = 1;
322     }
323
324     reserved 2;
325
326     // Minimum hash ring size. The larger the ring is (that is, the more hashes there are for each
327     // provided host) the better the request distribution will reflect the desired weights. Defaults
328     // to 1024 entries, and limited to 8M entries. See also
329     // :ref:`maximum_ring_size<envoy_api_field_Cluster.RingHashLbConfig.maximum_ring_size>`.
330     google.protobuf.UInt64Value minimum_ring_size = 1 [(validate.rules).uint64 = {lte: 8388608}];
331
332     // The hash function used to hash hosts onto the ketama ring. The value defaults to
333     // :ref:`XX_HASH<envoy_api_enum_value_Cluster.RingHashLbConfig.HashFunction.XX_HASH>`.
334     HashFunction hash_function = 3 [(validate.rules).enum = {defined_only: true}];
335
336     // Maximum hash ring size. Defaults to 8M entries, and limited to 8M entries, but can be lowered
337     // to further constrain resource use. See also
338     // :ref:`minimum_ring_size<envoy_api_field_Cluster.RingHashLbConfig.minimum_ring_size>`.
339     google.protobuf.UInt64Value maximum_ring_size = 4 [(validate.rules).uint64 = {lte: 8388608}];
340 }
341
342 // Specific configuration for the
343 // :ref:`Original Destination <arch_overview_load_balancing_types_original_destination>`
344 // load balancing policy.
345 message OriginalDstLbConfig {
346     // When true, :ref:`x-envoy-original-dst-host`
347     // <config_http_conn_man_headers_x-envoy-original-dst-host>` can be used to override destination
348     // address.
349     //
350     // .. attention::
351     //
352     // This header isn't sanitized by default, so enabling this feature allows HTTP clients to
353     // route traffic to arbitrary hosts and/or ports, which may have serious security
354     // consequences.
355     //
356     // .. note::
357     //
358     // If the header appears multiple times only the first value is used.
359     bool use_http_header = 1;
360 }
361
362 // Common configuration for all load balancer implementations.
363 // [#next-free-field: 8]
364 message CommonLbConfig {
365     // Configuration for :ref:`zone aware routing`
366     // <arch_overview_load_balancing_zone_aware_routing>`.
367     message ZoneAwareLbConfig {
368         // Configures percentage of requests that will be considered for zone aware routing
369         // if zone aware routing is configured. If not specified, the default is 100%.
370         // * :ref:`runtime values <config_cluster_manager_cluster_runtime_zone_routing>`.
371         // * :ref:`Zone aware routing support <arch_overview_load_balancing_zone_aware_routing>`.
372         type.Percent routing_enabled = 1;

```

```

373
374 // Configures minimum upstream cluster size required for zone aware routing
375 // If upstream cluster size is less than specified, zone aware routing is not performed
376 // even if zone aware routing is configured. If not specified, the default is 6.
377 // * :ref:`runtime values <config_cluster_manager_cluster_runtime_zone_routing>`.
378 // * :ref:`Zone aware routing support <arch_overview_load_balancing_zone_aware_routing>`.
379 google.protobuf.UInt64Value min_cluster_size = 2;
380
381 // If set to true, Envoy will not consider any hosts when the cluster is in :ref:`panic
382 // mode<arch_overview_load_balancing_panic_threshold>`. Instead, the cluster will fail all
383 // requests as if all hosts are unhealthy. This can help avoid potentially overwhelming a
384 // failing service.
385 bool fail_traffic_on_panic = 3;
386 }
387
388 // Configuration for :ref:`locality weighted load balancing
389 // <arch_overview_load_balancing_locality_weighted_lb>`
390 message LocalityWeightedLbConfig {
391 }
392
393 // Common Configuration for all consistent hashing load balancers (MaglevLb, RingHashLb, etc.)
394 message ConsistentHashingLbConfig {
395 // If set to `true`, the cluster will use hostname instead of the resolved
396 // address as the key to consistently hash to an upstream host. Only valid for StrictDNS clusters with hostnames which resolve to a single IP address.
397 bool use_hostname_for_hashing = 1;
398 }
399
400 // Configures the :ref:`healthy panic threshold <arch_overview_load_balancing_panic_threshold>`.
401 // If not specified, the default is 50%.
402 // To disable panic mode, set to 0%.
403 //
404 // .. note::
405 // The specified percent will be truncated to the nearest 1%.
406 type.Percent healthy_panic_threshold = 1;
407
408 oneof locality_config_specifier {
409     ZoneAwareLbConfig zone_aware_lb_config = 2;
410
411     LocalityWeightedLbConfig locality_weighted_lb_config = 3;
412 }
413
414 // If set, all health check/weight/metadata updates that happen within this duration will be
415 // merged and delivered in one shot when the duration expires. The start of the duration is when
416 // the first update happens. This is useful for big clusters, with potentially noisy deploys
417 // that might trigger excessive CPU usage due to a constant stream of healthcheck state changes
418 // or metadata updates. The first set of updates to be seen apply immediately (e.g.: a new
419 // cluster). Please always keep in mind that the use of sandbox technologies may change this
420 // behavior.
421 //
422 // If this is not set, we default to a merge window of 1000ms. To disable it, set the merge
423 // window to 0.
424 //
425 // Note: merging does not apply to cluster membership changes (e.g.: adds/removes); this is
426 // because merging those updates isn't currently safe. See
427 // https://github.com/envoyproxy/envoy/pull/3941.
428 google.protobuf.Duration update_merge_window = 4;
429
430 // If set to true, Envoy will not consider new hosts when computing load balancing weights until
431 // they have been health checked for the first time. This will have no effect unless
432 // active health checking is also configured.
433 //
434 // Ignoring a host means that for any load balancing calculations that adjust weights based
435 // on the ratio of eligible hosts and total hosts (priority spillover, locality weighting and
436 // panic mode) Envoy will exclude these hosts in the denominator.
437 //
438 // For example, with hosts in two priorities P0 and P1, where P0 looks like
439 // {healthy, unhealthy (new), unhealthy (new)}
440 // and where P1 looks like
441 // {healthy, healthy}
442 // all traffic will still hit P0, as 1 / (3 - 2) = 1.
443 //
444 // Enabling this will allow scaling up the number of hosts for a given cluster without entering
445 // panic mode or triggering priority spillover, assuming the hosts pass the first health check.
446 //
447 // If panic mode is triggered, new hosts are still eligible for traffic; they simply do not
448 // contribute to the calculation when deciding whether panic mode is enabled or not.
449 bool ignore_new_hosts_until_first_hc = 5;
450
451 // If set to `true`, the cluster manager will drain all existing
452 // connections to upstream hosts whenever hosts are added or removed from the cluster.
453 bool close_connections_on_host_set_change = 6;
454
455 // Common Configuration for all consistent hashing load balancers (MaglevLb, RingHashLb, etc.)
456 ConsistentHashingLbConfig consistent_hashing_lb_config = 7;
457 }
458
459 message RefreshRate {
460 // Specifies the base interval between refreshes. This parameter is required and must be greater
461 // than zero and less than
462 // :ref:`max_interval <envoy_api_field_Cluster.RefreshRate.max_interval>`.
463 google.protobuf.Duration base_interval = 1 [(validate.rules).duration = {
464     required: true
465     gt {nanos: 1000000}
466 }];
467
468 // Specifies the maximum interval between refreshes. This parameter is optional, but must be
469 // greater than or equal to the
470 // :ref:`base_interval <envoy_api_field_Cluster.RefreshRate.base_interval>` if set. The default

```

```

471 // is 10 times the :ref:`base_interval` <envoy_api_field_Cluster.RefreshRate.base_interval>`.
472 google.protobuf.Duration max_interval = 2 [(validate.rules).duration = {gt {nanos: 1000000}}];
473 }
474
475 reserved 12, 15;
476
477 // Configuration to use different transport sockets for different endpoints.
478 // The entry of *envoy.transport_socket_match* in the
479 // :ref:`LbEndpoint.Metadata` <envoy_api_field_LbEndpoint.Metadata>`
480 // is used to match against the transport sockets as they appear in the list. The first
481 // :ref:`match` <envoy_api_msg_Cluster.TransportSocketMatch>` is used.
482 // For example, with the following match
483 //
484 // .. code-block:: yaml
485 //
486 //   transport_socket_matches:
487 //   - name: "enableMTLS"
488 //     match:
489 //       acceptMTLS: true
490 //     transport_socket:
491 //       name: envoy.transport_sockets.tls
492 //       config: { ... } # tls socket configuration
493 //   - name: "defaultToPlaintext"
494 //     match: {}
495 //     transport_socket:
496 //       name: envoy.transport_sockets.raw_buffer
497 //
498 // Connections to the endpoints whose metadata value under *envoy.transport_socket_match*
499 // having "acceptMTLS"/"true" key/value pair use the "enableMTLS" socket configuration.
500 //
501 // If a :ref:`socket match` <envoy_api_msg_Cluster.TransportSocketMatch>` with empty match
502 // criteria is provided, that always match any endpoint. For example, the "defaultToPlaintext"
503 // socket match in case above.
504 //
505 // If an endpoint metadata's value under *envoy.transport_socket_match* does not match any
506 // *TransportSocketMatch*, socket configuration fallbacks to use the *tls_context* or
507 // *transport_socket* specified in this cluster.
508 //
509 // This field allows gradual and flexible transport socket configuration changes.
510 //
511 // The metadata of endpoints in EDS can indicate transport socket capabilities. For example,
512 // an endpoint's metadata can have two key value pairs as "acceptMTLS": "true",
513 // "acceptPlaintext": "true". While some other endpoints, only accepting plaintext traffic
514 // has "acceptPlaintext": "true" metadata information.
515 //
516 // Then the xDS server can configure the CDS to a client, Envoy A, to send mutual TLS
517 // traffic for endpoints with "acceptMTLS": "true", by adding a corresponding
518 // *TransportSocketMatch* in this field. Other client Envoy's receive CDS without
519 // *transport_socket_match* set, and still send plain text traffic to the same cluster.
520 //
521 // [#comment:TODO(incfly): add a detailed architecture doc on intended usage.]
522 repeated TransportSocketMatch transport_socket_matches = 43;
523
524 // Supplies the name of the cluster which must be unique across all clusters.
525 // The cluster name is used when emitting
526 // :ref:`statistics` <config_cluster_manager_cluster_stats>` if :ref:`alt_stat_name`
527 // <envoy_api_field_Cluster.alt_stat_name>` is not provided.
528 // Any ``:`` in the cluster name will be converted to ``_`` when emitting statistics.
529 string name = 1 [(validate.rules).string = {min_bytes: 1}];
530
531 // An optional alternative to the cluster name to be used while emitting stats.
532 // Any ``:`` in the name will be converted to ``_`` when emitting statistics. This should not be
533 // confused with :ref:`Router Filter Header`
534 // <config_http_filters_router_x-envoy-upstream-alt-stat-name>`.
535 string alt_stat_name = 28;
536
537 oneof cluster_discovery_type {
538 // The :ref:`service discovery type` <arch_overview_service_discovery_types>`
539 // to use for resolving the cluster.
540 DiscoveryType type = 2 [(validate.rules).enum = {defined_only: true}];
541
542 // The custom cluster type.
543 CustomClusterType cluster_type = 38;
544 }
545
546 // Configuration to use for EDS updates for the Cluster.
547 EdsClusterConfig eds_cluster_config = 3;
548
549 // The timeout for new network connections to hosts in the cluster.
550 google.protobuf.Duration connect_timeout = 4 [(validate.rules).duration = {gt {}}];
551
552 // Soft limit on size of the cluster's connections read and write buffers. If
553 // unspecified, an implementation defined default is applied (1MiB).
554 google.protobuf.UInt32Value per_connection_buffer_limit_bytes = 5;
555
556 // The :ref:`load balancer type` <arch_overview_load_balancing_types>` to use
557 // when picking a host in the cluster.
558 LbPolicy lb_policy = 6 [(validate.rules).enum = {defined_only: true}];
559
560 // If the service discovery type is
561 // :ref:`STATIC` <envoy_api_enum_value_Cluster.DiscoveryType.STATIC>`,
562 // :ref:`STRICT_DNS` <envoy_api_enum_value_Cluster.DiscoveryType.STRICT_DNS>`
563 // or :ref:`LOGICAL_DNS` <envoy_api_enum_value_Cluster.DiscoveryType.LOGICAL_DNS>`,
564 // then hosts is required.
565 //
566 // .. attention::
567 //
568 //   **This field is deprecated**. Set the

```

```

569 // :ref:'load_assignment<envoy_api_field_Cluster.load_assignment>' field instead.
570 //
571 repeated core.Address hosts = 7 [deprecated = true];
572
573 // Setting this is required for specifying members of
574 // :ref:'STATIC<envoy_api_enum_value_Cluster.DiscoveryType.STATIC>',
575 // :ref:'STRICT_DNS<envoy_api_enum_value_Cluster.DiscoveryType.STRICT_DNS>'
576 // or :ref:'LOGICAL_DNS<envoy_api_enum_value_Cluster.DiscoveryType.LOGICAL_DNS>' clusters.
577 // This field supersedes the *hosts* field in the v2 API.
578 //
579 // .. attention::
580 //
581 // Setting this allows non-EDS cluster types to contain embedded EDS equivalent
582 // :ref:'endpoint assignments<envoy_api_msg_ClusterLoadAssignment>'.
583 //
584 ClusterLoadAssignment load_assignment = 33;
585
586 // Optional :ref:'active health checking <arch_overview_health_checking>'
587 // configuration for the cluster. If no
588 // configuration is specified no health checking will be done and all cluster
589 // members will be considered healthy at all times.
590 repeated core.HealthCheck health_checks = 8;
591
592 // Optional maximum requests for a single upstream connection. This parameter
593 // is respected by both the HTTP/1.1 and HTTP/2 connection pool
594 // implementations. If not specified, there is no limit. Setting this
595 // parameter to 1 will effectively disable keep alive.
596 google.protobuf.UInt32Value max_requests_per_connection = 9;
597
598 // Optional :ref:'circuit breaking <arch_overview_circuit_break>' for the cluster.
599 cluster.CircuitBreakers circuit_breakers = 10;
600
601 // The TLS configuration for connections to the upstream cluster.
602 //
603 // .. attention::
604 //
605 // **This field is deprecated**. Use 'transport_socket' with name 'tls' instead. If both are
606 // set, 'transport_socket' takes priority.
607 auth.UpstreamTlsContext tls_context = 11
608     [deprecated = true, (envoy.annotations.disallowed_by_default) = true];
609
610 // HTTP protocol options that are applied only to upstream HTTP connections.
611 // These options apply to all HTTP versions.
612 core.UpstreamHttpProtocolOptions upstream_http_protocol_options = 46;
613
614 // Additional options when handling HTTP requests upstream. These options will be applicable to
615 // both HTTP1 and HTTP2 requests.
616 core.HttpProtocolOptions common_http_protocol_options = 29;
617
618 // Additional options when handling HTTP1 requests.
619 core.Http1ProtocolOptions http_protocol_options = 13;
620
621 // Even if default HTTP2 protocol options are desired, this field must be
622 // set so that Envoy will assume that the upstream supports HTTP/2 when
623 // making new HTTP connection pool connections. Currently, Envoy only
624 // supports prior knowledge for upstream connections. Even if TLS is used
625 // with ALPN, 'http2_protocol_options' must be specified. As an aside this allows HTTP/2
626 // connections to happen over plain text.
627 core.Http2ProtocolOptions http2_protocol_options = 14;
628
629 // The extension_protocol_options field is used to provide extension-specific protocol options
630 // for upstream connections. The key should match the extension filter name, such as
631 // "envoy.filters.network.thrift_proxy". See the extension's documentation for details on
632 // specific options.
633 map<string, google.protobuf.Struct> extension_protocol_options = 35
634     [deprecated = true, (envoy.annotations.disallowed_by_default) = true];
635
636 // The extension_protocol_options field is used to provide extension-specific protocol options
637 // for upstream connections. The key should match the extension filter name, such as
638 // "envoy.filters.network.thrift_proxy". See the extension's documentation for details on
639 // specific options.
640 map<string, google.protobuf.Any> typed_extension_protocol_options = 36;
641
642 // If the DNS refresh rate is specified and the cluster type is either
643 // :ref:'STRICT_DNS<envoy_api_enum_value_Cluster.DiscoveryType.STRICT_DNS>',
644 // or :ref:'LOGICAL_DNS<envoy_api_enum_value_Cluster.DiscoveryType.LOGICAL_DNS>',
645 // this value is used as the cluster's DNS refresh
646 // rate. The value configured must be at least 1ms. If this setting is not specified, the
647 // value defaults to 5000ms. For cluster types other than
648 // :ref:'STRICT_DNS<envoy_api_enum_value_Cluster.DiscoveryType.STRICT_DNS>'
649 // and :ref:'LOGICAL_DNS<envoy_api_enum_value_Cluster.DiscoveryType.LOGICAL_DNS>'
650 // this setting is ignored.
651 google.protobuf.Duration dns_refresh_rate = 16
652     [(validate.rules).duration = {gt {nanos: 1000000}}];
653
654 // If the DNS failure refresh rate is specified and the cluster type is either
655 // :ref:'STRICT_DNS<envoy_api_enum_value_Cluster.DiscoveryType.STRICT_DNS>',
656 // or :ref:'LOGICAL_DNS<envoy_api_enum_value_Cluster.DiscoveryType.LOGICAL_DNS>',
657 // this is used as the cluster's DNS refresh rate when requests are failing. If this setting is
658 // not specified, the failure refresh rate defaults to the DNS refresh rate. For cluster types
659 // other than :ref:'STRICT_DNS<envoy_api_enum_value_Cluster.DiscoveryType.STRICT_DNS>' and
660 // :ref:'LOGICAL_DNS<envoy_api_enum_value_Cluster.DiscoveryType.LOGICAL_DNS>' this setting is
661 // ignored.
662 RefreshRate dns_failure_refresh_rate = 44;
663
664 // Optional configuration for setting cluster's DNS refresh rate. If the value is set to true,
665 // cluster's DNS refresh rate will be set to resource record's TTL which comes from DNS
666 // resolution.

```

```

667 bool respect_dns_ttl = 39;
668
669 // The DNS IP address resolution policy. If this setting is not specified, the
670 // value defaults to
671 // :ref:`AUTO<envoy_api_enum_value_Cluster.DnsLookupFamily.AUTO>`.
672 DnsLookupFamily dns_lookup_family = 17 [(validate.rules).enum = {defined_only: true}];
673
674 // If DNS resolvers are specified and the cluster type is either
675 // :ref:`STRICT_DNS<envoy_api_enum_value_Cluster.DiscoveryType.STRICT_DNS>`,
676 // or :ref:`LOGICAL_DNS<envoy_api_enum_value_Cluster.DiscoveryType.LOGICAL_DNS>`,
677 // this value is used to specify the cluster's dns resolvers.
678 // If this setting is not specified, the value defaults to the default
679 // resolver, which uses /etc/resolv.conf for configuration. For cluster types
680 // other than
681 // :ref:`STRICT_DNS<envoy_api_enum_value_Cluster.DiscoveryType.STRICT_DNS>`
682 // and :ref:`LOGICAL_DNS<envoy_api_enum_value_Cluster.DiscoveryType.LOGICAL_DNS>`
683 // this setting is ignored.
684 // Setting this value causes failure if the
685 // ``envoy.restart_features.use_apple_api_for_dns_lookups`` runtime value is true during
686 // server startup. Apple's API only allows overriding DNS resolvers via system settings.
687 repeated core.Address dns_resolvers = 18;
688
689 // [#next-major-version: Reconcile DNS options in a single message.]
690 // Always use TCP queries instead of UDP queries for DNS lookups.
691 // Setting this value causes failure if the
692 // ``envoy.restart_features.use_apple_api_for_dns_lookups`` runtime value is true during
693 // server startup. Apple's API only uses UDP for DNS resolution.
694 bool use_tcp_for_dns_lookups = 45;
695
696 // If specified, outlier detection will be enabled for this upstream cluster.
697 // Each of the configuration values can be overridden via
698 // :ref:`runtime values <config_cluster_manager_cluster_runtime_outlier_detections>`.
699 cluster.OutlierDetection outlier_detection = 19;
700
701 // The interval for removing stale hosts from a cluster type
702 // :ref:`ORIGINAL_DST<envoy_api_enum_value_Cluster.DiscoveryType.ORIGINAL_DST>`.
703 // Hosts are considered stale if they have not been used
704 // as upstream destinations during this interval. New hosts are added
705 // to original destination clusters on demand as new connections are
706 // redirected to Envoy, causing the number of hosts in the cluster to
707 // grow over time. Hosts that are not stale (they are actively used as
708 // destinations) are kept in the cluster, which allows connections to
709 // them remain open, saving the latency that would otherwise be spent
710 // on opening new connections. If this setting is not specified, the
711 // value defaults to 5000ms. For cluster types other than
712 // :ref:`ORIGINAL_DST<envoy_api_enum_value_Cluster.DiscoveryType.ORIGINAL_DST>`
713 // this setting is ignored.
714 google.protobuf.Duration cleanup_interval = 20 [(validate.rules).duration = {gt {}}];
715
716 // Optional configuration used to bind newly established upstream connections.
717 // This overrides any bind_config specified in the bootstrap proto.
718 // If the address and port are empty, no bind will be performed.
719 core.BindConfig upstream_bind_config = 21;
720
721 // Configuration for load balancing subsetting.
722 lbSubsetConfig lb_subset_config = 22;
723
724 // Optional configuration for the load balancing algorithm selected by
725 // LbPolicy. Currently only
726 // :ref:`RING_HASH<envoy_api_enum_value_Cluster.LbPolicy.RING_HASH>` and
727 // :ref:`LEAST_REQUEST<envoy_api_enum_value_Cluster.LbPolicy.LEAST_REQUEST>`
728 // has additional configuration options.
729 // Specifying ring_hash_lb_config or least_request_lb_config without setting the corresponding
730 // LbPolicy will generate an error at runtime.
731 oneof lb_config {
732     // Optional configuration for the Ring Hash load balancing policy.
733     RingHashLbConfig ring_hash_lb_config = 23;
734
735     // Optional configuration for the Original Destination load balancing policy.
736     OriginalDstLbConfig original_dst_lb_config = 34;
737
738     // Optional configuration for the LeastRequest load balancing policy.
739     LeastRequestLbConfig least_request_lb_config = 37;
740 }
741
742 // Common configuration for all load balancer implementations.
743 CommonLbConfig common_lb_config = 27;
744
745 // Optional custom transport socket implementation to use for upstream connections.
746 // To setup TLS, set a transport socket with name `tls` and
747 // :ref:`UpstreamTlsContexts <envoy_api_msg_auth.UpstreamTlsContext>` in the `typed_config`.
748 // If no transport socket configuration is specified, new connections
749 // will be set up with plaintext.
750 core.TransportSocket transport_socket = 24;
751
752 // The Metadata field can be used to provide additional information about the
753 // cluster. It can be used for stats, logging, and varying filter behavior.
754 // Fields should use reverse DNS notation to denote which entity within Envoy
755 // will need the information. For instance, if the metadata is intended for
756 // the Router filter, the filter name should be specified as `envoy.filters.http.router`.
757 core.Metadata metadata = 25;
758
759 // Determines how Envoy selects the protocol used to speak to upstream hosts.
760 ClusterProtocolSelection protocol_selection = 26;
761
762 // Optional options for upstream connections.
763 UpstreamConnectionOptions upstream_connection_options = 30;
764

```



```

765 // If an upstream host becomes unhealthy (as determined by the configured health checks
766 // or outlier detection), immediately close all connections to the failed host.
767 //
768 // .. note::
769 //
770 // This is currently only supported for connections created by tcp_proxy.
771 //
772 // .. note::
773 //
774 // The current implementation of this feature closes all connections immediately when
775 // the unhealthy status is detected. If there are a large number of connections open
776 // to an upstream host that becomes unhealthy, Envoy may spend a substantial amount of
777 // time exclusively closing these connections, and not processing any other traffic.
778 bool close_connections_on_host_health_failure = 31;
779
780 // If set to true, Envoy will ignore the health value of a host when processing its removal
781 // from service discovery. This means that if active health checking is used, Envoy will *not*
782 // wait for the endpoint to go unhealthy before removing it.
783 bool drain_connections_on_host_removal = 32
784   [(udpa.annotations.field_migrate).rename = "ignore_health_on_host_removal"];
785
786 // An (optional) network filter chain, listed in the order the filters should be applied.
787 // The chain will be applied to all outgoing connections that Envoy makes to the upstream
788 // servers of this cluster.
789 repeated cluster.Filter filters = 40;
790
791 // [#not-implemented-hide:] New mechanism for LB policy configuration. Used only if the
792 // :ref:`lb_policy<envoy_api_field_Cluster.lb_policy>` field has the value
793 // :ref:`LOAD_BALANCING_POLICY_CONFIG<envoy_api_enum_value_Cluster.LbPolicy.LOAD_BALANCING_POLICY_CONFIG>`.
794 LoadBalancingPolicy load_balancing_policy = 41;
795
796 // [#not-implemented-hide:]
797 // If present, tells the client where to send load reports via LRS. If not present, the
798 // client will fall back to a client-side default, which may be either (a) don't send any
799 // load reports or (b) send load reports for all clusters to a single default server
800 // (which may be configured in the bootstrap file).
801 //
802 // Note that if multiple clusters point to the same LRS server, the client may choose to
803 // create a separate stream for each cluster or it may choose to coalesce the data for
804 // multiple clusters onto a single stream. Either way, the client must make sure to send
805 // the data for any given cluster on no more than one stream.
806 //
807 // [#next-major-version: In the v3 API, we should consider restructuring this somehow,
808 // maybe by allowing LRS to go on the ADS stream, or maybe by moving some of the negotiation
809 // from the LRS stream here.]
810 core.ConfigSource lrs_server = 42;
811
812 // If track_timeout_budgets is true, the :ref:`timeout budget histograms
813 // <config_cluster_manager_cluster_stats_timeout_budgets>` will be published for each
814 // request. These show what percentage of a request's per try and global timeout was used. A value
815 // of 0 would indicate that none of the timeout was used or that the timeout was infinite. A value
816 // of 100 would indicate that the request took the entirety of the timeout given to it.
817 bool track_timeout_budgets = 47;
818 }
819
820 // [#not-implemented-hide:] Extensible load balancing policy configuration.
821 //
822 // Every LB policy defined via this mechanism will be identified via a unique name using reverse
823 // DNS notation. If the policy needs configuration parameters, it must define a message for its
824 // own configuration, which will be stored in the config field. The name of the policy will tell
825 // clients which type of message they should expect to see in the config field.
826 //
827 // Note that there are cases where it is useful to be able to independently select LB policies
828 // for choosing a locality and for choosing an endpoint within that locality. For example, a
829 // given deployment may always use the same policy to choose the locality, but for choosing the
830 // endpoint within the locality, some clusters may use weighted-round-robin, while others may
831 // use some sort of session-based balancing.
832 //
833 // This can be accomplished via hierarchical LB policies, where the parent LB policy creates a
834 // child LB policy for each locality. For each request, the parent chooses the locality and then
835 // delegates to the child policy for that locality to choose the endpoint within the locality.
836 //
837 // To facilitate this, the config message for the top-level LB policy may include a field of
838 // type LoadBalancingPolicy that specifies the child policy.
839 message LoadBalancingPolicy {
840   message Policy {
841     // Required. The name of the LB policy.
842     string name = 1;
843
844     // Optional config for the LB policy.
845     // No more than one of these two fields may be populated.
846     google.protobuf.Struct config = 2 [deprecated = true];
847
848     google.protobuf.Any typed_config = 3;
849   }
850
851   // Each client will iterate over the list in order and stop at the first policy that it
852   // supports. This provides a mechanism for starting to use new LB policies that are not yet
853   // supported by all clients.
854   repeated Policy policies = 1;
855 }
856
857 // An extensible structure containing the address Envoy should bind to when
858 // establishing upstream connections.
859 message UpstreamBindConfig {
860   // The address Envoy should bind to when establishing upstream connections.
861   core.Address source_address = 1;
862 }

```

```
863 |
864 | message UpstreamConnectionOptions {
865 |     // If set then set SO_KEEPALIVE on the socket to enable TCP Keepalives.
866 |     core.TcpKeepalive tcp_keepalive = 1;
867 | }
```