

Heap buffer overflow caused by rounding

Low mihairaruseac published GHSA-jfp7-4j67-8r3q on May 12, 2021

Package

tensorflow, tensorflow-cpu, tensorflow-gpu (pip)

Affected versions

< 2.5.0

Patched versions

2.1.4, 2.2.3, 2.3.3, 2.4.2

Description

Impact

An attacker can trigger a heap buffer overflow in `tf.raw_ops.QuantizedResizeBilinear` by manipulating input values so that float rounding results in off-by-one error in accessing image elements:

```
import tensorflow as tf

l = [256, 328, 361, 17, 361, 361, 361, 361, 361, 361, 361, 361, 361, 361, 384]
images = tf.constant(l, shape=[1, 1, 15, 1], dtype=tf.qint32)
size = tf.constant([12, 6], shape=[2], dtype=tf.int32)
min = tf.constant(80.22522735595783)
max = tf.constant(80.3921585830078)

tf.raw_ops.QuantizedResizeBilinear(images=images, size=size, min=min, max=max,
                                  align_corners=True, half_pixel_centers=True)
```

This is because the [implementation](#) computes two integers (representing the upper and lower bounds for interpolation) by ceiling and flooring a floating point value:

```
const float in_f = std::floor(in);
interpolation->lower[i] = std::max(static_cast<int64>(in_f), static_cast<int64>(0));
interpolation->upper[i] = std::min(static_cast<int64>(std::ceil(in)), in_size - 1);
```

For some values of `in`, `interpolation->upper[i]` might be smaller than `interpolation->lower[i]`. This is an issue if `interpolation->upper[i]` is capped at `in_size-1` as it means that `interpolation->lower[i]` points outside of the image. Then, in the [interpolation code](#), this would result in heap buffer overflow:

```
template<int RESOLUTION, typename T, typename T_SCALE, typename T_CALC>
inline void OutputLerpForChannels(const InterpolationCache<T_SCALE>& xs,
                                const int64 x, const T_SCALE ys_ilerp,
                                const int channels, const float min,
                                const float max, const T* ys_input_lower_ptr,
                                const T* ys_input_upper_ptr,
                                T* output_y_ptr) {
    const int64 xs_lower = xs.lower[x];
    ...
    for (int c = 0; c < channels; ++c) {
        const T top_left = ys_input_lower_ptr[xs_lower + c];
        ...
    }
}
```

For the other cases where `interpolation->upper[i]` is smaller than `interpolation->lower[i]`, we can set them to be equal without affecting the output.

Patches

We have patched the issue in GitHub commit [f851613f8f0fb0c838d160ced13c134f778e3ce7](#).

The fix will be included in TensorFlow 2.5.0. We will also cherry-pick this commit on TensorFlow 2.4.2, TensorFlow 2.3.3, TensorFlow 2.2.3 and TensorFlow 2.1.4, as these are also affected and still in supported range.

For more information

Please consult [our security guide](#) for more information regarding the security model and how to contact us with issues and questions.

Attribution

This vulnerability has been reported by Ying Wang and Yakun Zhang of Baidu X-Team.

Severity

Low

CVE ID

CVE-2021-29529

Weaknesses

No CVEs