

Bug 27896 (CVE-2021-33574) - mq_notify does not handle separately allocated thread attributes (CVE-2021-33574)

Status: RESOLVED FIXED

Alias: CVE-2021-33574

Product: glibc

Component: nptl (show other bugs)

Version: 2.34

Importance: P2 normal

Target Milestone: 2.34

Assignee: Not yet assigned to anyone

URL:

Keywords:

Depends on:

Blocks:

Reported: 2021-05-21 08:54 UTC by Florian Weimer

Modified: 2021-08-02 00:51 UTC (History)

CC List: 7 users (show)

See Also:

Host:

Target:

Build:

Last reconfirmed:

Flags: siddhesh: security+

Attachments	
adapt for 2.28 (1.10 KB, application/mbox) 2021-06-17 02:00 UTC, liqingqing	Details
Add an attachment (proposed patch, testcase, etc.) View All	

Note
You need to [log in](#) before you can comment on or make changes to this bug.

Florian Weimer2021-05-21 08:54:12 UTC

Description

mq_notify makes a shallow copy of pthread_attr_t here:

```
if (notification->sigev_notify_attributes != NULL)
{
    /* The thread attribute has to be allocated separately. */
    data.attr = (pthread_attr_t *) malloc (sizeof (pthread_attr_t));
    if (data.attr == NULL)
        return -1;

    memcpy (data.attr, notification->sigev_notify_attributes,
           sizeof (pthread_attr_t));
}
```

This introduces a potential for a use-after-free bug because the affinity mask has been separately allocated, since before the addition of mq_notify. (A caller of mq_notify can call pthread_attr_destroy immediately after mq_notify returns and before the new thread is created.)

Found through code inspection. No known application impact.

Siddhesh Poyarekar2021-05-31 06:42:58 UTC

Comment 1

Sorry I missed this in my Mitre CVE report: the only use-after-free indirection is through the extensions member of struct pthread_attr (see sysdeps/nptl/internals.h) and it got introduced in glibc-2.32. As a result, only glibc-2.32 and glibc-2.33 have a use-after-free.

Siddhesh Poyarekar2021-05-31 07:24:46 UTC

Comment 2

(In reply to Siddhesh Poyarekar from [comment #1](#))
> Sorry I missed this in my Mitre CVE report: the only use-after-free
> indirection is through the extensions member of struct pthread_attr (see
> sysdeps/nptl/internals.h) and it got introduced in glibc-2.32. As a
> result, only glibc-2.32 and glibc-2.33 have a use-after-free.

Sorry again, this does in fact affect all versions of glibc because even though extensions were introduced in 2.32, the cpuset before that were also an additional indirection and hence would result in a similar use-after-free.

Andreas Schwab2021-06-01 15:14:23 UTC

Comment 3

Fixed in 2.34.

Siddhesh Poyarekar2021-06-03 05:54:38 UTC

Comment 4

Fixed commits:

<https://sourceware.org/git/?p=glibc.git;a=commit;h=42d359350510506b87101cf77202fefcbfc790cb>
<https://sourceware.org/git/?p=glibc.git;a=commit;h=217b6dc298156bdb0d6aea9ea93e7e394a5ff091>

A note on the security impact based on my analysis of the bug. In order to mount a minimal attack using this flaw, an attacker needs many pre-requisites to be able to even crash a program using this mq_notify bug:

1. The program call to mq_notify needs to be controlled by the attacker
2. The program must provide attributes to control creation of the notification thread in mq_notify
3. The program must have the race condition where it may potentially destroy the notification thread attributes before the notification thread is created
4. The program must set CPU affinity or signal mask of the notification thread to actually cause the use-after-free dereference

There are no known applications in distributions that have *all* these pre-requisites and it's quite rare to have all of those conditions together, so I reckon the attack complexity is very high.

In the worst (or theoretical) case that such an application exists, an attacker would at best be able to control on which CPU the notification thread runs; the signal mask, even if set is overwritten (i.e. all signals unblocked) before the notification function is called. The change in scheduling should not have an impact on confidentiality or integrity of a compliant application.

liqingqing2021-06-04 03:56:21 UTC

Comment 5

hi all, what about the attr->stack? I think if one thread can destroy cpuset, means that it can also free the other memory.

Siddhesh Poyarekar2021-06-04 04:19:41 UTC

Comment 6

The thread stack memory is not freed with pthread_attr destroy; it continues to be reachable and valid since the thread would need it. Just like with regular threads, it is the responsibility of the application to ensure that the stack memory remains reachable and valid for the duration of the notification thread.

manojh3012 2021-06-16 17:43:27 UTC

[Comment 7](#)

The suggested patch doesn't work for glibc 2.28 since `__pthread_attr_copy` is not available in that version. Any suggestions/fixes?

Florian Weimer 2021-06-16 18:18:43 UTC

[Comment 8](#)

(In reply to manojh3012 from [comment #7](#))
> The suggested patch doesn't work for glibc 2.28 since `__pthread_attr_copy`
> is not available in that version. Any suggestions/fixes?

You need to backport this additional commit:

commit 331c6e8a184167dd21a9f0b3fcl65aeefea6eeca
Author: Florian Weimer <fweimer@redhat.com>
Date: Tue May 19 12:32:39 2020 +0200

nptl: Add __pthread_attr_copy for copying pthread_attr_t objects

It should be self-contained.

manojh3012 2021-06-17 01:26:12 UTC

[Comment 9](#)

Thanks but that commit uses __pthread_attr_setaffinity_np which is not available in 2.28 as well. I have to backport other commits looks like to bring in __pthread_attr_setaffinity_np. Any pointers on what is needed for that?

Also, is it possible to create a patch suitable for older glibc versions that don't have __pthread_attr_copy()?

liqingqing 2021-06-17 02:00:54 UTC

[Comment 10](#)

Created [attachment 13497](#) [[details](#)]
adapt for 2.28

hi, all, how about this two patches? there are all from glibc upstream, and I had do some modification for the second patch.

```
diff --git a/sysdeps/unix/sysv/linux/mq_notify.c
b/sysdeps/unix/sysv/linux/mq_notify.c
index c4091169..76963567 100644
--- a/sysdeps/unix/sysv/linux/mq_notify.c
+++ b/sysdeps/unix/sysv/linux/mq_notify.c
@@ -260,7 +260,34 @@ mq_notify (mqd_t mqdes, const struct sigevent *notification)
     if (data.attr == NULL)
         return -1;

-    __pthread_attr_copy (data.attr, notification->sigev_notify_attributes);
+    memcpy (data.attr, notification->sigev_notify_attributes,
+           sizeof (pthread_attr_t));
+
+    struct pthread_attr *source =
+        (struct pthread_attr *) (notification->sigev_notify_attributes);
+    struct pthread_attr *target = (struct pthread_attr *) (data.attr);
+    cpu_set_t *newp;
+    cpu_set_t *cpuset = source->cpuset;
+    size_t cpusetsize = source->cpusetsize;
+
+    /* alloc a new memory for cpuset to avoid use after free */
+    if (cpuset != NULL && cpusetsize > 0)
+    {
+        newp = (cpu_set_t *) malloc (cpusetsize);
+        if (newp == NULL)
+        {
+            free (data.attr);
+            return -1;
+        }
+        memcpy (newp, cpuset, cpusetsize);
+        target->cpuset = newp;
+    }
+    else
+    {
+        target->cpuset = NULL;
+        target->cpusetsize = 0;
+    }
+}

/* Construct the new request. */
@@ -273,7 +300,7 @@ mq_notify (mqd_t mqdes, const struct sigevent *notification)
int retval = INLINE_SYSCALL (mq_notify, 2, mqdes, &se);

/* If it failed, free the allocated memory. */
- if (__glibc_unlikely (retval != 0))
+ if (retval != 0 && data.attr != NULL)
 {
     pthread_attr_destroy (data.attr);
     free (data.attr);
--
```

details:
https://sourceware.org/git/?p=openEuler/glibc/blob/openEuler-20.03-LTS-SP2/backport-CVE-2021-33574-0002-Fix-mq_notify-bug-27896.patch

liqingqing 2021-06-17 02:02:55 UTC

[Comment 11](#)

(In reply to liqingqing from [comment #10](#))
> Created [attachment 13497](#) [[details](#)]
> adapt for 2.28
>
> hi, all, how about this two patches? there are all from glibc upstream, and
> I had do some modification for the second patch.
>
>
> diff --git a/sysdeps/unix/sysv/linux/mq_notify.c
> b/sysdeps/unix/sysv/linux/mq_notify.c
> index c4091169..76963567 100644
> --- a/sysdeps/unix/sysv/linux/mq_notify.c
> +++ b/sysdeps/unix/sysv/linux/mq_notify.c
> @@ -260,7 +260,34 @@ mq_notify (mqd_t mqdes, const struct sigevent
> *notification)
> if (data.attr == NULL)
> return -1;
>
> - __pthread_attr_copy (data.attr,
> notification->sigev_notify_attributes);
> + memcpy (data.attr, notification->sigev_notify_attributes,
> + sizeof (pthread_attr_t));
> +
> + struct pthread_attr *source =
> + (struct pthread_attr *) (notification->sigev_notify_attributes);
> + struct pthread_attr *target = (struct pthread_attr *) (data.attr);

```

> +     cpu_set_t *newp;
> +     cpu_set_t *cpuset = source->cpuset;
> +     siz_t Cpusetsize = source->cpusetsize;
> +
> +     /* alloc a new memory for cpuset to avoid use after free */
> +     if (cpuset != NULL && cpusetsize > 0)
> +     {
> +         newp = (cpu_set_t *) malloc (cpusetsize);
> +         if (newp == NULL)
> +         {
> +             free(data.attr);
> +             return -1;
> +         }
> +         memcpy (newp, cpuset, cpusetsize);
> +         target->cpuset = newp;
> +     }
> +     else
> +     {
> +         target->cpuset = NULL;
> +         target->cpusetsize = 0;
> +     }
> + }
> +
> + /* Construct the new request. */
> + @ -273,7 +300,7 @@ mq_notify (mqd_t mqdes, const struct sigevent
> + *notification)
> + {
> +     int retval = INLINE_SYSCALL (mq_notify, 2, mqdes, &se);
> +
> +     /* If it failed, free the allocated memory. */
> +     if (__glibc_unlikely (retval != 0))
> +     if (retval != 0 && data.attr != NULL)
> +     {
> +         pthread_attr_destroy (data.attr);
> +         free (data.attr);
> +     }
> + }
> +
> + details:
> + https://gitee.com/src-openeuler/glibc/blob/openEuler-20.03-LTS-SP2/backport-
> + CVE-2021-33574-0001-Fix-mq\_notify-bug-27896.patch
> +
> + the first one is commit: 42d359350510506b87101cf77202fefcbfc790cb :
> + https://gitee.com/src-openeuler/glibc/blob/openEuler-20.03-LTS-SP2/backport-
> + CVE-2021-33574-0001-Fix-mq\_notify-bug-27896.patch

```

Liming Liu 2021-08-01 15:54:44 UTC

[Comment 12](#)

when will the 2.28 in buster be patched? thanks.

Carlos O'Donnell 2021-08-02 00:51:52 UTC

[Comment 13](#)

(In reply to Liming Liu from [comment #12](#))
> when will the 2.28 in buster be patched? thanks.

This is the upstream glibc bug tracker. Please consider reporting your issue with the Debian glibc team. It looks like this has already been reported here: <https://bugs.debian.org/cgi-bin/bugreport.cgi?bug=989147>, I would follow up there with Debian.