

btrfs: fix deadlock with concurrent chunk allocations involving syste...

Browse files

_m chunks

When a task attempting to allocate a new chunk verifies that there is not currently enough free space in the system space_info and there is another task that allocated a new system chunk but it did not finish yet the creation of the respective block group, it waits for that other task to finish creating the block group. This is to avoid exhaustion of the system chunk array in the superblock, which is limited, when we have a thundering herd of tasks allocating new chunks. This problem was described and fixed by commit [eafa4fd](#) ("btrfs: fix exhaustion of the system chunk array due to concurrent allocations").

However there are two very similar scenarios where this can lead to a deadlock:

- 1) Task B allocated a new system chunk and task A is waiting on task B to finish creation of the respective system block group. However before task B ends its transaction handle and finishes the creation of the system block group, it attempts to allocate another chunk (like a data chunk for an fallocate operation for a very large range). Task B will be unable to progress and allocate the new chunk, because task A set space_info->chunk_alloc to 1 and therefore it loops at btrfs_chunk_alloc() waiting for task A to finish its chunk allocation and set space_info->chunk_alloc to 0, but task A is waiting on task B to finish creation of the new system block group, therefore resulting in a deadlock;
- 2) Task B allocated a new system chunk and task A is waiting on task B to finish creation of the respective system block group. By the time that task B enter the final phase of block group allocation, which happens at btrfs_create_pending_block_groups(), when it modifies the extent tree, the device tree or the chunk tree to insert the items for some new block group, it needs to allocate a new chunk, so it ends up at btrfs_chunk_alloc() and keeps looping there because task A has set space_info->chunk_alloc to 1, but task A is waiting for task B to finish creation of the new system block group and release the reserved system space, therefore resulting in a deadlock.

In short, the problem is if a task B needs to allocate a new chunk after it previously allocated a new system chunk and if another task A is currently waiting for task B to complete the allocation of the new system chunk.

Unfortunately this deadlock scenario introduced by the previous fix for the system chunk array exhaustion problem does not have a simple and short fix, and requires a big change to rework the chunk allocation code so that chunk btree updates are all made in the first phase of chunk allocation. And since this deadlock regression is being frequently hit on zoned filesystems and the system chunk array exhaustion problem is triggered in more extreme cases (originally observed on PowerPC with a node size of 64K when running the fallocate tests from stress-ng), revert the changes from that commit. The next patch in the series, with a subject of "btrfs: rework chunk allocation to avoid exhaustion of the system chunk array" does the necessary changes to fix the system chunk array exhaustion problem.

Reported-by: Naohiro Aota <naohiro.aota@wdc.com>
Link: <https://lore.kernel.org/linux-btrfs/20210621015922.ewgbffxuawia7liz@naota-xeon/>
Fixes: [eafa4fd](#) ("btrfs: fix exhaustion of the system chunk array due to concurrent allocations")
CC: stable@vger.kernel.org # 5.12+
Tested-by: Shin'ichiro Kawasaki <shinichiro.kawasaki@wdc.com>
Tested-by: Naohiro Aota <naohiro.aota@wdc.com>
Signed-off-by: Filipe Manana <fdmanana@suse.com>
Tested-by: David Sterba <dsterba@suse.com>
Signed-off-by: David Sterba <dsterba@suse.com>

👤 master
📦 v6.1 ... v5.14-rc2

👤 fdmanana authored and kdave committed on Jul 7, 2021 parent 5f93e77 commit 1cb3db1cf383a3c7dbda1aa0ce748b0958759947

Showing 3 changed files with 1 addition and 69 deletions.

Split Unified

```
58 fs/btrfs/block-group.c
3377 3377 */
3378 3378 void check_system_chunk(struct btrfs_trans_handle *trans, u64 type)
3379 3379 {
3380 3380 - struct btrfs_transaction *cur_trans = trans->transaction;
3381 3380 struct btrfs_fs_info *fs_info = trans->fs_info;
3382 3381 struct btrfs_space_info *info;
3383 3382 u64 left;
3392 3391 lockdep_assert_held(&fs_info->chunk_mutex);
3393 3392
3394 3393 info = btrfs_find_space_info(fs_info, BTRFS_BLOCK_GROUP_SYSTEM);
3395 3393 - again:
3396 3394 spin_lock(&info->lock);
3397 3395 left = info->total_bytes - btrfs_space_info_used(info, true);
3398 3396 spin_unlock(&info->lock);
3411 3409
3412 3410 if (left < thresh) {
3413 3411 u64 flags = btrfs_system_alloc_profile(fs_info);
```

```

3414         u64 reserved = atomic64_read(&cur_trans->chunk_bytes_reserved);
3415
3416     /*
3417     * If there's not available space for the chunk tree (system
3418     * space) and there are other tasks that reserved space for
3419     * creating a new system block group, wait for them to complete
3420     * the creation of their system block group and release excess
3421     * reserved space. We do this because:
3422     *
3423     * *) We can end up allocating more system chunks than necessary
3424     * when there are multiple tasks that are concurrently
3425     * allocating block groups, which can lead to exhaustion of
3426     * the system array in the superblock;
3427     *
3428     * *) If we allocate extra and unnecessary system block groups,
3429     * despite being empty for a long time, and possibly forever,
3430     * they end not being added to the list of unused block groups
3431     * because that typically happens only when deallocating the
3432     * last extent from a block group - which never happens since
3433     * we never allocate from them in the first place. The few
3434     * exceptions are when mounting a filesystem or running scrub,
3435     * which add unused block groups to the list of unused block
3436     * groups, to be deleted by the cleaner kthread.
3437     * And even when they are added to the list of unused block
3438     * groups, it can take a long time until they get deleted,
3439     * since the cleaner kthread might be sleeping or busy with
3440     * other work (deleting subvolumes, running delayed inputs,
3441     * defrag scheduling, etc);
3442     *
3443     * This is rare in practice, but can happen when too many tasks
3444     * are allocating blocks groups in parallel (via fallocation())
3445     * and before the one that reserved space for a new system block
3446     * group finishes the block group creation and releases the space
3447     * reserved in excess (at btrfs_create_pending_block_groups()),
3448     * other tasks end up here and see free system space temporarily
3449     * not enough for updating the chunk tree.
3450     *
3451     * We unlock the chunk mutex before waiting for such tasks and
3452     * lock it again after the wait, otherwise we would deadlock.
3453     * It is safe to do so because allocating a system chunk is the
3454     * first thing done while allocating a new block group.
3455     */
3456     if (reserved > trans->chunk_bytes_reserved) {
3457         const u64 min_needed = reserved - thresh;
3458
3459         mutex_unlock(&fs_info->chunk_mutex);
3460         wait_event(cur_trans->chunk_reserve_wait,
3461             atomic64_read(&cur_trans->chunk_bytes_reserved) <=
3462             min_needed);
3463         mutex_lock(&fs_info->chunk_mutex);
3464         goto again;
3465     }
3466
3467     /*
3468     * Ignore failure to create system chunk. We might end up not
3469     * ret = btrfs_block_rsv_add(fs_info->chunk_root,
3470         &fs_info->chunk_block_rsv,
3471         thresh, BTRFS_RESERVE_NO_FLUSH);
3472
3473     if (!ret) {
3474         atomic64_add(thresh, &cur_trans->chunk_bytes_reserved);
3475     }
3476     if (!ret)
3477         trans->chunk_bytes_reserved += thresh;
3478 }
3479
3480 }
3481
3482 }
3483
3484 }
3485
3486 }

```

fs/btrfs/transaction.c

```

260 void btrfs_trans_release_chunk_metadata(struct btrfs_trans_handle *trans)
261 {
262     struct btrfs_fs_info *fs_info = trans->fs_info;
263     struct btrfs_transaction *cur_trans = trans->transaction;
264
265     if (!trans->chunk_bytes_reserved)
266         return;
267
268     btrfs_block_rsv_release(fs_info, &fs_info->chunk_block_rsv,
269         trans->chunk_bytes_reserved, NULL);
270
271     atomic64_sub(trans->chunk_bytes_reserved, &cur_trans->chunk_bytes_reserved);
272     cond_wake_up(&cur_trans->chunk_reserve_wait);
273     trans->chunk_bytes_reserved = 0;
274 }
275
276
277
278
279
280
281
282
283 spin_lock_init(&cur_trans->dropped_roots_lock);
284 INIT_LIST_HEAD(&cur_trans->releasing_ebs);
285 spin_lock_init(&cur_trans->releasing_ebs_lock);
286 atomic64_set(&cur_trans->chunk_bytes_reserved, 0);
287 init_waitqueue_head(&cur_trans->chunk_reserve_wait);
288 list_add_tail(&cur_trans->list, &fs_info->trans_list);
289 extent_io_tree_init(fs_info, &cur_trans->dirty_pages,
290     IO_TREE_TRANS_DIRTY_PAGES, fs_info->btree_inode);

```

fs/btrfs/transaction.h

```

96
97 spinlock_t releasing_ebs_lock;

```

```
98 98 struct list_head releasing_ebs;
99 -
100 - /*
101 -  * The number of bytes currently reserved, by all transaction handles
102 -  * attached to this transaction, for metadata extents of the chunk tree.
103 -  */
104 - atomic64_t chunk_bytes_reserved;
105 - wait_queue_head_t chunk_reserve_wait;
106 99 };
107 100
108 101 #define __TRANS_FREEZABLE (1U << 0)
```

0 comments on commit [1cb3db1](#)

Please [sign in](#) to comment.