

git\_connect\_git(): forbid newlines in host and path

Browse files

When we connect to a git:// server, we send an initial request that looks something like:

```
002dgit-upload-pack repo.git@host=example.com
```

If the repo path contains a newline, then it's included literally, and we get:

```
002egit-upload-pack repo
.git@host=example.com
```

This works fine if you really do have a newline in your repository name; the server side uses the pktline framing to parse the string, not newlines. However, there are many \_other\_ protocols in the wild that do parse on newlines, such as HTTP. So a carefully constructed git:// URL can actually turn into a valid HTTP request. For example:

```
git://localhost:1234/%0d%0a%0d%0aGET%20/%20HTTP/1.1 %0d%0aHost:localhost%0d%0a%0d%0a
```

becomes:

```
0050git-upload-pack /
GET / HTTP/1.1
Host:localhost

host=localhost:1234
```

on the wire. Again, this isn't a problem for a real Git server, but it does mean that feeding a malicious URL to Git (e.g., through a submodule) can cause it to make unexpected cross-protocol requests. Since repository names with newlines are presumably quite rare (and indeed, we already disallow them in git-over-http), let's just disallow them over this protocol.

Hostnames could likewise inject a newline, but this is unlikely a problem in practice; we'd try resolving the hostname with a newline in it, which wouldn't work. Still, it doesn't hurt to err on the side of caution there, since we would not expect them to work in the first place.

The ssh and local code paths are unaffected by this patch. In both cases we're trying to run upload-pack via a shell, and will quote the newline so that it makes it intact. An attacker can point an ssh url at an arbitrary port, of course, but unless there's an actual ssh server there, we'd never get as far as sending our shell command anyway. We \_could\_ similarly restrict newlines in those protocols out of caution, but there seems little benefit to doing so.

The new test here is run alongside the git-daemon tests, which cover the same protocol, but it shouldn't actually contact the daemon at all. In theory we could make the test more robust by setting up an actual repository with a newline in it (so that our clone would succeed if our new check didn't kick in). But a repo directory with newline in it is likely not portable across all filesystems. Likewise, we could check git-daemon's log that it was not contacted at all, but we do not currently record the log (and anyway, it would make the test racy with the daemon's log write). We'll just check the client-side stderr to make sure we hit the expected code path.

Reported-by: Harold Kim <h.kim@flatt.tech>  
Signed-off-by: Jeff King <peff@peff.net>  
Signed-off-by: Junio C Hamano <gitster@pobox.com>

🔗 master  
🔄 v2.39.0 ... v2.30.1

👤 peff authored and gitster committed on Jan 7, 2021 parent 041bc65 commit a02ea577174ab8ed18f847cf1693f213e0b9c473

Showing 2 changed files with 7 additions and 0 deletions.

Split Unified

🔗 2 📄 connect.c 📄

1063	1063	target_host = xstrdup(hostandport);
1064	1064	
1065	1065	transport_check_allowed("git");
	1066 +	if (strchr(target_host, '\n')    strchr(path, '\n'))
	1067 +	die_("newline is forbidden in git:// hosts and repo paths");
1066	1068	
1067	1069	/*
1068	1070	* These underlying connection commands die() if they

🔗 5 📄 t/t5570-git-daemon.sh 📄

102	102	)
103	103	'
104	104	
	105 +	test_expect_success 'client refuses to ask for repo with newline' '
	106 +	test_must_fail git clone "\$GIT_DAEMON_URL/repo\$LF.git" dst 2>stderr &&
	107 +	test_i18ngrep newline.is.forbidden stderr
	108 +	'
	109 +	

105	110	test_remote_error()
106	111	{
107	112	do_export=VesPlease

0 comments on commit `a02ea57`

Please [sign in](#) to comment.