

[New issue](#)
[Jump to bottom](#)

Out-of-bounds memory access in DXF loader (path identification) #4037

✓ Closed eldstal opened this issue on Jan 5 · 9 comments

eldstal commented on Jan 5

Contributor

Summary

A DXF-format drawing with particular (not necessarily malformed!) properties may cause an out-of-bounds memory access when imported using `import()`.

Vulnerable versions

- OpenSCAD Linux (commit [eedf370](#))
- OpenSCAD Windows x64 (2021.01)

Steps to reproduce

Two Proof-of-concept files are provided. `oob_dxfdata_505` is larger, but more reliably reproduces the fault. `oob_dxfdata_k5` is my attempt at a minimal working example, but only reliably crashes the linux build from the latest commit on the `main` branch (i.e. it does not crash in the windows release). Both illustrate the same crash, so I expect the `_min` example to be more helpful in testing.

[oob_dxfdata_505.zip, the larger messier p-o-c](#)
[oob_dxfdata_k5.zip, the smaller p-o-c](#)

1. Unzip one of the provided proof of concept files
2. Open the `.scad` file in OpenSCAD
3. Render with F6
4. Observe the application crashing

Alternatively, for headless operation:

1. Unzip one of the provided proof of concept files
2. `openscad --export-format stl -o /dev/null oob_dxfdata_k5.scad`

3. Observe the segmentation fault

Screenshot

```

Program received signal SIGSEGV, Segmentation fault.
DxfData::DxfData (this=0x7fffffff9f0, fn=0, fs=2, fa=12, filename="/home/albin/fuzz/opencad/pocs_manual/oob_dxfdata_k5.dxf", layername=""
=, xorigin=0, yorigin=0, scale=1) at /home/albin/fuzz/opencad/opencad.git/src/dxfdata.cc:470
470         if (grid.eq(ref_point[0], ref_point[1], this->points[lines[k].idx[0]][0], this->points[lines[k].
idx[0]][1])) {
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA

[ REGISTERS ]
RAX 0x555556a85528
RBX 0x555556a21a60 -> 0x555556a24040 -> 0x1
RCX 0x0
RDX 0x8
RDI 0x7fffffff110 -> 0x555556a85520
RSI 0x1
R8 0x555556a21ab0 -> 0x500000004
R9 0x555556a23fe0 -> 0x555556a23f20 -> 0x555556a21ad0 -> 0x555556a23f80 -> 0x555556a23f50 -> ...
R10 0x1
R11 0x7ffff54eebe0 (main_arena+96) -> 0x555556a38c10 -> 0x0
R12 0x555556a208e0 -> 0x5555569c7c78 -> 0x555556b50bf8 (SourceFile::~SourceFile()) -> endbr64
R13 0x1
R14 0x0
R15 0x0
RBP 0x7fffffff9a0 -> 0x7fffffffca80 -> 0x7fffffffcb10 -> 0x7fffffffcb40 -> 0x7fffffffcb80 -> ...
RSP 0x7fffffff170 -> 0x4014660000000000
RIP 0x555556f26 -> movsd xmm5, qword ptr [rax]

[ DISASM ]
> 0x555556f26 movsd xmm5, qword ptr [rax]
0x555556f2a movsd qword ptr [rbp - 0x820], xmm5
0x555556f32 mov eax, dword ptr [rbp - 0x760]
0x555556f38 movsxd rdx, eax
0x555556f3b lea rax, [rbp - 0x600]
0x555556f42 mov rsi, rdx
0x555556f45 mov rdi, rax
0x555556f48 call 0x555556ff728
0x555556f4d mov eax, dword ptr [rax]
0x555556f4f movsxd rdx, eax
0x555556f52 mov rax, qword ptr [rbp - 0x7d8]

[ SOURCE (CODE) ]
In file: /home/albin/fuzz/opencad/opencad.git/src/dxfdata.cc
465 enabled_lines.erase(current_line);
466 auto lv = grid.data(ref_point[0], ref_point[1]);
467 for (size_t ki = 0; ki < lv.size(); ++ki) {
468     int k = lv.at(ki);
469     if (lines[k].disabled) continue;
470     if (grid.eq(ref_point[0], ref_point[1], this->points[lines[k].idx[0]][0], this->points[lines[k].idx[0]][
1])) {
471         current_line = k;
472         current_point = 0;
473         goto found_next_line_in_open_path;
474     }
475     if (grid.eq(ref_point[0], ref_point[1], this->points[lines[k].idx[1]][0], this->points[lines[k].idx[1]][
1])) {
[ STACK ]
00:0000 rsp 0x7fffffff170 -> 0x4014660000000000
01:0008 0x7fffffff178 -> 0x4014660000000000
02:0010 0x7fffffff180 -> 0x4018660000000000
03:0018 0x7fffffff188 -> 0x3ff0000000000000
04:0020 0x7fffffff190 -> 0x0
05:0028 0x7fffffff198 -> 0x0
06:0030 0x7fffffff1a0 -> 0x555556a21c68 -> 0x555556a21c78 -> 0x0
07:0038 0x7fffffff1a8 -> 0x555556a21c48 -> 0x555556a22080 -> '/home/albin/fuzz/opencad/pocs_manual/oob_dxfdata_k5.dxf'

[ BACKTRACE ]
> f 0 0x555556f26
f 1 0x555556f124a
f 2 0x555556f36f49
f 3 0x555556f192d5
f 4 0x555556f277d
f 5 0x555556f244d
f 6 0x555556b38168
f 7 0x555556b3822c

pwndbg> bt
#0 DxfData::DxfData (this=0x7fffffff9f0, fn=0, fs=2, fa=12, filename="/home/albin/fuzz/opencad/pocs_manual/oob_dxfdata_k5.dxf", layer
name="", xorigin=0, yorigin=0, scale=1) at /home/albin/fuzz/opencad/opencad.git/src/dxfdata.cc:470
#1 0x0000555556f124a in ImportNode::createGeometry (this=0x555556a21c10) at /home/albin/fuzz/opencad/opencad.git/src/import.cc:198
#2 0x0000555556f36f49 in GeometryEvaluator::visit (this=0x7fffffff130, state=..., node=...) at /home/albin/fuzz/opencad/opencad.git/
src/GeometryEvaluator.cc:607
#3 0x0000555556f192d5 in NodeVisitor::visit (this=0x7fffffff130, state=..., node=...) at /home/albin/fuzz/opencad/opencad.git/src/No
deVisitor.h:72
#4 0x0000555556f277d in BaseVisitable::acceptImpl<ImportNode> (state=..., node=..., visitor=...) at /home/albin/fuzz/opencad/opencad
.git/src/BaseVisitable.h:30
#5 0x0000555556f244d in ImportNode::accept (this=0x555556a21c10, state=..., visitor=...) at /home/albin/fuzz/opencad/opencad.git/src
/importnode.h:22
#6 0x0000555556b38168 in NodeVisitor::traverse (this=0x7fffffff130, node=..., state=...) at /home/albin/fuzz/opencad/opencad.git/src
/NodeVisitor.cc:14
#7 0x0000555556b3822c in NodeVisitor::traverse (this=0x7fffffff130, node=..., state=...) at /home/albin/fuzz/opencad/opencad.git/src
/NodeVisitor.cc:20
#8 0x0000555556b3822c in NodeVisitor::traverse (this=0x7fffffff130, node=..., state=...) at /home/albin/fuzz/opencad/opencad.git/src
/NodeVisitor.cc:20
#9 0x0000555556f33c68 in GeometryEvaluator::evaluateGeometry (this=0x7fffffff130, node=..., allownef=true) at /home/albin/fuzz/opensca
d/opencad.git/src/GeometryEvaluator.cc:63
#10 0x0000555556d585 in do_export (cmd=..., render_variables=..., curFormat=FileFormat::ASCIISTL, root_file=0x555556a208e0) at /home/a

```

```
lbin/fuzz/openscad/openscad.git/src/openscad.cc:564
#11 0x0000555555560c0e8 in cmdline (cmd=...) at /home/albin/fuzz/openscad/openscad.git/src/openscad.cc:432
#12 0x000055555556125da in main (argc=6, argv=0x7fffffffe068) at /home/albin/fuzz/openscad/openscad.git/src/openscad.cc:1230
#13 0x00007ffff532a0b3 in __libc_start_main (main=0x555555560f718 <main(int, char**)>, argc=6, argv=0x7fffffffe068, init=<optimized out>,
fini=<optimized out>, rtld_fini=<optimized out>, stack_end=0x7fffffffe058) at ../csu/libc-start.c:308
#14 0x00005555555f6fae in _start ()
pwndbg>
```

Cause

I haven't quite been able to wrap my head around the DXF parser yet, but the OOB access occurs in `src/dxfdata.cc` on one of the lines [470](#), [475](#), [505](#), or [510](#) depending on the input file.

It appears to be related to the fact that if multiple line segments share points in common, they are merged into contiguous paths. As a part of this, The `ADD_LINE` macro manipulates the `grid.data<>` and `lines<>` structures. By creating lines either in the `ENTITIES` section of the DXF file or outside of it, an attacker is able to ensure that the `grid.data` entry created on line [108](#) points to an index in `lines` which never becomes valid.

On line 470 (and the others listed above), this value (`k`) is used as an index into `lines` , out of bounds.

Impact

It appears that the out-of-bounds data can only be read, and not written. An out-of-bounds read does not expose a security vulnerability on its own, but can be used to bypass automatic security features such as stack canaries and pointer encryption.

Proposed mitigation

The algorithm employed in `DxfData::DxfData` needs to be revised to prevent this type of aliasing. Without more insight into the DXF format and how it is used by OpenSCAD, I cannot say for sure where the actual flaw lies.

It seems likely that the `ADD_LINE` macro on lines 108-109 is the culprit, since it inserts values that will be used for indexing even though the values are not yet valid indices.

ChrisCoxArt commented on Jan 15 • edited ▼

Contributor

I see lots of warnings, but no crash on MacOS with the current development code.

I even tried with guard malloc, guard edges, and zombie (use after free) detection enabled.

Ah, but got it to fail with malloc_scribble enabled!

It's reading one entry over the end of the lines vector (size = 10, index = 10). Because of the way vectors double allocations (so this one has 16 values allocated), the address is valid, but the data is bogus, and the data is used as an index to another array, which then gets the access violation. Under normal debugging situations, the memory was probably zero-ed and leading to a valid index.

ChrisCoxArt commented on Jan 15 • edited ▼

Contributor

And there are at least 3 places in the code with the same fragility.

OK, probably best to skip any bad indices, and log a warning about the bad index in the DXF file.

I hope this doesn't break any import tests... Good, it doesn't break any of the tests. I still want to test this with more files, though.

ChrisCoxArt commented on Jan 16 • edited ▼

Contributor

I thought maybe we could catch the error earlier in the process (when adding lines), but so far I haven't figured out a way to do it.

This DXF code is more than a little fragile.

BTW - the second example actually hits the same sort of error in a different place, so thanks for providing both!

kintel commented on Jan 16

Member

OpenSCAD's DXF support is really a rather large piece of technical debt. At some point it's worth digging further into replacing it with an existing library, see <https://github.com/openscad/openscad/wiki/Project%3A-Improve-DXF-import-and-export>

ChrisCoxArt commented on Jan 16

Contributor

Ok, I have a fix that is safe, if not optimal. Trying to catch the error earlier runs into a lot of code that I cannot decipher easily (it's a huge state machine, with many inter-dependent arrays of data) -- so I cannot say that any changes there would be safe.

eldstal commented on Feb 4

Contributor

Author

Thanks for taking a look at this issue! Has the fix been merged?

This vulnerability has been assigned [CVE-2022-0496](#).

ChrisCoxArt commented on Feb 4

Contributor

No, the fix has not been merged. It's waiting on other pull requests to be merged (which have been sitting for 3 weeks).

 **ChrisCoxArt** added a commit to ChrisCoxArt/openscad that referenced this issue on Feb 4



add safety to line lookups in DXF import, fixes [openscad#4037](#) ...

6beaca5

ChrisCoxArt commented on Feb 4

Contributor

Now it has been merged! yay!

t-paul commented on Feb 4

Member

Hold the horses, it will be in a bit :)



t-paul closed this as completed in [770e323](#) on Feb 4



t-paul added a commit that referenced this issue on Feb 4



Merge pull request [#4090](#) from openscad/dxf-updates ...

✓ 52ce0f7



t-paul added a commit that referenced this issue on Feb 5



[CVE-2022-0496](#) Out-of-bounds memory access in DXF loader. ...

00a4692

Assignees

No one assigned

Labels

None yet

Projects

None yet

Milestone

No milestone

Development

No branches or pull requests

4 participants

