



author Jan Kara <jack@suse.cz> 2022-01-17 18:22:13 +0100
committer Jan Kara <jack@suse.cz> 2022-01-24 14:45:02 +0100
commit 7fc3b7c2981bbd1047916ade327beccb90994eee (patch)
tree 410cae0a02add6911b15af81d8472738e50d0d8b
parent dd81e1c7d5fb126e5fbc5c9e334d7b3ec29a16a0 (diff)
download linux-7fc3b7c2981bbd1047916ade327beccb90994eee.tar.gz

diff options

context: 3 ▾

space: include ▾

mode: unified ▾

udf: Fix NULL ptr deref when converting from inline format

udf_expand_file_adinicb() calls directly ->writepage to write data expanded into a page. This however misses to setup inode for writeback properly and so we can crash on inode->i_wb dereference when submitting page for IO like:

```
BUG: kernel NULL pointer dereference, address: 00000000000000158
#PF: supervisor read access in kernel mode
```

```
...
```

```
<TASK>
__folio_start_writeback+0x2ac/0x350
__block_write_full_page+0x37d/0x490
udf_expand_file_adinicb+0x255/0x400 [udf]
udf_file_write_iter+0xbe/0x1b0 [udf]
new_sync_write+0x125/0x1c0
vfs_write+0x28e/0x400
```

Fix the problem by marking the page dirty and going through the standard writeback path to write the page. Strictly speaking we would not even have to write the page but we want to catch e.g. ENOSPC errors early.

Reported-by: butt3rflyh4ck <butterflyhuangxx@gmail.com>

CC: stable@vger.kernel.org

Fixes: 52ebea749aae ("writeback: make backing_dev_info host cgroup-specific bdi_writebacks")

Reviewed-by: Christoph Hellwig <hch@lst.de>

Signed-off-by: Jan Kara <jack@suse.cz>

Diffstat

```
-rw-r--r-- fs/udf/inode.c 8
```

1 files changed, 3 insertions, 5 deletions

```
diff --git a/fs/udf/inode.c b/fs/udf/inode.c
```

```
index 1d6b7a50736ba..d6aa506b6b584 100644
```

```
--- a/fs/udf/inode.c
```

```
+++ b/fs/udf/inode.c
```

```
@@ -258,10 +258,6 @@ int udf_expand_file_adinicb(struct inode *inode)
     char *kaddr;
     struct udf_inode_info *iinfo = UDF_I(inode);
     int err;

-    struct writeback_control udf_wbc = {
-        .sync_mode = WB_SYNC_NONE,
-        .nr_to_write = 1,
-    };
```

```
WARN_ON_ONCE(!inode_is_locked(inode));
```

```

        if (!iinfo->i_lenAlloc) {
@@ -305,8 +301,10 @@ int udf_expand_file_adinicb(struct inode *inode)
            iinfo->i_alloc_type = ICBTAG_FLAG_AD_LONG;
            /* from now on we have normal address_space methods */
            inode->i_data.a_ops = &udf_aops;
+       set_page_dirty(page);
+       unlock_page(page);
            up_write(&iinfo->i_data_sem);
-       err = inode->i_data.a_ops->writepage(page, &udf_wbc);
+       err = filemap_fdatawrite(inode->i_mapping);
+       if (err) {
            /* Restore everything back so that we don't lose data... */
            lock_page(page);

```