



# index : kernel/git/torvalds/linux.git

master

Linux kernel source tree

Linus Torvalds

[about](#) [summary](#) [refs](#) [log](#) [tree](#) [commit](#) [diff](#) [stats](#)

path: [root/kernel/dma/swiotlb.c](#)

author Halil Pasic <pasic@linux.ibm.com> 2022-03-05 18:07:14 +0100  
committer Linus Torvalds <torvalds@linux-foundation.org> 2022-03-07 11:26:02 -0800  
commit aa6f8dcbab473f3a3c7454b74caa46d36cdc5d13 (patch)  
tree f4ff6004c13374d66d37db7221e120f8d409799c /kernel/dma/swiotlb.c  
parent ffb217a13a2eaf6d5bd974fc83036a53ca69f1e2 (diff)  
download [linux-aa6f8dcbab473f3a3c7454b74caa46d36cdc5d13.tar.gz](#)

## diff options

context:    
space:    
mode:

## swiotlb: rework "fix info leak with DMA\_FROM\_DEVICE"

Unfortunately, we ended up merging an old version of the patch "fix info leak with DMA\_FROM\_DEVICE" instead of merging the latest one. Christoph (the swiotlb maintainer), he asked me to create an incremental fix (after I have pointed this out the mix up, and asked him for guidance). So here we go.

The main differences between what we got and what was agreed are:

- \* swiotlb\_sync\_single\_for\_device is also required to do an extra bounce
- \* We decided not to introduce DMA\_ATTR\_OVERWRITE until we have exploiters
- \* The implantation of DMA\_ATTR\_OVERWRITE is flawed: DMA\_ATTR\_OVERWRITE must take precedence over DMA\_ATTR\_SKIP\_CPU\_SYNC

Thus this patch removes DMA\_ATTR\_OVERWRITE, and makes swiotlb\_sync\_single\_for\_device() bounce unconditionally (that is, also when dir == DMA\_TO\_DEVICE) in order do avoid synchronising back stale data from the swiotlb buffer.

Let me note, that if the size used with dma\_sync\_\* API is less than the size used with dma\_[un]map\_\*, under certain circumstances we may still end up with swiotlb not being transparent. In that sense, this is no perfect fix either.

To get this bullet proof, we would have to bounce the entire mapping/bounce buffer. For that we would have to figure out the starting address, and the size of the mapping in swiotlb\_sync\_single\_for\_device(). While this does seem possible, there seems to be no firm consensus on how things are supposed to work.

Signed-off-by: Halil Pasic <pasic@linux.ibm.com>  
Fixes: ddbd89deb7d3 ("swiotlb: fix info leak with DMA\_FROM\_DEVICE")  
Cc: stable@vger.kernel.org  
Reviewed-by: Christoph Hellwig <hch@lst.de>  
Signed-off-by: Linus Torvalds <torvalds@linux-foundation.org>

## Diffstat (limited to 'kernel/dma/swiotlb.c')

```
-rw-r--r-- kernel/dma/swiotlb.c 23
```

1 files changed, 15 insertions, 8 deletions

```
diff --git a/kernel/dma/swiotlb.c b/kernel/dma/swiotlb.c
index bfc56cb217059..6db1c475ec827 100644
--- a/kernel/dma/swiotlb.c
+++ b/kernel/dma/swiotlb.c
@@ -627,10 +627,14 @@ phys_addr_t swiotlb_tlb_map_single(struct device *dev, phys_addr_t orig_addr,
     for (i = 0; i < nr_slots(alloc_size + offset); i++)
         mem->slots[index + i].orig_addr = slot_addr(orig_addr, i);
     tlb_addr = slot_addr(mem->start, index) + offset;
-    if (!(attrs & DMA_ATTR_SKIP_CPU_SYNC) &&
-        (!(attrs & DMA_ATTR_OVERWRITE) || dir == DMA_TO_DEVICE ||
```

```

-         dir == DMA_BIDIRECTIONAL))
-         swiotlb_bounce(dev, tlb_addr, mapping_size, DMA_TO_DEVICE);
+     /*
+      * When dir == DMA_FROM_DEVICE we could omit the copy from the orig
+      * to the tlb buffer, if we knew for sure the device will
+      * overwrite the entire current content. But we don't. Thus
+      * unconditional bounce may prevent leaking swiotlb content (i.e.
+      * kernel memory) to user-space.
+      */
+     swiotlb_bounce(dev, tlb_addr, mapping_size, DMA_TO_DEVICE);
+     return tlb_addr;
+ }

@@ -697,10 +701,13 @@ void swiotlb_tlb_unmap_single(struct device *dev, phys_addr_t tlb_addr,
void swiotlb_sync_single_for_device(struct device *dev, phys_addr_t tlb_addr,
size_t size, enum dma_data_direction dir)
{
-     if (dir == DMA_TO_DEVICE || dir == DMA_BIDIRECTIONAL)
-         swiotlb_bounce(dev, tlb_addr, size, DMA_TO_DEVICE);
-     else
-         BUG_ON(dir != DMA_FROM_DEVICE);
+     /*
+      * Unconditional bounce is necessary to avoid corruption on
+      * sync * _for_cpu or dma_ummap_* when the device didn't overwrite
+      * the whole length of the bounce buffer.
+      */
+     swiotlb_bounce(dev, tlb_addr, size, DMA_TO_DEVICE);
+     BUG_ON(!valid_dma_direction(dir));
+ }

void swiotlb_sync_single_for_cpu(struct device *dev, phys_addr_t tlb_addr,

```