



index : kernel/git/stable/linux.git

linux-2.6.11.y

Linux kernel stable tree

Stable Group

[about](#) [summary](#) [refs](#) [log](#) [tree](#) [commit](#) [diff](#) [stats](#)

path: [root/fs/io_uring.c](#)

author Pavel Begunkov <asml.silence@gmail.com> 2022-04-14 08:50:50 +0100
committer Greg Kroah-Hartman <gregkh@linuxfoundation.org> 2022-04-15 14:18:41 +0200
commit 1a623d361ffe5cecd4244a02f449528416360038 (patch)
tree b99709904f01723c8475b87507637b4427b9c94c /fs/io_uring.c
parent 33fcb359a64213bea0510770abc761de9780ca89 (diff)
download linux-1a623d361ffe5cecd4244a02f449528416360038.tar.gz

diff options

context:
space:
mode:

io_uring: fix fs->users overflow

There is a bunch of cases where we can grab req->fs but not put it, this can be used to cause a controllable overflow with further implications. Release req->fs in the request free path and make sure we zero the field to be sure we don't do it twice.

Fixes: cac68d12c531 ("io_uring: grab ->fs as part of async offload")
Reported-by: Bing-Jhong Billy Jheng <billy@starlabs.sg>
Signed-off-by: Pavel Begunkov <asml.silence@gmail.com>
Signed-off-by: Greg Kroah-Hartman <gregkh@linuxfoundation.org>

Diffstat (limited to 'fs/io_uring.c')

-rw-r--r-- fs/io_uring.c 28

1 files changed, 18 insertions, 10 deletions

diff --git a/fs/io_uring.c b/fs/io_uring.c
index 478df7e10767a..e73969fa96bcb 100644

--- a/fs/io_uring.c
+++ b/fs/io_uring.c

```
@@ -438,6 +438,22 @@ static struct io_ring_ctx *io_ring_ctx_alloc(struct io_uring_params *p)
    return ctx;
}

+static void io_req_put_fs(struct io_kiocb *req)
+{
+    struct fs_struct *fs = req->fs;
+
+    if (!fs)
+        return;
+
+    spin_lock(&req->fs->lock);
+    if (--fs->users)
+        fs = NULL;
+    spin_unlock(&req->fs->lock);
+    if (fs)
+        free_fs_struct(fs);
+    req->fs = NULL;
+}
+
static inline bool __io_sequence_defer(struct io_ring_ctx *ctx,
                                       struct io_kiocb *req)
{
@@ -695,6 +711,7 @@ static void io_free_req_many(struct io_ring_ctx *ctx, void **reqs, int *nr)

static void __io_free_req(struct io_kiocb *req)
{
+    io_req_put_fs(req);
    if (req->file && !(req->flags & REQ_F_FIXED_FILE))
        fput(req->file);
    percpu_ref_put(&req->ctx->refs);
@@ -1701,16 +1718,7 @@ static int io_send_recvmmsg(struct io_kiocb *req, const struct io_uring_sqe *sqe,
```

```
        ret = -EINTR;
    }

-   if (req->fs) {
-       struct fs_struct *fs = req->fs;
-
-       spin_lock(&req->fs->lock);
-       if (--fs->users)
-           fs = NULL;
-       spin_unlock(&req->fs->lock);
-       if (fs)
-           free_fs_struct(fs);
-   }
+   io_req_put_fs(req);
+   io_cqring_add_event(req->ctx, sqe->user_data, ret);
+   io_put_req(req);
+   return 0;
```