

From: kerneljasonxing@gmail.com
To: edumazet@google.com, davem@davemloft.net, kuznet@ms2.inr.ac.ru,
yoshfuji@linux-ipv6.org
Cc: netdev@vger.kernel.org, kerneljasonxing@gmail.com,
linux-kernel@vger.kernel.org, liweishi@kuaishou.com,
lishujin@kuaishou.com
Subject: [PATCH] tcp: fix TCP socks unreleased in BBR mode
Date: Tue, 2 Jun 2020 16:04:25 +0800 [\[thread overview\]](#)
Message-ID: <20200602080425.93712-1-kerneljasonxing@gmail.com> ([raw](#))

From: Jason Xing <kerneljasonxing@gmail.com>

TCP socks cannot be released because of the sock_hold() increasing the sk_refcnt in the manner of tcp_internal_pacing() when RTO happens. Therefore, this situation could increase the slab memory and then trigger the OOM if the machine has been running for a long time. This issue, however, can happen on some machine only running a few days.

We add one exception case to avoid unneeded use of sock_hold if the pacing_timer is enqueued.

Reproduce procedure:

- 0) cat /proc/slabinfo | grep TCP
- 1) switch net.ipv4.tcp_congestion_control to bbr
- 2) using wrk tool something like that to send packages
- 3) using tc to increase the delay in the dev to simulate the busy case.
- 4) cat /proc/slabinfo | grep TCP
- 5) kill the wrk command and observe the number of objects and slabs in TCP.
- 6) at last, you could notice that the number would not decrease.

Signed-off-by: Jason Xing <kerneljasonxing@gmail.com>
Signed-off-by: liweishi <liweishi@kuaishou.com>
Signed-off-by: Shujin Li <lishujin@kuaishou.com>

```
---
net/ipv4/tcp_output.c | 3 ++-
1 file changed, 2 insertions(+), 1 deletion(-)

diff --git a/net/ipv4/tcp_output.c b/net/ipv4/tcp_output.c
index cc4ba42..5cf63d9 100644
--- a/net/ipv4/tcp_output.c
+++ b/net/ipv4/tcp_output.c
@@ -969,7 +969,8 @@ static void tcp_internal_pacing(struct sock *sk, const struct sk_buff
 *skb)
     u64 len_ns;
     u32 rate;

-    if (!tcp_needs_internal_pacing(sk))
+    if (!tcp_needs_internal_pacing(sk) ||
+        hrtimer_is_queued(&tcp_sk(sk)->pacing_timer))
         return;
     rate = sk->sk_pacing_rate;
     if (!rate || rate == ~0U)
--
1.8.3.1
```

[next](#) [reply](#) [other threads](#): [[~2020-06-02 8:05 UTC](#)|[newest](#)]

Thread overview: 22+ messages / [expand](#)[[flat](#)|[nested](#)] [mbox.gz](#) [Atom feed](#) [top](#)
2020-06-02 8:04 [kerneljasonxing](#) [[this message](#)]

```
2020-06-02 13:05 ` [PATCH] tcp: fix TCP socks unreleased in BBR mode Eric Dumazet
2020-06-03 1:53 ` Jason Xing
2020-06-03 2:14 ` David Miller
2020-06-03 2:29 ` Eric Dumazet
2020-06-03 2:41 ` Jason Xing
2020-06-03 2:43 ` Eric Dumazet
2020-06-03 2:48 ` Jason Xing
2020-06-03 5:05 ` Jason Xing
2020-06-03 5:44 ` Eric Dumazet
2020-06-03 6:32 ` Jason Xing
2020-06-03 12:02 ` Neal Cardwell
2020-06-03 13:51 ` Jason Xing
2020-06-03 13:54 ` Eric Dumazet
2020-06-03 14:07 ` Neal Cardwell
2020-06-04 8:40 ` Jason Xing
2020-06-04 9:00 ` [PATCH v2 4.19] " kerneljasonxing
2020-06-04 13:10 ` Eric Dumazet
2020-06-04 13:47 ` Jason Xing
2020-08-11 10:37 ` Jason Xing
2020-08-11 15:32 ` Eric Dumazet
2021-04-30 10:51 ` Jason Xing
```

find likely ancestor, descendant, or conflicting patches for [this message](#):

dfblob:cc4ba42 dfblob:5cf63d9

search

([help](#))

Reply instructions:

You may reply publicly to [this message](#) via plain-text email using any one of the following methods:

- * Save the following mbox file, import it into your mail client, and reply-to-all from there: [mbox](#)

Avoid top-posting and favor interleaved quoting:

https://en.wikipedia.org/wiki/Posting_style#Interleaved_style

- * Reply using the **--to**, **--cc**, and **--in-reply-to** switches of `git-send-email(1)`:

```
git send-email \
  --in-reply-to=20200602080425.93712-1-kerneljasonxing@gmail.com \
  --to=kerneljasonxing@gmail.com \
  --cc=davem@davemloft.net \
  --cc=edumazet@google.com \
  --cc=kuznet@ms2.inr.ac.ru \
  --cc=linux-kernel@vger.kernel.org \
  --cc=lishujin@kuaishou.com \
  --cc=liweishi@kuaishou.com \
  --cc=netdev@vger.kernel.org \
  --cc=yoshfuji@linux-ipv6.org \
  /path/to/YOUR_REPLY
```

<https://kernel.org/pub/software/scm/git/docs/git-send-email.html>

- * If your mail client supports setting the **In-Reply-To** header via `mailto:` links, try the `mailto:` [link](#)

Be sure your reply has a **Subject:** header at the top and a blank line before the message body.

This is an external index of several public inboxes, see [mirroring instructions](#) on how to clone and mirror all data and code used by this external index.

