

## bpf.vger.kernel.org archive mirror

[help](#) / [color](#) / [mirror](#) / [Atom feed](#)

From: Daniel Borkmann <daniel@iogearbox.net>  
To: ast@kernel.org  
Cc: andrii@kernel.org, bpf@vger.kernel.org,  
Daniel Borkmann <daniel@iogearbox.net>,  
Hsin-Wei Hung <hsinweih@uci.edu>,  
Shung-Hsi Yu <shung-hsi.yu@suse.com>  
Subject: [\[PATCH bpf\] bpf: Don't use tnum\\_range on array range checking for poke descriptors](#)  
Date: Thu, 25 Aug 2022 23:26:47 +0200 [\[thread overview\]](#)  
Message-ID: <984b37f9fdf7ac36831d2137415a4a915744c1b6.1661462653.git.daniel@iogearbox.net>  
([raw](#))

Hsin-Wei reported a KASAN splat triggered by their BPF runtime fuzzer which is based on a customized syzkaller:

```
BUG: KASAN: slab-out-of-bounds in bpf_int_jit_compile+0x1257/0x13f0
Read of size 8 at addr ffff888004e90b58 by task syz-executor.0/1489
CPU: 1 PID: 1489 Comm: syz-executor.0 Not tainted 5.19.0 #1
Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS
1.13.0-1ubuntu1.1 04/01/2014
Call Trace:
<TASK>
dump_stack_lvl+0x9c/0xc9
print_address_description.constprop.0+0x1f/0x1f0
? bpf_int_jit_compile+0x1257/0x13f0
kasan_report.cold+0xeb/0x197
? kvmalloc_node+0x170/0x200
? bpf_int_jit_compile+0x1257/0x13f0
bpf_int_jit_compile+0x1257/0x13f0
? arch_prepare_bpf_dispatcher+0xd0/0xd0
? rcu_read_lock_sched_held+0x43/0x70
bpf_prog_select_runtime+0x3e8/0x640
? bpf_obj_name_cpy+0x149/0x1b0
bpf_prog_load+0x102f/0x2220
? __bpf_prog_put.constprop.0+0x220/0x220
? find_held_lock+0x2c/0x110
? __might_fault+0xd6/0x180
? lock_downgrade+0x6e0/0x6e0
? lock_is_held_type+0xa6/0x120
? __might_fault+0x147/0x180
__sys_bpf+0x137b/0x6070
? bpf_perf_link_attach+0x530/0x530
? new_sync_read+0x600/0x600
? __fget_files+0x255/0x450
? lock_downgrade+0x6e0/0x6e0
? fput+0x30/0x1a0
? ksys_write+0x1a8/0x260
__x64_sys_bpf+0x7a/0xc0
? syscall_enter_from_user_mode+0x21/0x70
do_syscall_64+0x3b/0x90
entry_SYSCALL_64_after_hwframe+0x63/0xcd
RIP: 0033:0x7f917c4e2c2d
```

The problem here is that a range of `tnum_range(0, map->max_entries - 1)` has limited ability to represent the concrete tight range with the `tnum` as the set of resulting states from `value + mask` can result in a superset of the actual intended range, and as such a `tnum_in(range, reg->var_off)` check may yield true when it shouldn't, for example `tnum_range(0, 2)` would result in `00XX -> v = 0000, m = 0011` such that the intended set of `{0, 1, 2}` is here represented by a less precise superset of `{0, 1, 2, 3}`. As the register is known const scalar, really just use the concrete `reg->var_off.value` for the upper index check.

Fixes: d2e4c1e6c294 ("bpf: Constant map key tracking for prog array pokes")

Reported-by: Hsin-Wei Hung <hsinwei@uci.edu>

Signed-off-by: Daniel Borkmann <daniel@iogearbox.net>

Cc: Shung-Hsi Yu <shung-hsi.yu@suse.com>

---

kernel/bpf/verifier.c | 10 ++++-----

1 file changed, 4 insertions(+), 6 deletions(-)

diff --git a/kernel/bpf/verifier.c b/kernel/bpf/verifier.c

index 30c6eebce146..3eadb14e090b 100644

--- a/kernel/bpf/verifier.c

+++ b/kernel/bpf/verifier.c

@@ -7033,8 +7033,7 @@ record\_func\_key(struct bpf\_verifier\_env \*env, struct bpf\_call\_arg\_meta \*meta,

struct bpf\_insn\_aux\_data \*aux = &env->insn\_aux\_data[insn\_idx];

struct bpf\_reg\_state \*regs = cur\_regs(env), \*reg;

struct bpf\_map \*map = meta->map\_ptr;

- struct tnum range;

- u64 val;

+ u64 val, max;

int err;

if (func\_id != BPF\_FUNC\_tail\_call)

@@ -7044,10 +7043,11 @@ record\_func\_key(struct bpf\_verifier\_env \*env, struct bpf\_call\_arg\_meta \*meta,

return -EINVAL;

}

- range = tnum\_range(0, map->max\_entries - 1);

reg = &regs[BPF\_REG\_3];

+ val = reg->var\_off.value;

+ max = map->max\_entries;

- if (!register\_is\_const(reg) || !tnum\_in(range, reg->var\_off)) {

+ if (!(register\_is\_const(reg) && val < max)) {

bpf\_map\_key\_store(aux, BPF\_MAP\_KEY\_POISON);

return 0;

}

@@ -7055,8 +7055,6 @@ record\_func\_key(struct bpf\_verifier\_env \*env, struct bpf\_call\_arg\_meta \*meta,

err = mark\_chain\_precision(env, BPF\_REG\_3);

if (err)

return err;

-

- val = reg->var\_off.value;

if (bpf\_map\_key\_unseen(aux))

bpf\_map\_key\_store(aux, val);

else if (!bpf\_map\_key\_poisoned(aux) &&

--

2.21.0

---

next                    reply   other threads:[~2022-08-25 21:27 UTC|newest]

**Thread overview:** 3+ messages / expand[flat|nested]   mbox.gz   Atom feed   top

**2022-08-25 21:26 Daniel Borkmann [this message]**

2022-08-25 21:53 ` [PATCH bpf] bpf: Don't use tnum\_range on array range checking for poke descriptors John Fastabend

2022-08-25 22:00 ` patchwork-bot+netdevbpf

find likely ancestor, descendant, or conflicting patches for [this message](#):

dfblob:30c6eebce14 dfblob:3eadb14e090

---

**Reply instructions:**

You may reply publicly to [this message](#) via plain-text email using any one of the following methods:

- \* Save the following mbox file, import it into your mail client, and reply-to-all from there: [mbox](#)

Avoid top-posting and favor interleaved quoting:

[https://en.wikipedia.org/wiki/Posting\\_style#Interleaved\\_style](https://en.wikipedia.org/wiki/Posting_style#Interleaved_style)

- \* Reply using the **--to**, **--cc**, and **--in-reply-to** switches of `git-send-email(1)`:

```
git send-email \
  --in-reply-
to=984b37f9fdf7ac36831d2137415a4a915744c1b6.1661462653.git.daniel@iogearbox.net \
  --to=daniel@iogearbox.net \
  --cc=andrii@kernel.org \
  --cc=ast@kernel.org \
  --cc=bpf@vger.kernel.org \
  --cc=hsinweih@uci.edu \
  --cc=shung-hsi.yu@suse.com \
  /path/to/YOUR_REPLY
```

<https://kernel.org/pub/software/scm/git/docs/git-send-email.html>

- \* If your mail client supports setting the **In-Reply-To** header via `mailto:` links, try the `mailto: link`

Be sure your reply has a **Subject:** header at the top and a blank line before the message body.

---

This is a public inbox, see [mirroring instructions](#) for how to clone and mirror all data and code used for this inbox; as well as URLs for NNTP newsgroup(s).