

5100e359ae ▾

...


tensorflow / tensorflow / core / kernels / fractional\_avg\_pool\_op.cc



jpienaar Rename to underlying type rather than alias ... ✓

History

7 contributors



374 lines (333 sloc) | 15.8 KB

...

```

1  /* Copyright 2016 The TensorFlow Authors. All Rights Reserved.
2
3  Licensed under the Apache License, Version 2.0 (the "License");
4  you may not use this file except in compliance with the License.
5  You may obtain a copy of the License at
6
7      http://www.apache.org/licenses/LICENSE-2.0
8
9  Unless required by applicable law or agreed to in writing, software
10 distributed under the License is distributed on an "AS IS" BASIS,
11 WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
12 See the License for the specific language governing permissions and
13 limitations under the License.
14 =====*/
15 #define EIGEN_USE_THREADS
16
17 #include <algorithm>
18 #include <cmath>
19 #include <random>
20 #include <vector>
21
22 #include "tensorflow/core/kernels/fractional_pool_common.h"
23
24 #include "third_party/eigen3/unsupported/Eigen/CXX11/Tensor"
25 #include "tensorflow/core/framework/numeric_op.h"
26 #include "tensorflow/core/framework/op_kernel.h"
27 #include "tensorflow/core/lib/random/random.h"
28 #include "tensorflow/core/platform/logging.h"
29 #include "tensorflow/core/platform/mutex.h"

```

```

30 #include "tensorflow/core/util/guarded_philox_random.h"
31
32 namespace tensorflow {
33 typedef Eigen::ThreadPoolDevice CPUDevice;
34
35 template <typename T>
36 class FractionalAvgPoolOp : public OpKernel {
37 public:
38     explicit FractionalAvgPoolOp(OpKernelConstruction* context)
39         : OpKernel(context) {
40         OP_REQUIRES_OK(context, context->GetAttr("pooling_ratio", &pooling_ratio_));
41         OP_REQUIRES_OK(context, context->GetAttr("pseudo_random", &pseudo_random_));
42         OP_REQUIRES_OK(context, context->GetAttr("overlapping", &overlapping_));
43         OP_REQUIRES(context, pooling_ratio_.size() == 4,
44             errors::InvalidArgument(
45                 "pooling_ratio field must specify 4 dimensions"));
46         OP_REQUIRES(
47             context, pooling_ratio_[0] == 1 || pooling_ratio_[3] == 1,
48             errors::Unimplemented("Fractional average pooling is not yet "
49                 "supported on the batch nor channel dimension.));
50         OP_REQUIRES_OK(context, context->GetAttr("deterministic", &deterministic_));
51         OP_REQUIRES_OK(context, context->GetAttr("seed", &seed_));
52         OP_REQUIRES_OK(context, context->GetAttr("seed2", &seed2_));
53         if (deterministic_) {
54             // If both seeds are not set when deterministic_ is true, force set seeds.
55             if ((seed_ == 0) && (seed2_ == 0)) {
56                 seed_ = random::New64();
57                 seed2_ = random::New64();
58             }
59         } else {
60             OP_REQUIRES(
61                 context, (seed_ == 0) && (seed2_ == 0),
62                 errors::InvalidArgument(
63                     "Both seed and seed2 should be 0 if deterministic is false.));
64         }
65     }
66
67     void Compute(OpKernelContext* context) override {
68         typedef Eigen::Map<const Eigen::Matrix<T, Eigen::Dynamic, Eigen::Dynamic>>
69             ConstEigenMatrixMap;
70         typedef Eigen::Map<Eigen::Matrix<T, Eigen::Dynamic, Eigen::Dynamic>>
71             EigenMatrixMap;
72
73         constexpr int tensor_in_and_out_dims = 4;
74
75         const Tensor& tensor_in = context->input(0);
76         OP_REQUIRES(context, tensor_in.dims() == tensor_in_and_out_dims,
77             errors::InvalidArgument("tensor_in must be 4-dimensional"));
78

```

```

79     std::vector<int> input_size(tensor_in_and_out_dims);
80     std::vector<int> output_size(tensor_in_and_out_dims);
81     for (int i = 0; i < tensor_in_and_out_dims; ++i) {
82         input_size[i] = tensor_in.dim_size(i);
83         OP_REQUIRES(
84             context, pooling_ratio[i] <= input_size[i],
85             errors::InvalidArgument(
86                 "Pooling ratio cannot be bigger than input tensor dim size."));
87     }
88     // Output size.
89     for (int i = 0; i < tensor_in_and_out_dims; ++i) {
90         output_size[i] =
91             static_cast<int>(std::floor(input_size[i] / pooling_ratio[i]));
92         DCHECK_GT(output_size[i], 0);
93     }
94
95     // Generate pooling sequence.
96     std::vector<int64_t> row_cum_seq;
97     std::vector<int64_t> col_cum_seq;
98     GuardedPhiloxRandom generator;
99     generator.Init(seed_, seed2_);
100    row_cum_seq = GeneratePoolingSequence(input_size[1], output_size[1],
101                                         &generator, pseudo_random_);
102    col_cum_seq = GeneratePoolingSequence(input_size[2], output_size[2],
103                                         &generator, pseudo_random_);
104
105    // Prepare output.
106    Tensor* output_tensor = nullptr;
107    OP_REQUIRES_OK(context, context->allocate_output(
108        0,
109        TensorShape({output_size[0], output_size[1],
110                     output_size[2], output_size[3]}),
111        &output_tensor));
112    Tensor* output_row_seq_tensor = nullptr;
113    OP_REQUIRES_OK(
114        context, context->allocate_output(
115            1, TensorShape({static_cast<int64_t>(row_cum_seq.size())}),
116            &output_row_seq_tensor));
117    Tensor* output_col_seq_tensor = nullptr;
118    OP_REQUIRES_OK(
119        context, context->allocate_output(
120            2, TensorShape({static_cast<int64_t>(col_cum_seq.size())}),
121            &output_col_seq_tensor));
122
123    ConstEigenMatrixMap in_mat(tensor_in.flat<T>().data(), input_size[3],
124                               input_size[2] * input_size[1] * input_size[0]);
125
126    EigenMatrixMap out_mat(output_tensor->flat<T>().data(), output_size[3],
127                           output_size[2] * output_size[1] * output_size[0]);

```

```

128 // out_count corresponds to number of elements in each pooling cell.
129 Eigen::Matrix<T, Eigen::Dynamic, 1> out_count(out_mat.cols());
130
131 // Initializes the output tensor and out_count with 0.
132 out_mat.setZero();
133 out_count.setZero();
134
135 auto output_row_seq_flat = output_row_seq_tensor->flat<int64_t>();
136 auto output_col_seq_flat = output_col_seq_tensor->flat<int64_t>();
137
138 // Set output tensors.
139 for (int i = 0; i < row_cum_seq.size(); ++i) {
140     output_row_seq_flat(i) = row_cum_seq[i];
141 }
142
143 for (int i = 0; i < col_cum_seq.size(); ++i) {
144     output_col_seq_flat(i) = col_cum_seq[i];
145 }
146
147 // For both input and output,
148 // 0: batch
149 // 1: row / row
150 // 2: col / col
151 // 3: depth / channel
152 const int64_t row_max = input_size[1] - 1;
153 const int64_t col_max = input_size[2] - 1;
154 for (int64_t b = 0; b < input_size[0]; ++b) {
155     // row sequence.
156     for (int64_t hs = 0; hs < row_cum_seq.size() - 1; ++hs) {
157         // row start and end.
158         const int64_t row_start = row_cum_seq[hs];
159         int64_t row_end =
160             overlapping_ ? row_cum_seq[hs + 1] : row_cum_seq[hs + 1] - 1;
161         row_end = std::min(row_end, row_max);
162
163         // col sequence.
164         for (int64_t ws = 0; ws < col_cum_seq.size() - 1; ++ws) {
165             const int64_t out_offset =
166                 (b * output_size[1] + hs) * output_size[2] + ws;
167             // col start and end.
168             const int64_t col_start = col_cum_seq[ws];
169             int64_t col_end =
170                 overlapping_ ? col_cum_seq[ws + 1] : col_cum_seq[ws + 1] - 1;
171             col_end = std::min(col_end, col_max);
172             for (int64_t h = row_start; h <= row_end; ++h) {
173                 for (int64_t w = col_start; w <= col_end; ++w) {
174                     const int64_t in_offset =
175                         (b * input_size[1] + h) * input_size[2] + w;
176                     out_mat.col(out_offset) += in_mat.col(in_offset);

```

```

177         out_count(out_offset)++;
178     }
179 }
180 }
181 }
182 }
183 DCHECK_GT(out_count.minCoeff(), 0);
184 out_mat.array().rowwise() /= out_count.transpose().array();
185 }
186
187 private:
188     bool deterministic_;
189     int64_t seed_;
190     int64_t seed2_;
191     std::vector<float> pooling_ratio_;
192     bool pseudo_random_;
193     bool overlapping_;
194 };
195
196 #define REGISTER_FRACTIONALAVGPPOOL(type) \
197     REGISTER_KERNEL_BUILDER( \
198         Name("FractionalAvgPool").Device(DEVICE_CPU).TypeConstraint<type>("T"), \
199         FractionalAvgPoolOp<type>)
200
201 REGISTER_FRACTIONALAVGPPOOL(int32);
202 REGISTER_FRACTIONALAVGPPOOL(int64_t);
203 REGISTER_FRACTIONALAVGPPOOL(float);
204 REGISTER_FRACTIONALAVGPPOOL(double);
205
206 #undef REGISTER_FRACTIONALAVGPPOOL
207
208 template <class T>
209 class FractionalAvgPoolGradOp : public OpKernel {
210 public:
211     explicit FractionalAvgPoolGradOp(OpKernelConstruction* context)
212         : OpKernel(context) {
213         OP_REQUIRES_OK(context, context->GetAttr("overlapping", &overlapping_));
214     }
215
216     void Compute(OpKernelContext* context) override {
217         // Here's the basic idea:
218         // Batch and depth dimension are independent from row and col dimension. And
219         // because FractionalAvgPool currently only support pooling along row and
220         // col, we can basically think of this 4D tensor backpropagation as
221         // operation of a series of 2D planes.
222         //
223         // For each element of a 'slice' (2D plane) of output_backprop, we need to
224         // figure out its contributors when doing FractionalAvgPool operation. This
225         // can be done based on row_pooling_sequence, col_pooling_seq and

```

```

226 // overlapping.
227 // Once we figure out the original contributors, we just need to evenly
228 // divide the value of this element among these contributors.
229 //
230 // Internally, we divide the out_backprop tensor and store it in a temporary
231 // tensor of double type. And cast it to the corresponding type.
232 typedef Eigen::Map<const Eigen::Matrix<T, Eigen::Dynamic, Eigen::Dynamic>>
233     ConstEigenMatrixMap;
234 typedef Eigen::Map<Eigen::Matrix<double, Eigen::Dynamic, Eigen::Dynamic>>
235     EigenDoubleMatrixMap;
236
237 // Grab the inputs.
238 const Tensor& orig_input_tensor_shape = context->input(0);
239 OP_REQUIRES(context,
240             orig_input_tensor_shape.dims() == 1 &&
241             orig_input_tensor_shape.NumElements() == 4,
242             errors::InvalidArgument("original input tensor shape must be "
243                                     "1-dimensional and 4 elements"));
244 const Tensor& out_backprop = context->input(1);
245 const Tensor& row_seq_tensor = context->input(2);
246 const Tensor& col_seq_tensor = context->input(3);
247
248 const int64_t out_batch = out_backprop.dim_size(0);
249 const int64_t out_rows = out_backprop.dim_size(1);
250 const int64_t out_cols = out_backprop.dim_size(2);
251 const int64_t out_depth = out_backprop.dim_size(3);
252
253 OP_REQUIRES(context, row_seq_tensor.NumElements() > out_rows,
254             errors::InvalidArgument("Given out_backprop shape ",
255                                     out_backprop.shape().DebugString(),
256                                     ", row_seq_tensor must have at least ",
257                                     out_rows + 1, " elements, but got ",
258                                     row_seq_tensor.NumElements()));
259 OP_REQUIRES(context, col_seq_tensor.NumElements() > out_cols,
260             errors::InvalidArgument("Given out_backprop shape ",
261                                     out_backprop.shape().DebugString(),
262                                     ", col_seq_tensor must have at least ",
263                                     out_cols + 1, " elements, but got ",
264                                     col_seq_tensor.NumElements()));
265
266 auto row_seq_tensor_flat = row_seq_tensor.flat<int64_t>();
267 auto col_seq_tensor_flat = col_seq_tensor.flat<int64_t>();
268 auto orig_input_tensor_shape_flat = orig_input_tensor_shape.flat<int64_t>();
269
270 const int64_t in_batch = orig_input_tensor_shape_flat(0);
271 const int64_t in_rows = orig_input_tensor_shape_flat(1);
272 const int64_t in_cols = orig_input_tensor_shape_flat(2);
273 const int64_t in_depth = orig_input_tensor_shape_flat(3);
274 OP_REQUIRES(

```

```

275     context, in_batch != 0,
276     errors::InvalidArgument("Batch dimension of input must not be 0"));
277 OP_REQUIRES(
278     context, in_rows != 0,
279     errors::InvalidArgument("Rows dimension of input must not be 0"));
280 OP_REQUIRES(
281     context, in_cols != 0,
282     errors::InvalidArgument("Columns dimension of input must not be 0"));
283 OP_REQUIRES(
284     context, in_depth != 0,
285     errors::InvalidArgument("Depth dimension of input must not be 0"));
286
287 constexpr int tensor_in_and_out_dims = 4;
288 // Transform orig_input_tensor_shape into TensorShape
289 TensorShape in_shape;
290 for (auto i = 0; i < tensor_in_and_out_dims; ++i) {
291     in_shape.AddDim(orig_input_tensor_shape_flat(i));
292 }
293
294 // Create intermediate in_backprop.
295 Tensor in_backprop_tensor_temp;
296 OP_REQUIRES_OK(context, context->forward_input_or_allocate_temp(
297     {0}, DataTypeToEnum<double>::v(), in_shape,
298     &in_backprop_tensor_temp));
299 in_backprop_tensor_temp.flat<double>().setZero();
300 // Transform 4D tensor to 2D matrix.
301 EigenDoubleMatrixMap in_backprop_tensor_temp_mat(
302     in_backprop_tensor_temp.flat<double>().data(), in_depth,
303     in_cols * in_rows * in_batch);
304 ConstEigenMatrixMap out_backprop_mat(out_backprop.flat<T>().data(),
305     out_depth,
306     out_cols * out_rows * out_batch);
307 // Loop through each element of out_backprop and evenly distribute the
308 // element to the corresponding pooling cell.
309 const int64_t in_max_row_index = in_rows - 1;
310 const int64_t in_max_col_index = in_cols - 1;
311 for (int64_t b = 0; b < out_batch; ++b) {
312     for (int64_t r = 0; r < out_rows; ++r) {
313         const int64_t in_row_start = row_seq_tensor_flat(r);
314         int64_t in_row_end = overlapping_ ? row_seq_tensor_flat(r + 1)
315             : row_seq_tensor_flat(r + 1) - 1;
316         in_row_end = std::min(in_row_end, in_max_row_index);
317         for (int64_t c = 0; c < out_cols; ++c) {
318             const int64_t in_col_start = col_seq_tensor_flat(c);
319             int64_t in_col_end = overlapping_ ? col_seq_tensor_flat(c + 1)
320                 : col_seq_tensor_flat(c + 1) - 1;
321             in_col_end = std::min(in_col_end, in_max_col_index);
322
323             const int64_t num_elements_in_pooling_cell =

```

```

324         (in_row_end - in_row_start + 1) * (in_col_end - in_col_start + 1);
325         const int64_t out_index = (b * out_rows + r) * out_cols + c;
326         // Now we can evenly distribute out_backprop(b, h, w, *) to
327         // in_backprop(b, hs:he, ws:we, *).
328         for (int64_t in_r = in_row_start; in_r <= in_row_end; ++in_r) {
329             for (int64_t in_c = in_col_start; in_c <= in_col_end; ++in_c) {
330                 const int64_t in_index = (b * in_rows + in_r) * in_cols + in_c;
331                 // Walk through each channel (depth).
332                 for (int64_t d = 0; d < out_depth; ++d) {
333                     const double out_backprop_element = static_cast<double>(
334                         out_backprop_mat.coeffRef(d, out_index));
335                     double& in_backprop_ref =
336                         in_backprop_tensor_temp_mat.coeffRef(d, in_index);
337                     in_backprop_ref +=
338                         out_backprop_element / num_elements_in_pooling_cell;
339                 }
340             }
341         }
342     }
343 }
344 }
345
346 // Depending on the type, cast double to type T.
347 Tensor* in_backprop_tensor = nullptr;
348 OP_REQUIRES_OK(context, context->forward_input_or_allocate_output(
349     {0}, 0, in_shape, &in_backprop_tensor));
350 auto in_backprop_tensor_flat = in_backprop_tensor->flat<T>();
351 auto in_backprop_tensor_temp_flat = in_backprop_tensor_temp.flat<double>();
352 for (int64_t i = 0; i < in_backprop_tensor_flat.size(); ++i) {
353     in_backprop_tensor_flat(i) =
354         static_cast<T>(in_backprop_tensor_temp_flat(i));
355 }
356 }
357
358 private:
359     bool overlapping_;
360 };
361
362 #define REGISTER_FRACTIONALAVGPPOOLGRAD(type) \
363     REGISTER_KERNEL_BUILDER(Name("FractionalAvgPoolGrad") \
364         .Device(DEVICE_CPU) \
365         .TypeConstraint<type>("T"), \
366         FractionalAvgPoolGradOp<type>)
367
368 REGISTER_FRACTIONALAVGPPOOLGRAD(int32);
369 REGISTER_FRACTIONALAVGPPOOLGRAD(int64_t);
370 REGISTER_FRACTIONALAVGPPOOLGRAD(float);
371 REGISTER_FRACTIONALAVGPPOOLGRAD(double);
372

```



```
373 #undef REGISTER_FRACTIONALAVGPPOOLGRAD
374 } // namespace tensorflow
```