

5100e359ae ▾

...

tensorflow / tensorflow / core / kernels / data / experimental / threadpool\_dataset\_op.cc



tensorflow-gardener [tf.data] Change Cardinality() implementation to read c... ... ✖

History

9 contributors



582 lines (498 sloc) | 21.2 KB

...

```

1  /* Copyright 2017 The TensorFlow Authors. All Rights Reserved.
2
3  Licensed under the Apache License, Version 2.0 (the "License");
4  you may not use this file except in compliance with the License.
5  You may obtain a copy of the License at
6
7      http://www.apache.org/licenses/LICENSE-2.0
8
9  Unless required by applicable law or agreed to in writing, software
10 distributed under the License is distributed on an "AS IS" BASIS,
11 WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
12 See the License for the specific language governing permissions and
13 limitations under the License.
14 =====*/
15 #include "tensorflow/core/kernels/data/experimental/threadpool_dataset_op.h"
16
17 #include <memory>
18
19 #include "tensorflow/core/data/dataset_utils.h"
20 #include "tensorflow/core/framework/dataset.h"
21 #include "tensorflow/core/framework/op_kernel.h"
22 #include "tensorflow/core/framework/resource_mgr.h"
23 #include "tensorflow/core/lib/core/refcount.h"
24 #include "tensorflow/core/lib/core/threadpool.h"
25 #include "tensorflow/core/platform/cpu_info.h"
26 #include "tensorflow/core/platform/stringprintf.h"
27 #include "tensorflow/core/platform/thread_annotations.h"
28 #include "tensorflow/core/util/work_sharder.h"
29

```

```

30 namespace tensorflow {
31 namespace data {
32 namespace experimental {
33
34 /* static */ constexpr const char* const
35     MaxIntraOpParallelismDatasetOp::kDatasetType;
36 /* static */ constexpr const char* const
37     MaxIntraOpParallelismDatasetOp::kDatasetOp;
38 /* static */ constexpr const char* const
39     PrivateThreadPoolDatasetOp::kDatasetType;
40 /* static */ constexpr const char* const PrivateThreadPoolDatasetOp::kDatasetOp;
41
42 class ThreadPoolResource : public ResourceBase {
43 public:
44     ThreadPoolResource(Env* env, const ThreadOptions& thread_options,
45                       const string& name, int num_threads, bool low_latency_hint,
46                       int max_intra_op_parallelism)
47         : thread_pool_(env, thread_options, name, num_threads, low_latency_hint),
48           max_intra_op_parallelism_(max_intra_op_parallelism) {}
49
50     // Schedules fn() for execution in the pool of threads.
51     void Schedule(std::function<void()> fn) {
52         if (max_intra_op_parallelism_ < 0) {
53             thread_pool_.Schedule(std::move(fn));
54         } else {
55             thread_pool_.Schedule(std::bind(
56                 [this](std::function<void()> bound_fn) {
57                     // TODO(mrry): Consider moving this thread-local configuration to
58                     // the threads themselves.
59                     ScopedPerThreadMaxParallelism scope(max_intra_op_parallelism_);
60                     bound_fn();
61                 },
62                 std::move(fn)));
63         }
64     }
65
66     int32 NumThreads() { return thread_pool_.NumThreads(); }
67
68     string DebugString() const override { return "ThreadPoolResource"; }
69
70 private:
71     thread::ThreadPool thread_pool_;
72     const int max_intra_op_parallelism_;
73 };
74
75 // Creates a handle to a ThreadPool resource. Note that we don't use
76 // ResourceOpKernel here because the ThreadPoolResource constructor requires
77 // access to `OpKernelContext::env()`, which isn't provided by
78 // `ResourceOpKernel<T>::CreateResource()`.

```

```

79 class ThreadPoolHandleOp : public OpKernel {
80 public:
81     explicit ThreadPoolHandleOp(OpKernelConstruction* ctx) : OpKernel(ctx) {
82         OP_REQUIRES_OK(ctx, ctx->GetAttr("display_name", &display_name_));
83         OP_REQUIRES_OK(ctx, ctx->GetAttr("num_threads", &num_threads_));
84         OP_REQUIRES_OK(ctx, ctx->GetAttr("max_intra_op_parallelism",
85                                         &max_intra_op_parallelism_));
86         OP_REQUIRES(
87             ctx, num_threads_ > 0,
88             errors::InvalidArgument("`num_threads` must be greater than zero."));
89     }
90
91     // The resource is deleted from the resource manager only when it is private
92     // to kernel. Ideally the resource should be deleted when it is no longer held
93     // by anyone, but it would break backward compatibility.
94     ~ThreadPoolHandleOp() override {
95         if (cinfo_.resource_is_private_to_kernel()) {
96             if (!cinfo_.resource_manager()
97                 ->Delete<ThreadPoolResource>(cinfo_.container(), cinfo_.name())
98                 .ok()) {
99                 // Do nothing; the resource can have been deleted by session resets.
100             }
101         }
102     }
103
104     void Compute(OpKernelContext* ctx) override TF_LOCKS_EXCLUDED(mu_) {
105         mutex_lock l(mu_);
106         if (!initialized_) {
107             ResourceMgr* mgr = ctx->resource_manager();
108             OP_REQUIRES_OK(ctx, cinfo_.Init(mgr, def()));
109             ThreadPoolResource* resource;
110             OP_REQUIRES_OK(ctx, mgr->LookupOrCreate<ThreadPoolResource>(
111                 cinfo_.container(), cinfo_.name(), &resource,
112                 [this, ctx](ThreadPoolResource** ret)
113                     TF_EXCLUSIVE_LOCKS_REQUIRED(mu_) {
114                         *ret = new ThreadPoolResource(
115                             ctx->env(), {}, display_name_,
116                             num_threads_,
117                             /*low_latency_hint=*/false,
118                             max_intra_op_parallelism_);
119                         return Status::OK();
120                     }));
121             initialized_ = true;
122         }
123         OP_REQUIRES_OK(ctx, MakeResourceHandleToOutput(
124             ctx, 0, cinfo_.container(), cinfo_.name(),
125             TypeIndex::Make<ThreadPoolResource>()));
126     }
127

```

```

128 private:
129     mutex mu_;
130     ContainerInfo cinfo_ TF_GUARDED_BY(mu_);
131     bool initialized_ TF_GUARDED_BY(mu_) = false;
132     string display_name_;
133     int num_threads_;
134     int max_intra_op_parallelism_;
135 };
136
137 class ThreadPoolDatasetOp : public UnaryDatasetOpKernel {
138 public:
139     explicit ThreadPoolDatasetOp(OpKernelConstruction* ctx)
140         : UnaryDatasetOpKernel(ctx) {}
141
142     void MakeDataset(OpKernelContext* ctx, DatasetBase* input,
143                     DatasetBase** output) override {
144         core::RefCountPtr<ThreadPoolResource> threadpool_resource;
145         OP_REQUIRES_OK(ctx, LookupResource(ctx, HandleFromInput(ctx, 1),
146                                     &threadpool_resource));
147         *output = new Dataset(ctx, input, ctx->input(1), threadpool_resource.get());
148     }
149
150 private:
151     class Dataset : public DatasetBase {
152     public:
153         Dataset(OpKernelContext* ctx, const DatasetBase* input,
154                 const Tensor& resource_handle, ThreadPoolResource* threadpool)
155             : DatasetBase(DatasetContext(ctx)),
156               input_(input),
157               resource_handle_(resource_handle),
158               threadpool_(threadpool) {
159             input_->Ref();
160             threadpool_->Ref();
161         }
162
163         ~Dataset() override {
164             input_->Unref();
165             threadpool_->Unref();
166         }
167
168         std::unique_ptr<IteratorBase> MakeIteratorInternal(
169             const string& prefix) const override {
170             return absl::make_unique<Iterator>(
171                 Iterator::Params{this, strings::StrCat(prefix, "::ThreadPool")});
172         }
173
174         const DataTypeVector& output_dtypes() const override {
175             return input_->output_dtypes();
176         }

```

[illegible]

```

226         return input_impl_->GetNext(IteratorContext(CreateParams(ctx)),
227                                     out_tensors, end_of_sequence);
228     }
229
230 protected:
231     std::shared_ptr<model::Node> CreateNode(
232         IteratorContext* ctx, model::Node::Args args) const override {
233         return model::MakeKnownRatioNode(std::move(args),
234                                           /*ratio=*/1);
235     }
236
237     Status SaveInternal(SerializationContext* ctx,
238                        IteratorStateWriter* writer) override {
239         DCHECK(input_impl_ != nullptr);
240         TF_RETURN_IF_ERROR(SaveInput(ctx, writer, input_impl_));
241         return Status::OK();
242     }
243
244     Status RestoreInternal(IteratorContext* ctx,
245                           IteratorStateReader* reader) override {
246         TF_RETURN_IF_ERROR(RestoreInput(ctx, reader, input_impl_));
247         return Status::OK();
248     }
249
250 private:
251     IteratorContext::Params CreateParams(IteratorContext* ctx) {
252         ThreadPoolResource* pool = dataset()->threadpool_;
253         IteratorContext::Params params(ctx);
254         params.runner = [pool](std::function<void()> c) {
255             pool->Schedule(std::move(c));
256         };
257         params.runner_threadpool_size = pool->NumThreads();
258         return params;
259     }
260
261     std::unique_ptr<IteratorBase> input_impl_;
262 };
263
264 const DatasetBase* const input_;
265 const Tensor resource_handle_;
266 ThreadPoolResource* const threadpool_;
267 };
268 };
269
270 class MaxIntraOpParallelismDatasetOp::Dataset : public DatasetBase {
271 public:
272     Dataset(OpKernelContext* ctx, const DatasetBase* input,
273            int64_t max_intra_op_parallelism)
274         : Dataset(DatasetContext(ctx), input, max_intra_op_parallelism) {}

```

```

275
276 Dataset(DatasetContext&& ctx, const DatasetBase* input,
277         int64_t max_intra_op_parallelism)
278     : DatasetBase(std::move(ctx)),
279       input_(input),
280       max_intra_op_parallelism_(max_intra_op_parallelism),
281       traceme_metadata_(
282         {"parallelism",
283          strings::Printf("%lld", static_cast<long long>(
284             max_intra_op_parallelism_))}) {
285     input_>Ref();
286 }
287
288 ~Dataset() override { input_>Unref(); }
289
290 std::unique_ptr<IteratorBase> MakeIteratorInternal(
291     const string& prefix) const override {
292     return absl::make_unique<Iterator>(Iterator::Params{
293         this, strings::StrCat(prefix, ":", MaxIntraOpParallelism)});
294 }
295
296 const DataTypeVector& output_dtypes() const override {
297     return input_>output_dtypes();
298 }
299 const std::vector<PartialTensorShape>& output_shapes() const override {
300     return input_>output_shapes();
301 }
302
303 string DebugString() const override {
304     return "MaxIntraOpParallelismDatasetOp::Dataset";
305 }
306
307 int64_t CardinalityInternal() const override { return input_>Cardinality(); }
308
309 Status InputDatasets(std::vector<const DatasetBase*>* inputs) const override {
310     inputs->clear();
311     inputs->push_back(input_);
312     return Status::OK();
313 }
314
315 Status CheckExternalState() const override {
316     return input_>CheckExternalState();
317 }
318
319 protected:
320 Status AsGraphDefInternal(SerializationContext* ctx,
321                           DatasetGraphDefBuilder* b,
322                           Node** output) const override {
323     Node* input_graph_node = nullptr;

```

```

324     TF_RETURN_IF_ERROR(b->AddInputDataset(ctx, input_, &input_graph_node));
325     Node* max_intra_op_parallelism_node = nullptr;
326     TF_RETURN_IF_ERROR(b->AddScalar(max_intra_op_parallelism_,
327                                     &max_intra_op_parallelism_node));
328     TF_RETURN_IF_ERROR(b->AddDataset(
329         this, {input_graph_node, max_intra_op_parallelism_node}, output));
330     return Status::OK();
331 }
332
333 private:
334     class Iterator : public DatasetIterator<Dataset> {
335     public:
336         explicit Iterator(const Params& params)
337             : DatasetIterator<Dataset>(params) {}
338
339         Status Initialize(IteratorContext* ctx) override {
340             return dataset()->input_->MakeIterator(ctx, this, prefix(), &input_impl_);
341         }
342
343         Status GetNextInternal(IteratorContext* ctx,
344                               std::vector<Tensor>* out_tensors,
345                               bool* end_of_sequence) override {
346             IteratorContext::Params params(ctx);
347             auto max_parallelism = dataset()->max_intra_op_parallelism_;
348             params.runner = RunnerWithMaxParallelism(*ctx->runner(), max_parallelism);
349             return input_impl_->GetNext(IteratorContext{std::move(params)},
350                                         out_tensors, end_of_sequence);
351         }
352
353     protected:
354         std::shared_ptr<model::Node> CreateNode(
355             IteratorContext* ctx, model::Node::Args args) const override {
356             return model::MakeKnownRatioNode(std::move(args), /*ratio=*/1);
357         }
358
359         Status SaveInternal(SerializationContext* ctx,
360                             IteratorStateWriter* writer) override {
361             DCHECK(input_impl_ != nullptr);
362             TF_RETURN_IF_ERROR(SaveInput(ctx, writer, input_impl_));
363             return Status::OK();
364         }
365
366         Status RestoreInternal(IteratorContext* ctx,
367                               IteratorStateReader* reader) override {
368             TF_RETURN_IF_ERROR(RestoreInput(ctx, reader, input_impl_));
369             return Status::OK();
370         }
371
372         TraceMeMetadata GetTraceMeMetadata() const override {

```



```

373     return dataset()->traceme_metadata_;
374 }
375
376 private:
377     std::unique_ptr<IteratorBase> input_impl_;
378 };
379
380     const DatasetBase* const input_;
381     const int64_t max_intra_op_parallelism_;
382     const TraceMeMetadata traceme_metadata_;
383 };
384
385 /* static */
386 void MaxIntraOpParallelismDatasetOp::MakeDatasetFromOptions(
387     OpKernelContext* ctx, DatasetBase* input, int32_t max_intra_op_parallelism,
388     DatasetBase** output) {
389     OP_REQUIRES(
390         ctx, max_intra_op_parallelism >= 0,
391         errors::InvalidArgument("`max_intra_op_parallelism` must be >= 0"));
392     *output = new Dataset(DatasetContext(DatasetContext::Params(
393         {MaxIntraOpParallelismDatasetOp::kDatasetType,
394         MaxIntraOpParallelismDatasetOp::kDatasetOp})),
395         input, max_intra_op_parallelism);
396 }
397
398 void MaxIntraOpParallelismDatasetOp::MakeDataset(OpKernelContext* ctx,
399     DatasetBase* input,
400     DatasetBase** output) {
401     int64_t max_intra_op_parallelism;
402     OP_REQUIRES_OK(ctx,
403         ParseScalarArgument<int64_t>(ctx, "max_intra_op_parallelism",
404         &max_intra_op_parallelism));
405     OP_REQUIRES(
406         ctx, max_intra_op_parallelism >= 0,
407         errors::InvalidArgument("`max_intra_op_parallelism` must be >= 0"));
408     *output = new Dataset(ctx, input, max_intra_op_parallelism);
409 }
410
411 class PrivateThreadPoolDatasetOp::Dataset : public DatasetBase {
412 public:
413     Dataset(OpKernelContext* ctx, const DatasetBase* input, int num_threads)
414         : Dataset(ctx, DatasetContext(ctx), input, num_threads) {}
415
416     Dataset(OpKernelContext* ctx, DatasetContext&& dataset_ctx,
417         const DatasetBase* input, int num_threads)
418         : DatasetBase(std::move(dataset_ctx)),
419         input_(input),
420         num_threads_(num_threads == 0 ? port::MaxParallelism() : num_threads),
421         traceme_metadata_(

```

```

422         {"num_threads",
423          strings::Printf("%lld", static_cast<long long>(num_threads_))}}) {
424     thread_pool_ = absl::make_unique<thread::ThreadPool>(
425         ctx->env(), ThreadOptions{}, "data_private_threadpool", num_threads_);
426     input_->Ref();
427 }
428
429 ~Dataset() override { input_->Unref(); }
430
431 std::unique_ptr<IteratorBase> MakeIteratorInternal(
432     const string& prefix) const override {
433     return absl::make_unique<Iterator>(
434         Iterator::Params{this, strings::StrCat(prefix, "::PrivateThreadPool")});
435 }
436
437 const DataTypeVector& output_dtypes() const override {
438     return input_->output_dtypes();
439 }
440 const std::vector<PartialTensorShape>& output_shapes() const override {
441     return input_->output_shapes();
442 }
443
444 string DebugString() const override {
445     return "PrivateThreadPoolDatasetOp::Dataset";
446 }
447
448 int64_t CardinalityInternal() const override { return input_->Cardinality(); }
449
450 Status InputDatasets(std::vector<const DatasetBase*>* inputs) const override {
451     inputs->clear();
452     inputs->push_back(input_);
453     return Status::OK();
454 }
455
456 Status CheckExternalState() const override {
457     return input_->CheckExternalState();
458 }
459
460 protected:
461 Status AsGraphDefInternal(SerializationContext* ctx,
462                           DatasetGraphDefBuilder* b,
463                           Node** output) const override {
464     Node* input_graph_node = nullptr;
465     TF_RETURN_IF_ERROR(b->AddInputDataset(ctx, input_, &input_graph_node));
466     Node* num_threads_node = nullptr;
467     TF_RETURN_IF_ERROR(b->AddScalar(num_threads_, &num_threads_node));
468     TF_RETURN_IF_ERROR(
469         b->AddDataset(this, {input_graph_node, num_threads_node}, output));
470     return Status::OK();

```

```

471     }
472
473 private:
474     class Iterator : public DatasetIterator<Dataset> {
475     public:
476         explicit Iterator(const Params& params)
477             : DatasetIterator<Dataset>(params) {}
478
479         Status Initialize(IteratorContext* ctx) override {
480             return dataset()->input_->MakeIterator(ctx, this, prefix(), &input_impl_);
481         }
482
483         Status GetNextInternal(IteratorContext* ctx,
484                               std::vector<Tensor>* out_tensors,
485                               bool* end_of_sequence) override {
486             thread::ThreadPool* pool = dataset()->thread_pool_.get();
487             IteratorContext::Params params(ctx);
488             params.runner = [pool](std::function<void()> c) {
489                 pool->Schedule(std::move(c));
490             };
491             params.runner_threadpool_size = dataset()->num_threads_;
492             return input_impl_->GetNext(IteratorContext{std::move(params)},
493                                         out_tensors, end_of_sequence);
494         }
495
496     protected:
497         std::shared_ptr<model::Node> CreateNode(
498             IteratorContext* ctx, model::Node::Args args) const override {
499             return model::MakeKnownRatioNode(std::move(args), /*ratio=*/1);
500         }
501
502         Status SaveInternal(SerializationContext* ctx,
503                             IteratorStateWriter* writer) override {
504             DCHECK(input_impl_ != nullptr);
505             TF_RETURN_IF_ERROR(SaveInput(ctx, writer, input_impl_));
506             return Status::OK();
507         }
508
509         Status RestoreInternal(IteratorContext* ctx,
510                               IteratorStateReader* reader) override {
511             TF_RETURN_IF_ERROR(RestoreInput(ctx, reader, input_impl_));
512             return Status::OK();
513         }
514
515         TraceMeMetadata GetTraceMeMetadata() const override {
516             return dataset()->traceme_metadata_;
517         }
518
519     private:

```

```

520     std::unique_ptr<IteratorBase> input_impl_;
521 };
522
523     const DatasetBase* const input_;
524     const int64_t num_threads_;
525     const TraceMeMetadata traceme_metadata_;
526     std::unique_ptr<thread::ThreadPool> thread_pool_;
527 };
528
529 /* static */
530 void PrivateThreadPoolDatasetOp::MakeDatasetFromOptions(OpKernelContext* ctx,
531                                                         DatasetBase* input,
532                                                         int32_t num_threads,
533                                                         DatasetBase** output) {
534     OP_REQUIRES(ctx, num_threads >= 0,
535                 errors::InvalidArgument("`num_threads` must be >= 0"));
536     *output = new Dataset(ctx,
537                           DatasetContext(DatasetContext::Params(
538                               {PrivateThreadPoolDatasetOp::kDatasetType,
539                                PrivateThreadPoolDatasetOp::kDatasetOp})),
540                           input, num_threads);
541 }
542
543 void PrivateThreadPoolDatasetOp::MakeDataset(OpKernelContext* ctx,
544                                              DatasetBase* input,
545                                              DatasetBase** output) {
546     int64_t num_threads = 0;
547     OP_REQUIRES_OK(
548         ctx, ParseScalarArgument<int64_t>(ctx, "num_threads", &num_threads));
549     OP_REQUIRES(ctx, num_threads >= 0,
550                 errors::InvalidArgument("`num_threads` must be >= 0"));
551     *output = new Dataset(ctx, input, num_threads);
552 }
553
554 namespace {
555
556 REGISTER_KERNEL_BUILDER(Name("MaxIntraOpParallelismDataset").Device(DEVICE_CPU),
557                         MaxIntraOpParallelismDatasetOp);
558 REGISTER_KERNEL_BUILDER(
559     Name("ExperimentalMaxIntraOpParallelismDataset").Device(DEVICE_CPU),
560     MaxIntraOpParallelismDatasetOp);
561
562 REGISTER_KERNEL_BUILDER(Name("PrivateThreadPoolDataset").Device(DEVICE_CPU),
563                         PrivateThreadPoolDatasetOp);
564 REGISTER_KERNEL_BUILDER(
565     Name("ExperimentalPrivateThreadPoolDataset").Device(DEVICE_CPU),
566     PrivateThreadPoolDatasetOp);
567
568 REGISTER_KERNEL_BUILDER(Name("ThreadPoolHandle").Device(DEVICE_CPU),

```

```
569         ThreadPoolHandleOp);
570 REGISTER_KERNEL_BUILDER(Name("ExperimentalThreadPoolHandle").Device(DEVICE_CPU),
571         ThreadPoolHandleOp);
572
573 REGISTER_KERNEL_BUILDER(Name("ThreadPoolDataset").Device(DEVICE_CPU),
574         ThreadPoolDatasetOp);
575 REGISTER_KERNEL_BUILDER(
576     Name("ExperimentalThreadPoolDataset").Device(DEVICE_CPU),
577     ThreadPoolDatasetOp);
578
579 } // namespace
580 } // namespace experimental
581 } // namespace data
582 } // namespace tensorflow
```