

5100e359ae ▾

...

tensorflow / tensorflow / core / kernels / map_stage_op.cc



jpienaar Rename to underlying type rather than alias ... ✓

History

7 contributors



808 lines (644 sloc) | 25.7 KB

...

```

1  /* Copyright 2017 The TensorFlow Authors. All Rights Reserved.
2
3  Licensed under the Apache License, Version 2.0 (the "License");
4  you may not use this file except in compliance with the License.
5  You may obtain a copy of the License at
6
7      http://www.apache.org/licenses/LICENSE-2.0
8
9  Unless required by applicable law or agreed to in writing, software
10 distributed under the License is distributed on an "AS IS" BASIS,
11 WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
12 See the License for the specific language governing permissions and
13 limitations under the License.
14 =====*/
15
16 #include <cstdint>
17 #include <functional>
18 #include <map>
19 #include <mutex>
20 #include <numeric>
21 #include <unordered_map>
22 #include <vector>
23
24 #include "tensorflow/core/framework/op_kernel.h"
25 #include "tensorflow/core/framework/resource_mgr.h"
26 #include "tensorflow/core/framework/tensor.h"
27 #include "tensorflow/core/framework/tensor_shape.h"
28 #include "tensorflow/core/lib/gtl/optional.h"
29 #include "tensorflow/core/lib/strings/strcat.h"

```

```

30 #include "tensorflow/core/platform/env.h"
31 #include "tensorflow/core/platform/mutex.h"
32 #include "tensorflow/core/platform/thread_annotations.h"
33
34 namespace tensorflow {
35 namespace {
36
37 // Partial Ordering Comparator for Tensor keys containing scalar int64's
38 struct KeyTensorLess {
39     bool operator()(const Tensor& lhs, const Tensor& rhs) const {
40         return std::less<int64_t>{}(lhs.scalar<int64_t>()(),
41                                     rhs.scalar<int64_t>()());
42     }
43 };
44
45 // Key Equality operator for Tensor keys containing scalar int64's
46 struct KeyTensorEqual {
47     bool operator()(const Tensor& lhs, const Tensor& rhs) const {
48         return std::equal_to<int64_t>{}(lhs.scalar<int64_t>()(),
49                                         rhs.scalar<int64_t>()());
50     }
51 };
52
53 // Hash for Tensor keys containing scalar int64's
54 struct KeyTensorHash {
55     std::size_t operator()(const Tensor& key) const {
56         return std::hash<int64_t>{}(key.scalar<int64_t>()());
57     }
58 };
59
60 // Primary template.
61 template <bool Ordered, typename Data>
62 struct MapTraits;
63
64 // Partial specialization for ordered.
65 template <typename Data>
66 struct MapTraits<true, Data> {
67     using KeyType = Tensor;
68     using DataType = Data;
69     using MapType = std::map<KeyType, Data, KeyTensorLess>;
70 };
71
72 // Partial specialization for unordered.
73 template <typename Data>
74 struct MapTraits<false, Data> {
75     using KeyType = Tensor;
76     using DataType = Data;
77     using MapType =
78         std::unordered_map<KeyType, Data, KeyTensorHash, KeyTensorEqual>;

```

```

79 };
80
81 // Wrapper around map/unordered_map.
82 template <bool Ordered>
83 class StagingMap : public ResourceBase {
84 public:
85     // Public typedefs
86     using Tuple = std::vector<Tensor>;
87     using OptionalTensor = gtl::optional<Tensor>;
88     using OptionalTuple = std::vector<OptionalTensor>;
89
90     using MapType = typename MapTraits<Ordered, OptionalTuple>::MapType;
91     using KeyType = typename MapTraits<Ordered, OptionalTuple>::KeyType;
92
93     using IncompleteType = typename MapTraits<false, OptionalTuple>::MapType;
94
95 private:
96     // Private variables
97     DataTypeVector dtypes_ TF_GUARDED_BY(mu_);
98     std::size_t capacity_ TF_GUARDED_BY(mu_);
99     std::size_t memory_limit_ TF_GUARDED_BY(mu_);
100     std::size_t current_bytes_ TF_GUARDED_BY(mu_);
101     tensorflow::mutex mu_;
102     tensorflow::condition_variable not_empty_;
103     tensorflow::condition_variable full_;
104     IncompleteType incomplete_ TF_GUARDED_BY(mu_);
105     MapType map_ TF_GUARDED_BY(mu_);
106
107 private:
108     // private methods
109
110     // If map is configured for bounded capacity, notify
111     // waiting inserters that space is now available
112     void notify_inserters_if_bounded() TF_EXCLUSIVE_LOCKS_REQUIRED(mu_) {
113         if (has_capacity() || has_memory_limit()) {
114             // Notify all inserters. The removal of an element
115             // may make memory available for many inserters
116             // to insert new elements
117             full_.notify_all();
118         }
119     }
120
121     // Notify all removers waiting to extract values
122     // that data is now available
123     void notify_removers() {
124         // Notify all removers. This is because they are
125         // waiting for specific keys to appear in the map
126         // so we don't know which one to wake up.
127         not_empty_.notify_all();

```

```

128     }
129
130     bool has_capacity() const TF_EXCLUSIVE_LOCKS_REQUIRED(mu_) {
131         return capacity_ > 0;
132     }
133
134     bool has_memory_limit() const TF_EXCLUSIVE_LOCKS_REQUIRED(mu_) {
135         return memory_limit_ > 0;
136     }
137
138     bool would_exceed_memory_limit(std::size_t bytes) const
139         TF_EXCLUSIVE_LOCKS_REQUIRED(mu_) {
140         return has_memory_limit() && bytes + current_bytes_ > memory_limit_;
141     }
142
143     bool is_capacity_full() const TF_EXCLUSIVE_LOCKS_REQUIRED(mu_) {
144         return has_capacity() && map_.size() >= capacity_;
145     }
146
147     // Get number of bytes in the tuple
148     std::size_t get_tuple_bytes(const Tuple& tuple) {
149         return std::accumulate(tuple.begin(), tuple.end(),
150                                static_cast<std::size_t>(0),
151                                [](const std::size_t& lhs, const Tensor& rhs) {
152                                    return lhs + rhs.TotalBytes();
153                                });
154     }
155
156     // Get number of bytes in the incomplete tuple
157     std::size_t get_tuple_bytes(const OptionalTuple& tuple) {
158         return std::accumulate(
159             tuple.begin(), tuple.end(), static_cast<std::size_t>(0),
160             [](const std::size_t& lhs, const OptionalTensor& rhs) {
161                 return (lhs + rhs.has_value()) ? rhs.value().TotalBytes() : 0;
162             });
163     }
164
165     // Check that the index is within bounds
166     Status check_index(const Tensor& key, std::size_t index)
167         TF_EXCLUSIVE_LOCKS_REQUIRED(mu_) {
168         if (index >= dtypes_.size()) {
169             return Status(errors::InvalidArgument(
170                 "Index '", index, "' for key '", key.scalar<int64_t>()(),
171                 "' was out of bounds '", dtypes_.size(), "'."));
172         }
173
174         return Status::OK();
175     }
176

```

```

177 Status copy_or_move_tensors(OptionalTuple* map_tuple, const Tensor& key,
178                             const Tensor& indices, Tuple* output,
179                             bool copy = false)
180     TF_EXCLUSIVE_LOCKS_REQUIRED(mu_) {
181     auto findices = indices.flat<int>();
182
183     // Return values at specified indices
184     for (std::size_t i = 0; i < findices.dimension(0); ++i) {
185         std::size_t index = findices(i);
186
187         TF_RETURN_IF_ERROR(check_index(key, index));
188
189         // Insist on a value present at the specified index
190         if (!(*map_tuple)[index].has_value()) {
191             return Status(errors::InvalidArgument(
192                 "Tensor at index '", index, "' for key '", key.scalar<int64_t>()(),
193                 "' has already been removed."));
194         }
195
196         // Copy the contained tensor and
197         // remove from the OptionalTuple
198         output->push_back((*map_tuple)[index].value());
199
200         // Clear out the entry if we're not copying (moving)
201         if (!copy) {
202             (*map_tuple)[index].reset();
203         }
204     }
205
206     return Status::OK();
207 }
208
209 // Check that the optional value at the specified index
210 // is uninitialized
211 Status check_index_uninitialized(const Tensor& key, std::size_t index,
212                                 const OptionalTuple& tuple)
213     TF_EXCLUSIVE_LOCKS_REQUIRED(mu_) {
214     if (tuple[index].has_value()) {
215         return errors::InvalidArgument("The tensor for index '", index,
216                                         "' for key '", key.scalar<int64_t>()(),
217                                         "' was already initialized '",
218                                         dtypes_.size(), "'.");
219     }
220
221     return Status::OK();
222 }
223
224 // Check that the indices are strictly ordered
225 Status check_index_ordering(const Tensor& indices) {

```

```

226     if (indices.NumElements() == 0) {
227         return errors::InvalidArgument("Indices are empty");
228     }
229
230     auto findices = indices.flat<int>();
231
232     for (std::size_t i = 0; i < findices.dimension(0) - 1; ++i) {
233         if (findices(i) < findices(i + 1)) {
234             continue;
235         }
236
237         return errors::InvalidArgument("Indices are not strictly ordered");
238     }
239
240     return Status::OK();
241 }
242
243 // Check bytes are within memory limits memory limits
244 Status check_memory_limit(std::size_t bytes)
245     TF_EXCLUSIVE_LOCKS_REQUIRED(mu_) {
246     if (has_memory_limit() && bytes > memory_limit_) {
247         return errors::ResourceExhausted(
248             "Attempted to insert tensors with combined size of '", bytes,
249             "' bytes into Staging Area with a memory limit of '", memory_limit_,
250             "'.");
251     }
252
253     return Status::OK();
254 }
255
256 // Insert incomplete data into the Barrier
257 Status put_incomplete(const KeyType& key, const Tensor& indices,
258                     OptionalTuple* tuple, tensorflow::mutex_lock* lock)
259     TF_EXCLUSIVE_LOCKS_REQUIRED(mu_) {
260     auto findices = indices.flat<int>();
261
262     // Search for the key in our incomplete set
263     auto it = incomplete_.find(key);
264
265     // Check that the tuple fits within the memory limit
266     std::size_t tuple_bytes = get_tuple_bytes(*tuple);
267     TF_RETURN_IF_ERROR(check_memory_limit(tuple_bytes));
268
269     // Wait until we don't exceed the memory limit
270     while (would_exceed_memory_limit(tuple_bytes)) {
271         full_.wait(*lock);
272     }
273
274     // This key isn't present in the incomplete set

```

```

275 // Create OptionalTuple and insert
276 if (it == incomplete_.end()) {
277     OptionalTuple empty(dtypes_.size());
278
279     // Initialize empty tuple with given dta
280     for (std::size_t i = 0; i < findindices.dimension(0); ++i) {
281         std::size_t index = findindices(i);
282         TF_RETURN_IF_ERROR(check_index(key, index));
283
284         // Assign tuple at this index
285         empty[index] = std::move((*tuple)[i]);
286     }
287
288     // Insert into incomplete map
289     incomplete_.insert({key, std::move(empty)});
290
291     // Increment size
292     current_bytes_ += tuple_bytes;
293 }
294 // Found an entry in the incomplete index
295 // Update with given data and insert complete entries
296 // into the main map
297 else {
298     // Reference existing incomplete tuple
299     OptionalTuple& present = it->second;
300
301     // Assign given data
302     for (std::size_t i = 0; i < findindices.dimension(0); ++i) {
303         std::size_t index = findindices(i);
304         TF_RETURN_IF_ERROR(check_index(key, index));
305         TF_RETURN_IF_ERROR(check_index_uninitialized(key, index, present));
306
307         // Assign tuple at this index
308         present[index] = std::move((*tuple)[i]);
309     }
310
311     // Increment size
312     current_bytes_ += tuple_bytes;
313
314     // Do we have values at all tuple elements?
315     bool complete =
316         std::all_of(present.begin(), present.end(),
317             [](const OptionalTensor& v) { return v.has_value(); });
318
319     // If so, put the tuple in the actual map
320     if (complete) {
321         OptionalTuple insert_tuple = std::move(it->second);
322
323         // Remove from incomplete

```

```

324         incomplete_.erase(it);
325
326         TF_RETURN_IF_ERROR(put_complete(key, &insert_tuple));
327     }
328 }
329
330     return Status::OK();
331 }
332
333 // Does the insertion into the actual staging area
334 Status put_complete(const KeyType& key, OptionalTuple* tuple)
335     TF_EXCLUSIVE_LOCKS_REQUIRED(mu_) {
336     // Insert key and tuples into the map
337     map_.insert({key, std::move(*tuple)});
338
339     notify_removers();
340
341     return Status::OK();
342 }
343
344 public:
345     // public methods
346     explicit StagingMap(const DataTypeVector& dtypes, std::size_t capacity,
347                        std::size_t memory_limit)
348         : dtypes_(dtypes),
349           capacity_(capacity),
350           memory_limit_(memory_limit),
351           current_bytes_(0) {}
352
353     Status put(KeyType* key, const Tensor* indices, OptionalTuple* tuple) {
354         tensorflow::mutex_lock lock(mu_);
355
356         // Sanity check the indices
357         TF_RETURN_IF_ERROR(check_index_ordering(*indices));
358
359         // Handle incomplete inserts
360         if (indices->NumElements() != dtypes_.size()) {
361             return put_incomplete(*key, *indices, tuple, &lock);
362         }
363
364         std::size_t tuple_bytes = get_tuple_bytes(*tuple);
365         // Check that tuple_bytes fits within the memory limit
366         TF_RETURN_IF_ERROR(check_memory_limit(tuple_bytes));
367
368         // Wait until there's space for insertion.
369         while (would_exceed_memory_limit(tuple_bytes) || is_capacity_full()) {
370             full_.wait(lock);
371         }
372

```



```

373     // Do the put operation
374     TF_RETURN_IF_ERROR(put_complete(*key, tuple));
375
376     // Update the current size
377     current_bytes_ += tuple_bytes;
378
379     return Status::OK();
380 }
381
382 Status get(const KeyType* key, const Tensor* indices, Tuple* tuple) {
383     tensorflow::mutex_lock lock(mu_);
384
385     // Sanity check the indices
386     TF_RETURN_IF_ERROR(check_index_ordering(*indices));
387
388     typename MapType::iterator it;
389
390     // Wait until the element with the requested key is present
391     while ((it = map_.find(*key)) == map_.end()) {
392         not_empty_.wait(lock);
393     }
394
395     TF_RETURN_IF_ERROR(
396         copy_or_move_tensors(&it->second, *key, *indices, tuple, true));
397
398     // Update bytes in the Staging Area
399     current_bytes_ -= get_tuple_bytes(*tuple);
400
401     return Status::OK();
402 }
403
404 Status pop(const KeyType* key, const Tensor* indices, Tuple* tuple) {
405     tensorflow::mutex_lock lock(mu_);
406
407     // Sanity check the indices
408     TF_RETURN_IF_ERROR(check_index_ordering(*indices));
409
410     typename MapType::iterator it;
411
412     // Wait until the element with the requested key is present
413     while ((it = map_.find(*key)) == map_.end()) {
414         not_empty_.wait(lock);
415     }
416
417     TF_RETURN_IF_ERROR(
418         copy_or_move_tensors(&it->second, *key, *indices, tuple));
419
420     // Remove entry if all the values have been consumed
421     if (!std::any_of(

```

```

422         it->second.begin(), it->second.end(),
423         [](const OptionalTensor& tensor) { return tensor.has_value(); })) {
424     map_.erase(it);
425 }
426
427 // Update bytes in the Staging Area
428 current_bytes_ -= get_tuple_bytes(*tuple);
429
430 notify_inserters_if_bounded();
431
432 return Status::OK();
433 }
434
435 Status popitem(KeyType* key, const Tensor* indices, Tuple* tuple) {
436     tensorflow::mutex_lock lock(mu_);
437
438     // Sanity check the indices
439     TF_RETURN_IF_ERROR(check_index_ordering(*indices));
440
441     // Wait until map is not empty
442     while (this->map_.empty()) {
443         not_empty_.wait(lock);
444     }
445
446     // Move from the first element and erase it
447
448     auto it = map_.begin();
449
450     TF_RETURN_IF_ERROR(
451         copy_or_move_tensors(&it->second, *key, *indices, tuple));
452
453     *key = it->first;
454
455     // Remove entry if all the values have been consumed
456     if (!std::any_of(
457         it->second.begin(), it->second.end(),
458         [](const OptionalTensor& tensor) { return tensor.has_value(); })) {
459         map_.erase(it);
460     }
461
462     // Update bytes in the Staging Area
463     current_bytes_ -= get_tuple_bytes(*tuple);
464
465     notify_inserters_if_bounded();
466
467     return Status::OK();
468 }
469
470 Status clear() {

```

```

471     tensorflow::mutex_lock lock(mu_);
472     map_.clear();
473     incomplete_.clear();
474     current_bytes_ = 0;
475
476     notify_inserters_if_bounded();
477
478     return Status::OK();
479 }
480
481 std::size_t incomplete_size() {
482     tensorflow::mutex_lock lock(mu_);
483     return incomplete_.size();
484 }
485
486 std::size_t size() {
487     tensorflow::mutex_lock lock(mu_);
488     return map_.size();
489 }
490
491 string DebugString() const override { return "StagingMap"; }
492 };
493
494 template <bool Ordered>
495 Status GetStagingMap(OpKernelContext* ctx, const NodeDef& ndef,
496                    StagingMap<Ordered>** map) {
497     auto rm = ctx->resource_manager();
498     ContainerInfo cinfo;
499
500     // Lambda for creating the Staging Area
501     auto create_fn = [&ndef](StagingMap<Ordered>** ret) -> Status {
502         DataTypeVector dtypes;
503         int64_t capacity;
504         int64_t memory_limit;
505         TF_RETURN_IF_ERROR(GetNodeAttr(ndef, "dtypes", &dtypes));
506         TF_RETURN_IF_ERROR(GetNodeAttr(ndef, "capacity", &capacity));
507         TF_RETURN_IF_ERROR(GetNodeAttr(ndef, "memory_limit", &memory_limit));
508         *ret = new StagingMap<Ordered>(dtypes, capacity, memory_limit);
509         return Status::OK();
510     };
511
512     TF_RETURN_IF_ERROR(cinfo.Init(rm, ndef, true /* use name() */));
513     TF_RETURN_IF_ERROR(rm->LookupOrCreate<StagingMap<Ordered>>(
514         cinfo.container(), cinfo.name(), map, create_fn));
515     return Status::OK();
516 }
517
518 template <bool Ordered>
519 class MapStageOp : public OpKernel {

```

```

520 public:
521     explicit MapStageOp(OpKernelConstruction* ctx) : OpKernel(ctx) {}
522
523     void Compute(OpKernelContext* ctx) override {
524         StagingMap<Ordered>* map = nullptr;
525         OP_REQUIRES_OK(ctx, GetStagingMap(ctx, def(), &map));
526         core::ScopedUnref scope(map);
527         typename StagingMap<Ordered>::OptionalTuple tuple;
528
529         const Tensor* key_tensor;
530         const Tensor* indices_tensor;
531         OpInputList values_tensor;
532
533         OP_REQUIRES_OK(ctx, ctx->input("key", &key_tensor));
534         OP_REQUIRES_OK(ctx, ctx->input("indices", &indices_tensor));
535         OP_REQUIRES_OK(ctx, ctx->input_list("values", &values_tensor));
536         OP_REQUIRES(ctx, key_tensor->NumElements() > 0,
537                     errors::InvalidArgument("key must not be empty"));
538
539         // Create copy for insertion into Staging Area
540         Tensor key(*key_tensor);
541
542         // Create the tuple to store
543         for (std::size_t i = 0; i < values_tensor.size(); ++i) {
544             tuple.push_back(values_tensor[i]);
545         }
546
547         // Store the tuple in the map
548         OP_REQUIRES_OK(ctx, map->put(&key, indices_tensor, &tuple));
549     }
550 };
551
552 REGISTER_KERNEL_BUILDER(Name("MapStage").Device(DEVICE_CPU), MapStageOp<false>);
553 REGISTER_KERNEL_BUILDER(Name("OrderedMapStage").Device(DEVICE_CPU),
554                         MapStageOp<true>);
555
556 #if GOOGLE_CUDA || TENSORFLOW_USE_ROCM
557 REGISTER_KERNEL_BUILDER(
558     Name("MapStage").HostMemory("key").HostMemory("indices").Device(DEVICE_GPU),
559     MapStageOp<false>);
560 REGISTER_KERNEL_BUILDER(Name("OrderedMapStage")
561                         .HostMemory("key")
562                         .HostMemory("indices")
563                         .Device(DEVICE_GPU),
564                         MapStageOp<true>);
565 #endif // GOOGLE_CUDA || TENSORFLOW_USE_ROCM
566
567
568 template <bool Ordered>

```

```

569 class MapUnstageOp : public OpKernel {
570 public:
571     explicit MapUnstageOp(OpKernelConstruction* ctx) : OpKernel(ctx) {}
572
573     // Using this op in such a way that it blocks forever
574     // is an error. As such cancellation is not handled.
575     void Compute(OpKernelContext* ctx) override {
576         StagingMap<Ordered>* map = nullptr;
577         OP_REQUIRES_OK(ctx, GetStagingMap(ctx, def(), &map));
578         core::ScopedUnref scope(map);
579         typename StagingMap<Ordered>::Tuple tuple;
580
581         const Tensor* key_tensor;
582         const Tensor* indices_tensor;
583
584         OP_REQUIRES_OK(ctx, ctx->input("key", &key_tensor));
585         OP_REQUIRES_OK(ctx, ctx->input("indices", &indices_tensor));
586         OP_REQUIRES_OK(ctx, map->pop(key_tensor, indices_tensor, &tuple));
587
588         OP_REQUIRES(
589             ctx, tuple.size() == indices_tensor->NumElements(),
590             errors::InvalidArgument("output/indices size mismatch: ", tuple.size(),
591                                     " vs. ", indices_tensor->NumElements()));
592
593         for (std::size_t i = 0; i < tuple.size(); ++i) {
594             ctx->set_output(i, tuple[i]);
595         }
596     }
597 };
598
599 REGISTER_KERNEL_BUILDER(Name("MapUnstage").Device(DEVICE_CPU),
600                         MapUnstageOp<false>);
601 REGISTER_KERNEL_BUILDER(Name("OrderedMapUnstage").Device(DEVICE_CPU),
602                         MapUnstageOp<true>);
603
604 #if GOOGLE_CUDA || TENSORFLOW_USE_ROCM
605 REGISTER_KERNEL_BUILDER(Name("MapUnstage")
606                         .HostMemory("key")
607                         .HostMemory("indices")
608                         .Device(DEVICE_GPU),
609                         MapUnstageOp<false>);
610 REGISTER_KERNEL_BUILDER(Name("OrderedMapUnstage")
611                         .HostMemory("key")
612                         .HostMemory("indices")
613                         .Device(DEVICE_GPU),
614                         MapUnstageOp<true>);
615 #endif
616
617 template <bool Ordered>

```

```

618 class MapPeekOp : public OpKernel {
619 public:
620     explicit MapPeekOp(OpKernelConstruction* ctx) : OpKernel(ctx) {}
621
622     // Using this op in such a way that it blocks forever
623     // is an error. As such cancellation is not handled.
624     void Compute(OpKernelContext* ctx) override {
625         StagingMap<Ordered>* map = nullptr;
626         OP_REQUIRES_OK(ctx, GetStagingMap(ctx, def(), &map));
627         core::ScopedUnref scope(map);
628         typename StagingMap<Ordered>::Tuple tuple;
629
630         const Tensor* key_tensor;
631         const Tensor* indices_tensor;
632
633         OP_REQUIRES_OK(ctx, ctx->input("key", &key_tensor));
634         OP_REQUIRES_OK(ctx, ctx->input("indices", &indices_tensor));
635         OP_REQUIRES_OK(ctx, map->get(key_tensor, indices_tensor, &tuple));
636
637         OP_REQUIRES(
638             ctx, tuple.size() == indices_tensor->NumElements(),
639             errors::InvalidArgument("output/indices size mismatch: ", tuple.size(),
640                                     " vs. ", indices_tensor->NumElements()));
641
642         for (std::size_t i = 0; i < tuple.size(); ++i) {
643             ctx->set_output(i, tuple[i]);
644         }
645     }
646 };
647
648 REGISTER_KERNEL_BUILDER(Name("MapPeek").Device(DEVICE_CPU), MapPeekOp<false>);
649 REGISTER_KERNEL_BUILDER(Name("OrderedMapPeek").Device(DEVICE_CPU),
650                         MapPeekOp<true>);
651
652 #if GOOGLE_CUDA || TENSORFLOW_USE_ROCM
653 REGISTER_KERNEL_BUILDER(
654     Name("MapPeek").HostMemory("key").HostMemory("indices").Device(DEVICE_GPU),
655     MapPeekOp<false>);
656 REGISTER_KERNEL_BUILDER(Name("OrderedMapPeek")
657                         .HostMemory("key")
658                         .HostMemory("indices")
659                         .Device(DEVICE_GPU),
660                         MapPeekOp<true>);
661 #endif
662
663
664 template <bool Ordered>
665 class MapUnstageNoKeyOp : public OpKernel {
666 public:

```



```

716 #endif
717
718
719 template <bool Ordered>
720 class MapSizeOp : public OpKernel {
721 public:
722     explicit MapSizeOp(OpKernelConstruction* ctx) : OpKernel(ctx) {}
723
724     void Compute(OpKernelContext* ctx) override {
725         StagingMap<Ordered>* map = nullptr;
726         OP_REQUIRES_OK(ctx, GetStagingMap(ctx, def(), &map));
727         core::ScopedUnref scope(map);
728
729         // Allocate size output tensor
730         Tensor* size = nullptr;
731         OP_REQUIRES_OK(ctx, ctx->allocate_output(0, TensorShape({}), &size));
732
733         // Set it to the actual size
734         size->scalar<int32>().setConstant(map->size());
735     }
736 };
737
738 REGISTER_KERNEL_BUILDER(Name("MapSize").Device(DEVICE_CPU), MapSizeOp<false>);
739 REGISTER_KERNEL_BUILDER(Name("OrderedMapSize").Device(DEVICE_CPU),
740                         MapSizeOp<true>);
741
742 #if GOOGLE_CUDA || TENSORFLOW_USE_ROCM
743 REGISTER_KERNEL_BUILDER(Name("MapSize").Device(DEVICE_GPU).HostMemory("size"),
744                         MapSizeOp<false>);
745 REGISTER_KERNEL_BUILDER(
746     Name("OrderedMapSize").Device(DEVICE_GPU).HostMemory("size"),
747     MapSizeOp<true>);
748 #endif
749
750 template <bool Ordered>
751 class MapIncompleteSizeOp : public OpKernel {
752 public:
753     explicit MapIncompleteSizeOp(OpKernelConstruction* ctx) : OpKernel(ctx) {}
754
755     void Compute(OpKernelContext* ctx) override {
756         StagingMap<Ordered>* map = nullptr;
757         OP_REQUIRES_OK(ctx, GetStagingMap(ctx, def(), &map));
758         core::ScopedUnref scope(map);
759
760         // Allocate size output tensor
761         Tensor* size = nullptr;
762         OP_REQUIRES_OK(ctx, ctx->allocate_output(0, TensorShape({}), &size));
763
764         // Set it to the actual size

```



```

765     size->scalar<int32>().setConstant(map->incomplete_size());
766 }
767 };
768
769 REGISTER_KERNEL_BUILDER(Name("MapIncompleteSize").Device(DEVICE_CPU),
770     MapIncompleteSizeOp<false>);
771 REGISTER_KERNEL_BUILDER(Name("OrderedMapIncompleteSize").Device(DEVICE_CPU),
772     MapIncompleteSizeOp<true>);
773
774 #if GOOGLE_CUDA || TENSORFLOW_USE_ROCM
775 REGISTER_KERNEL_BUILDER(
776     Name("MapIncompleteSize").Device(DEVICE_GPU).HostMemory("size"),
777     MapIncompleteSizeOp<false>);
778 REGISTER_KERNEL_BUILDER(
779     Name("OrderedMapIncompleteSize").Device(DEVICE_GPU).HostMemory("size"),
780     MapIncompleteSizeOp<true>);
781 #endif
782
783 template <bool Ordered>
784 class MapClearOp : public OpKernel {
785 public:
786     explicit MapClearOp(OpKernelConstruction* ctx) : OpKernel(ctx) {}
787
788     void Compute(OpKernelContext* ctx) override {
789         StagingMap<Ordered>* map = nullptr;
790         OP_REQUIRES_OK(ctx, GetStagingMap(ctx, def(), &map));
791         core::ScopedUnref scope(map);
792
793         OP_REQUIRES_OK(ctx, map->clear());
794     }
795 };
796
797 REGISTER_KERNEL_BUILDER(Name("MapClear").Device(DEVICE_CPU), MapClearOp<false>);
798 REGISTER_KERNEL_BUILDER(Name("OrderedMapClear").Device(DEVICE_CPU),
799     MapClearOp<true>);
800
801 #if GOOGLE_CUDA || TENSORFLOW_USE_ROCM
802 REGISTER_KERNEL_BUILDER(Name("MapClear").Device(DEVICE_GPU), MapClearOp<false>);
803 REGISTER_KERNEL_BUILDER(Name("OrderedMapClear").Device(DEVICE_GPU),
804     MapClearOp<true>);
805 #endif
806
807 } // namespace
808 } // namespace tensorflow

```