

5100e359ae ▾

...

tensorflow / tensorflow / core / kernels / fractional_max_pool_op.cc



jpienaar Rename to underlying type rather than alias ... ✓

History

3 contributors



390 lines (343 sloc) | 15.9 KB

...

```

1  /* Copyright 2016 The TensorFlow Authors. All Rights Reserved.
2
3  Licensed under the Apache License, Version 2.0 (the "License");
4  you may not use this file except in compliance with the License.
5  You may obtain a copy of the License at
6
7      http://www.apache.org/licenses/LICENSE-2.0
8
9  Unless required by applicable law or agreed to in writing, software
10 distributed under the License is distributed on an "AS IS" BASIS,
11 WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
12 See the License for the specific language governing permissions and
13 limitations under the License.
14 =====*/
15 #define EIGEN_USE_THREADS
16
17 #include <algorithm>
18 #include <cmath>
19 #include <random>
20 #include <vector>
21
22 #include "tensorflow/core/kernels/fractional_pool_common.h"
23
24 #include "third_party/eigen3/unsupported/Eigen/CXX11/Tensor"
25 #include "tensorflow/core/framework/numeric_op.h"
26 #include "tensorflow/core/framework/op_kernel.h"
27 #include "tensorflow/core/lib/random/random.h"
28 #include "tensorflow/core/platform/logging.h"
29 #include "tensorflow/core/platform/mutex.h"

```

```

30 #include "tensorflow/core/util/guarded_philox_random.h"
31
32 namespace tensorflow {
33 typedef Eigen::ThreadPoolDevice CPUDevice;
34
35 template <typename T>
36 class FractionalMaxPoolOp : public OpKernel {
37 public:
38     explicit FractionalMaxPoolOp(OpKernelConstruction* context)
39         : OpKernel(context) {
40         OP_REQUIRES_OK(context, context->GetAttr("pooling_ratio", &pooling_ratio_));
41         OP_REQUIRES_OK(context, context->GetAttr("pseudo_random", &pseudo_random_));
42         OP_REQUIRES_OK(context, context->GetAttr("overlapping", &overlapping_));
43
44         OP_REQUIRES(context, pooling_ratio_.size() == 4,
45                     errors::InvalidArgument("pooling_ratio field must "
46                                             "specify 4 dimensions"));
47
48         OP_REQUIRES(
49             context, pooling_ratio_[0] == 1 || pooling_ratio_[3] == 1,
50             errors::Unimplemented("Fractional max pooling is not yet "
51                                   "supported on the batch nor channel dimension.));
52
53         OP_REQUIRES_OK(context, context->GetAttr("deterministic", &deterministic_));
54         OP_REQUIRES_OK(context, context->GetAttr("seed", &seed_));
55         OP_REQUIRES_OK(context, context->GetAttr("seed2", &seed2_));
56         if (deterministic_) {
57             // If both seeds are not set when deterministic_ is true, force set seeds.
58             if ((seed_ == 0) && (seed2_ == 0)) {
59                 seed_ = random::New64();
60                 seed2_ = random::New64();
61             }
62         } else {
63             OP_REQUIRES(
64                 context, (seed_ == 0) && (seed2_ == 0),
65                 errors::InvalidArgument(
66                     "Both seed and seed2 should be 0 if deterministic is false.));
67         }
68     }
69
70     void Compute(OpKernelContext* context) override {
71         typedef Eigen::Map<const Eigen::Matrix<T, Eigen::Dynamic, Eigen::Dynamic>>
72             ConstEigenMatrixMap;
73         typedef Eigen::Map<Eigen::Matrix<T, Eigen::Dynamic, Eigen::Dynamic>>
74             EigenMatrixMap;
75
76         constexpr int tensor_in_and_out_dims = 4;
77
78         const Tensor& tensor_in = context->input(0);

```

```

79     OP_REQUIRES(context, tensor_in.dims() == tensor_in_and_out_dims,
80                 errors::InvalidArgument("tensor_in must be 4-dimensional"));
81
82     std::vector<int> input_size(tensor_in_and_out_dims);
83     std::vector<int> output_size(tensor_in_and_out_dims);
84     for (int i = 0; i < tensor_in_and_out_dims; ++i) {
85         input_size[i] = tensor_in.dim_size(i);
86     }
87     // Output size.
88     for (int i = 0; i < tensor_in_and_out_dims; ++i) {
89         // This must match the same logic in the shape function in
90         // core/ops/nn_ops.cc.
91         output_size[i] =
92             static_cast<int>(std::floor(input_size[i] / pooling_ratio_[i]));
93         DCHECK_GT(output_size[i], 0);
94     }
95
96     // Generate pooling sequence.
97     std::vector<int64_t> height_cum_seq;
98     std::vector<int64_t> width_cum_seq;
99     GuardedPhiloxRandom generator;
100    generator.Init(seed_, seed2_);
101    height_cum_seq = GeneratePoolingSequence(input_size[1], output_size[1],
102                                            &generator, pseudo_random_);
103    width_cum_seq = GeneratePoolingSequence(input_size[2], output_size[2],
104                                           &generator, pseudo_random_);
105
106    // Prepare output.
107    Tensor* output_tensor = nullptr;
108    OP_REQUIRES_OK(context, context->allocate_output(
109        0,
110        TensorShape({output_size[0], output_size[1],
111                    output_size[2], output_size[3]}),
112        &output_tensor));
113    Tensor* output_height_seq_tensor = nullptr;
114    OP_REQUIRES_OK(
115        context,
116        context->allocate_output(
117            1, TensorShape({static_cast<int64_t>(height_cum_seq.size())}),
118            &output_height_seq_tensor));
119    Tensor* output_width_seq_tensor = nullptr;
120    OP_REQUIRES_OK(
121        context,
122        context->allocate_output(
123            2, TensorShape({static_cast<int64_t>(width_cum_seq.size())}),
124            &output_width_seq_tensor));
125
126    ConstEigenMatrixMap in_mat(tensor_in.flat<T>().data(), input_size[3],
127                               input_size[2] * input_size[1] * input_size[0]);

```

```

128
129     EigenMatrixMap out_mat(output_tensor->flat<T>().data(), output_size[3],
130                             output_size[2] * output_size[1] * output_size[0]);
131
132     // Initializes the output tensor with MIN<T>.
133     output_tensor->flat<T>().setConstant(Eigen::NumTraits<T>::lowest());
134
135     auto output_height_seq_flat = output_height_seq_tensor->flat<int64_t>();
136     auto output_width_seq_flat = output_width_seq_tensor->flat<int64_t>();
137
138     // Set output tensors.
139     for (int i = 0; i < height_cum_seq.size(); ++i) {
140         output_height_seq_flat(i) = height_cum_seq[i];
141     }
142
143     for (int i = 0; i < width_cum_seq.size(); ++i) {
144         output_width_seq_flat(i) = width_cum_seq[i];
145     }
146
147     // For both input and output,
148     // 0: batch
149     // 1: height / row
150     // 2: width / col
151     // 3: depth / channel
152     const int64_t height_max = input_size[1] - 1;
153     const int64_t width_max = input_size[2] - 1;
154     for (int64_t b = 0; b < input_size[0]; ++b) {
155         // height sequence.
156         for (int64_t hs = 0; hs < height_cum_seq.size() - 1; ++hs) {
157             // height start and end.
158             const int64_t height_start = height_cum_seq[hs];
159             int64_t height_end =
160                 overlapping_ ? height_cum_seq[hs + 1] : height_cum_seq[hs + 1] - 1;
161             height_end = std::min(height_end, height_max);
162
163             // width sequence.
164             for (int64_t ws = 0; ws < width_cum_seq.size() - 1; ++ws) {
165                 const int64_t out_offset =
166                     (b * output_size[1] + hs) * output_size[2] + ws;
167                 // width start and end.
168                 const int64_t width_start = width_cum_seq[ws];
169                 int64_t width_end =
170                     overlapping_ ? width_cum_seq[ws + 1] : width_cum_seq[ws + 1] - 1;
171                 width_end = std::min(width_end, width_max);
172                 for (int64_t h = height_start; h <= height_end; ++h) {
173                     for (int64_t w = width_start; w <= width_end; ++w) {
174                         const int64_t in_offset =
175                             (b * input_size[1] + h) * input_size[2] + w;
176                         out_mat.col(out_offset) =

```

```

177         out_mat.col(out_offset).cwiseMax(in_mat.col(in_offset));
178     }
179 }
180 }
181 }
182 }
183 }
184
185 private:
186     bool deterministic_;
187     int64_t seed_;
188     int64_t seed2_;
189     std::vector<float> pooling_ratio_;
190     bool pseudo_random_;
191     bool overlapping_;
192 };
193
194 #define REGISTER_FRACTIONALMAXPOOL(type) \
195     REGISTER_KERNEL_BUILDER( \
196         Name("FractionalMaxPool").Device(DEVICE_CPU).TypeConstraint<type>("T"), \
197         FractionalMaxPoolOp<type>)
198
199 REGISTER_FRACTIONALMAXPOOL(int32);
200 REGISTER_FRACTIONALMAXPOOL(int64_t);
201 REGISTER_FRACTIONALMAXPOOL(float);
202 REGISTER_FRACTIONALMAXPOOL(double);
203
204 #undef REGISTER_FRACTIONALMAXPOOL
205
206 static const int kInvalidMaxPoolingIndex = -1;
207
208 template <class T>
209 class FractionalMaxPoolGradOp : public OpKernel {
210 public:
211     explicit FractionalMaxPoolGradOp(OpKernelConstruction* context)
212         : OpKernel(context) {
213         OP_REQUIRES_OK(context, context->GetAttr("overlapping", &overlapping_));
214     }
215
216     void Compute(OpKernelContext* context) override {
217         // There are two steps when calculating gradient for FractionalMaxPool.
218         // 1) Walk through the process of calculating fractional pooling given
219         //    pooling region; however, in the process, keep track of where the max
220         //    element comes from. (arg_max)
221         // 2) Populate the value of out_backprop to where arg_max indicates. If
222         //    we support overlapping, it is likely to have multiple out_backprop[i]
223         //    propagates back to the same arg_max value.
224         typedef Eigen::Map<const Eigen::Matrix<T, Eigen::Dynamic, Eigen::Dynamic>>
225             ConstEigenMatrixMap;

```

```

226     typedef Eigen::Map<Eigen::Matrix<T, Eigen::Dynamic, Eigen::Dynamic>>
227         EigenMatrixMap;
228     typedef Eigen::Map<Eigen::Matrix<int64, Eigen::Dynamic, Eigen::Dynamic>>
229         EigenIndexMatrixMap;
230
231     const Tensor& tensor_in = context->input(0);
232     const Tensor& tensor_out = context->input(1);
233     const Tensor& out_backprop = context->input(2);
234     const Tensor& height_seq_tensor = context->input(3);
235     const Tensor& width_seq_tensor = context->input(4);
236
237     // Just to make it similar to FractionalMaxPoolOp.
238     constexpr int tensor_in_and_out_dims = 4;
239     OP_REQUIRES(
240         context, tensor_in.dims() == tensor_in_and_out_dims,
241         errors::InvalidArgument("orig_input should be a tensor of rank 4, got ",
242                                 tensor_in.DebugString()));
243     OP_REQUIRES(context, tensor_in.NumElements() > 0,
244                 errors::InvalidArgument("orig_input must not be empty, got ",
245                                         tensor_in.DebugString()));
246     OP_REQUIRES(context, tensor_out.dims() == tensor_in_and_out_dims,
247                 errors::InvalidArgument(
248                     "orig_output should be a tensor of rank 4, got ",
249                     tensor_out.DebugString()));
250     OP_REQUIRES(context, tensor_out.NumElements() > 0,
251                 errors::InvalidArgument("orig_output must not be empty, got ",
252                                         tensor_out.DebugString()));
253     std::vector<int64_t> input_size(tensor_in_and_out_dims);
254     std::vector<int64_t> output_size(tensor_in_and_out_dims);
255     for (int i = 0; i < tensor_in_and_out_dims; ++i) {
256         input_size[i] = tensor_in.dim_size(i);
257     }
258     for (int i = 0; i < tensor_in_and_out_dims; ++i) {
259         output_size[i] = tensor_out.dim_size(i);
260     }
261
262     // -----
263     // Step 1
264     // -----
265     Tensor tensor_out_dup;
266     OP_REQUIRES_OK(context, context->forward_input_or_allocate_temp(
267         {1}, DataTypeToEnum<T>::v(), tensor_out.shape(),
268         &tensor_out_dup));
269     Tensor tensor_out_arg_max;
270     OP_REQUIRES_OK(context, context->allocate_temp(DataTypeToEnum<int64_t>::v(),
271         tensor_out.shape(),
272         &tensor_out_arg_max));
273     // Find arg_max for each tensor_out
274     ConstEigenMatrixMap tensor_in_mat(

```

```

275     tensor_in.flat<T>().data(), input_size[3],
276     input_size[2] * input_size[1] * input_size[0]);
277 EigenMatrixMap tensor_out_dup_mat(
278     tensor_out_dup.flat<T>().data(), output_size[3],
279     output_size[2] * output_size[1] * output_size[0]);
280 EigenIndexMatrixMap tensor_out_arg_max_mat(
281     tensor_out_arg_max.flat<int64_t>().data(), output_size[3],
282     output_size[2] * output_size[1] * output_size[0]);
283
284 tensor_out_arg_max.flat<int64_t>().setConstant(kInvalidMaxPoolingIndex);
285 // Initializes the duplicate output tensor with MIN<T>.
286 tensor_out_dup.flat<T>().setConstant(Eigen::NumTraits<T>::lowest());
287
288 auto height_seq_tensor_flat = height_seq_tensor.flat<int64_t>();
289 auto width_seq_tensor_flat = width_seq_tensor.flat<int64_t>();
290
291 // Now walk through the process of fractional max pooling again.
292 // For both input and output,
293 // 0: batch
294 // 1: height / row
295 // 2: width / col
296 // 3: depth / channel
297 const int64_t height_max = input_size[1] - 1;
298 const int64_t width_max = input_size[2] - 1;
299 for (int64_t b = 0; b < input_size[0]; ++b) {
300     // height sequence.
301     for (int64_t hs = 0; hs < height_seq_tensor.dim_size(0) - 1; ++hs) {
302         // height start and end.
303         const int64_t height_start = height_seq_tensor_flat(hs);
304         int64_t height_end = overlapping_ ? height_seq_tensor_flat(hs + 1)
305                                         : height_seq_tensor_flat(hs + 1) - 1;
306         height_end = std::min(height_end, height_max);
307
308         // width sequence.
309         for (int64_t ws = 0; ws < width_seq_tensor.dim_size(0) - 1; ++ws) {
310             const int64_t out_index =
311                 (b * output_size[1] + hs) * output_size[2] + ws;
312             // width start and end.
313             const int64_t width_start = width_seq_tensor_flat(ws);
314             int64_t width_end = overlapping_ ? width_seq_tensor_flat(ws + 1)
315                                             : width_seq_tensor_flat(ws + 1) - 1;
316             width_end = std::min(width_end, width_max);
317             for (int64_t h = height_start; h <= height_end; ++h) {
318                 for (int64_t w = width_start; w <= width_end; ++w) {
319                     const int64_t in_index =
320                         (b * input_size[1] + h) * input_size[2] + w;
321                     // Walk through each channel (depth).
322                     for (int64_t d = 0; d < input_size[3]; ++d) {
323                         const T& input_ref = tensor_in_mat.coeffRef(d, in_index);

```

```

324         T& output_ref = tensor_out_dup_mat.coeffRef(d, out_index);
325         int64_t& out_arg_max_ref =
326             tensor_out_arg_max_mat.coeffRef(d, out_index);
327         if (output_ref < input_ref ||
328             out_arg_max_ref == kInvalidMaxPoolingIndex) {
329             output_ref = input_ref;
330             int input_offset = in_index * input_size[3] + d;
331             out_arg_max_ref = input_offset;
332         }
333     }
334 }
335 }
336 }
337 }
338 }
339
340 // Check tensor_out_dup is the same as tensor_out.
341 ConstEigenMatrixMap tensor_out_mat(
342     tensor_out.flat<T>().data(), output_size[3],
343     output_size[2] * output_size[1] * output_size[0]);
344 const int64_t num_resaped_cols =
345     output_size[2] * output_size[1] * output_size[0];
346 for (int64_t i = 0; i < num_resaped_cols; ++i) {
347     for (int64_t j = 0; j < output_size[3]; ++j) {
348         DCHECK_EQ(tensor_out_dup_mat(j, i), tensor_out_mat(j, i));
349     }
350 }
351
352 Tensor* output = nullptr;
353 OP_REQUIRES_OK(context, context->forward_input_or_allocate_output(
354     {0}, 0, tensor_in.shape(), &output));
355 output->flat<T>().setZero();
356
357 auto out_backprop_flat = out_backprop.flat<T>();
358 auto input_backprop_flat = output->flat<T>();
359 auto out_arg_max_flat = tensor_out_arg_max.flat<int64_t>();
360 int num_total_outputs = out_backprop_flat.size();
361 int num_total_inputs = input_backprop_flat.size();
362
363 for (int index = 0; index < num_total_outputs; ++index) {
364     int input_backprop_index = out_arg_max_flat(index);
365     // According to maxpooling_op.cc, the performance impact below is small.
366     CHECK(input_backprop_index >= 0 &&
367         input_backprop_index < num_total_inputs)
368         << "Invalid input backprop index: " << input_backprop_index << ", "
369         << num_total_inputs;
370     input_backprop_flat(input_backprop_index) += out_backprop_flat(index);
371 }
372 }

```



```

373
374     private:
375         bool overlapping_;
376 };
377
378 #define REGISTER_FRACTIONALMAXPOOLGRAD(type) \
379     REGISTER_KERNEL_BUILDER(Name("FractionalMaxPoolGrad") \
380                             .Device(DEVICE_CPU) \
381                             .TypeConstraint<type>("T"), \
382                             FractionalMaxPoolGradOp<type>)
383
384 REGISTER_FRACTIONALMAXPOOLGRAD(int32);
385 REGISTER_FRACTIONALMAXPOOLGRAD(int64_t);
386 REGISTER_FRACTIONALMAXPOOLGRAD(float);
387 REGISTER_FRACTIONALMAXPOOLGRAD(double);
388
389 #undef REGISTER_FRACTIONALMAXPOOLGRAD
390 } // namespace tensorflow

```