

5100e359ae ▾

...


tensorflow / tensorflow / core / kernels / data / sparse_tensor_slice_dataset_op.cc



tensorflow-gardener [tf.data] Change Cardinality() implementation to read c... ... ✖

History

8 contributors



305 lines (266 sloc) | 11.8 KB

...

```

1  /* Copyright 2017 The TensorFlow Authors. All Rights Reserved.
2
3  Licensed under the Apache License, Version 2.0 (the "License");
4  you may not use this file except in compliance with the License.
5  You may obtain a copy of the License at
6
7      http://www.apache.org/licenses/LICENSE-2.0
8
9  Unless required by applicable law or agreed to in writing, software
10 distributed under the License is distributed on an "AS IS" BASIS,
11 WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
12 See the License for the specific language governing permissions and
13 limitations under the License.
14 =====*/
15 #include <numeric>
16
17 #include "tensorflow/core/framework/dataset.h"
18 #include "tensorflow/core/framework/partial_tensor_shape.h"
19 #include "tensorflow/core/framework/register_types.h"
20 #include "tensorflow/core/framework/tensor.h"
21 #include "tensorflow/core/util/sparse/sparse_tensor.h"
22
23 namespace tensorflow {
24 namespace data {
25 namespace {
26
27 // See documentation in ../../ops/dataset_ops.cc for a high-level
28 // description of the following op.
29

```

```

30 template <typename T>
31 class Dataset : public DatasetBase {
32 public:
33     explicit Dataset(OpKernelContext* ctx,
34                     const sparse::SparseTensor& sparse_tensor)
35         : DatasetBase(DatasetContext(ctx)),
36           sparse_tensor_(sparse_tensor),
37           dtypes_({DT_INT64, sparse_tensor.dtype(), DT_INT64}),
38           shapes_({{-1, sparse_tensor.dims() - 1},
39                  {-1},
40                  {sparse_tensor.dims() - 1}}) {}
41
42     std::unique_ptr<IteratorBase> MakeIteratorInternal(
43         const string& prefix) const override {
44         return absl::make_unique<Iterator><typename Iterator::Params{
45             this, strings::StrCat(prefix, "::SparseTensorSlice")}>();
46     }
47
48     const DataTypeVector& output_dtypes() const override { return dtypes_; }
49     const std::vector<PartialTensorShape>& output_shapes() const override {
50         return shapes_;
51     }
52
53     string DebugString() const override {
54         return "SparseTensorSliceDatasetOp::Dataset";
55     }
56
57     int64_t CardinalityInternal() const override {
58         return sparse_tensor_.shape()[0];
59     }
60
61     Status InputDatasets(std::vector<const DatasetBase*>* inputs) const override {
62         return Status::OK();
63     }
64
65     Status CheckExternalState() const override { return Status::OK(); }
66
67 protected:
68     Status AsGraphDefInternal(SerializationContext* ctx,
69                             DatasetGraphDefBuilder* b,
70                             Node** output) const override {
71         Node* indices_node;
72         TF_RETURN_IF_ERROR(b->AddTensor(sparse_tensor_.indices(), &indices_node));
73         Node* value_node;
74         TF_RETURN_IF_ERROR(b->AddTensor(sparse_tensor_.values(), &value_node));
75         Node* dense_shape_node;
76         std::vector<int64_t> dense_shape;
77         dense_shape.reserve(sparse_tensor_.shape().size());
78         for (int i = 0; i < sparse_tensor_.shape().size(); i++)

```

```

79     dense_shape.emplace_back(sparse_tensor_.shape()[i]);
80     TF_RETURN_IF_ERROR(b->AddVector(dense_shape, &dense_shape_node));
81     AttrValue val_dtype;
82     b->BuildAttrValue(sparse_tensor_.dtype(), &val_dtype);
83     TF_RETURN_IF_ERROR(
84         b->AddDataset(this, {indices_node, value_node, dense_shape_node},
85             {"Tvalues", val_dtype}, output));
86     return Status::OK();
87 }
88
89 private:
90     class Iterator : public DatasetIterator<Dataset<T>> {
91     public:
92         explicit Iterator(const typename Iterator::Params& params)
93             : DatasetIterator<Dataset<T>>(params),
94             num_elements_(params.dataset->sparse_tensor_.shape()[0]),
95             dense_shape_(DT_INT64, {params.dataset->sparse_tensor_.dims() - 1}),
96             group_iterable_(params.dataset->sparse_tensor_.group({0})),
97             iter_(group_iterable_.begin()) {
98             for (size_t i = 0; i < dense_shape_.NumElements(); ++i) {
99                 dense_shape_.vec<int64_t>()(i) =
100                     params.dataset->sparse_tensor_.shape()[i + 1];
101             }
102         }
103
104         Status GetNextInternal(IteratorContext* ctx,
105             std::vector<Tensor>* out_tensors,
106             bool* end_of_sequence) override {
107             mutex_lock l(mu_);
108             if (i_ == num_elements_) {
109                 *end_of_sequence = true;
110                 return Status::OK();
111             }
112
113             out_tensors->clear();
114             out_tensors->reserve(3);
115             const int rank = Iterator::dataset()->sparse_tensor_.dims();
116
117             if (i_ > next_non_empty_i_ && iter_ != group_iterable_.end()) {
118                 // We still have elements to consume from `group_iterable`
119                 // and we have emitted all elements up to and including the
120                 // current position.
121                 sparse::Group group = *iter_;
122                 const auto indices = group.indices();
123                 const auto values = group.values<T>();
124                 const int64_t num_entries = values.size();
125                 next_non_empty_i_ = indices(0, 0);
126
127                 next_indices_ = Tensor(DT_INT64, {num_entries, rank - 1});

```

```

128     next_values_ = Tensor(DataTypeToEnum<T>::value, {num_entries});
129
130     auto next_indices_t = next_indices_.matrix<int64_t>();
131     auto next_values_t = next_values_.vec<T>();
132
133     for (int64_t i = 0; i < num_entries; ++i) {
134         for (int d = 1; d < rank; ++d) {
135             next_indices_t(i, d - 1) = indices(i, d);
136         }
137         next_values_t(i) = values(i);
138     }
139
140     ++iter_;
141 }
142 if (i_ == next_non_empty_i_) {
143     // The current position is non-empty in the input
144     // `SparseTensor`, and we have already read the value from the
145     // `GroupIterable`.
146     out_tensors->push_back(std::move(next_indices_));
147     out_tensors->push_back(std::move(next_values_));
148     out_tensors->push_back(dense_shape_);
149     next_non_empty_i_ = kNextNonEmptyUnknown;
150 } else {
151     DCHECK(i_ < next_non_empty_i_ || iter_ == group_iterable_.end());
152     // The current position is empty in the input `SparseTensor`,
153     // so emit empty indices and values.
154     out_tensors->push_back(Tensor(DT_INT64, TensorShape({0, rank - 1})));
155     out_tensors->push_back(Tensor(DataTypeToEnum<T>::value, {0}));
156     out_tensors->push_back(dense_shape_);
157 }
158
159 ++i_;
160 *end_of_sequence = false;
161 return Status::OK();
162 }
163
164 protected:
165     std::shared_ptr<model::Node> CreateNode(
166         IteratorContext* ctx, model::Node::Args args) const override {
167         return model::MakeSourceNode(std::move(args));
168     }
169
170     Status SaveInternal(SerializationContext* ctx,
171         IteratorStateWriter* writer) override {
172         mutex_lock l(mu_);
173         TF_RETURN_IF_ERROR(writer->WriteScalar(Iterator::full_name("i"), i_));
174         TF_RETURN_IF_ERROR(
175             writer->WriteScalar(Iterator::full_name("iter_loc"), iter_.loc()));
176         TF_RETURN_IF_ERROR(writer->WriteScalar(

```

```

177         Iterator::full_name("next_non_empty_i_"), next_non_empty_i_));
178     if (i_ <= next_non_empty_i_) {
179         TF_RETURN_IF_ERROR(writer->WriteTensor(
180             Iterator::full_name("next_indices_"), next_indices_));
181         TF_RETURN_IF_ERROR(writer->WriteTensor(
182             Iterator::full_name("next_values_"), next_values_));
183     }
184     return Status::OK();
185 }
186
187 Status RestoreInternal(IteratorContext* ctx,
188                       IteratorStateReader* reader) override {
189     mutex_lock l(mu_);
190     TF_RETURN_IF_ERROR(reader->ReadScalar(Iterator::full_name("i"), &i_));
191     int64_t iter_loc;
192     TF_RETURN_IF_ERROR(
193         reader->ReadScalar(Iterator::full_name("iter_loc"), &iter_loc));
194     iter_ = group_iterable_.at(iter_loc);
195     TF_RETURN_IF_ERROR(reader->ReadScalar(
196         Iterator::full_name("next_non_empty_i_"), &next_non_empty_i_));
197     if (i_ <= next_non_empty_i_) {
198         TF_RETURN_IF_ERROR(reader->ReadTensor(
199             Iterator::full_name("next_indices_"), &next_indices_));
200         TF_RETURN_IF_ERROR(reader->ReadTensor(
201             Iterator::full_name("next_values_"), &next_values_));
202     }
203     return Status::OK();
204 }
205
206 private:
207     const int64_t num_elements_;
208
209     Tensor dense_shape_;
210
211     mutex mu_;
212     sparse::GroupIterable group_iterable_ TF_GUARDED_BY(mu_);
213     sparse::GroupIterable::IteratorStep iter_ TF_GUARDED_BY(mu_);
214     int64_t i_ TF_GUARDED_BY(mu_) = 0;
215     const int64_t kNextNonEmptyUnknown = -1;
216     int64_t next_non_empty_i_ TF_GUARDED_BY(mu_) = kNextNonEmptyUnknown;
217     Tensor next_indices_ TF_GUARDED_BY(mu_);
218     Tensor next_values_ TF_GUARDED_BY(mu_);
219 };
220
221 const sparse::SparseTensor sparse_tensor_;
222 const DataTypeVector dtypes_;
223 const std::vector<PartialTensorShape> shapes_;
224 };
225

```

```

226 template <typename T>
227 class SparseTensorSliceDatasetOp : public DatasetOpKernel {
228 public:
229     explicit SparseTensorSliceDatasetOp(OpKernelConstruction* ctx)
230         : DatasetOpKernel(ctx) {}
231
232     void MakeDataset(OpKernelContext* ctx, DatasetBase** output) override {
233         // Create a new SparseTensorSliceDatasetOp::Dataset, insert it in
234         // the step container, and return it as the output.
235         const Tensor* indices;
236         OP_REQUIRES_OK(ctx, ctx->input("indices", &indices));
237         const Tensor* values;
238         OP_REQUIRES_OK(ctx, ctx->input("values", &values));
239         const Tensor* dense_shape;
240         OP_REQUIRES_OK(ctx, ctx->input("dense_shape", &dense_shape));
241
242         OP_REQUIRES(ctx, TensorShapeUtils::IsMatrix(indices->shape()),
243             errors::InvalidArgument(
244                 "Input indices should be a matrix but received shape ",
245                 indices->shape().DebugString()));
246
247         const auto num_indices = indices->NumElements();
248         const auto num_values = values->NumElements();
249         if (num_indices == 0 || num_values == 0) {
250             OP_REQUIRES(ctx, num_indices == num_values,
251                 errors::InvalidArgument(
252                     "If indices or values are empty, the other one must also "
253                     "be. Got indices of shape ",
254                     indices->shape().DebugString(), " and values of shape ",
255                     values->shape().DebugString()));
256         }
257         OP_REQUIRES(ctx, TensorShapeUtils::IsVector(values->shape()),
258             errors::InvalidArgument(
259                 "Input values should be a vector but received shape ",
260                 indices->shape().DebugString()));
261         OP_REQUIRES(ctx, TensorShapeUtils::IsVector(dense_shape->shape()),
262             errors::InvalidArgument(
263                 "Input shape should be a vector but received shape ",
264                 dense_shape->shape().DebugString()));
265
266         // We currently ensure that `sparse_tensor` is ordered in the
267         // batch dimension.
268         // TODO(mrry): Investigate ways to avoid this unconditional check
269         // if we can be sure that the sparse tensor was produced in an
270         // appropriate order (e.g. by `tf.parse_example()` or a Dataset
271         // that batches elements into rows of a SparseTensor).
272         int64_t previous_batch_index = -1;
273         for (int64_t i = 0; i < indices->dim_size(0); ++i) {
274             int64_t next_batch_index = indices->matrix<int64_t>()(i, 0);

```

```

275     OP_REQUIRES(
276         ctx, next_batch_index >= previous_batch_index,
277         errors::Unimplemented("The SparseTensor must be ordered in the batch "
278             "dimension; handling arbitrarily ordered input "
279             "is not currently supported."));
280     previous_batch_index = next_batch_index;
281 }
282 gtl::InlinedVector<int64_t, 8> std_order(dense_shape->NumElements(), 0);
283 sparse::SparseTensor tensor;
284 OP_REQUIRES_OK(
285     ctx, sparse::SparseTensor::Create(
286         *indices, *values, TensorShape(dense_shape->vec<int64_t>()),
287         std_order, &tensor));
288     *output = new Dataset<T>(ctx, std::move(tensor));
289 }
290
291 private:
292 };
293
294 #define REGISTER_DATASET_KERNEL(type) \
295     REGISTER_KERNEL_BUILDER(Name("SparseTensorSliceDataset") \
296         .Device(DEVICE_CPU) \
297         .TypeConstraint<type>("Tvalues"), \
298         SparseTensorSliceDatasetOp<type>);
299
300 TF_CALL_DATASET_TYPES(REGISTER_DATASET_KERNEL);
301 #undef REGISTER_DATASET_KERNEL
302
303 } // namespace
304 } // namespace data
305 } // namespace tensorflow

```