

5100e359ae ▾

...


tensorflow / tensorflow / core / kernels / quantized_pooling_ops.cc



jpienaar Rename to underlying type rather than alias ... ✖

History

6 contributors



150 lines (125 sloc) | 5.92 KB

...

```

1  /* Copyright 2015 The TensorFlow Authors. All Rights Reserved.
2
3  Licensed under the Apache License, Version 2.0 (the "License");
4  you may not use this file except in compliance with the License.
5  You may obtain a copy of the License at
6
7      http://www.apache.org/licenses/LICENSE-2.0
8
9  Unless required by applicable law or agreed to in writing, software
10 distributed under the License is distributed on an "AS IS" BASIS,
11 WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
12 See the License for the specific language governing permissions and
13 limitations under the License.
14 =====*/
15
16 // See docs in ../ops/nn_ops.cc.
17
18 #define EIGEN_USE_THREADS
19
20 #include "third_party/eigen3/unsupported/Eigen/CXX11/Tensor"
21 #include "tensorflow/core/framework/numeric_op.h"
22 #include "tensorflow/core/framework/op_kernel.h"
23 #include "tensorflow/core/framework/tensor.h"
24 #include "tensorflow/core/framework/tensor_shape.h"
25 #include "tensorflow/core/kernels/ops_util.h"
26 #include "tensorflow/core/kernels/pooling_ops_common.h"
27 #include "tensorflow/core/lib/core/errors.h"
28 #include "tensorflow/core/platform/logging.h"
29 #include "tensorflow/core/util/padding.h"

```

```

30 #include "tensorflow/core/util/tensor_format.h"
31
32 namespace tensorflow {
33
34 typedef Eigen::ThreadPoolDevice CPUDevice;
35
36 template <typename Device, typename T>
37 class QuantizedAvgPoolingOp : public OpKernel {
38 public:
39     explicit QuantizedAvgPoolingOp(OpKernelConstruction* context)
40         : OpKernel(context) {
41         OP_REQUIRES_OK(context, context->GetAttr("ksize", &ksize_));
42         OP_REQUIRES(context, ksize_.size() == 4,
43             errors::InvalidArgument("Sliding window ksize field must "
44                                     "specify 4 dimensions"));
45         OP_REQUIRES_OK(context, context->GetAttr("strides", &stride_));
46         OP_REQUIRES(context, stride_.size() == 4,
47             errors::InvalidArgument("Sliding window strides field must "
48                                     "specify 4 dimensions"));
49         OP_REQUIRES_OK(context, context->GetAttr("padding", &padding_));
50         OP_REQUIRES(context, ksize_[0] == 1 && stride_[0] == 1,
51             errors::Unimplemented(
52                 "Pooling is not yet supported on the batch dimension.));
53     }
54
55     void Compute(OpKernelContext* context) override {
56         const Tensor& tensor_in = context->input(0);
57         PoolParameters params{context,
58                               ksize_,
59                               stride_,
60                               padding_,
61                               /*explicit_paddings=*/{},
62                               FORMAT_NHWC,
63                               tensor_in.shape()};
64         if (!context->status().ok()) {
65             return;
66         }
67
68         const float min_input = context->input(1).flat<float>()(0);
69         const float max_input = context->input(2).flat<float>()(0);
70
71         OP_REQUIRES(context, params.depth_window == 1,
72             errors::Unimplemented("Non-spatial pooling is not "
73                                     "yet supported. Volunteers? :));
74
75         OP_REQUIRES(context, tensor_in.dims() == 4,
76             errors::InvalidArgument("tensor_in must be 4-dimensional));
77
78         Tensor* output = nullptr;

```

```

79     OP_REQUIRES_OK(context, context->allocate_output(
80         0, params.forward_output_shape(), &output));
81     const int32_t highest = static_cast<int32>(Eigen::NumTraits<T>::highest());
82     const int32_t lowest = static_cast<int32>(Eigen::NumTraits<T>::lowest());
83
84     // TODO(vrv): Switch this to the Eigen::Tensor version of
85     // SpatialAvgPooling once that version is running quickly.
86     Tensor int32_output(DT_INT32, params.forward_output_shape());
87     // Cast input to int32 tensor and call SpatialAvgPool.
88     Tensor int32_input(DT_INT32, tensor_in.shape());
89     int32_input.flat<int32>() = tensor_in.flat<T>().template cast<int32>();
90     SpatialAvgPool<Device, int32>(context, &int32_output, int32_input, params,
91         padding_);
92
93     // Clamp the int32 output back into quantized space.
94     output->flat<T>() = int32_output.flat<int32>()
95         .cwiseMax(lowest)
96         .cwiseMin(highest)
97         .template cast<T>();
98
99     Tensor* output_min = nullptr;
100    OP_REQUIRES_OK(context, context->allocate_output(1, {}, &output_min));
101    output_min->flat<float>()(0) = min_input;
102    Tensor* output_max = nullptr;
103    OP_REQUIRES_OK(context, context->allocate_output(2, {}, &output_max));
104    output_max->flat<float>()(0) = max_input;
105 }
106
107 private:
108     std::vector<int32> ksize_;
109     std::vector<int32> stride_;
110     Padding padding_;
111 };
112
113 template <typename Device, typename T>
114 class QuantizedMaxPoolingOp : public MaxPoolingOp<Device, T> {
115 public:
116     explicit QuantizedMaxPoolingOp(OpKernelConstruction* context)
117         : MaxPoolingOp<Device, T>(context) {}
118
119     void Compute(OpKernelContext* context) override {
120         const float min_input = context->input(1).flat<float>()(0);
121         const float max_input = context->input(2).flat<float>()(0);
122         MaxPoolingOp<Device, T>::Compute(context);
123         Tensor* output_min = nullptr;
124         OP_REQUIRES_OK(context, context->allocate_output(1, {}, &output_min));
125         output_min->flat<float>()(0) = min_input;
126         Tensor* output_max = nullptr;
127         OP_REQUIRES_OK(context, context->allocate_output(2, {}, &output_max));

```

```
128     output_max->flat<float>()(0) = max_input;
129 }
130 };
131
132 REGISTER_KERNEL_BUILDER(
133     Name("QuantizedAvgPool").Device(DDEVICE_CPU).TypeConstraint<quint8>("T"),
134     QuantizedAvgPoolingOp<CPUDevice, quint8>);
135
136 REGISTER_KERNEL_BUILDER(
137     Name("QuantizedMaxPool").Device(DDEVICE_CPU).TypeConstraint<quint8>("T"),
138     QuantizedMaxPoolingOp<CPUDevice, quint8>);
139
140 #ifdef INTEL_MKL
141 REGISTER_KERNEL_BUILDER(
142     Name("QuantizedAvgPool").Device(DDEVICE_CPU).TypeConstraint<qint8>("T"),
143     QuantizedAvgPoolingOp<CPUDevice, qint8>);
144
145 REGISTER_KERNEL_BUILDER(
146     Name("QuantizedMaxPool").Device(DDEVICE_CPU).TypeConstraint<qint8>("T"),
147     QuantizedMaxPoolingOp<CPUDevice, qint8>);
148 #endif
149
150 } // namespace tensorflow
```