

5100e359ae ▾

...

tensorflow / tensorflow / lite / kernels / internal / common.h



tensorflow-gardener Merge pull request #45342 from Tessil:toupstream/tabl... ... ✓

History

23 contributors



1083 lines (974 sloc) | 42.8 KB

...

```

1  /* Copyright 2017 The TensorFlow Authors. All Rights Reserved.
2
3  Licensed under the Apache License, Version 2.0 (the "License");
4  you may not use this file except in compliance with the License.
5  You may obtain a copy of the License at
6
7      http://www.apache.org/licenses/LICENSE-2.0
8
9  Unless required by applicable law or agreed to in writing, software
10 distributed under the License is distributed on an "AS IS" BASIS,
11 WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
12 See the License for the specific language governing permissions and
13 limitations under the License.
14 =====*/
15 #ifndef TENSORFLOW_LITE_KERNELS_INTERNAL_COMMON_H_
16 #define TENSORFLOW_LITE_KERNELS_INTERNAL_COMMON_H_
17
18 #ifndef ALLOW_SLOW_GENERIC_DEPTHWISECONV_FALLBACK
19 #ifndef GEMMLOWP_ALLOW_SLOW_SCALAR_FALLBACK
20 #define ALLOW_SLOW_GENERIC_DEPTHWISECONV_FALLBACK
21 #endif
22 #endif
23
24 #include <functional>
25
26 #include "fixedpoint/fixedpoint.h"
27 #include "tensorflow/lite/kernels/internal/cppmath.h"
28 #include "tensorflow/lite/kernels/internal/optimized/neon_check.h"
29 #include "tensorflow/lite/kernels/internal/types.h"

```

```

30
31 namespace tfLite {
32
33 constexpr int kReverseShift = -1;
34
35 inline void GetActivationMinMax(FusedActivationFunctionType ac,
36                               float* output_activation_min,
37                               float* output_activation_max) {
38     switch (ac) {
39         case FusedActivationFunctionType::kNone:
40             *output_activation_min = std::numeric_limits<float>::lowest();
41             *output_activation_max = std::numeric_limits<float>::max();
42             break;
43         case FusedActivationFunctionType::kRelu:
44             *output_activation_min = 0.f;
45             *output_activation_max = std::numeric_limits<float>::max();
46             break;
47         case FusedActivationFunctionType::kRelu1:
48             *output_activation_min = -1.f;
49             *output_activation_max = 1.f;
50             break;
51         case FusedActivationFunctionType::kRelu6:
52             *output_activation_min = 0.f;
53             *output_activation_max = 6.f;
54             break;
55     }
56 }
57
58 template <typename T>
59 inline T ActivationFunctionWithMinMax(T x, T output_activation_min,
60                                     T output_activation_max) {
61     using std::max;
62     using std::min;
63     return min(max(x, output_activation_min), output_activation_max);
64 }
65
66 // Legacy function, left for compatibility only.
67 template <FusedActivationFunctionType Ac>
68 float ActivationFunction(float x) {
69     float output_activation_min, output_activation_max;
70     GetActivationMinMax(Ac, &output_activation_min, &output_activation_max);
71     return ActivationFunctionWithMinMax(x, output_activation_min,
72                                       output_activation_max);
73 }
74
75 inline void BiasAndClamp(float clamp_min, float clamp_max, int bias_size,
76                         const float* bias_data, int array_size,
77                         float* array_data) {
78     // Note: see b/132215220: in May 2019 we thought it would be OK to replace

```

[illegible]

```

128     }
129 }
130 #else // not NEON
131     for (int array_offset = 0; array_offset < array_size;
132          array_offset += bias_size) {
133         for (int i = 0; i < bias_size; i++) {
134             array_data[array_offset + i] = ActivationFunctionWithMinMax(
135                 array_data[array_offset + i] + bias_data[i], clamp_min, clamp_max);
136         }
137     }
138 #endif
139 }
140
141 inline int32_t MultiplyByQuantizedMultiplierSmallerThanOneExp(
142     int32_t x, int32_t quantized_multiplier, int left_shift) {
143     using gemmlowp::RoundingDivideByPOT;
144     using gemmlowp::SaturatingRoundingDoublingHighMul;
145     return RoundingDivideByPOT(
146         SaturatingRoundingDoublingHighMul(x, quantized_multiplier), -left_shift);
147 }
148
149 inline int32_t MultiplyByQuantizedMultiplierGreaterThanOne(
150     int32_t x, int32_t quantized_multiplier, int left_shift) {
151     using gemmlowp::SaturatingRoundingDoublingHighMul;
152     return SaturatingRoundingDoublingHighMul(x * (1 << left_shift),
153         quantized_multiplier);
154 }
155
156 inline int32_t MultiplyByQuantizedMultiplier(int32_t x,
157     int32_t quantized_multiplier,
158     int shift) {
159     using gemmlowp::RoundingDivideByPOT;
160     using gemmlowp::SaturatingRoundingDoublingHighMul;
161     int left_shift = shift > 0 ? shift : 0;
162     int right_shift = shift > 0 ? 0 : -shift;
163     return RoundingDivideByPOT(SaturatingRoundingDoublingHighMul(
164         x * (1 << left_shift), quantized_multiplier),
165         right_shift);
166 }
167
168 inline int32_t MultiplyByQuantizedMultiplier(int64_t x,
169     int32_t quantized_multiplier,
170     int shift) {
171     // Inputs:
172     // - quantized_multiplier has fixed point at bit 31
173     // - shift is -31 to +7 (negative for right shift)
174     //
175     // Assumptions: The following input ranges are assumed
176     // - quantize_scale >= 0 (the usual range is (1 << 30) to (1 >> 31) - 1)

```

```

177 // - scaling is chosen so final scaled result fits in int32_t
178 // - input x is in the range -(1<<47) <= x < (1<<47)
179 assert(quantized_multiplier >= 0);
180 assert(shift >= -31 && shift < 8);
181 assert(x >= -(static_cast<int64_t>(1) << 47) &&
182        x < (static_cast<int64_t>(1) << 47));
183
184 int32_t reduced_multiplier = (quantized_multiplier < 0x7FFF0000)
185                             ? ((quantized_multiplier + (1 << 15)) >> 16)
186                             : 0x7FFF;
187 int total_shift = 15 - shift;
188 x = (x * (int64_t)reduced_multiplier) + ((int64_t)1 << (total_shift - 1));
189 int32_t result = x >> total_shift;
190 return result;
191 }
192
193 #ifdef USE_NEON
194 // Round uses ARM's rounding shift right.
195 inline int32x4x4_t MultiplyByQuantizedMultiplier4Rows(
196     int32x4x4_t input_val, int32_t quantized_multiplier, int shift) {
197     const int left_shift = std::max(shift, 0);
198     const int right_shift = std::min(shift, 0);
199     int32x4x4_t result;
200
201     int32x4_t multiplier_dup = vdupq_n_s32(quantized_multiplier);
202     int32x4_t left_shift_dup = vdupq_n_s32(left_shift);
203     int32x4_t right_shift_dup = vdupq_n_s32(right_shift);
204
205     result.val[0] =
206         vrshlq_s32(vqrdmulhq_s32(vshlq_s32(input_val.val[0], left_shift_dup),
207                                         multiplier_dup),
208                   right_shift_dup);
209
210     result.val[1] =
211         vrshlq_s32(vqrdmulhq_s32(vshlq_s32(input_val.val[1], left_shift_dup),
212                                         multiplier_dup),
213                   right_shift_dup);
214
215     result.val[2] =
216         vrshlq_s32(vqrdmulhq_s32(vshlq_s32(input_val.val[2], left_shift_dup),
217                                         multiplier_dup),
218                   right_shift_dup);
219
220     result.val[3] =
221         vrshlq_s32(vqrdmulhq_s32(vshlq_s32(input_val.val[3], left_shift_dup),
222                                         multiplier_dup),
223                   right_shift_dup);
224
225     return result;

```

```

226 }
227 #endif
228
229 template <typename T>
230 int CountLeadingZeros(T integer_input) {
231     static_assert(std::is_unsigned<T>::value,
232         "Only unsigned integer types handled.");
233     #if defined(__GNUC__)
234         return integer_input ? __builtin_clz(integer_input)
235             : std::numeric_limits<T>::digits;
236     #else
237         if (integer_input == 0) {
238             return std::numeric_limits<T>::digits;
239         }
240
241         const T one_in_leading_positive = static_cast<T>(1)
242             << (std::numeric_limits<T>::digits - 1);
243         int leading_zeros = 0;
244         while (integer_input < one_in_leading_positive) {
245             integer_input <= 1;
246             ++leading_zeros;
247         }
248         return leading_zeros;
249     #endif
250 }
251
252 template <typename T>
253 inline int CountLeadingSignBits(T integer_input) {
254     static_assert(std::is_signed<T>::value, "Only signed integer types handled.");
255     #if defined(__GNUC__) && !defined(__clang__)
256         return integer_input ? __builtin_clrsb(integer_input)
257             : std::numeric_limits<T>::digits;
258     #else
259         using U = typename std::make_unsigned<T>::type;
260         return integer_input >= 0
261             ? CountLeadingZeros(static_cast<U>(integer_input)) - 1
262             : integer_input != std::numeric_limits<T>::min()
263                 ? CountLeadingZeros(2 * static_cast<U>(-integer_input) - 1)
264                 : 0;
265     #endif
266 }
267
268 // Use "count leading zeros" helper functions to do a fast Floor(log2(x)).
269 template <typename Integer>
270 inline Integer FloorLog2(Integer n) {
271     static_assert(std::is_integral<Integer>::value, "");
272     static_assert(std::is_signed<Integer>::value, "");
273     static_assert(sizeof(Integer) == 4 || sizeof(Integer) == 8, "");
274     TFLITE_CHECK_GT(n, 0);

```

```

275     if (sizeof(Integer) == 4) {
276         return 30 - CountLeadingSignBits(n);
277     } else {
278         return 62 - CountLeadingSignBits(n);
279     }
280 }
281
282 // The size of the LUT depends on the type of input. For int8 inputs a simple
283 // 256 entries LUT is used. For int16 inputs the high 9 bits are used for
284 // indexing and the 7 remaining bits are used for interpolation. We thus use a
285 // 513-entries LUT for int16 cases, 512 for the 9-bit indexing and 1 extra entry
286 // to interpolate the last value.
287 template <typename LutInT>
288 constexpr int lut_size() {
289     static_assert(std::is_same<LutInT, int8_t>::value ||
290                 std::is_same<LutInT, int16_t>::value,
291                 "Only LUTs with int8 or int16 inputs are supported.");
292     return std::is_same<LutInT, int8_t>::value ? 256 : 513;
293 }
294
295 // Generate a LUT for 'func' which can be used to approximate functions like
296 // exp, log, ...
297 //
298 // - func: the function to build the LUT for (e.g exp(x))
299 // - input_min, input_max: range of the func inputs
300 // - output_min, output_max: range of the func outputs
301 // - lut: pointer to the LUT table to fill, the table must be of size
302 // lut_size<LutInT>()
303 template <typename FloatT, typename LutInT, typename LutOutT>
304 inline void gen_lut(FloatT (*func)(FloatT), FloatT input_min, FloatT input_max,
305                   FloatT output_min, FloatT output_max, LutOutT* lut) {
306     static_assert(std::is_same<LutInT, int8_t>::value ||
307                 std::is_same<LutInT, int16_t>::value,
308                 "Only LUTs with int8 or int16 inputs are supported.");
309     static_assert(std::is_same<LutOutT, int8_t>::value ||
310                 std::is_same<LutOutT, int16_t>::value,
311                 "Only LUTs with int8 or int16 outputs are supported.");
312     static_assert(std::is_floating_point<FloatT>::value,
313                 "FloatT must be a floating-point type.");
314
315     const int nb_steps = std::is_same<LutInT, int8_t>::value ? 256 : 512;
316     const FloatT step = (input_max - input_min) / nb_steps;
317     const FloatT half_step = step / 2;
318     const FloatT output_scaling_inv =
319         static_cast<FloatT>(std::numeric_limits<LutOutT>::max() -
320                             std::numeric_limits<LutOutT>::min() + 1) /
321         (output_max - output_min);
322     const FloatT table_min =
323         static_cast<FloatT>(std::numeric_limits<LutOutT>::min());

```

```

324     const FloatT table_max =
325         static_cast<FloatT>(std::numeric_limits<LutOutT>::max());
326
327     for (int i = 0; i < nb_steps; i++) {
328         const FloatT val = func(input_min + i * step);
329         const FloatT val_midpoint = func(input_min + i * step + half_step);
330         const FloatT val_next = func(input_min + (i + 1) * step);
331
332         const FloatT sample_val = TfliteRound(val * output_scaling_inv);
333         const FloatT midpoint_interp_val =
334             TfliteRound((val_next * output_scaling_inv +
335                 TfliteRound(val * output_scaling_inv)) /
336                 2);
337         const FloatT midpoint_val = TfliteRound(val_midpoint * output_scaling_inv);
338         const FloatT midpoint_err = midpoint_interp_val - midpoint_val;
339         const FloatT bias = TfliteRound(midpoint_err / 2);
340
341         lut[i] = static_cast<LutOutT>(std::min<FloatT>(
342             std::max<FloatT>(sample_val - bias, table_min), table_max));
343     }
344
345     const bool with_extra_interpolation_value =
346         std::is_same<LutInT, int16_t>::value;
347     if (with_extra_interpolation_value) {
348         lut[nb_steps] = static_cast<LutOutT>(std::min<FloatT>(
349             std::max<FloatT>(TfliteRound(func(input_max) * output_scaling_inv),
350                 table_min),
351             table_max));
352     }
353 }
354
355 // LUT must have 513 values
356 template <typename LutOutT>
357 inline LutOutT lut_lookup_with_interpolation(int16_t value,
358                                             const LutOutT* lut) {
359     static_assert(std::is_same<LutOutT, int8_t>::value ||
360         std::is_same<LutOutT, int16_t>::value,
361         "Only LUTs with int8 or int16 outputs are supported.");
362     // 512 base values, lut[513] is only used to calculate the slope
363     const uint16_t index = static_cast<uint16_t>(256 + (value >> 7));
364     assert(index < 512 && "LUT index out of range.");
365     const int16_t offset = value & 0x7f;
366
367     // Base and slope are Q0.x
368     const LutOutT base = lut[index];
369     const LutOutT slope = lut[index + 1] - lut[index];
370
371     // Q0.x * Q0.7 = Q0.(x + 7)
372     // Round and convert from Q0.(x + 7) to Q0.x

```



```

373     const int delta = (slope * offset + 64) >> 7;
374
375     // Q0.15 + Q0.15
376     return static_cast<LutOutT>(base + delta);
377 }
378
379 // int16_t -> int16_t table lookup with interpolation
380 // LUT must have 513 values
381 inline int16_t lut_lookup(int16_t value, const int16_t* lut) {
382     return lut_lookup_with_interpolation(value, lut);
383 }
384
385 // int16_t -> int8_t table lookup with interpolation
386 // LUT must have 513 values
387 inline int8_t lut_lookup(int16_t value, const int8_t* lut) {
388     return lut_lookup_with_interpolation(value, lut);
389 }
390
391 // int8_t -> int8_t table lookup without interpolation
392 // LUT must have 256 values
393 inline int8_t lut_lookup(int8_t value, const int8_t* lut) {
394     return lut[128 + value];
395 }
396
397 // int8_t -> int16_t table lookup without interpolation
398 // LUT must have 256 values
399 inline int16_t lut_lookup(int8_t value, const int16_t* lut) {
400     return lut[128 + value];
401 }
402
403 // Table of sigmoid(i/24) at 0.16 format - 256 elements.
404
405 // We use combined sigmoid and tanh look-up table, since
406 // tanh(x) = 2*sigmoid(2*x) -1.
407 // Both functions are symmetric, so the LUT table is only needed
408 // for the absolute value of the input.
409 static const uint16_t sigmoid_table_uint16[256] = {
410     32768, 33451, 34133, 34813, 35493, 36169, 36843, 37513, 38180, 38841, 39498,
411     40149, 40794, 41432, 42064, 42688, 43304, 43912, 44511, 45102, 45683, 46255,
412     46817, 47369, 47911, 48443, 48964, 49475, 49975, 50464, 50942, 51409, 51865,
413     52311, 52745, 53169, 53581, 53983, 54374, 54755, 55125, 55485, 55834, 56174,
414     56503, 56823, 57133, 57433, 57724, 58007, 58280, 58544, 58800, 59048, 59288,
415     59519, 59743, 59959, 60168, 60370, 60565, 60753, 60935, 61110, 61279, 61441,
416     61599, 61750, 61896, 62036, 62172, 62302, 62428, 62549, 62666, 62778, 62886,
417     62990, 63090, 63186, 63279, 63368, 63454, 63536, 63615, 63691, 63765, 63835,
418     63903, 63968, 64030, 64090, 64148, 64204, 64257, 64308, 64357, 64405, 64450,
419     64494, 64536, 64576, 64614, 64652, 64687, 64721, 64754, 64786, 64816, 64845,
420     64873, 64900, 64926, 64950, 64974, 64997, 65019, 65039, 65060, 65079, 65097,
421     65115, 65132, 65149, 65164, 65179, 65194, 65208, 65221, 65234, 65246, 65258,

```

```

422     65269, 65280, 65291, 65301, 65310, 65319, 65328, 65337, 65345, 65352, 65360,
423     65367, 65374, 65381, 65387, 65393, 65399, 65404, 65410, 65415, 65420, 65425,
424     65429, 65433, 65438, 65442, 65445, 65449, 65453, 65456, 65459, 65462, 65465,
425     65468, 65471, 65474, 65476, 65479, 65481, 65483, 65485, 65488, 65489, 65491,
426     65493, 65495, 65497, 65498, 65500, 65501, 65503, 65504, 65505, 65507, 65508,
427     65509, 65510, 65511, 65512, 65513, 65514, 65515, 65516, 65517, 65517, 65518,
428     65519, 65520, 65520, 65521, 65522, 65522, 65523, 65523, 65524, 65524, 65525,
429     65525, 65526, 65526, 65526, 65527, 65527, 65528, 65528, 65528, 65529, 65529,
430     65529, 65529, 65530, 65530, 65530, 65530, 65531, 65531, 65531, 65531, 65531,
431     65532, 65532, 65532, 65532, 65532, 65532, 65533, 65533, 65533, 65533, 65533,
432     65533, 65533, 65533, 65534, 65534, 65534, 65534, 65534, 65534, 65534,
433     65534, 65534, 65535};
434
435 // TODO(b/77858996): Add these to gemmlowp.
436 template <typename IntegerType>
437 IntegerType SaturatingAddNonGemmlowp(IntegerType a, IntegerType b) {
438     static_assert(std::is_same<IntegerType, void>::value, "unimplemented");
439     return a;
440 }
441
442 template <>
443 inline std::int32_t SaturatingAddNonGemmlowp(std::int32_t a, std::int32_t b) {
444     std::int64_t a64 = a;
445     std::int64_t b64 = b;
446     std::int64_t sum = a64 + b64;
447     return static_cast<std::int32_t>(std::min(
448         static_cast<std::int64_t>(std::numeric_limits<std::int32_t>::max()),
449         std::max(
450             static_cast<std::int64_t>(std::numeric_limits<std::int32_t>::min()),
451             sum)));
452 }
453
454 template <typename tRawType, int tIntegerBits>
455 gemmlowp::FixedPoint<tRawType, tIntegerBits> SaturatingAddNonGemmlowp(
456     gemmlowp::FixedPoint<tRawType, tIntegerBits> a,
457     gemmlowp::FixedPoint<tRawType, tIntegerBits> b) {
458     return gemmlowp::FixedPoint<tRawType, tIntegerBits>::FromRaw(
459         SaturatingAddNonGemmlowp(a.raw(), b.raw()));
460 }
461
462 template <typename IntegerType>
463 IntegerType SaturatingSub(IntegerType a, IntegerType b) {
464     static_assert(std::is_same<IntegerType, void>::value, "unimplemented");
465     return a;
466 }
467
468 template <>
469 inline std::int16_t SaturatingSub(std::int16_t a, std::int16_t b) {
470     std::int32_t a32 = a;

```

```

471     std::int32_t b32 = b;
472     std::int32_t diff = a32 - b32;
473     return static_cast<std::int16_t>(
474         std::min(static_cast<int32_t>(32767),
475             std::max(static_cast<int32_t>(-32768), diff)));
476 }
477
478 template <>
479 inline std::int32_t SaturatingSub(std::int32_t a, std::int32_t b) {
480     std::int64_t a64 = a;
481     std::int64_t b64 = b;
482     std::int64_t diff = a64 - b64;
483     return static_cast<std::int32_t>(std::min(
484         static_cast<std::int64_t>(std::numeric_limits<std::int32_t>::max()),
485         std::max(
486             static_cast<std::int64_t>(std::numeric_limits<std::int32_t>::min()),
487             diff)));
488 }
489
490 template <typename tRawType, int tIntegerBits>
491 gemmlowp::FixedPoint<tRawType, tIntegerBits> SaturatingSub(
492     gemmlowp::FixedPoint<tRawType, tIntegerBits> a,
493     gemmlowp::FixedPoint<tRawType, tIntegerBits> b) {
494     return gemmlowp::FixedPoint<tRawType, tIntegerBits>::FromRow(
495         SaturatingSub(a.raw(), b.raw()));
496 }
497 // End section to be moved to gemmlowp.
498
499 template <typename IntegerType>
500 IntegerType SaturatingRoundingMultiplyByPOTParam(IntegerType x, int exponent) {
501     if (exponent == 0) {
502         return x;
503     }
504     using ScalarIntegerType =
505         typename gemmlowp::FixedPointRawTypeTraits<IntegerType>::ScalarRawType;
506     const IntegerType min =
507         gemmlowp::Dup<IntegerType>(std::numeric_limits<ScalarIntegerType>::min());
508     const IntegerType max =
509         gemmlowp::Dup<IntegerType>(std::numeric_limits<ScalarIntegerType>::max());
510     const int ScalarIntegerTypeBits = 8 * sizeof(ScalarIntegerType);
511
512     const std::int32_t threshold =
513         ((1 << (ScalarIntegerTypeBits - 1 - exponent)) - 1);
514     const IntegerType positive_mask =
515         gemmlowp::MaskIfGreaterThan(x, gemmlowp::Dup<IntegerType>(threshold));
516     const IntegerType negative_mask =
517         gemmlowp::MaskIfLessThan(x, gemmlowp::Dup<IntegerType>(-threshold));
518
519     IntegerType result = gemmlowp::ShiftLeft(x, exponent);

```

```

520     result = gemmlowp::SelectUsingMask(positive_mask, max, result);
521     result = gemmlowp::SelectUsingMask(negative_mask, min, result);
522     return result;
523 }
524
525 // If we want to leave IntegerBits fixed, then multiplication
526 // by a power of two has to be saturating/rounding, not exact anymore.
527 template <typename tRawType, int tIntegerBits>
528 gemmlowp::FixedPoint<tRawType, tIntegerBits>
529 SaturatingRoundingMultiplyByPOTParam(
530     gemmlowp::FixedPoint<tRawType, tIntegerBits> a, int exponent) {
531     return gemmlowp::FixedPoint<tRawType, tIntegerBits>::FromRaw(
532         SaturatingRoundingMultiplyByPOTParam(a.raw(), exponent));
533 }
534
535 // Convert int32_t multiplier to int16_t with rounding.
536 inline void DownScaleInt32ToInt16Multiplier(int32_t multiplier_int32_t,
537                                             int16_t* multiplier_int16_t) {
538     TFLITE_DCHECK_GE(multiplier_int32_t, 0);
539     static constexpr int32_t kRoundingOffset = 1 << 15;
540     if (multiplier_int32_t >=
541         std::numeric_limits<int32_t>::max() - kRoundingOffset) {
542         *multiplier_int16_t = std::numeric_limits<int16_t>::max();
543         return;
544     }
545     const int32_t result = (multiplier_int32_t + kRoundingOffset) >> 16;
546     TFLITE_DCHECK_LE(result << 16, multiplier_int32_t + kRoundingOffset);
547     TFLITE_DCHECK_GT(result << 16, multiplier_int32_t - kRoundingOffset);
548     *multiplier_int16_t = result;
549     TFLITE_DCHECK_EQ(*multiplier_int16_t, result);
550 }
551
552 // Minimum output bits to accommodate log of maximum input range. It actually
553 // does not matter if one considers, say, [-64,64] or [-64,64).
554 //
555 // For example, run this through Octave:
556 // [0:127; ...
557 //   ceil(log(abs( log(2.^(0:127))+1 ))/log(2)); ...
558 //   ceil(log(abs( log(2.^(0:127))+1 ))/log(2))]
559 constexpr int min_log_x_output_bits(int input_bits) {
560     return input_bits > 90 ? 7
561         : input_bits > 44 ? 6
562         : input_bits > 21 ? 5
563         : input_bits > 10 ? 4
564         : input_bits > 4  ? 3
565         : input_bits > 1  ? 2
566         : 1;
567 }
568

```

```

569 // Although currently the name of this function says that it cannot handle
570 // values less than 1, in practice it can handle as low as 1/x_max, where
571 // x_max is the largest representable input. In other words, the output range
572 // is symmetric.
573 template <int OutputIntegerBits, int InputIntegerBits>
574 inline gemmlowp::FixedPoint<int32_t, OutputIntegerBits>
575 log_x_for_x_greater_than_or_equal_to_1_impl(
576     gemmlowp::FixedPoint<int32_t, InputIntegerBits> input_val) {
577     // assert(__builtin_clz(0u) >= std::numeric_limits<uint32_t>::digits - 1);
578     // assert(__builtin_clz(0u) <= std::numeric_limits<uint32_t>::digits);
579     using FixedPoint0 = gemmlowp::FixedPoint<int32_t, 0>;
580     // The reason for accumulating the result with an extra bit of headroom is
581     // that z_pow_2_adj * log_2 might be saturated, and adding num_scaled *
582     // recip_denom will otherwise introduce an error.
583     static constexpr int kAccumIntegerBits = OutputIntegerBits + 1;
584     using FixedPointAccum = gemmlowp::FixedPoint<int32_t, kAccumIntegerBits>;
585
586     const FixedPoint0 log_2 = GEMMLOWP_CHECKED_FIXEDPOINT_CONSTANT(
587         FixedPoint0, 1488522236, std::log(2.0));
588     const FixedPoint0 sqrt_sqrt_half = GEMMLOWP_CHECKED_FIXEDPOINT_CONSTANT(
589         FixedPoint0, 1805811301, std::sqrt(std::sqrt(0.5)));
590     const FixedPoint0 sqrt_half = GEMMLOWP_CHECKED_FIXEDPOINT_CONSTANT(
591         FixedPoint0, 1518500250, std::sqrt(0.5));
592     const FixedPoint0 one_quarter =
593         GEMMLOWP_CHECKED_FIXEDPOINT_CONSTANT(FixedPoint0, 536870912, 1.0 / 4.0);
594
595     const FixedPoint0 alpha_n = GEMMLOWP_CHECKED_FIXEDPOINT_CONSTANT(
596         FixedPoint0, 117049297, 11.0 / 240.0 * std::sqrt(std::sqrt(2.0)));
597     const FixedPoint0 alpha_d = GEMMLOWP_CHECKED_FIXEDPOINT_CONSTANT(
598         FixedPoint0, 127690142, 1.0 / 20.0 * std::sqrt(std::sqrt(2.0)));
599     const FixedPoint0 alpha_i = GEMMLOWP_CHECKED_FIXEDPOINT_CONSTANT(
600         FixedPoint0, 1057819769,
601         2.0 / std::sqrt(std::sqrt(2.0)) - std::sqrt(std::sqrt(2.0)));
602     const FixedPoint0 alpha_f = GEMMLOWP_CHECKED_FIXEDPOINT_CONSTANT(
603         FixedPoint0, 638450708, 1.0 / 4.0 * std::sqrt(std::sqrt(2.0)));
604
605     const FixedPointAccum shifted_quarter =
606         gemmlowp::Rescale<kAccumIntegerBits>(one_quarter);
607
608     // Reinterpret the input value as Q0.31, because we will figure out the
609     // required shift "ourselves" instead of using, say, Rescale.
610     FixedPoint0 z_a = FixedPoint0::FromRaw(input_val.raw());
611     // z_a_pow_2 = input_integer_bits - z_a_headroom;
612     int z_a_headroom_plus_1 = CountLeadingZeros(static_cast<uint32_t>(z_a.raw()));
613     FixedPoint0 r_a_tmp =
614         SaturatingRoundingMultiplyByPOTParam(z_a, (z_a_headroom_plus_1 - 1));
615     const int32_t r_a_raw =
616         SaturatingRoundingMultiplyByPOTParam((r_a_tmp * sqrt_half).raw(), 1);
617     // z_pow_2_adj = max(z_pow_2_a - 0.75, z_pow_2_b - 0.25);

```

```

618 // z_pow_2_adj = max(InputIntegerBits - z_a_headroom_plus_1 + 0.25,
619 //                    InputIntegerBits - z_b_headroom - 0.25);
620 const FixedPointAccum z_a_pow_2_adj = SaturatingAddNonGemmlowp(
621     FixedPointAccum::FromRaw(SaturatingRoundingMultiplyByPOTParam(
622         static_cast<int32_t>(InputIntegerBits - z_a_headroom_plus_1),
623         31 - kAccumIntegerBits)),
624     shifted_quarter);
625
626 // z_b is treated like z_a, but premultiplying by sqrt(0.5).
627 FixedPoint0 z_b = z_a * sqrt_half;
628 int z_b_headroom = CountLeadingZeros(static_cast<uint32_t>(z_b.raw())) - 1;
629 const int32_t r_b_raw =
630     SaturatingRoundingMultiplyByPOTParam(z_a.raw(), z_b_headroom);
631 const FixedPointAccum z_b_pow_2_adj = SaturatingSub(
632     FixedPointAccum::FromRaw(SaturatingRoundingMultiplyByPOTParam(
633         static_cast<int32_t>(InputIntegerBits - z_b_headroom),
634         31 - kAccumIntegerBits)),
635     shifted_quarter);
636
637 const FixedPoint0 r = FixedPoint0::FromRaw(std::min(r_a_raw, r_b_raw));
638 const FixedPointAccum z_pow_2_adj = FixedPointAccum::FromRaw(
639     std::max(z_a_pow_2_adj.raw(), z_b_pow_2_adj.raw()));
640
641 const FixedPoint0 p = gemmlowp::RoundingHalfSum(r, sqrt_sqrt_half);
642 FixedPoint0 q = r - sqrt_sqrt_half;
643 q = q + q;
644
645 const FixedPoint0 common_sq = q * q;
646 const FixedPoint0 num = q * r + q * common_sq * alpha_n;
647 const FixedPoint0 denom_minus_one_0 =
648     p * (alpha_i + q + alpha_d * common_sq) + alpha_f * q;
649 const FixedPoint0 recip_denom =
650     one_over_one_plus_x_for_x_in_0_1(denom_minus_one_0);
651
652 const FixedPointAccum num_scaled = gemmlowp::Rescale<kAccumIntegerBits>(num);
653 return gemmlowp::Rescale<OutputIntegerBits>(z_pow_2_adj * log_2 +
654     num_scaled * recip_denom);
655 }
656
657 template <int OutputIntegerBits, int InputIntegerBits>
658 inline gemmlowp::FixedPoint<int32_t, OutputIntegerBits>
659 log_x_for_x_greater_than_or_equal_to_1(
660     gemmlowp::FixedPoint<int32_t, InputIntegerBits> input_val) {
661     static_assert(
662         OutputIntegerBits >= min_log_x_output_bits(InputIntegerBits),
663         "Output integer bits must be sufficient to accommodate logs of inputs.");
664     return log_x_for_x_greater_than_or_equal_to_1_impl<OutputIntegerBits,
665         InputIntegerBits>(
666         input_val);

```

```

667 }
668
669 inline int32_t GetReciprocal(int32_t x, int x_integer_digits,
670                             int* num_bits_over_unit) {
671     int headroom_plus_one = CountLeadingZeros(static_cast<uint32_t>(x));
672     // This is the number of bits to the left of the binary point above 1.0.
673     // Consider x=1.25. In that case shifted_scale=0.8 and
674     // no later adjustment will be needed.
675     *num_bits_over_unit = x_integer_digits - headroom_plus_one;
676     const int32_t shifted_sum_minus_one =
677         static_cast<int32_t>((static_cast<uint32_t>(x) << headroom_plus_one) -
678                             (static_cast<uint32_t>(1) << 31));
679
680     gemmlowp::FixedPoint<int32_t, 0> shifted_scale =
681         gemmlowp::one_over_one_plus_x_for_x_in_0_1(
682             gemmlowp::FixedPoint<int32_t, 0>::FromRaw(shifted_sum_minus_one));
683     return shifted_scale.raw();
684 }
685
686 inline void GetInvSqrtQuantizedMultiplierExp(int32_t input, int reverse_shift,
687                                               int32_t* output_inv_sqrt,
688                                               int* output_shift) {
689     TFLITE_DCHECK_GE(input, 0);
690     if (input <= 1) {
691         // Handle the input value 1 separately to avoid overflow in that case
692         // in the general computation below (b/143972021). Also handle 0 as if it
693         // were a 1. 0 is an invalid input here (divide by zero) and 1 is a valid
694         // but rare/unrealistic input value. We can expect both to occur in some
695         // incompletely trained models, but probably not in fully trained models.
696         *output_inv_sqrt = std::numeric_limits<std::int32_t>::max();
697         *output_shift = 0;
698         return;
699     }
700     TFLITE_DCHECK_GT(input, 1);
701     *output_shift = 11;
702     while (input >= (1 << 29)) {
703         input /= 4;
704         ++*output_shift;
705     }
706     const unsigned max_left_shift_bits =
707         CountLeadingZeros(static_cast<uint32_t>(input)) - 1;
708     const unsigned max_left_shift_bit_pairs = max_left_shift_bits / 2;
709     const unsigned left_shift_bit_pairs = max_left_shift_bit_pairs - 1;
710     *output_shift -= left_shift_bit_pairs;
711     input <= 2 * left_shift_bit_pairs;
712     TFLITE_DCHECK_GE(input, (1 << 27));
713     TFLITE_DCHECK_LT(input, (1 << 29));
714     using gemmlowp::FixedPoint;
715     using gemmlowp::Rescale;

```

```

716     using gemmlowp::SaturatingRoundingMultiplyByPOT;
717     // Using 3 integer bits gives us enough room for the internal arithmetic in
718     // this Newton-Raphson iteration.
719     using F3 = FixedPoint<int32_t, 3>;
720     using F0 = FixedPoint<int32_t, 0>;
721     const F3 fixedpoint_input = F3::FromRaw(input >> 1);
722     const F3 fixedpoint_half_input =
723         SaturatingRoundingMultiplyByPOT<-1>(fixedpoint_input);
724     const F3 fixedpoint_half_three =
725         GEMMLOWP_CHECKED_FIXEDPOINT_CONSTANT(F3, (1 << 28) + (1 << 27), 1.5);
726     // Newton-Raphson iteration
727     // Naive unoptimized starting guess: x = 1
728     F3 x = F3::One();
729     // Naive unoptimized number of iterations: 5
730     for (int i = 0; i < 5; i++) {
731         const F3 x3 = Rescale<3>(x * x * x);
732         x = Rescale<3>(fixedpoint_half_three * x - fixedpoint_half_input * x3);
733     }
734     const F0 fixedpoint_half_sqrt_2 =
735         GEMMLOWP_CHECKED_FIXEDPOINT_CONSTANT(F0, 1518500250, std::sqrt(2.) / 2.);
736     x = x * fixedpoint_half_sqrt_2;
737     *output_inv_sqrt = x.raw();
738     if (*output_shift < 0) {
739         *output_inv_sqrt <<= -*output_shift;
740         *output_shift = 0;
741     }
742     // Convert right shift (right is positive) to left shift.
743     *output_shift *= reverse_shift;
744 }
745
746 // DO NOT USE THIS STRUCT FOR NEW FUNCTIONALITY BEYOND IMPLEMENTING
747 // BROADCASTING.
748 //
749 // NdArrayDesc<N> describes the shape and memory layout of an N-dimensional
750 // rectangular array of numbers.
751 //
752 // NdArrayDesc<N> is basically identical to Dims<N> defined in types.h.
753 // However, as Dims<N> is to be deprecated, this class exists as an adaptor
754 // to enable simple unoptimized implementations of element-wise broadcasting
755 // operations.
756 template <int N>
757 struct NdArrayDesc {
758     // The "extent" of each dimension. Indices along dimension d must be in the
759     // half-open interval [0, extents[d]).
760     int extents[N];
761
762     // The number of *elements* (not bytes) between consecutive indices of each
763     // dimension.
764     int strides[N];

```



```

765 };
766
767 // DO NOT USE THIS FUNCTION FOR NEW FUNCTIONALITY BEYOND IMPLEMENTING
768 // BROADCASTING.
769 //
770 // Same as Offset(), except takes as NdArrayDesc<N> instead of Dims<N>.
771 inline int SubscriptToIndex(const NdArrayDesc<4>& desc, int i0, int i1, int i2,
772                             int i3) {
773     TFLITE_DCHECK(i0 >= 0 && i0 < desc.extents[0]);
774     TFLITE_DCHECK(i1 >= 0 && i1 < desc.extents[1]);
775     TFLITE_DCHECK(i2 >= 0 && i2 < desc.extents[2]);
776     TFLITE_DCHECK(i3 >= 0 && i3 < desc.extents[3]);
777     return i0 * desc.strides[0] + i1 * desc.strides[1] + i2 * desc.strides[2] +
778           i3 * desc.strides[3];
779 }
780
781 inline int SubscriptToIndex(const NdArrayDesc<5>& desc, int indexes[5]) {
782     return indexes[0] * desc.strides[0] + indexes[1] * desc.strides[1] +
783           indexes[2] * desc.strides[2] + indexes[3] * desc.strides[3] +
784           indexes[4] * desc.strides[4];
785 }
786
787 inline int SubscriptToIndex(const NdArrayDesc<8>& desc, int indexes[8]) {
788     return indexes[0] * desc.strides[0] + indexes[1] * desc.strides[1] +
789           indexes[2] * desc.strides[2] + indexes[3] * desc.strides[3] +
790           indexes[4] * desc.strides[4] + indexes[5] * desc.strides[5] +
791           indexes[6] * desc.strides[6] + indexes[7] * desc.strides[7];
792 }
793
794 // Given the dimensions of the operands for an element-wise binary broadcast,
795 // adjusts them so that they can be directly iterated over with simple loops.
796 // Returns the adjusted dims as instances of NdArrayDesc in 'desc0_out' and
797 // 'desc1_out'. 'desc0_out' and 'desc1_out' cannot be nullptr.
798 //
799 // This function assumes that the two input shapes are compatible up to
800 // broadcasting and the shorter one has already been prepended with 1s to be the
801 // same length. E.g., if shape0 is (1, 16, 16, 64) and shape1 is (1, 64),
802 // shape1 must already have been prepended to be (1, 1, 1, 64). Recall that
803 // Dims<N> refer to shapes in reverse order. In this case, input0_dims will be
804 // (64, 16, 16, 1) and input1_dims will be (64, 1, 1, 1).
805 //
806 // When two shapes are compatible up to broadcasting, for each dimension d,
807 // the input extents are either equal, or one of them is 1.
808 //
809 // This function performs the following for each dimension d:
810 // - If the extents are equal, then do nothing since the loop that walks over
811 //   both of the input arrays is correct.
812 // - Otherwise, one (and only one) of the extents must be 1. Say extent0 is 1
813 //   and extent1 is e1. Then set extent0 to e1 and stride0 *to 0*. This allows

```

```

814 // array0 to be referenced *at any index* in dimension d and still access the
815 // same slice.
816 template <int N>
817 inline void NdArrayDescsForElementwiseBroadcast(const Dims<N>& input0_dims,
818                                                  const Dims<N>& input1_dims,
819                                                  NdArrayDesc<N>* desc0_out,
820                                                  NdArrayDesc<N>* desc1_out) {
821     TFLITE_DCHECK(desc0_out != nullptr);
822     TFLITE_DCHECK(desc1_out != nullptr);
823
824     // Copy dims to desc.
825     for (int i = 0; i < N; ++i) {
826         desc0_out->extents[i] = input0_dims.sizes[i];
827         desc0_out->strides[i] = input0_dims.strides[i];
828         desc1_out->extents[i] = input1_dims.sizes[i];
829         desc1_out->strides[i] = input1_dims.strides[i];
830     }
831
832     // Walk over each dimension. If the extents are equal do nothing.
833     // Otherwise, set the desc with extent 1 to have extent equal to the other and
834     // stride 0.
835     for (int i = 0; i < N; ++i) {
836         const int extent0 = ArraySize(input0_dims, i);
837         const int extent1 = ArraySize(input1_dims, i);
838         if (extent0 != extent1) {
839             if (extent0 == 1) {
840                 desc0_out->strides[i] = 0;
841                 desc0_out->extents[i] = extent1;
842             } else {
843                 TFLITE_DCHECK_EQ(extent1, 1);
844                 desc1_out->strides[i] = 0;
845                 desc1_out->extents[i] = extent0;
846             }
847         }
848     }
849 }
850
851 // Copies dims to desc, calculating strides.
852 template <int N>
853 inline void CopyDimsToDesc(const RuntimeShape& input_shape,
854                            NdArrayDesc<N>* desc_out) {
855     int desc_stride = 1;
856     for (int i = N - 1; i >= 0; --i) {
857         desc_out->extents[i] = input_shape.Dims(i);
858         desc_out->strides[i] = desc_stride;
859         desc_stride *= input_shape.Dims(i);
860     }
861 }
862

```

```

863 template <int N>
864 inline void NdArrayDescsForElementwiseBroadcast(
865     const RuntimeShape& input0_shape, const RuntimeShape& input1_shape,
866     NdArrayDesc<N>* desc0_out, NdArrayDesc<N>* desc1_out) {
867     TFLITE_DCHECK(desc0_out != nullptr);
868     TFLITE_DCHECK(desc1_out != nullptr);
869
870     auto extended_input0_shape = RuntimeShape::ExtendedShape(N, input0_shape);
871     auto extended_input1_shape = RuntimeShape::ExtendedShape(N, input1_shape);
872
873     // Copy dims to desc, calculating strides.
874     CopyDimsToDesc<N>(extended_input0_shape, desc0_out);
875     CopyDimsToDesc<N>(extended_input1_shape, desc1_out);
876
877     // Walk over each dimension. If the extents are equal do nothing.
878     // Otherwise, set the desc with extent 1 to have extent equal to the other and
879     // stride 0.
880     for (int i = 0; i < N; ++i) {
881         const int extent0 = extended_input0_shape.Dims(i);
882         const int extent1 = extended_input1_shape.Dims(i);
883         if (extent0 != extent1) {
884             if (extent0 == 1) {
885                 desc0_out->strides[i] = 0;
886                 desc0_out->extents[i] = extent1;
887             } else {
888                 TFLITE_DCHECK_EQ(extent1, 1);
889                 desc1_out->strides[i] = 0;
890                 desc1_out->extents[i] = extent0;
891             }
892         }
893     }
894 }
895
896 template <int N>
897 inline void NdArrayDescsForElementwiseBroadcast(
898     const RuntimeShape& input0_shape, const RuntimeShape& input1_shape,
899     const RuntimeShape& input2_shape, NdArrayDesc<N>* desc0_out,
900     NdArrayDesc<N>* desc1_out, NdArrayDesc<N>* desc2_out) {
901     TFLITE_DCHECK(desc0_out != nullptr);
902     TFLITE_DCHECK(desc1_out != nullptr);
903     TFLITE_DCHECK(desc2_out != nullptr);
904
905     auto extended_input0_shape = RuntimeShape::ExtendedShape(N, input0_shape);
906     auto extended_input1_shape = RuntimeShape::ExtendedShape(N, input1_shape);
907     auto extended_input2_shape = RuntimeShape::ExtendedShape(N, input2_shape);
908
909     // Copy dims to desc, calculating strides.
910     CopyDimsToDesc<N>(extended_input0_shape, desc0_out);
911     CopyDimsToDesc<N>(extended_input1_shape, desc1_out);

```

```

912 CopyDimsToDesc<N>(extended_input2_shape, desc2_out);
913
914 // Walk over each dimension. If the extents are equal do nothing.
915 // Otherwise, set the desc with extent 1 to have extent equal to the other and
916 // stride 0.
917 for (int i = 0; i < N; ++i) {
918     const int extent0 = extended_input0_shape.Dims(i);
919     const int extent1 = extended_input1_shape.Dims(i);
920     const int extent2 = extended_input2_shape.Dims(i);
921
922     int extent = extent0;
923     if (extent1 != 1) extent = extent1;
924     if (extent2 != 1) extent = extent2;
925
926     TFLITE_DCHECK(extent0 == 1 || extent0 == extent);
927     TFLITE_DCHECK(extent1 == 1 || extent1 == extent);
928     TFLITE_DCHECK(extent2 == 1 || extent2 == extent);
929
930     if (!(extent0 == extent1 && extent1 == extent2)) {
931         if (extent0 == 1) {
932             desc0_out->strides[i] = 0;
933             desc0_out->extents[i] = extent;
934         }
935         if (extent1 == 1) {
936             desc1_out->strides[i] = 0;
937             desc1_out->extents[i] = extent;
938         }
939         if (extent2 == 1) {
940             desc2_out->strides[i] = 0;
941             desc2_out->extents[i] = extent;
942         }
943     }
944 }
945 }
946
947 // Detailed implementation of NDOpsHelper, the indexes must be a zero array.
948 // This implementation is equivalent to N nested loops. Ex, if N=4, it can be
949 // re-written as:
950 // for (int b = 0; b < output.extents[0]; ++b) {
951 //     for (int y = 0; y < output.extents[1]; ++y) {
952 //         for (int x = 0; x < output.extents[2]; ++x) {
953 //             for (int c = 0; c < output.extents[3]; ++c) {
954 //                 calc({b,y,x,c});
955 //             }
956 //         }
957 //     }
958 // }
959 template <int N, int DIM, typename Calc>
960 typename std::enable_if<DIM != N - 1, void>::type NDOpsHelperImpl(

```

```

961     const NdArrayDesc<N>& output, const Calc& calc, int indexes[N]) {
962     for (indexes[DIM] = 0; indexes[DIM] < output.extents[DIM]; ++indexes[DIM]) {
963         NDOpsHelperImpl<N, DIM + 1, Calc>(output, calc, indexes);
964     }
965 }
966
967 template <int N, int DIM, typename Calc>
968 typename std::enable_if<DIM == N - 1, void>::type NDOpsHelperImpl(
969     const NdArrayDesc<N>& output, const Calc& calc, int indexes[N]) {
970     for (indexes[DIM] = 0; indexes[DIM] < output.extents[DIM]; ++indexes[DIM]) {
971         calc(indexes);
972     }
973 }
974
975 // Execute the calc function in the innermost iteration based on the shape of
976 // the output. The calc function should take a single argument of type int[N].
977 template <int N, typename Calc>
978 inline void NDOpsHelper(const NdArrayDesc<N>& output, const Calc& calc) {
979     int indexes[N] = {0};
980     NDOpsHelperImpl<N, 0, Calc>(output, calc, indexes);
981 }
982 // Copied from gemmlowp::RoundDown when we dropped direct dependency on
983 // gemmlowp.
984 //
985 // Returns the runtime argument rounded down to the nearest multiple of
986 // the fixed Modulus.
987 template <unsigned Modulus, typename Integer>
988 Integer RoundDown(Integer i) {
989     return i - (i % Modulus);
990 }
991
992 // Copied from gemmlowp::RoundUp when we dropped direct dependency on
993 // gemmlowp.
994 //
995 // Returns the runtime argument rounded up to the nearest multiple of
996 // the fixed Modulus.
997 template <unsigned Modulus, typename Integer>
998 Integer RoundUp(Integer i) {
999     return RoundDown<Modulus>(i + Modulus - 1);
1000 }
1001
1002 // Copied from gemmlowp::CeilQuotient when we dropped direct dependency on
1003 // gemmlowp.
1004 //
1005 // Returns the quotient a / b rounded up ('ceil') to the nearest integer.
1006 template <typename Integer>
1007 Integer CeilQuotient(Integer a, Integer b) {
1008     return (a + b - 1) / b;
1009 }

```

```

1010
1011 // This function is a copy of gemmlowp::HowManyThreads, copied when we dropped
1012 // the direct dependency of internal/optimized/ on gemmlowp.
1013 //
1014 // It computes a reasonable number of threads to use for a GEMM of shape
1015 // (rows, cols, depth).
1016 //
1017 // TODO(b/131910176): get rid of this function by switching each call site
1018 // to its own more sensible logic for its own workload.
1019 template <int KernelRows>
1020 inline int LegacyHowManyThreads(int max_num_threads, int rows, int cols,
1021                                int depth) {
1022     // Early-exit in the default case where multi-threading is disabled.
1023     if (max_num_threads == 1) {
1024         return 1;
1025     }
1026
1027     // Ensure that each thread has KernelRows rows to process, if at all possible.
1028     int thread_count = std::min(max_num_threads, rows / KernelRows);
1029
1030     // Limit the number of threads according to the overall size of the problem.
1031     if (thread_count > 1) {
1032         // Empirically determined value.
1033         static constexpr std::uint64_t min_cubic_size_per_thread = 64 * 1024;
1034
1035         // We can only multiply two out of three sizes without risking overflow
1036         const std::uint64_t cubic_size =
1037             std::uint64_t(rows) * std::uint64_t(cols) * std::uint64_t(depth);
1038
1039         thread_count = std::min(
1040             thread_count, static_cast<int>(cubic_size / min_cubic_size_per_thread));
1041     }
1042
1043     if (thread_count < 1) {
1044         thread_count = 1;
1045     }
1046
1047     assert(thread_count > 0 && thread_count <= max_num_threads);
1048     return thread_count;
1049 }
1050
1051 template <typename T>
1052 void optimized_ops_preload_l1_stream(const T* ptr) {
1053     #ifdef __GNUC__
1054         // builtin offered by GCC-compatible compilers including clang
1055         __builtin_prefetch(ptr, /* 0 means read */ 0, /* 0 means no locality */ 0);
1056     #else
1057         (void)ptr;
1058     #endif

```

```
1059 }
1060
1061 template <typename T>
1062 void optimized_ops_preload_l1_keep(const T* ptr) {
1063     #ifdef __GNUC__
1064         // builtin offered by GCC-compatible compilers including clang
1065         __builtin_prefetch(ptr, /* 0 means read */ 0, /* 3 means high locality */ 3);
1066     #else
1067         (void)ptr;
1068     #endif
1069 }
1070
1071 template <typename T>
1072 void optimized_ops_prefetch_write_l1_keep(const T* ptr) {
1073     #ifdef __GNUC__
1074         // builtin offered by GCC-compatible compilers including clang
1075         __builtin_prefetch(ptr, /* 1 means write */ 1, /* 3 means high locality */ 3);
1076     #else
1077         (void)ptr;
1078     #endif
1079 }
1080
1081 } // namespace tflite
1082
1083 #endif // TENSORFLOW_LITE_KERNELS_INTERNAL_COMMON_H_
```