

ca6f96b62a ▾

...


tensorflow / tensorflow / lite / kernels / embedding\_lookup\_sparse.cc



mihaimaruseac Fix a dangerous integer overflow and a malloc of negative size. ... ✖

History

6 contributors



269 lines (232 sloc) | 9.8 KB

...

```

1  /* Copyright 2017 The TensorFlow Authors. All Rights Reserved.
2
3  Licensed under the Apache License, Version 2.0 (the "License");
4  you may not use this file except in compliance with the License.
5  You may obtain a copy of the License at
6
7      http://www.apache.org/licenses/LICENSE-2.0
8
9  Unless required by applicable law or agreed to in writing, software
10 distributed under the License is distributed on an "AS IS" BASIS,
11 WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
12 See the License for the specific language governing permissions and
13 limitations under the License.
14 =====*/
15
16 // Op that looks up items from a sparse tensor in an embedding matrix.
17 // The sparse lookup tensor is represented by three individual tensors: lookup,
18 // indices, and dense_shape. The representation assume that the corresponding
19 // dense tensor would satisfy:
20 //   * dense.shape = dense_shape
21 //   * dense[tuple(indices[i])] = lookup[i]
22 //
23 // By convention, indices should be sorted.
24 //
25 // Options:
26 //   combiner: The reduction op (SUM, MEAN, SQRTN).
27 //   * SUM computes the weighted sum of the embedding results.
28 //   * MEAN is the weighted sum divided by the total weight.
29 //   * SQRTN is the weighted sum divided by the square root of the sum of the

```

```

30 //      squares of the weights.
31 //
32 // Input:
33 //      Tensor[0]: Ids to lookup, dim.size == 1, int32.
34 //      Tensor[1]: Indices, int32.
35 //      Tensor[2]: Dense shape, int32.
36 //      Tensor[3]: Weights to use for aggregation, float.
37 //      Tensor[4]: Params, a matrix of multi-dimensional items,
38 //                  dim.size >= 2, float.
39 //
40 // Output:
41 //      A (dense) tensor representing the combined embeddings for the sparse ids.
42 //      For each row in the sparse tensor represented by (lookup, indices, shape)
43 //      the op looks up the embeddings for all ids in that row, multiplies them by
44 //      the corresponding weight, and combines these embeddings as specified in the
45 //      last dimension.
46 //
47 //      Output.dim = [l0, ... , ln-1, e1, ..., em]
48 //      Where dense_shape == [l0, ..., ln] and Tensor[4].dim == [e0, e1, ..., em]
49 //
50 //      For instance, if params is a 10x20 matrix and ids, weights are:
51 //
52 //      [0, 0]: id 1, weight 2.0
53 //      [0, 1]: id 3, weight 0.5
54 //      [1, 0]: id 0, weight 1.0
55 //      [2, 3]: id 1, weight 3.0
56 //
57 //      with combiner=MEAN, then the output will be a (3, 20) tensor where:
58 //
59 //      output[0, :] = (params[1, :] * 2.0 + params[3, :] * 0.5) / (2.0 + 0.5)
60 //      output[1, :] = (params[0, :] * 1.0) / 1.0
61 //      output[2, :] = (params[1, :] * 3.0) / 3.0
62 //
63 //      When indices are out of bound, the op will not succeed.
64
65 #include <stdint.h>
66
67 #include <algorithm>
68 #include <cmath>
69
70 #include "tensorflow/lite/c/builtin_op_data.h"
71 #include "tensorflow/lite/c/common.h"
72 #include "tensorflow/lite/kernels/internal/tensor_ctypes.h"
73 #include "tensorflow/lite/kernels/internal/tensor_utils.h"
74 #include "tensorflow/lite/kernels/kernel_util.h"
75
76 namespace tflite {
77 namespace ops {
78 namespace builtin {

```

```

79
80 namespace {
81
82 TfliteStatus Prepare(TfLiteContext* context, TfLiteNode* node) {
83     TF_LITE_ENSURE_EQ(context, NumInputs(node), 5);
84     TF_LITE_ENSURE_EQ(context, NumOutputs(node), 1);
85
86     const TfLiteTensor* ids;
87     TF_LITE_ENSURE_OK(context, GetInputSafe(context, node, 0, &ids));
88     TF_LITE_ENSURE_EQ(context, NumDimensions(ids), 1);
89     TF_LITE_ENSURE_EQ(context, ids->type, kTfLiteInt32);
90
91     const TfLiteTensor* indices;
92     TF_LITE_ENSURE_OK(context, GetInputSafe(context, node, 1, &indices));
93     TF_LITE_ENSURE_EQ(context, NumDimensions(indices), 2);
94     TF_LITE_ENSURE_EQ(context, indices->type, kTfLiteInt32);
95
96     const TfLiteTensor* shape;
97     TF_LITE_ENSURE_OK(context, GetInputSafe(context, node, 2, &shape));
98     TF_LITE_ENSURE_EQ(context, NumDimensions(shape), 1);
99     TF_LITE_ENSURE_EQ(context, shape->type, kTfLiteInt32);
100
101     const TfLiteTensor* weights;
102     TF_LITE_ENSURE_OK(context, GetInputSafe(context, node, 3, &weights));
103     TF_LITE_ENSURE_EQ(context, NumDimensions(weights), 1);
104     TF_LITE_ENSURE_EQ(context, weights->type, kTfLiteFloat32);
105
106     TF_LITE_ENSURE_EQ(context, SizeOfDimension(indices, 0),
107         SizeOfDimension(ids, 0));
108     TF_LITE_ENSURE_EQ(context, SizeOfDimension(indices, 0),
109         SizeOfDimension(weights, 0));
110
111     const TfLiteTensor* value;
112     TF_LITE_ENSURE_OK(context, GetInputSafe(context, node, 4, &value));
113     TF_LITE_ENSURE(context, NumDimensions(value) >= 2);
114
115     // Mark the output as a dynamic tensor.
116     TfLiteTensor* output;
117     TF_LITE_ENSURE_OK(context, GetOutputSafe(context, node, 0, &output));
118     TF_LITE_ENSURE_TYPES_EQ(context, output->type, kTfLiteFloat32);
119     output->allocation_type = kTfLiteDynamic;
120
121     return kTfLiteOk;
122 }
123
124 void FinalizeAggregation(TfLiteCombinerType combiner, int num_elements,
125     float current_total_weight,
126     float current_squares_weight, int embedding_size,
127     float* output) {

```

```

128 if (combiner != kTfLiteCombinerTypeSum && num_elements > 0) {
129     float multiplier = 1.0;
130     switch (combiner) {
131         case kTfLiteCombinerTypeMean:
132             multiplier = current_total_weight;
133             break;
134         case kTfLiteCombinerTypeSqrtn:
135             multiplier = std::sqrt(current_squares_weight);
136             break;
137         default:
138             break;
139     }
140     for (int k = 0; k < embedding_size; k++) {
141         output[k] /= multiplier;
142     }
143 }
144 }
145
146 TfLiteStatus Eval(TfLiteContext* context, TfLiteNode* node) {
147     auto* params =
148         reinterpret_cast<TfLiteEmbeddingLookupSparseParams*>(node->builtin_data);
149     TfLiteTensor* output;
150     TF_LITE_ENSURE_OK(context, GetOutputSafe(context, node, 0, &output));
151     const TfLiteTensor* ids;
152     TF_LITE_ENSURE_OK(context, GetInputSafe(context, node, 0, &ids));
153     const TfLiteTensor* indices;
154     TF_LITE_ENSURE_OK(context, GetInputSafe(context, node, 1, &indices));
155     const TfLiteTensor* dense_shape;
156     TF_LITE_ENSURE_OK(context, GetInputSafe(context, node, 2, &dense_shape));
157     const TfLiteTensor* weights;
158     TF_LITE_ENSURE_OK(context, GetInputSafe(context, node, 3, &weights));
159     const TfLiteTensor* value;
160     TF_LITE_ENSURE_OK(context, GetInputSafe(context, node, 4, &value));
161
162     const int lookup_rank = SizeOfDimension(indices, 1);
163     const int embedding_rank = NumDimensions(value);
164     const int num_lookups = SizeOfDimension(ids, 0);
165     const int num_rows = SizeOfDimension(value, 0);
166
167     // The last dimension gets replaced by the embedding.
168     const int output_rank = (lookup_rank - 1) + (embedding_rank - 1);
169
170     // Make sure that the actual dense shape of the sparse tensor represented by
171     // (lookup, indices, dense_shape) is consistent.
172     TF_LITE_ENSURE_EQ(context, SizeOfDimension(dense_shape, 0), lookup_rank);
173
174     // Resize output tensor.
175     TfLiteIntArray* output_shape = TfLiteIntArrayCreate(output_rank);
176     TF_LITE_ENSURE(context, output_shape != nullptr);

```

```

177     int k = 0;
178     int embedding_size = 1;
179     int lookup_size = 1;
180     for (int i = 0; i < lookup_rank - 1; i++, k++) {
181         const int dim = dense_shape->data.i32[i];
182         lookup_size *= dim;
183         output_shape->data[k] = dim;
184     }
185     for (int i = 1; i < embedding_rank; i++, k++) {
186         const int dim = SizeOfDimension(value, i);
187         embedding_size *= dim;
188         output_shape->data[k] = dim;
189     }
190     TF_LITE_ENSURE_STATUS(context->ResizeTensor(context, output, output_shape));
191     const int output_size = lookup_size * embedding_size;
192     TfLiteTensorRealloc(output_size * sizeof(float), output);
193
194     float* output_ptr = GetTensorData<float>(output);
195     const float* weights_ptr = GetTensorData<float>(weights);
196     const float* value_ptr = GetTensorData<float>(value);
197
198     std::fill_n(output_ptr, output_size, 0.0f);
199
200     // Keep track of the current bucket for aggregation/combination.
201     int current_output_offset = 0;
202     float current_total_weight = 0.0;
203     float current_squares_weight = 0.0;
204     int num_elements = 0;
205
206     for (int i = 0; i < num_lookups; i++) {
207         int idx = ids->data.i32[i];
208         if (idx >= num_rows || idx < 0) {
209             context->ReportError(context,
210                                 "Embedding Lookup Sparse: index out of bounds. "
211                                 "Got %d, and bounds are [0, %d]",
212                                 idx, num_rows - 1);
213             return kTfLiteError;
214         }
215
216         // Check where we need to aggregate.
217         const int example_indices_offset = i * lookup_rank;
218         int output_bucket = 0;
219         int stride = 1;
220         for (int k = (lookup_rank - 1) - 1; k >= 0; k--) {
221             output_bucket += indices->data.i32[example_indices_offset + k] * stride;
222             stride *= dense_shape->data.i32[k];
223         }
224         const int output_offset = output_bucket * embedding_size;
225

```

```

226 // If we are in a new aggregation bucket and the combiner is not the sum,
227 // go back and finalize the result of the previous bucket.
228 if (output_offset != current_output_offset) {
229     FinalizeAggregation(params->combiner, num_elements, current_total_weight,
230                          current_squares_weight, embedding_size,
231                          &output_ptr[current_output_offset]);
232
233     // Track next bucket.
234     num_elements = 0;
235     current_total_weight = 0.0;
236     current_squares_weight = 0.0;
237     current_output_offset = output_offset;
238 }
239
240 // Add element to aggregation.
241 ++num_elements;
242 const int example_embedding_offset = idx * embedding_size;
243 const float w = weights_ptr[i];
244 current_squares_weight += w * w;
245 current_total_weight += w;
246 for (int k = 0; k < embedding_size; k++) {
247     output_ptr[current_output_offset + k] +=
248         value_ptr[example_embedding_offset + k] * w;
249 }
250 }
251
252 // Finalize last bucket.
253 FinalizeAggregation(params->combiner, num_elements, current_total_weight,
254                    current_squares_weight, embedding_size,
255                    &GetTensorData<float>(output)[current_output_offset]);
256
257 return kTfLiteOk;
258 }
259
260 } // namespace
261
262 TfLiteRegistration* Register_EMBEDDING_LOOKUP_SPARSE() {
263     static TfLiteRegistration r = {nullptr, nullptr, Prepare, Eval};
264     return &r;
265 }
266
267 } // namespace builtin
268 } // namespace ops
269 } // namespace tfLite

```