

5100e359ae ▾

...

tensorflow / tensorflow / core / kernels / sparse_tensors_map_ops.cc



jpienaar Rename to underlying type rather than alias ... ✓

History

5 contributors



519 lines (427 sloc) | 19.6 KB

...

```

1  /* Copyright 2015 The TensorFlow Authors. All Rights Reserved.
2
3  Licensed under the Apache License, Version 2.0 (the "License");
4  you may not use this file except in compliance with the License.
5  You may obtain a copy of the License at
6
7      http://www.apache.org/licenses/LICENSE-2.0
8
9  Unless required by applicable law or agreed to in writing, software
10 distributed under the License is distributed on an "AS IS" BASIS,
11 WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
12 See the License for the specific language governing permissions and
13 limitations under the License.
14 =====*/
15
16 #define EIGEN_USE_THREADS
17
18 #include <algorithm>
19 #include <numeric>
20 #include <unordered_map>
21 #include <utility>
22 #include <vector>
23
24 #include "tensorflow/core/framework/op_kernel.h"
25 #include "tensorflow/core/framework/register_types.h"
26 #include "tensorflow/core/framework/resource_mgr.h"
27 #include "tensorflow/core/framework/tensor.h"
28 #include "tensorflow/core/framework/tensor_util.h"
29 #include "tensorflow/core/framework/types.h"

```

```

30 #include "tensorflow/core/lib/gtl/inlined_vector.h"
31 #include "tensorflow/core/util/overflow.h"
32 #include "tensorflow/core/util/sparse/sparse_tensor.h"
33
34 namespace tensorflow {
35
36 typedef Eigen::ThreadPoolDevice CPUDevice;
37
38 using sparse::SparseTensor;
39
40 class SparseTensorsMap : public ResourceBase {
41 public:
42     explicit SparseTensorsMap(const string& name) : name_(name), counter_(0) {}
43
44     string DebugString() const override { return "A SparseTensorsMap"; }
45
46     typedef struct {
47         Tensor indices;
48         Tensor values;
49         gtl::InlinedVector<int64_t, 8> shape;
50     } PersistentSparseTensor;
51
52     Status AddSparseTensor(OpKernelContext* ctx, const SparseTensor& sp,
53                           int64_t* handle) {
54         Tensor ix;
55         TF_RETURN_IF_ERROR(
56             ctx->allocate_temp(sp.indices().dtype(), sp.indices().shape(), &ix));
57         ix = sp.indices();
58
59         Tensor values;
60         TF_RETURN_IF_ERROR(ctx->allocate_temp(sp.indices().dtype(),
61                                                sp.indices().shape(), &values));
62         values = sp.values();
63         {
64             mutex_lock l(mu_);
65             int64_t unique_st_handle = counter_++; // increment is guarded on purpose
66             sp_tensors_[unique_st_handle] = PersistentSparseTensor{
67                 ix, values,
68                 gtl::InlinedVector<int64_t, 8>(sp.shape().begin(), sp.shape().end())};
69             *handle = unique_st_handle;
70         }
71         return Status::OK();
72     }
73
74     Status RetrieveAndClearSparseTensors(
75         OpKernelContext* ctx, const TTypes<int64_t>::ConstVec& handles,
76         std::vector<SparseTensor>* sparse_tensors) {
77         sparse_tensors->clear();
78         sparse_tensors->reserve(handles.size());

```

```

79     {
80         mutex_lock l(mu_);
81         for (size_t i = 0; i < handles.size(); ++i) {
82             const int64_t handle = handles(i);
83             auto sp_iter = sp_tensors_.find(handle);
84             if (sp_iter == sp_tensors_.end()) {
85                 return errors::InvalidArgument(
86                     "Unable to find SparseTensor: ", handle, " in map: ", name_);
87             }
88             const Tensor* ix = &sp_iter->second.indices;
89             const Tensor* values = &sp_iter->second.values;
90             const auto& shape = sp_iter->second.shape;
91             SparseTensor tensor;
92             TF_RETURN_IF_ERROR(SparseTensor::Create(*ix, *values, shape, &tensor));
93             sparse_tensors->push_back(std::move(tensor));
94             sp_tensors_.erase(sp_iter);
95         }
96     }
97
98     return Status::OK();
99 }
100
101 protected:
102     ~SparseTensorsMap() override {}
103
104 private:
105     string name_;
106
107     mutex mu_;
108     int64_t counter_ TF_GUARDED_BY(mu_);
109     std::unordered_map<int64_t, PersistentSparseTensor> sp_tensors_
110         TF_GUARDED_BY(mu_);
111 };
112
113 class SparseTensorAccessingOp : public OpKernel {
114 public:
115     typedef std::function<Status(SparseTensorsMap**)> CreatorCallback;
116
117     explicit SparseTensorAccessingOp(OpKernelConstruction* context)
118         : OpKernel(context), sparse_tensors_map_(nullptr) {}
119
120 protected:
121     ~SparseTensorAccessingOp() override {
122         if (sparse_tensors_map_) sparse_tensors_map_->Unref();
123     }
124
125     Status GetMap(OpKernelContext* ctx, bool is_writing,
126                 SparseTensorsMap** sparse_tensors_map) {
127         mutex_lock l(mu_);

```

```

128
129     if (sparse_tensors_map_) {
130         *sparse_tensors_map = sparse_tensors_map_;
131         return Status::OK();
132     }
133
134     TF_RETURN_IF_ERROR(cinfo_.Init(ctx->resource_manager(), def(),
135                                     is_writing /* use_node_name_as_default */));
136
137     CreatorCallback sparse_tensors_map_creator = [this](SparseTensorsMap** c) {
138         SparseTensorsMap* map = new SparseTensorsMap(cinfo_.name());
139         *c = map;
140         return Status::OK();
141     };
142
143     TF_RETURN_IF_ERROR(
144         cinfo_.resource_manager()->LookupOrCreate<SparseTensorsMap>(
145             cinfo_.container(), cinfo_.name(), &sparse_tensors_map_,
146             sparse_tensors_map_creator));
147
148     *sparse_tensors_map = sparse_tensors_map_;
149     return Status::OK();
150 }
151
152 private:
153     ContainerInfo cinfo_;
154
155     mutex mu_;
156     SparseTensorsMap* sparse_tensors_map_ TF_PT_GUARDED_BY(mu_);
157 };
158
159 class AddSparseToTensorsMapOp : public SparseTensorAccessingOp {
160 public:
161     explicit AddSparseToTensorsMapOp(OpKernelConstruction* context)
162         : SparseTensorAccessingOp(context) {}
163
164     void Compute(OpKernelContext* context) override {
165         const Tensor* input_indices;
166         const Tensor* input_values;
167         const Tensor* input_shape;
168         SparseTensorsMap* map;
169
170         OP_REQUIRES_OK(context, context->input("sparse_indices", &input_indices));
171         OP_REQUIRES_OK(context, context->input("sparse_values", &input_values));
172         OP_REQUIRES_OK(context, context->input("sparse_shape", &input_shape));
173         OP_REQUIRES_OK(context, GetMap(context, true /* is_writing */, &map));
174
175         OP_REQUIRES(context, TensorShapeUtils::IsMatrix(input_indices->shape()),
176                     errors::InvalidArgument(

```

```

177         "Input indices should be a matrix but received shape ",
178         input_indices->shape().DebugString());
179
180     OP_REQUIRES(context, TensorShapeUtils::IsVector(input_values->shape()),
181         errors::InvalidArgument(
182             "Input values should be a vector but received shape ",
183             input_values->shape().DebugString()));
184
185     OP_REQUIRES(context, TensorShapeUtils::IsVector(input_shape->shape()),
186         errors::InvalidArgument(
187             "Input shape should be a vector but received shape ",
188             input_shape->shape().DebugString()));
189
190     TensorShape input_shape_object;
191     OP_REQUIRES_OK(
192         context, TensorShapeUtils::MakeShape(input_shape->vec<int64_t>().data(),
193             input_shape->NumElements(),
194             &input_shape_object));
195     SparseTensor st;
196     OP_REQUIRES_OK(context, SparseTensor::Create(*input_indices, *input_values,
197         input_shape_object, &st));
198     int64_t handle;
199     OP_REQUIRES_OK(context, map->AddSparseTensor(context, st, &handle));
200
201     Tensor sparse_handle(DT_INT64, TensorShape({}));
202     auto sparse_handle_t = sparse_handle.scalar<int64_t>();
203
204     sparse_handle_t() = handle;
205
206     context->set_output(0, sparse_handle);
207 }
208 };
209
210 REGISTER_KERNEL_BUILDER(Name("AddSparseToTensorsMap").Device(DEVICE_CPU),
211     AddSparseToTensorsMapOp);
212
213 template <typename T>
214 class AddManySparseToTensorsMapOp : public SparseTensorAccessingOp {
215 public:
216     explicit AddManySparseToTensorsMapOp(OpKernelConstruction* context)
217         : SparseTensorAccessingOp(context) {}
218
219     void Compute(OpKernelContext* context) override {
220         const Tensor* input_indices;
221         const Tensor* input_values;
222         const Tensor* input_shape;
223         SparseTensorsMap* map;
224
225         OP_REQUIRES_OK(context, context->input("sparse_indices", &input_indices));

```

[illegible]

```

275
276     const int64_t N = input_shape_vec(0);
277
278     Tensor sparse_handles(DT_INT64, TensorShape({N}));
279     auto sparse_handles_t = sparse_handles.vec<int64_t>();
280
281     OP_REQUIRES_OK(context, input_st.IndicesValid());
282
283     // We can generate the output shape proto string now, for all
284     // minibatch entries.
285     TensorShape output_shape;
286     OP_REQUIRES_OK(context, TensorShapeUtils::MakeShape(
287         input_shape_vec.data() + 1,
288         input_shape->NumElements() - 1, &output_shape));
289
290     // Get groups by minibatch dimension
291     std::unordered_set<int64_t> visited;
292     sparse::GroupIterable minibatch = input_st.group({0});
293     for (const auto& subset : minibatch) {
294         const int64_t b = subset.group()[0];
295         visited.insert(b);
296         OP_REQUIRES(
297             context, b > -1 && b < N,
298             errors::InvalidArgument(
299                 "Received unexpected column 0 value in input SparseTensor: ", b,
300                 " < 0 or >= N (= ", N, ")");
301
302         const auto indices = subset.indices();
303         const auto values = subset.values<T>();
304         const int64_t num_entries = values.size();
305
306         Tensor output_indices = Tensor(DT_INT64, {num_entries, rank - 1});
307         Tensor output_values = Tensor(DataTypeToEnum<T>::value, {num_entries});
308
309         auto output_indices_t = output_indices.matrix<int64_t>();
310         auto output_values_t = output_values.vec<T>();
311
312         for (int i = 0; i < num_entries; ++i) {
313             for (int d = 1; d < rank; ++d) {
314                 output_indices_t(i, d - 1) = indices(i, d);
315             }
316             output_values_t(i) = values(i);
317         }
318
319         SparseTensor st_i;
320         OP_REQUIRES_OK(context,
321             SparseTensor::Create(output_indices, output_values,
322                 output_shape, &st_i));
323         int64_t handle;

```

```

324     OP_REQUIRES_OK(context, map->AddSparseTensor(context, st_i, &handle));
325     sparse_handles_t(b) = handle;
326 }
327
328 // Fill in any gaps; we must provide an empty ST for batch entries
329 // the grouper didn't find.
330 if (visited.size() < N) {
331     Tensor empty_indices(DT_INT64, {0, rank - 1});
332     Tensor empty_values(DataTypeToEnum<T>::value, {0});
333     SparseTensor empty_st;
334     OP_REQUIRES_OK(context, SparseTensor::Create(empty_indices, empty_values,
335                                                    output_shape, &empty_st));
336
337     for (int64_t b = 0; b < N; ++b) {
338         // We skipped this batch entry.
339         if (visited.find(b) == visited.end()) {
340             int64_t handle;
341             OP_REQUIRES_OK(context,
342                             map->AddSparseTensor(context, empty_st, &handle));
343             sparse_handles_t(b) = handle;
344         }
345     }
346 }
347
348 context->set_output(0, sparse_handles);
349 }
350 };
351
352 #define REGISTER_KERNELS(type) \
353     REGISTER_KERNEL_BUILDER(Name("AddManySparseToTensorsMap") \
354                             .Device(DEVICE_CPU) \
355                             .TypeConstraint<type>("T"), \
356                             AddManySparseToTensorsMapOp<type>)
357
358 TF_CALL_ALL_TYPES(REGISTER_KERNELS);
359 #undef REGISTER_KERNELS
360
361 template <typename T>
362 class TakeManySparseFromTensorsMapOp : public SparseTensorAccessingOp {
363 public:
364     explicit TakeManySparseFromTensorsMapOp(OpKernelConstruction* context)
365         : SparseTensorAccessingOp(context) {}
366
367     void Compute(OpKernelContext* context) override {
368         SparseTensorsMap* map = nullptr;
369         OP_REQUIRES_OK(context, GetMap(context, false /* is_writing */, &map));
370
371         const Tensor& sparse_handles = context->input(0);
372

```



```

373 OP_REQUIRES(context, TensorShapeUtils::IsVector(sparse_handles.shape()),
374             errors::InvalidArgument(
375                 "sparse_handles should be a vector but received shape ",
376                 sparse_handles.shape().DebugString()));
377
378 int64_t N = sparse_handles.shape().dim_size(0);
379
380 OP_REQUIRES(
381     context, N > 0,
382     errors::InvalidArgument("Must have at least 1 serialized SparseTensor, "
383                             "but input matrix has 0 rows"));
384
385 std::vector<Tensor> indices_to_concat;
386 std::vector<Tensor> values_to_concat;
387 std::vector<TensorShape> shapes_to_concat;
388
389 const auto& sparse_handles_t = sparse_handles.vec<int64_t>();
390
391 std::vector<SparseTensor> sparse_tensors;
392
393 OP_REQUIRES_OK(context, map->RetrieveAndClearSparseTensors(
394     context, sparse_handles_t, &sparse_tensors));
395
396 for (int64_t i = 0; i < N; ++i) {
397     const SparseTensor& st = sparse_tensors[i];
398     const Tensor& output_indices = st.indices();
399     const Tensor& output_values = st.values();
400     const auto output_shape = st.shape();
401
402     OP_REQUIRES(context, TensorShapeUtils::IsMatrix(output_indices.shape()),
403                 errors::InvalidArgument(
404                     "Expected sparse_handles[" + i,
405                     "]" to represent an index matrix but received shape ",
406                     output_indices.shape().DebugString()));
407     OP_REQUIRES(context, TensorShapeUtils::IsVector(output_values.shape()),
408                 errors::InvalidArgument(
409                     "Expected sparse_handles[" + i,
410                     "]" to represent a values vector but received shape ",
411                     output_values.shape().DebugString()));
412     OP_REQUIRES(
413         context, DataTypeToEnum<T>::value == output_values.dtype(),
414         errors::InvalidArgument(
415             "Requested SparseTensor of type ",
416             DataTypeString(DataTypeToEnum<T>::value), " but SparseTensor[" + i,
417             "].values.dtype() == ", DataTypeString(output_values.dtype())));
418
419     int64_t num_entries = output_indices.dim_size(0);
420     OP_REQUIRES(context, num_entries == output_values.dim_size(0),
421                 errors::InvalidArgument(

```

```

422         "Expected row counts of SparseTensor[" , i,
423         "].indices and SparseTensor[" , i,
424         "].values to match but they do not: " , num_entries,
425         " vs. " , output_values.dim_size(0)));
426 int rank = output_indices.dim_size(1);
427 OP_REQUIRES(
428     context, rank == output_shape.size(),
429     errors::InvalidArgument("Expected column counts of SparseTensor[" , i,
430                             "].indices to match size of SparseTensor[" , i,
431                             "].shape "
432                             "but they do not: " ,
433                             rank, " vs. " , output_shape.size()));
434
435 // Now we expand each SparseTensors' indices and shape by
436 // prefixing a dimension
437 Tensor expanded_indices(
438     DT_INT64, TensorShape({num_entries, 1 + output_indices.dim_size(1)}));
439 Tensor expanded_shape(DT_INT64, TensorShape({1 + rank}));
440 const auto& output_indices_t = output_indices.matrix<int64_t>();
441 auto expanded_indices_t = expanded_indices.matrix<int64_t>();
442 auto expanded_shape_t = expanded_shape.vec<int64_t>();
443 expanded_indices_t.chip<1>(0).setZero();
444 Eigen::DSizes<Eigen::DenseIndex, 2> indices_start(0, 1);
445 Eigen::DSizes<Eigen::DenseIndex, 2> indices_sizes(num_entries, rank);
446 expanded_indices_t.slice(indices_start, indices_sizes) = output_indices_t;
447 expanded_shape_t(0) = 1;
448 // TODO: copy shape from TensorShape to &expanded_shape_t(1)
449 // std::copy_n(&output_shape_t(0), rank, &expanded_shape_t(1));
450 for (int i = 0; i < rank; ++i) {
451     expanded_shape_t(i + 1) = output_shape[i];
452 }
453 TensorShape expanded_tensor_shape(expanded_shape_t);
454
455 indices_to_concat.push_back(std::move(expanded_indices));
456 values_to_concat.push_back(output_values);
457 shapes_to_concat.push_back(std::move(expanded_tensor_shape));
458 }
459
460 int rank = -1;
461 for (int i = 0; i < N; ++i) {
462     if (rank < 0) rank = shapes_to_concat[i].dims();
463     OP_REQUIRES(context, rank == shapes_to_concat[i].dims(),
464                 errors::InvalidArgument(
465                     "Inconsistent rank across SparseTensors: rank prior to "
466                     "SparseTensor[" ,
467                     i, "]" was: " , rank, " but rank of SparseTensor[" , i,
468                     "]" is: " , shapes_to_concat[i].dims()));
469 }
470

```

```

471 // SparseTensor::Concat requires consistent shape for all but the
472 // primary order dimension (dimension 0 in this case). So we get
473 // the maximum value across all the input SparseTensors for each
474 // dimension and use that.
475 TensorShape preconcat_shape(shapes_to_concat[0]);
476 for (int i = 0; i < N; ++i) {
477     for (int d = 0; d < rank; ++d) {
478         preconcat_shape.set_dim(d, std::max(preconcat_shape.dim_size(d),
479                                             shapes_to_concat[i].dim_size(d)));
480     }
481 }
482
483 // Dimension 0 is the primary dimension.
484 gtl::InlinedVector<int64_t, 8> std_order(rank);
485 std::iota(std_order.begin(), std_order.end(), 0);
486
487 std::vector<SparseTensor> tensors_to_concat;
488 tensors_to_concat.reserve(N);
489 for (int i = 0; i < N; ++i) {
490     SparseTensor tensor;
491     OP_REQUIRES_OK(context,
492                     SparseTensor::Create(std::move(indices_to_concat[i]),
493                                         std::move(values_to_concat[i]),
494                                         preconcat_shape, std_order, &tensor));
495     tensors_to_concat.push_back(std::move(tensor));
496 }
497
498 auto output = SparseTensor::Concat<T>(tensors_to_concat);
499 Tensor final_output_shape(DT_INT64, TensorShape({output.dims()}));
500
501 std::copy_n(output.shape().data(), output.dims(),
502             final_output_shape.vec<int64_t>().data());
503
504 context->set_output(0, output.indices());
505 context->set_output(1, output.values());
506 context->set_output(2, final_output_shape);
507 }
508 };
509
510 #define REGISTER_KERNELS(type) \
511     REGISTER_KERNEL_BUILDER(Name("TakeManySparseFromTensorsMap") \
512                             .Device(DEVICE_CPU) \
513                             .TypeConstraint<type>("dtype"), \
514                             TakeManySparseFromTensorsMapOp<type>)
515
516 TF_CALL_ALL_TYPES(REGISTER_KERNELS);
517 #undef REGISTER_KERNELS
518
519 } // namespace tensorflow

```

