

a1320ec1ea ▾

...


tensorflow / tensorflow / core / common\_runtime / immutable\_executor\_state.cc



qcfish Fix a NPE issue in invalid Exit op. Now it will report an error inste... .. ✖

History

6 contributors



380 lines (336 sloc) | 13 KB

...

```

1  /* Copyright 2015 The TensorFlow Authors. All Rights Reserved.
2
3  Licensed under the Apache License, Version 2.0 (the "License");
4  you may not use this file except in compliance with the License.
5  You may obtain a copy of the License at
6
7      http://www.apache.org/licenses/LICENSE-2.0
8
9  Unless required by applicable law or agreed to in writing, software
10 distributed under the License is distributed on an "AS IS" BASIS,
11 WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
12 See the License for the specific language governing permissions and
13 limitations under the License.
14 =====*/
15
16 #include "tensorflow/core/common_runtime/immutable_executor_state.h"
17
18 #include "absl/memory/memory.h"
19 #include "tensorflow/core/framework/function.h"
20 #include "tensorflow/core/framework/metrics.h"
21 #include "tensorflow/core/framework/node_def_util.h"
22 #include "tensorflow/core/graph/edgeset.h"
23 #include "tensorflow/core/graph/graph.h"
24 #include "tensorflow/core/graph/graph_node_util.h"
25 #include "tensorflow/core/platform/errors.h"
26 #include "tensorflow/core/platform/logging.h"
27
28 namespace tensorflow {
29
```

```

30 namespace {
31 bool IsInitializationOp(const Node* node) {
32     return node->op_def().allows_uninitialized_input();
33 }
34 } // namespace
35
36 ImmutableExecutorState::~ImmutableExecutorState() {
37     for (int32_t i = 0; i < gview_.num_nodes(); i++) {
38         NodeItem* item = gview_.node(i);
39         if (item != nullptr) {
40             params_.delete_kernel(item->kernel);
41         }
42     }
43 }
44
45 namespace {
46 void GetMaxPendingCounts(const Node* n, size_t* max_pending,
47                          size_t* max_dead_count) {
48     const size_t num_in_edges = n->in_edges().size();
49     size_t initial_count;
50     if (IsMerge(n)) {
51         // merge waits all control inputs so we initialize the pending
52         // count to be the number of control edges.
53         int32_t num_control_edges = 0;
54         for (const Edge* edge : n->in_edges()) {
55             if (edge->IsControlEdge()) {
56                 num_control_edges++;
57             }
58         }
59         // Use bit 0 to indicate if we are waiting for a ready live data input.
60         initial_count = 1 + (num_control_edges << 1);
61     } else {
62         initial_count = num_in_edges;
63     }
64
65     *max_pending = initial_count;
66     *max_dead_count = num_in_edges;
67 }
68 } // namespace
69
70 ImmutableExecutorState::FrameInfo* ImmutableExecutorState::EnsureFrameInfo(
71     const string& fname) {
72     auto iter = frame_info_.find(fname);
73     if (iter != frame_info_.end()) {
74         return iter->second.get();
75     } else {
76         auto frame_info = absl::make_unique<FrameInfo>(fname);
77         absl::string_view fname_view = frame_info->name;
78         auto emplace_result =

```

```

79     frame_info_.emplace(fname_view, std::move(frame_info));
80     return emplace_result.first->second.get();
81 }
82 }
83
84 Status ImmutableExecutorState::Initialize(const Graph& graph) {
85     TF_RETURN_IF_ERROR(gview_.Initialize(&graph));
86
87     // Build the information about frames in this subgraph.
88     ControlFlowInfo cf_info;
89     TF_RETURN_IF_ERROR(BuildControlFlowInfo(&graph, &cf_info));
90
91     for (auto& it : cf_info.unique_frame_names) {
92         EnsureFrameInfo(it)->nodes =
93             absl::make_unique<std::vector<const NodeItem*>>();
94     }
95     root_frame_info_ = frame_info_[""].get();
96
97     pending_ids_.resize(gview_.num_nodes());
98
99     // Preprocess every node in the graph to create an instance of op
100    // kernel for each node.
101    requires_control_flow_ = false;
102    for (const Node* n : graph.nodes()) {
103        if (IsSink(n)) continue;
104        if (IsSwitch(n) || IsMerge(n) || IsEnter(n) || IsExit(n)) {
105            requires_control_flow_ = true;
106        } else if (IsRecv(n)) {
107            // A Recv node from a different device may produce dead tensors from
108            // non-local control-flow nodes.
109            //
110            // TODO(mrry): Track whether control flow was present in the
111            // pre-partitioned graph, and enable the caller (e.g.
112            // `DirectSession`) to relax this constraint.
113            string send_device;
114            string recv_device;
115            TF_RETURN_IF_ERROR(GetNodeAttr(n->attrs(), "send_device", &send_device));
116            TF_RETURN_IF_ERROR(GetNodeAttr(n->attrs(), "recv_device", &recv_device));
117            if (send_device != recv_device) {
118                requires_control_flow_ = true;
119            }
120        }
121
122        const int id = n->id();
123        const string& frame_name = cf_info.frame_names[id];
124        FrameInfo* frame_info = EnsureFrameInfo(frame_name);
125
126        NodeItem* item = gview_.node(id);
127        item->node_id = id;

```

```

128
129     item->input_start = frame_info->total_inputs;
130     frame_info->total_inputs += n->num_inputs();
131
132     Status s = params_.create_kernel(n->properties(), &item->kernel);
133     if (!s.ok()) {
134         item->kernel = nullptr;
135         s = AttachDef(s, *n);
136         return s;
137     }
138     CHECK(item->kernel);
139     item->kernel_is_async = (item->kernel->AsAsync() != nullptr);
140     item->is_merge = IsMerge(n);
141     item->is_any_consumer_merge_or_control_trigger = false;
142     for (const Node* consumer : n->out_nodes()) {
143         if (IsMerge(consumer) || IsControlTrigger(consumer)) {
144             item->is_any_consumer_merge_or_control_trigger = true;
145             break;
146         }
147     }
148     const Tensor* const_tensor = item->kernel->const_tensor();
149     if (const_tensor) {
150         // Hold onto a shallow copy of the constant tensor in `*this` so that the
151         // reference count does not drop to 1. This prevents the constant tensor
152         // from being forwarded, and its buffer reused.
153         const_tensors_.emplace_back(*const_tensor);
154     }
155     item->const_tensor = const_tensor;
156     item->is_noop = (item->kernel->type_string_view() == "NoOp");
157     item->is_enter = IsEnter(n);
158     if (item->is_enter) {
159         bool is_constant_enter;
160         TF_RETURN_IF_ERROR(
161             GetNodeAttr(n->attrs(), "is_constant", &is_constant_enter));
162         item->is_constant_enter = is_constant_enter;
163
164         string frame_name;
165         TF_RETURN_IF_ERROR(GetNodeAttr(n->attrs(), "frame_name", &frame_name));
166         FrameInfo* frame_info = frame_info_[frame_name].get();
167
168         int parallel_iterations;
169         TF_RETURN_IF_ERROR(
170             GetNodeAttr(n->attrs(), "parallel_iterations", &parallel_iterations));
171
172         if (frame_info->parallel_iterations == -1) {
173             frame_info->parallel_iterations = parallel_iterations;
174         } else if (frame_info->parallel_iterations != parallel_iterations) {
175             LOG(WARNING) << "Loop frame \"" << frame_name
176                 << "\" had two different values for parallel_iterations: "

```

```

177         << frame_info->parallel_iterations << " vs. "
178         << parallel_iterations << ".";
179     }
180
181     if (enter_frame_info_.size() <= id) {
182         enter_frame_info_.resize(id + 1);
183     }
184     enter_frame_info_[id] = frame_info;
185 } else {
186     item->is_constant_enter = false;
187 }
188 item->is_exit = IsExit(n);
189 item->is_control_trigger = IsControlTrigger(n);
190 item->is_source = IsSource(n);
191 item->is_enter_exit_or_next_iter =
192     (IsEnter(n) || IsExit(n) || IsNextIteration(n));
193 item->is_transfer_node = IsTransferNode(n);
194 item->is_initialization_op = IsInitializationOp(n);
195 item->is_recv_or_switch = IsRecv(n) || IsSwitch(n);
196 item->is_next_iteration = IsNextIteration(n);
197 item->is_distributed_communication = IsDistributedCommunication(n);
198
199 // Compute the maximum values we'll store for this node in the
200 // pending counts data structure, and allocate a handle in
201 // that frame's pending counts data structure that has enough
202 // space to store these maximal count values.
203 size_t max_pending, max_dead;
204 GetMaxPendingCounts(n, &max_pending, &max_dead);
205 pending_ids_[id] =
206     frame_info->pending_counts_layout.CreateHandle(max_pending, max_dead);
207
208 // See if this node is a root node, and if so, add item to root_nodes_.
209 if (n->in_edges().empty()) {
210     root_nodes_.push_back(item);
211 }
212
213 // Initialize static information about the frames in the graph.
214 frame_info->nodes->push_back(item);
215 if (item->is_enter) {
216     string enter_name;
217     TF_RETURN_IF_ERROR(GetNodeAttr(n->attrs(), "frame_name", &enter_name));
218     EnsureFrameInfo(enter_name)->input_count++;
219 }
220
221 // Record information about whether each output of the op is used.
222 std::unique_ptr<bool[]> outputs_required(new bool[n->num_outputs()]);
223 std::fill(&outputs_required[0], &outputs_required[n->num_outputs()], false);
224 int32_t unused_outputs = n->num_outputs();
225 for (const Edge* e : n->out_edges()) {

```

```

226     if (IsSink(e->dst())) continue;
227     if (e->src_output() >= 0) {
228         if (!outputs_required[e->src_output()]) {
229             --unused_outputs;
230             outputs_required[e->src_output()] = true;
231         }
232     }
233 }
234 if (unused_outputs > 0) {
235     for (int i = 0; i < n->num_outputs(); ++i) {
236         if (!outputs_required[i]) {
237             metrics::RecordUnusedOutput(n->type_string());
238         }
239     }
240     item->outputs_required = std::move(outputs_required);
241 }
242 }
243
244 // Rewrite each `EdgeInfo::input_slot` member to refer directly to the input
245 // location.
246 for (const Node* n : graph.nodes()) {
247     if (IsSink(n)) continue;
248     const int id = n->id();
249     NodeItem* item = gview_.node(id);
250
251     for (EdgeInfo& e : item->mutable_output_edges()) {
252         const int dst_id = e.dst_id;
253         NodeItem* dst_item = gview_.node(dst_id);
254         e.input_slot += dst_item->input_start;
255     }
256 }
257
258 // Initialize PendingCounts only after pending_ids_[node.id] is initialized
259 // for all nodes.
260 InitializePending(&graph, cf_info);
261 return gview_.SetAllocAttrs(&graph, params_.device);
262 }
263
264 namespace {
265 // If a Node has been marked to use a ScopedAllocator x for output i, then
266 // sc_attr will contain the subsequence (i, x) at an even offset. This function
267 // extracts and transfers that ScopedAllocator id to alloc_attr. For now, we
268 // only allow one ScopedAllocator use per Node.
269 bool ExtractScopedAllocatorAttr(const std::vector<int>& sc_attr,
270                                int output_index,
271                                AllocatorAttributes* alloc_attr) {
272     DCHECK_LE(2, sc_attr.size());
273     for (int i = 0; i < sc_attr.size(); i += 2) {
274         if (sc_attr[i] == output_index) {

```

```

275     CHECK_EQ(alloc_attr->scope_id, 0);
276     alloc_attr->scope_id = sc_attr[i + 1];
277     return true;
278 }
279 }
280 return false;
281 }
282 } // namespace
283
284 Status ImmutableExecutorState::BuildControlFlowInfo(const Graph* g,
285                                                    ControlFlowInfo* cf_info) {
286     const int num_nodes = g->num_node_ids();
287     cf_info->frame_names.resize(num_nodes);
288     std::vector<Node*> parent_nodes;
289     parent_nodes.resize(num_nodes);
290     std::vector<bool> visited;
291     visited.resize(num_nodes);
292
293     string frame_name;
294     std::deque<Node*> ready;
295
296     // Initialize with the root nodes.
297     for (Node* n : g->nodes()) {
298         if (n->in_edges().empty()) {
299             visited[n->id()] = true;
300             cf_info->unique_frame_names.insert(frame_name);
301             ready.push_back(n);
302         }
303     }
304
305     while (!ready.empty()) {
306         Node* curr_node = ready.front();
307         int curr_id = curr_node->id();
308         ready.pop_front();
309
310         Node* parent = nullptr;
311         if (IsEnter(curr_node)) {
312             // Enter a child frame.
313             TF_RETURN_IF_ERROR(
314                 GetNodeAttr(curr_node->attrs(), "frame_name", &frame_name));
315             parent = curr_node;
316         } else if (IsExit(curr_node)) {
317             // Exit to the parent frame.
318             parent = parent_nodes[curr_id];
319             if (!parent) {
320                 return errors::InvalidArgument(
321                     "Invalid Exit op: Cannot find a corresponding Enter op.");
322             }
323             frame_name = cf_info->frame_names[parent->id()];

```

```

324     parent = parent_nodes[parent->id()];
325 } else {
326     parent = parent_nodes[curr_id];
327     frame_name = cf_info->frame_names[curr_id];
328 }
329
330 for (const Edge* out_edge : curr_node->out_edges()) {
331     Node* out = out_edge->dst();
332     if (IsSink(out)) continue;
333     const int out_id = out->id();
334
335     // Add to ready queue if not visited.
336     bool is_visited = visited[out_id];
337     if (!is_visited) {
338         ready.push_back(out);
339         visited[out_id] = true;
340
341         // Process the node 'out'.
342         cf_info->frame_names[out_id] = frame_name;
343         parent_nodes[out_id] = parent;
344         cf_info->unique_frame_names.insert(frame_name);
345     }
346 }
347 }
348
349 return Status::OK();
350 }
351
352 void ImmutableExecutorState::InitializePending(const Graph* graph,
353                                               const ControlFlowInfo& cf_info) {
354     for (auto& it : cf_info.unique_frame_names) {
355         FrameInfo* finfo = EnsureFrameInfo(it);
356         DCHECK_EQ(finfo->pending_counts.get(), nullptr);
357         finfo->pending_counts =
358             absl::make_unique<PendingCounts>(finfo->pending_counts_layout);
359     }
360
361     if (!requires_control_flow_) {
362         atomic_pending_counts_.reset(new std::atomic<int32>[gview_.num_nodes()]);
363         std::fill(atomic_pending_counts_.get(),
364                 atomic_pending_counts_.get() + gview_.num_nodes(), 0);
365     }
366
367     for (const Node* n : graph->nodes()) {
368         if (IsSink(n)) continue;
369         const int id = n->id();
370         const string& name = cf_info.frame_names[id];
371         size_t max_pending, max_dead;
372         GetMaxPendingCounts(n, &max_pending, &max_dead);

```



```
373     auto& counts = EnsureFrameInfo(name)->pending_counts;
374     counts->set_initial_count(pending_ids_[id], max_pending);
375     if (!requires_control_flow_) {
376         atomic_pending_counts_[id] = max_pending;
377     }
378 }
379 }
380 } // namespace tensorflow
```