

a1320ec1ea ▾

...

tensorflow / tensorflow / core / grappler / optimizers / dependency\_optimizer.cc



jpienaar Add default feedback in preparation for removal ... ✓

History

9 contributors



791 lines (736 sloc) | 30.7 KB

...

```

1  /* Copyright 2017 The TensorFlow Authors. All Rights Reserved.
2
3  Licensed under the Apache License, Version 2.0 (the "License");
4  you may not use this file except in compliance with the License.
5  You may obtain a copy of the License at
6
7      http://www.apache.org/licenses/LICENSE-2.0
8
9  Unless required by applicable law or agreed to in writing, software
10 distributed under the License is distributed on an "AS IS" BASIS,
11 WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
12 See the License for the specific language governing permissions and
13 limitations under the License.
14 =====*/
15
16 #include "tensorflow/core/grappler/optimizers/dependency_optimizer.h"
17
18 #include <unordered_set>
19
20 #include "absl/container/flat_hash_map.h"
21 #include "tensorflow/core/framework/node_def.pb.h"
22 #include "tensorflow/core/framework/node_def_util.h"
23 #include "tensorflow/core/framework/op.h"
24 #include "tensorflow/core/grappler/costs/graph_properties.h"
25 #include "tensorflow/core/grappler/grappler_item.h"
26 #include "tensorflow/core/grappler/op_types.h"
27 #include "tensorflow/core/grappler/optimizers/constant_folding.h"
28 #include "tensorflow/core/grappler/utils.h"
29 #include "tensorflow/core/grappler/utils/topological_sort.h"

```

```

30 #include "tensorflow/core/lib/core/errors.h"
31 #include "tensorflow/core/lib/core/stringpiece.h"
32 #include "tensorflow/core/lib/gtl/inlined_vector.h"
33 #include "tensorflow/core/lib/strings/str_util.h"
34 #include "tensorflow/core/lib/strings/strcat.h"
35 #include "tensorflow/core/util/device_name_utils.h"
36
37 namespace tensorflow {
38 namespace grappler {
39
40 namespace {
41
42 bool RemoveControlInput(NodeDef* node, const string& control_input_to_remove,
43                           NodeMap* node_map) {
44     for (int pos = node->input_size() - 1; pos >= 0; --pos) {
45         const string& input = node->input(pos);
46         if (input[0] != '^') break;
47         if (input == control_input_to_remove) {
48             node->mutable_input()->SwapElements(pos, node->input_size() - 1);
49             node->mutable_input()->RemoveLast();
50             node_map->RemoveOutput(NodeName(input), node->name());
51             return true;
52         }
53     }
54     return false;
55 }
56
57 } // namespace
58

```

...

```

59 bool DependencyOptimizer::SafeToRemoveIdentity(const NodeDef& node) const {
60     if (!IsIdentity(node) && !IsIdentityN(node)) {
61         return true;
62     }
63
64     if (nodes_to_preserve_.find(node.name()) != nodes_to_preserve_.end()) {
65         return false;
66     }
67     if (!fetch_nodes_known_) {
68         // The output values of this node may be needed.
69         return false;
70     }
71
72     if (node.input_size() < 1) {
73         // Node lacks input, is invalid
74         return false;
75     }
76
77     const NodeDef* input = node_map->GetNode(NodeName(node.input(0)));
78     CHECK(input != nullptr) << "node = " << node.name()

```

```

79         << " input = " << node.input(0);
80     // Don't remove Identity nodes corresponding to Variable reads or following
81     // Recv.
82     if (IsVariable(*input) || IsRecv(*input)) {
83         return false;
84     }
85     for (const auto& consumer : node_map_->GetOutputs(node.name())) {
86         if (node.input_size() > 1 && (IsRetVal(*consumer) || IsMerge(*consumer))) {
87             return false;
88         }
89         if (IsSwitch(*input)) {
90             for (const string& consumer_input : consumer->input()) {
91                 if (consumer_input == AsControlDependency(node.name())) {
92                     return false;
93                 }
94             }
95         }
96     }
97     return true;
98 }
99
100 bool DependencyOptimizer::SafeToConvertToNoOp(const NodeDef& node) const {
101     if (HasRegularOutputs(node, *node_map_)) {
102         // The output values of this node may be needed.
103         VLOG(3) << "Not safe to convert '" << node.name()
104             << " to NoOp. Node has outputs.";
105         return false;
106     }
107     if (!fetch_nodes_known_) {
108         VLOG(3) << "Not safe to convert '" << node.name()
109             << " to NoOp. Fetches unknown.";
110         return false;
111     }
112     if (nodes_to_preserve_.find(node.name()) != nodes_to_preserve_.end()) {
113         VLOG(3) << "Not safe to convert to NoOp: " << node.name()
114             << " is in preserve set.";
115         return false;
116     }
117     if (IsMerge(node) || IsSwitch(node) || ModifiesFrameInfo(node)) {
118         VLOG(3) << "Not safe to convert '" << node.name()
119             << " to NoOp. Node modifies frame info.";
120         return false;
121     }
122     // Ops reading variables are marked as stateful, but are safe to remove if
123     // redundant.
124     static const absl::flat_hash_set<string>* gather_ops =
125         new absl::flat_hash_set<string>{"Gather", "GatherV2", "GatherNd",
126             "ResourceGather", "ResourceGatherNd"};
127     const bool is_variable_read =

```

```

128     IsReadVariableOp(node) || IsReadVariablesOp(node) ||
129     gather_ops->find(node.op()) != gather_ops->end();
130 if (!is_variable_read && !IsFreeOfSideEffect(node)) {
131     VLOG(3) << "Not safe to convert '" << node.name()
132         << " to NoOp. Node has side effect.";
133     return false;
134 }
135 if (node.op().rfind("Submodel", 0) == 0) {
136     return false;
137 }
138 const OpDef* op_def = nullptr;
139 Status status = OpRegistry::Global()->LookUpOpDef(node.op(), &op_def);
140 if (!status.ok() || op_def->output_arg_size() == 0) {
141     return false;
142 }
143 const std::unordered_set<string> do_not_rewrite_ops{
144     "Assert",      "CheckNumerics",      "_Retval",
145     "_Arg",        "_ParallelConcatUpdate", "TPUExecute",
146     "TPUCompile", "ControlTrigger"};
147 if (do_not_rewrite_ops.find(node.op()) != do_not_rewrite_ops.end()) {
148     return false;
149 }
150 if (!SafeToRemoveIdentity(node)) {
151     return false;
152 }
153 return true;
154 }
155
156 int DependencyOptimizer::NumEdgesIfBypassed(
157     const NodeDef& node, const std::vector<NodeDef*>& output_nodes) const {
158     const bool is_multi_input_identity_n =
159         IsIdentityN(node) && !IsIdentityNSingleInput(node);
160     const int num_outputs = output_nodes.size();
161     const int num_inputs = node.input_size();
162
163     if (is_multi_input_identity_n) {
164         // multi-input identity_n with input/output control dependencies will likely
165         // increase number of edges after optimization.
166         int num_edges_if_bypassed(0);
167         for (const string& input_node_name : node.input()) {
168             if (IsControlInput(input_node_name)) {
169                 num_edges_if_bypassed += num_outputs;
170             } else {
171                 ++num_edges_if_bypassed;
172             }
173         }
174
175         for (auto consumer : output_nodes) {
176             for (int j = 0; j < consumer->input_size(); ++j) {

```

```

177     const TensorId consumer_input = ParseTensorName(consumer->input(j));
178     if (consumer_input.node() == node.name()) {
179         if (IsControlInput(consumer_input)) {
180             num_edges_if_bypassed += num_inputs;
181         } else {
182             ++num_edges_if_bypassed;
183         }
184     }
185 }
186 }
187 return num_edges_if_bypassed;
188 } else {
189     return num_inputs * num_outputs;
190 }
191 }
192
193 bool DependencyOptimizer::BypassingNodeIsBeneficial(
194     const NodeDef& node, const std::vector<NodeDef*>& input_nodes,
195     const std::vector<NodeDef*>& output_nodes) const {
196     const bool is_identity = IsIdentity(node) || IsIdentityNSingleInput(node);
197     const bool is_multi_input_identity_n =
198         IsIdentityN(node) && !IsIdentityNSingleInput(node);
199     const int num_outputs = output_nodes.size();
200     const int num_inputs = node.input_size();
201
202     if (NumEdgesIfBypassed(node, output_nodes) > num_inputs + num_outputs) {
203         return false;
204     }
205
206     // Make sure that we don't increase the number of edges that cross
207     // device boundaries.
208     if ((num_inputs == 1 && num_outputs > 1 &&
209         input_nodes[0]->device() != node.device()) ||
210         (num_inputs > 1 && num_outputs == 1 &&
211         output_nodes[0]->device() != node.device())) {
212         return false;
213     }
214
215     // TODO(rmlarsen): Not all device crossings are equally expensive.
216     // Assign a cost to each based on device affinity and compute a
217     // cost before and after.
218     const string& node_dev = node.device();
219     int num_cross_in = 0;
220     for (NodeDef* input_node : input_nodes) {
221         num_cross_in += static_cast<int>(input_node->device() != node_dev);
222     }
223     int num_cross_out = 0;
224     for (NodeDef* output_node : output_nodes) {
225         num_cross_out += static_cast<int>(output_node->device() != node_dev);

```

```

226     }
227
228     // Make sure we do not increase the number of device crossings.
229     const int num_cross_before = num_cross_in + num_cross_out;
230     int num_cross_after = 0;
231     for (NodeDef* input_node : input_nodes) {
232         for (NodeDef* output_node : output_nodes) {
233             num_cross_after +=
234                 static_cast<int>(input_node->device() != output_node->device());
235         }
236     }
237     if (num_cross_after > num_cross_before) {
238         return false;
239     }
240
241     if ((is_identity || is_multi_input_identity_n) && num_cross_in > 0 &&
242         num_cross_out > 0 && num_cross_after > 0) {
243         // This identity node follows a device crossing, so it might be
244         // following a _Recv node after partitioning. Do not remove such nodes,
245         // unless they only have consumers on the same device as themselves.
246         return false;
247     }
248
249     return true;
250 }
251
252 void DependencyOptimizer::OptimizeNode(int node_idx,
253                                       SetVector<int>* nodes_to_simplify,
254                                       std::set<int>* nodes_to_delete) {
255     NodeDef* node = optimized_graph->mutable_node(node_idx);
256     const bool is_noop = IsNoOp(*node);
257     const bool is_identity = IsIdentity(*node) || IsIdentityNSingleInput(*node);
258     const bool is_multi_input_identity =
259         IsIdentityN(*node) && !IsIdentityNSingleInput(*node);
260     const string node_name = node->name();
261     // Constant nodes with no input control dependency are always executed early,
262     // so we can prune all their output control dependencies.
263     if (IsConstant(*node) && node->input_size() == 0) {
264         const auto output_nodes = node_map->GetOutputs(node_name);
265         for (NodeDef* fanout : output_nodes) {
266             bool optimize_fanout = false;
267             bool data_connection = false;
268             for (int i = fanout->input_size() - 1; i >= 0; --i) {
269                 const TensorId input_tensor = ParseTensorName(fanout->input(i));
270                 if (input_tensor.node() == node_name) {
271                     if (input_tensor.index() < 0) {
272                         fanout->mutable_input()->SwapElements(i, fanout->input_size() - 1);
273                         fanout->mutable_input()->RemoveLast();
274                         optimize_fanout = true;

```

```

275         } else {
276             data_connection = true;
277         }
278     }
279 }
280 if (optimize_fanout) {
281     nodes_to_simplify->PushBack(node_to_idx_[fanout]);
282     if (!data_connection) {
283         node_map->RemoveOutput(node_name, fanout->name());
284     }
285 }
286 }
287 if (node_map->GetOutputs(node_name).empty() && fetch_nodes_known_ &&
288     nodes_to_preserve_.find(node_name) == nodes_to_preserve_.end()) {
289     // Mark the node for deletion.
290     nodes_to_delete->insert(node_to_idx_[node]);
291 }
292 return;
293 }
294
295 // Change ops that only have control dependencies as outputs to NoOps.
296 if (!is_noop && SafeToConvertToNoOp(*node)) {
297     VLOG(2) << "***** Replacing " << node_name << " (" << node->op()
298         << ") with NoOp.";
299     // The outputs of this node are not consumed. Replace its inputs with
300     // control dependencies and replace the op itself with the NoOp op.
301     std::unordered_set<string> ctrl_inputs;
302     int pos = 0;
303     while (pos < node->input_size()) {
304         const string old_input = node->input(pos);
305         if (IsControlInput(old_input)) {
306             if (!ctrl_inputs.insert(old_input).second) {
307                 // We found a duplicate control input. Remove it.
308                 node->mutable_input()->SwapElements(pos, node->input_size() - 1);
309                 node->mutable_input()->RemoveLast();
310             } else {
311                 ++pos;
312             }
313             continue;
314         }
315         // Replace a normal input with a control input.
316         const string ctrl_input = ConstantFolding::AddControlDependency(
317             old_input, optimized_graph_, node_map_.get());
318         ctrl_inputs.insert(ctrl_input);
319         node->set_input(pos, ctrl_input);
320         node_map->UpdateInput(node_name, old_input, ctrl_input);
321         const NodeDef* old_input_node = node_map->GetNode(old_input);
322         nodes_to_simplify->PushBack(node_to_idx_[old_input_node]);
323         ++pos;

```

```

324     }
325     node->set_op("NoOp");
326     EraseRegularNodeAttributes(node);
327     DedupControlInputs(node);
328     nodes_to_simplify->PushBack(node_to_idx_[node]);
329     return;
330 }
331
332 // Remove NoOp nodes if the product of their fan-in and fan-out is less than
333 // or equal to the sum of the fan-in and fan-out. The non-trivial rewrites
334 // take the following form:
335 //
336 // Case a)
337 //   x --^> +-----+          x --^> +---+
338 //   y --^> | NoOp | --^> a  ==> y --^> | a |
339 //   ...   |      |          ...   |   |
340 //   z --^> +-----+          z --^> +---+
341 //
342 // Case b)
343 //           +-----+ --^> a      +---+ --^> a
344 //   x --^> | NoOp | --^> b  ==>  | x | --^> b
345 //           |      | ...          |   | ...
346 //           +-----+ --^> c      +---+ --^> c
347 // Case c)
348 //           +-----+          x ---^> a
349 //   x --^> | NoOp | --^> a  ==>  \ /
350 //   y --^> |      | --^> b          /\
351 //           +-----+          y ---^> b
352 //
353 // We only apply this optimization if we don't increase the number of control
354 // edges across device boundaries, e.g. in cases a) and b) if NoOp and
355 // a and x, respectively, are on the same device. Control edges across device
356 // boundaries require inter-device communication (Send/Recv pairs to be
357 // inserted in the graph), which is very costly.
358 //
359 // We also remove identity nodes, subject to the same constraints on number of
360 // resulting control edges and device boundary crossings:
361 //
362 // Case a)
363 //           +-----+ ---> a      +---+ ---> a
364 //   x --> | Identity | --^> b  ==>  | x | --^> b
365 //           |      | ...          |   | ...
366 //           +-----+ --^> c      +---+ --^> c
367 //
368 // Case b)
369 //   x ---> +-----+ ---> a      x ---> +---+
370 //   y --^> | Identity |          ==> y --^> | a |
371 //   ...   |      |          ...   |   |
372 //   z --^> +-----+          z --^> +---+

```





```

422         node_map_->UpdateInput(consumer->name(),
423                                string(old_input.node()), new_input);
424         consumer->set_input(j, new_input);
425     } else if (old_input.index() == -1) {
426         // Control dependency
427         new_input = AsControlDependency(NodeName(input_to_forward));
428         node_map_->UpdateInput(consumer->name(),
429                                string(old_input.node()), new_input);
430         consumer->set_input(j, new_input);
431     }
432 }
433 }
434 updated_consumer = true;
435 } else {
436     // Forward dependency from input to consumer if it doesn't already
437     // depend on it.
438     if (node_map_->GetOutputs(input->name()).count(consumer) == 0) {
439         consumer->add_input(AsControlDependency(input->name()));
440         node_map_->AddOutput(input->name(), consumer->name());
441         nodes_to_simplify->PushBack(node_to_idx_[input]);
442         updated_consumer = true;
443     }
444 }
445 }
446 updated_consumer |= RemoveControlInput(
447     consumer, AsControlDependency(node_name), node_map_.get());
448 if (updated_consumer) {
449     nodes_to_simplify->PushBack(node_to_idx_[consumer]);
450 }
451 VLOG(2) << "consumer after:\n" << consumer->DebugString();
452 }
453 node_map_->RemoveOutputs(node_name);
454 if (fetch_nodes_known_ &&
455     nodes_to_preserve_.find(node_name) == nodes_to_preserve_.end()) {
456     // Mark the node for deletion.
457     nodes_to_delete->insert(node_idx);
458
459     // Disconnect the node from its inputs to enable further optimizations.
460     node_map_->RemoveInputs(node_name);
461     node->clear_input();
462 }
463 }
464 }
465
466 void DependencyOptimizer::CleanControlInputs() {
467     for (int i = 0; i < optimized_graph_->node_size(); ++i) {
468         DedupControlInputs(optimized_graph_->mutable_node(i));
469     }
470 }

```

```

471
472 Status DependencyOptimizer::OptimizeDependencies() {
473     SetVector<int> nodes_to_simplify;
474     std::set<int> nodes_to_delete;
475     for (int i = 0; i < optimized_graph->node_size(); ++i) {
476         const NodeDef& node = optimized_graph->node(i);
477         if (IsNoOp(node) || IsIdentity(node) || IsIdentityN(node) ||
478             IsConstant(node) || SafeToConvertToNoOp(node)) {
479             nodes_to_simplify.PushBack(i);
480         }
481     }
482     while (!nodes_to_simplify.Empty()) {
483         int node_to_simplify = nodes_to_simplify.PopBack();
484         // Discard nodes that were marked for deletion already.
485         while (nodes_to_delete.find(node_to_simplify) != nodes_to_delete.end()) {
486             node_to_simplify = nodes_to_simplify.PopBack();
487         }
488         OptimizeNode(node_to_simplify, &nodes_to_simplify, &nodes_to_delete);
489     }
490
491     if (fetch_nodes_known_) {
492         VLOG(1) << "Deleted " << nodes_to_delete.size() << " out of "
493             << optimized_graph->node_size() << " nodes.";
494         EraseNodesFromGraph(nodes_to_delete, optimized_graph_);
495         node_map_.reset(new NodeMap(optimized_graph_));
496         BuildNodeToIdx();
497     }
498     return Status::OK();
499 }
500
501 namespace {
502
503 enum DistanceFromSource : uint8 { ZERO = 0, ONE = 1, TWO_OR_GREATER = 2 };
504
505 void LongestPathsLowerBounds(
506     int source, const std::pair<int, int>& target_range,
507     const std::vector<std::vector<int>>& outputs,
508     std::vector<DistanceFromSource>* longest_distance) {
509     std::deque<int> queue;
510     queue.emplace_front(source);
511     while (!queue.empty()) {
512         int node = queue.front();
513         queue.pop_front();
514         for (int fanout : outputs[node]) {
515             // 1) Only nodes in the target range can be on paths from source to one of
516             //     its control outputs.
517             // 2) Since we only need a lower bound on the longest distance, we can
518             //     skip nodes for which we have already proven have a path of
519             //     length > 1 from the source.

```

```

520     if (fanout >= target_range.first && fanout <= target_range.second &&
521         (*longest_distance)[fanout] != TWO_OR_GREATER) {
522         (*longest_distance)[fanout] =
523             (*longest_distance)[fanout] == ZERO ? ONE : TWO_OR_GREATER;
524         queue.emplace_front(fanout);
525     }
526 }
527 }
528 }
529
530 } // namespace
531
532 Status DependencyOptimizer::TransitiveReduction() {
533     // PRECONDITION: optimized_graph_ must be sorted topologically.
534     const int num_nodes = optimized_graph_>node_size();
535     // Set up a compressed version of the graph to save a constant factor in the
536     // expensive algorithm below. Also cache the set of control outputs and the
537     // highest index of a target of any control output from each node.
538     int num_controls = 0;
539     std::vector<std::vector<int>> outputs(num_nodes);
540     std::vector<gtl::InlinedVector<std::pair<int, int>, 2>> control_outputs(
541         num_nodes);
542     // target_range[i] contains the range of node indices for which to compute
543     // longest paths starting from node i.
544     std::vector<std::pair<int, int>> target_range(num_nodes, {num_nodes, -1});
545     for (int node_idx = 0; node_idx < num_nodes; ++node_idx) {
546         const NodeDef& node = optimized_graph_>node(node_idx);
547         if (ModifiesFrameInfo(node) || !HasOpDef(node)) {
548             // Ignore function nodes and nodes that modify frame info.
549             continue;
550         }
551         for (int input_slot = 0; input_slot < node.input_size(); ++input_slot) {
552             const string& input = node.input(input_slot);
553             const NodeDef* input_node = node_map_>GetNode(input);
554             if (ModifiesFrameInfo(*input_node) || IsMerge(*input_node)) {
555                 // Ignore edges from nodes that modify frame info and from Merge nodes,
556                 // because we cannot know which of it's input paths executes.
557                 continue;
558             }
559             const int input_node_idx = node_to_idx_[input_node];
560             outputs[input_node_idx].push_back(node_idx);
561             target_range[input_node_idx].first =
562                 std::min(target_range[input_node_idx].first, node_idx);
563             if (IsControlInput(input)) {
564                 ++num_controls;
565                 control_outputs[input_node_idx].emplace_back(node_idx, input_slot);
566                 target_range[input_node_idx].second =
567                     std::max(target_range[input_node_idx].second, node_idx);
568             }

```

```

569     }
570 }
571
572 // Run the longest path in DAG algorithm for each source node that has control
573 // outputs. If, for any target node of a control output, there exists a path
574 // of length > 1, we can drop that control dependency.
575 int num_controls_removed = 0;
576 std::vector<DistanceFromSource> longest_distance(num_nodes);
577 // Map from target_index -> set of (input_slot, source_index), representing
578 // the control edges to remove. We sort them in reverse order by input slot,
579 // such that when we swap them out so we don't clobber the
580 // node(target).input() repeated field.
581 typedef std::pair<int, int> InputSlotAndSource;
582 absl::flat_hash_map<
583     int, std::set<InputSlotAndSource, std::greater<InputSlotAndSource>>>
584     control_edges_to_remove;
585 for (int source = 0; source < num_nodes; ++source) {
586     if (target_range[source].first >= target_range[source].second ||
587         target_range[source].second <= source) {
588         continue;
589     }
590     // Compute the set of nodes in the transitive fanout of source with
591     // topological sort index in [target_range.first : target_range.second]]
592     // to which there exists a path of length 2 or more from source.
593     std::fill(longest_distance.begin() + target_range[source].first,
594               longest_distance.begin() + target_range[source].second + 1, ZERO);
595     LongestPathsLowerBounds(source, target_range[source], outputs,
596                             &longest_distance);
597
598     // If the longest path from source to target of a control dependency is
599     // longer than 1, there exists an alternate path, and we can eliminate the
600     // redundant direct control dependency.
601     for (const auto& control_output : control_outputs[source]) {
602         const int target = control_output.first;
603         if (longest_distance[target] == TWO_OR_GREATER) {
604             const int input_slot = control_output.second;
605             control_edges_to_remove[target].emplace(input_slot, source);
606         }
607     }
608 }
609 for (const auto& it : control_edges_to_remove) {
610     const int target = it.first;
611     NodeDef* target_node = optimized_graph->mutable_node(target);
612     for (const InputSlotAndSource& slot_and_source : it.second) {
613         const int input_slot = slot_and_source.first;
614         const int source = slot_and_source.second;
615         const NodeDef& source_node = optimized_graph->node(source);
616         CHECK_LT(input_slot, target_node->input_size());
617         target_node->mutable_input()->SwapElements(input_slot,

```

```

618         target_node->input_size() - 1);
619     node_map_->RemoveOutput(source_node.name(), target_node->name());
620     target_node->mutable_input()->RemoveLast();
621     ++num_controls_removed;
622 }
623 }
624 VLOG(1) << "Removed " << num_controls_removed << " out of " << num_controls
625     << " control dependencies";
626 return Status::OK();
627 }
628
629 void DependencyOptimizer::BuildNodeToIdx() {
630     // Set up &node -> index map.
631     node_to_idx_.clear();
632     for (int i = 0; i < optimized_graph->node_size(); ++i) {
633         const NodeDef& node = optimized_graph->node(i);
634         node_to_idx_[&node] = i;
635     }
636 }
637
638 // Suppose there are cross-device control inputs to node C from multiple nodes
639 // that are located on another device, e.g., we have control edges:
640 // A->C, B->C
641 // where A and B are on device X and C is on device Y.
642 // We can reduce cross-device communication by introducing an intermediate
643 // NoOp node C' on device X and rewriting the control edges to:
644 // A->C', B->C', C' -> C
645 void DependencyOptimizer::GroupCrossDeviceControlEdges(bool host_granularity) {
646     VLOG(1)
647         << "DependencyOptimizer::GroupCrossDeviceControlEdges host_granularity="
648         << host_granularity;
649     const int num_nodes = optimized_graph->node_size();
650     for (int i = 0; i < num_nodes; ++i) {
651         NodeDef* node = optimized_graph->mutable_node(i);
652         if (node->device().empty()) continue;
653         string rest, node_device = node->device();
654         if (host_granularity) {
655             DeviceNameUtils::SplitDeviceName(node->device(), &node_device, &rest);
656         }
657
658         // Creates new noop nodes for devices on which multiple control inputs are
659         // located.
660
661         // Map keyed by device name to the newly introduced Noop node for that
662         // device. A nullptr value means that we have only seen a single node on
663         // that device.
664         std::map<string, NodeDef*> noops;
665         int num_noops = 0;
666         for (int j = 0; j < node->input_size(); ++j) {

```

```

667     if (IsControlInput(node->input(j))) {
668         const NodeDef* input = node_map_->GetNode(node->input(j));
669         if (input == nullptr || input->device().empty()) continue;
670         string input_device = input->device();
671         if (host_granularity) {
672             DeviceNameUtils::SplitDeviceName(input->device(), &input_device,
673                                               &rest);
674         }
675         if (input_device != node_device) {
676             VLOG(2) << "Cross-device " << node->name() << " " << input->device()
677                 << " -> " << node->device();
678             auto emplace_result = noops.emplace(input_device, nullptr);
679             if (!emplace_result.second &&
680                 emplace_result.first->second == nullptr) {
681                 VLOG(2) << "Duplicate input device from " << node->name();
682                 // This is the second cross-device control input from the same
683                 // device. Creates an intermediate noop node on that device.
684                 string group_name;
685                 NodeDef* noop;
686                 // Creates a fresh node name; there may be conflicting names from
687                 // a previous iteration of the optimizer.
688                 do {
689                     group_name = AddPrefixToNodeName(
690                         node->name(),
691                         strings::StrCat("GroupCrossDeviceControlEdges_", num_noops));
692                     noop = node_map_->GetNode(group_name);
693                     ++num_noops;
694                 } while (noop != nullptr);
695                 noop = optimized_graph_->add_node();
696                 noop->set_name(group_name);
697                 noop->set_device(input->device());
698                 noop->set_op("NoOp");
699                 node_map_->AddNode(noop->name(), noop);
700                 emplace_result.first->second = noop;
701                 VLOG(1) << "GroupCrossDeviceControlEdges: Added "
702                     << SummarizeNodeDef(*noop);
703             }
704         }
705     }
706 }

707
708 // Reroute existing control edges to go via the newly introduced NoOp nodes.
709 int pos = 0;
710 while (pos < node->input_size()) {
711     const string& input_name = node->input(pos);
712     if (IsControlInput(input_name)) {
713         NodeDef* input = node_map_->GetNode(input_name);
714         if (input == nullptr) {
715             ++pos;

```

```

716     } else {
717         string input_device = input->device();
718         if (host_granularity) {
719             DeviceNameUtils::SplitDeviceName(input->device(), &input_device,
720                                             &rest);
721         }
722         auto it = noops.find(input_device);
723         if (it == noops.end() || it->second == nullptr) {
724             ++pos;
725         } else {
726             VLOG(2) << "Rewriting input from " << input_name;
727             node->mutable_input()->SwapElements(pos, node->input_size() - 1);
728             node->mutable_input()->RemoveLast();
729             it->second->add_input(AsControlDependency(*input));
730             node_map_->UpdateOutput(input_name, node->name(),
731                                   it->second->name());
732         }
733     }
734     } else {
735         ++pos;
736     }
737 }
738 for (const auto& entry : noops) {
739     if (entry.second) {
740         node->add_input(AsControlDependency(*entry.second));
741         node_map_->AddOutput(entry.second->name(), node->name());
742     }
743 }
744 }
745 }
746
747 Status DependencyOptimizer::Optimize(Cluster* cluster, const GrapplerItem& item,
748                                     GraphDef* optimized_graph) {
749     optimized_graph_ = optimized_graph;
750     *optimized_graph_ = item.graph;
751     nodes_to_preserve_ = item.NodesToPreserve();
752     fetch_nodes_known_ = !item.fetch.empty();
753     CleanControlInputs();
754
755     const int num_iterations = 2;
756     for (int iteration = 0; iteration < num_iterations; ++iteration) {
757         GRAPPLER_RETURN_IF_DEADLINE_EXCEEDED();
758         Status topo_sort_status;
759         // Perform topological sort to prepare the graph for transitive reduction.
760         topo_sort_status = TopologicalSort(optimized_graph_);
761         // Set up index-based graph datastructures to speed up analysis steps below.
762         node_map_.reset(new NodeMap(optimized_graph_));
763         BuildNodeToIdx();
764     }

```



```

765     if (topo_sort_status.ok()) {
766         // Remove redundant control dependencies.
767         TF_RETURN_IF_ERROR(TransitiveReduction());
768     } else {
769         LOG(ERROR) << "Iteration = " << iteration
770             << ", topological sort failed with message: "
771             << topo_sort_status.error_message();
772     }
773     // Turn nodes with only control outputs into NoOps, prune NoOp and Identity
774     // nodes.
775     TF_RETURN_IF_ERROR(OptimizeDependencies());
776
777     // Dedup control inputs.
778     CleanControlInputs();
779
780     // Merge multiple control edges from the same device.
781     GroupCrossDeviceControlEdges(/*host_granularity=*/false);
782
783     // Merge control edges from the same host to reduce RPC traffic.
784     GroupCrossDeviceControlEdges(/*host_granularity=*/true);
785 }
786
787 return Status::OK();
788 }
789
790 } // end namespace grappler
791 } // end namespace tensorflow

```