

274df9b023 ▾

...

tensorflow / tensorflow / compiler / jit / xla_platform_info.cc



jpienaar Rename to underlying type rather than alias ... ✖

History

3 contributors



183 lines (164 sloc) | 7.5 KB

...

```

1  /* Copyright 2020 The TensorFlow Authors. All Rights Reserved.
2
3  Licensed under the Apache License, Version 2.0 (the "License");
4  you may not use this file except in compliance with the License.
5  You may obtain a copy of the License at
6
7      http://www.apache.org/licenses/LICENSE-2.0
8
9  Unless required by applicable law or agreed to in writing, software
10 distributed under the License is distributed on an "AS IS" BASIS,
11 WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
12 See the License for the specific language governing permissions and
13 limitations under the License.
14 =====*/
15
16 #include "tensorflow/compiler/jit/xla_platform_info.h"
17
18 #include "tensorflow/compiler/xla/client/client_library.h"
19
20 namespace tensorflow {
21
22 xla::StatusOr<absl::optional<std::set<int>>> ParseVisibleDeviceList(
23     absl::string_view visible_device_list) {
24     std::set<int> gpu_ids;
25     if (visible_device_list.empty()) {
26         return {{absl::nullopt}};
27     }
28     const std::vector<string> visible_devices =
29         absl::StrSplit(visible_device_list, ',');
```

```

30     for (const string& platform_device_id_str : visible_devices) {
31         int32_t platform_device_id;
32         if (!absl::SimpleAtoi(platform_device_id_str, &platform_device_id)) {
33             return errors::InvalidArgument(
34                 "Could not parse entry in 'visible_device_list': '",
35                 platform_device_id_str,
36                 "'. visible_device_list = ", visible_device_list);
37         }
38         gpu_ids.insert(platform_device_id);
39     }
40     return {{gpu_ids}};
41 }
42
43 Status BuildXlaCompilationCache(DeviceBase* device, FunctionLibraryRuntime* flr,
44                                 const XlaPlatformInfo& platform_info,
45                                 XlaCompilationCache** cache) {
46     if (platform_info.xla_device_metadata()) {
47         *cache = new XlaCompilationCache(
48             platform_info.xla_device_metadata()->client(),
49             platform_info.xla_device_metadata()->jit_device_type());
50         return Status::OK();
51     }
52
53     auto platform =
54         se::MultiPlatformManager::PlatformWithId(platform_info.platform_id());
55     if (!platform.ok()) {
56         return platform.status();
57     }
58
59     StatusOr<xla::Compiler*> compiler_for_platform =
60         xla::Compiler::GetForPlatform(platform.ValueOrDie());
61     if (!compiler_for_platform.ok()) {
62         // In some rare cases (usually in unit tests with very small clusters) we
63         // may end up transforming an XLA cluster with at least one GPU operation
64         // (which would normally force the cluster to be compiled using XLA:GPU)
65         // into an XLA cluster with no GPU operations (i.e. containing only CPU
66         // operations). Such a cluster can fail compilation (in way that
67         // MarkForCompilation could not have detected) if the CPU JIT is not linked
68         // in.
69         //
70         // So bail out of _XlaCompile in this case, and let the executor handle the
71         // situation for us.
72         const Status& status = compiler_for_platform.status();
73         if (status.code() == error::NOT_FOUND) {
74             return errors::Unimplemented("Could not find compiler for platform ",
75                                         platform.ValueOrDie()->Name(), ": ",
76                                         status.ToString());
77         }
78     }

```

```

79
80 xla::LocalClientOptions client_options;
81 client_options.set_platform(platform.ValueOrDie());
82 client_options.set_intra_op_parallelism_threads(
83     device->tensorflow_cpu_worker_threads()->num_threads);
84
85 string allowed_gpus =
86     flr->config_proto()->gpu_options().visible_device_list();
87 TF_ASSIGN_OR_RETURN(absl::optional<std::set<int>> gpu_ids,
88     ParseVisibleDeviceList(allowed_gpus));
89 client_options.set_allowed_devices(gpu_ids);
90
91 auto client = xla::ClientLibrary::GetOrCreateLocalClient(client_options);
92 if (!client.ok()) {
93     return client.status();
94 }
95 const XlaOpRegistry::DeviceRegistration* registration;
96 if (!XlaOpRegistry::GetCompilationDevice(platform_info.device_type().type(),
97     &registration)) {
98     return errors::InvalidArgument("No JIT device registered for ",
99     platform_info.device_type().type());
100 }
101 *cache = new XlaCompilationCache(
102     client.ValueOrDie(), DeviceType(registration->compilation_device_name));
103 return Status::OK();
104 }
105
106 XlaPlatformInfo XlaPlatformInfoFromDevice(DeviceBase* device_base) {
107     auto device = static_cast<Device*>(device_base);
108     se::Platform::Id platform_id = nullptr;
109     const XlaDevice::Metadata* xla_device_metadata = nullptr;
110     std::shared_ptr<se::DeviceMemoryAllocator> custom_allocator;
111
112     if (device->device_type() == DEVICE_CPU) {
113         platform_id = se::host::kHostPlatformId;
114     } else if (device->device_type() == DEVICE_GPU) {
115         platform_id = device->tensorflow_gpu_device_info()
116             ->stream->parent()
117             ->platform()
118             ->id();
119     } else if (XlaDevice::GetMetadataFromDevice(device, &xla_device_metadata)
120         .ok()) {
121         // If we are on an XlaDevice, use the underlying XLA platform's allocator
122         // directly. We could use the StreamExecutor's allocator which may
123         // theoretically be more correct, but XLA returns a nice OOM message in a
124         // Status and StreamExecutor does not.
125         //
126         // Importantly we can't use ctx->device()->GetAllocator() as the allocator
127         // (which xla_allocator above uses) as on an XlaDevice, this is a dummy

```

```

128     // allocator that returns XlaTensor objects. The XlaCompiler needs a real
129     // allocator to allocate real buffers.
130     platform_id = xla_device_metadata->platform()->id();
131     custom_allocator =
132         xla_device_metadata->client()->backend().shared_memory_allocator();
133 }
134
135     return XlaPlatformInfo(DeviceType(device->device_type()), platform_id,
136         xla_device_metadata, custom_allocator);
137 }
138
139 std::shared_ptr<se::DeviceMemoryAllocator> GetAllocator(
140     DeviceBase* device, se::Stream* stream,
141     const XlaPlatformInfo& platform_info) {
142     if (platform_info.custom_allocator()) {
143         return platform_info.custom_allocator();
144     }
145     auto* alloc = device->GetAllocator({});
146     if (!stream) {
147         // Stream is not set for the host platform.
148         se::Platform* platform =
149             se::MultiPlatformManager::PlatformWithId(platform_info.platform_id())
150                 .ValueOrDie();
151         return std::make_shared<se::TfAllocatorAdapter>(alloc, platform);
152     }
153     return std::make_shared<se::TfAllocatorAdapter>(alloc, stream);
154 }
155
156 XlaCompiler::Options GenerateCompilerOptions(
157     const XlaCompilationCache& cache,
158     const FunctionLibraryRuntime& function_library, DeviceBase* device,
159     se::Stream* stream, const XlaPlatformInfo& platform_info,
160     bool has_ref_vars) {
161     XlaCompiler::Options options;
162     options.client = static_cast<xla::LocalClient*>(cache.client());
163     if (stream != nullptr) {
164         options.device_ordinal = stream->parent()->device_ordinal();
165     }
166     options.device_type = cache.device_type();
167     options.flib_def = function_library.GetFunctionLibraryDefinition();
168     options.graph_def_version = function_library.graph_def_version();
169     options.allow_cpu_custom_calls =
170         (platform_info.platform_id() == se::host::kHostPlatformId);
171     options.device_allocator = GetAllocator(device, stream, platform_info);
172     if (platform_info.xla_device_metadata()) {
173         options.shape_representation_fn =
174             platform_info.xla_device_metadata()->shape_representation_fn();
175     }
176     // If reference variables are not present in the graph, we can safely alias

```

```
177 // passthrough parameters without performing a copy.
178 options.alias_passthrough_params =
179     !has_ref_vars && !platform_info.is_on_xla_device();
180 return options;
181 }
182
183 } // namespace tensorflow
```