

Talos Vulnerability Report

TALOS-2022-1486

HDF5 Group libhdf5 gif2h5 out-of-bounds read vulnerability

AUGUST 16, 2022

CVE NUMBER

CVE-2022-25942

SUMMARY

An out-of-bounds read vulnerability exists in the gif2h5 functionality of HDF5 Group libhdf5 1.10.4. A specially-crafted GIF file can lead to code execution. An attacker can provide a malicious file to trigger this vulnerability.

CONFIRMED VULNERABLE VERSIONS

The versions below were either tested or verified to be vulnerable by Talos or confirmed to be vulnerable by the vendor.

HDF5 Group libhdf5 1.10.4

PRODUCT URLS

libhdf5 - <https://www.hdfgroup.org>

CVSSV3 SCORE

7.8 - CVSS:3.0/AV:L/AC:L/PR:N/UI:R/S:U/C:H/I:H/A:H

CWE

CWE-125 - Out-of-bounds Read

DETAILS

HDF5 is a file format that is maintained by a non-profit organization, the HDF Group. HDF5 is designed to store and organize large amounts of scientific data. It is used to exchange data structures between applications in industries (such as the GIS industry) via libraries such as GDAL, OGR or as part of software like ArcGIS.

The vulnerability exists in the HDF5 gif2h5 tool/library, specifically in ReadGifHeader() function while attempting to convert GIF files to their HDF5 format. The vulnerability exists due to their failure to check the file size compared to the Global Palette Size specified in the GIF file. This leads to repeated data filling the gifHead->HDFPalette and returning what should be a pointer to the GIF data from the original file. Instead, it points just past the palette data in the destination structure.

In gif2mem.c:

```
96      /* Read the GIF image file header information */
97      ReadGifHeader(gifHead, &MemGif);
```

We can see this occur in the ReadGifHeader() function within gifread.c. The size of the file isn't checked against the table size before attempting to read palette data. In this example, we specified a palette size of 0x100 (max) bytes with a file that is only 0x5e bytes long

```
79      /* Check if a Global Color Table is present */
80      if (GifHead->PackedField & 0x80) {
81          /* Read number of color table entries */
82          tableSize      = (GIFWORD)(1L << ((GifHead->PackedField & 0x07) +
1));
83          GifHead->TableSize = tableSize;
84
85          /* Read the Global Color Table */
86
87          /*
88           * There are some changes made here apart from just reading in the
89           * global color table as would seem intuitively obvious. The colors
90           * are stored in the bottom part of the palette as opposed to the top
91           */
92
93          for (i = 0; i < tableSize; i++) {
94              GifHead->HDFPalette[i][0] = *(*MemGif2)++;
95              GifHead->HDFPalette[i][1] = *(*MemGif2)++;
96              GifHead->HDFPalette[i][2] = *(*MemGif2)++;
97          }
98      }
```

The MemGif2 structure exists at 0x000055555a838dd, which points to the palette data in the original GIF.

When i == 0x0 (Before consuming Memory)

```
gef> hexdump *MemGif2
0x000055555a838dd    10 00 00 0a 03 00 f4 e4 2c 46 38 39 61 0a 00 03
.....,F89a...
0x000055555a838ed    00 80 76 00 ff fa 2c 00 00 00 00 00 00 03 00 ff
..V...,.....
0x000055555a838fd    fe ff ef 00 01 e4 2c 46 38 39 61 0a 00 03 00 80
.....,F89a.....
0x000055555a8390d    76 f4 fe fa 2c 00 00 00 00 00 00 03 00 ff ff 00
V...,.....
```

The GifHead structure exists at 0x000055555a83940, which is only 0x63 bytes away

```
gef> hexdump GifHead
0x000055555a83940    ff 00 00 00 00 00 00 00 00 00 01 00 00 00 00 00
.....
0x000055555a83950    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....
0x000055555a83960    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....
0x000055555a83970    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....
```

The loop will read for a total of 0x300 bytes, incrementing the MemGif2 pointer until it eventually points into the GifHead structure itself and begins copying bytes from the GifHead structure into the HDFPalette data.

We can see this occur after the loop by dumping the beginning of the HDFPalette data. The same data from GifHead structure has been copied into the HDFPalette member starting 0x63 (99) bytes into the HDFPalette data at address 0x000055555a839b7.

```

gef> hexdump GifHead->HDFPalette 128
0x00005555555a83954      10 00 00 0a 03 00 f4 e4 2c 46 38 39 61 0a 00 03
.....,F89a...
0x00005555555a83964      00 80 76 00 ff fa 2c 00 00 00 00 00 03 00 ff
..V....,.....
0x00005555555a83974      fe ff ef 00 01 e4 2c 46 38 39 61 0a 00 03 00 80
.....,F89a.....
0x00005555555a83984      76 f4 fe fa 2c 00 00 00 00 00 00 03 00 ff ff 00
V....,.....
0x00005555555a83994      00 2c 00 00 00 9c 01 00 00 01 00 00 02 ff 0e 64
.,.....d
0x00005555555a839a4      3b 00 00 00 00 00 00 00 00 00 00 31 03 00 00 00
;.....1....
0x00005555555a839b4      00 00 00 ff 00 00 00 00 00 00 00 01 00 00 00
..... // <-- Here
0x00005555555a839c4      00 00 00 00 00 00 00 10 00 00 0a 03 00 f4 e4 2c
.....,

// The GifHead structure itself
gef> p/x *GifHead
$37 = {
    PackedField = 0xff,
    TableSize = 0x100,
    ImageCount = 0x0,
    CommentCount = 0x0,
    ApplicationCount = 0x0,
    PlainTextCount = 0x0,
    HDFPalette = {{0x10, 0x0, 0x0}, {0xa, 0x3, 0x0}, {0xf4, 0xe4, 0x2c}, {0x46, 0x38,
0x39}, {0x61, 0xa, 0x0}, {0x3, 0x0, 0x80}, {0x76, 0x0, 0xff}, {0xfa, 0x2c, 0x0},
{0x0, 0x0, 0x0}, {0x0, 0x0, 0x3}, {0x0, 0xff, 0xfe}, {0xff, 0xef, 0x0}, {0x1, 0xe4,
0x2c}, {0x46, 0x38, 0x39}, {0x61, 0xa, 0x0}, {0x3, 0x0, 0x80}, {0x76, 0xf4, 0xfe},
{0xfa, 0x2c, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x3}, {0x0, 0xff, 0xff}, {0x0, 0x0,
0x2c}, {0x0, 0x0, 0x0}, {0x9c, 0x1, 0x0}, {0x0, 0x1, 0x0}, {0x0, 0x2, 0xff}, {0xe,
0x64, 0x3b}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x31, 0x3},
{0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0xff, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0,
0x0}, {0x1, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x10}, {0x0,
0x0, 0xa}, {0x3, 0x0, 0xf4}, {0xe4, 0x2c, 0x46}, {0x38, 0x39, 0x61}, {0xa, 0x0,
0x3}, {0x0, 0x80, 0x76}, {0x0, 0xff, 0xfa}, {0x2c, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0,
0x3, 0x0}, {0xff, 0xfe, 0xff}, {0xef, 0x0, 0x1}, {0xe4, 0x2c, 0x46}, {0x38, 0x39,
0x61}, {0xa, 0x0, 0x3}, {0x0, 0x80, 0x76}, {0xf4, 0xfe, 0xfa}, {0x2c, 0x0, 0x0},
{0x0, 0x0, 0x0}, {0x0, 0x3, 0x0}, {0xff, 0xff, 0x0}, {0x0, 0x2c, 0x0}, {0x0, 0x0,
0x9c}, {0x1, 0x0, 0x0}, {0x1, 0x0, 0x0}, {0x2, 0xff, 0xe}, {0x64, 0x3b, 0x0}, {0x0,
0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x31, 0x3, 0x0}, {0x0, 0x0, 0x0},
{0x0, 0x0, 0xff}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x1}, {0x0, 0x0,
0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x10, 0x0}, {0x0, 0xa, 0x3}, {0x0,
0xf4, 0xe4}, {0x2c, 0x46, 0x38}, {0x39, 0x61, 0xa}, {0x0, 0x3, 0x0}, {0x80, 0x76,
0x0}, {0xff, 0xfa, 0x2c}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x3, 0x0, 0xff}, {0xfe,
0xff, 0xef}, {0x0, 0x1, 0xe4}, {0x2c, 0x46, 0x38}, {0x39, 0x61, 0xa}, {0x0, 0x3,
0x0}, {0x80, 0x76, 0xf4}, {0xfe, 0xfa, 0x2c}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0},
{0x3, 0x0, 0xff}, {0xff, 0x0, 0x0}, {0x2c, 0x0, 0x0}, {0x0, 0x9c, 0x1}, {0x0, 0x0,
0x1}, {0x0, 0x0, 0x2}, {0xff, 0xe, 0x64}, {0x3b, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0,
0x0, 0x0}, {0x0, 0x0, 0x31}, {0x3, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0xff, 0x0},
{0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x1, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0},
{0x0, 0x0, 0x0}, {0x10, 0x0, 0x0}, {0xa, 0x3, 0x0}, {0xf4, 0xe4, 0x2c}, {0x46, 0x38,
0x39}, {0x61, 0xa, 0x0}, {0x3, 0x0, 0x80}, {0x76, 0x0, 0xff}, {0xfa, 0x2c, 0x0},
{0x0, 0x0, 0x0}, {0x0, 0x0, 0x3}, {0x0, 0xff, 0xfe}, {0xff, 0xef, 0x0}, {0x1, 0xe4,
0x2c}, {0x46, 0x38, 0x39}, {0x61, 0xa, 0x0}, {0x3, 0x0, 0x80}, {0x76, 0xf4, 0xfe},
{0xfa, 0x2c, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x3}, {0x0, 0xff, 0xff}, {0x0, 0x0,

```

```

0x2c}, {0x0, 0x0, 0x0}, {0x9c, 0x1, 0x0}, {0x0, 0x1, 0x0}, {0x0, 0x2, 0xff}, {0xe,
0x64, 0x3b}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x31, 0x3},
{0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0xff, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0,
0x0}, {0x1, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x10}, {0x0,
0x0, 0xa}, {0x3, 0x0, 0xf4}, {0xe4, 0x2c, 0x46}, {0x38, 0x39, 0x61}, {0xa, 0x0,
0x3}, {0x0, 0x80, 0x76}, {0x0, 0xff, 0xfa}, {0x2c, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0,
0x3, 0x0}, {0xff, 0xfe, 0xff}, {0xef, 0x0, 0x1}, {0xe4, 0x2c, 0x46}, {0x38, 0x39,
0x61}, {0xa, 0x0, 0x3}, {0x0, 0x80, 0x76}, {0xf4, 0xfe, 0xfa}, {0x2c, 0x0, 0x0},
{0x0, 0x0, 0x0}, {0x0, 0x3, 0x0}, {0xff, 0xff, 0x0}, {0x0, 0x2c, 0x0}, {0x0, 0x0,
0x9c}, {0x1, 0x0, 0x0}, {0x1, 0x0, 0x0}, {0x2, 0xff, 0xe}, {0x64, 0x3b, 0x0}, {0x0,
0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x31, 0x3, 0x0}, {0x0, 0x0, 0x0},
{0x0, 0x0, 0xff}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x1}, {0x0, 0x0,
0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x10, 0x0}, {0x0, 0xa, 0x3}, {0x0,
0xf4, 0xe4}, {0x2c, 0x46, 0x38}, {0x39, 0x61, 0xa}, {0x0, 0x3, 0x0}, {0x80, 0x76,
0x0}, {0xff, 0xfa, 0x2c}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x3, 0x0, 0xff}, {0xfe,
0xff, 0xef}, {0x0, 0x1, 0xe4}, {0x2c, 0x46, 0x38}, {0x39, 0x61, 0xa}, {0x0, 0x3,
0x0}, {0x80, 0x76, 0xf4}, {0xfe, 0xfa, 0x2c}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0},
{0x3, 0x0, 0xff}, {0xff, 0x0, 0x0}, {0x2c, 0x0, 0x0}, {0x0, 0x9c, 0x1}, {0x0, 0x0,
0x1}, {0x0, 0x0, 0x2}, {0xff, 0xe, 0x64}, {0x3b, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0,
0x0, 0x0}, {0x0, 0x0, 0x31}, {0x3, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0xff, 0x0},
{0x0, 0x0, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x1, 0x0}, {0x0, 0x0, 0x0}, {0x0, 0x0, 0x0},
{0x0, 0x0, 0x0}, {0x10, 0x0, 0x0}, {0xa, 0x3, 0x0}, {0xf4, 0xe4, 0x2c}, {0x46, 0x38,
0x39}, {0x61, 0xa, 0x0}, {0x3, 0x0, 0x80}, {0x76, 0x0, 0xff}, {0xfa, 0x2c, 0x0},
{0x0, 0x0, 0x0}, {0x0, 0x0, 0x3}, {0x0, 0xff, 0xfe}, {0xff, 0xef, 0x0}, {0x1, 0xe4,
0x2c}, {0x46, 0x38, 0x39}, {0x61, 0xa, 0x0}, {0x3, 0x0, 0x80}, {0x76, 0xf4, 0xfe},
{0xfa, 0x2c, 0x0}},
    HeaderDump = {0x47, 0x49, 0x46, 0x38, 0x39, 0xb4},
    LSDDump = {0xa, 0x80, 0x0, 0x2, 0xff, 0x76, 0x0}
}

```

When we return back to the Gif2Mem() function within gif2mem.c, parsing of the image data begins using the MemGif pointer used previously:

```

97     ReadGifHeader(gifHead, &MemGif);
    /* <snip> */
117    for (;;) {
118        Identifier = *MemGif++;
119
120        switch (Identifier) {
121            case 0x3B: /* Trailer */
122                /*
123                 * The counts are stored to make it easier while putting stuff
124                 * into the HDF file and then deallocating space.

```

However, MemGif instead points into the GifHead struct itself. We can see this at address 0x0000555555a83c54, which is the gifHead->HeaderDump member and so on, meaning that the gifHead data will be interpreted and parsed as image data.

```

gef> hexdump gifHead->HeaderDump
0x000055555a83c54      47 49 46 38 39 b4 0a 80 00 02 ff 76 00 00 00 00
GIF89.....v....
0x000055555a83c64      00 00 00 00 00 a1 03 02 00 00 00 00 00 00 00 00
.....
0x000055555a83c74      00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....
0x000055555a83c84      00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....

gef> hexdump MemGif 128
0x000055555a83bdd      00 00 00 00 00 03 00 ff ff 00 00 2c 00 00 00 9c
.....,....
0x000055555a83bed      01 00 00 01 00 00 02 ff 0e 64 3b 00 00 00 00 00
.....d;....
0x000055555a83bfd      00 00 00 00 00 31 03 00 00 00 00 00 00 00 ff 00
.....1.....
0x000055555a83c0d      00 00 00 00 00 00 01 00 00 00 00 00 00 00 00 00
.....
0x000055555a83c1d      00 10 00 00 0a 03 00 f4 e4 2c 46 38 39 61 0a 00
.....,F89a..
0x000055555a83c2d      03 00 80 76 00 ff fa 2c 00 00 00 00 00 00 03 00
...V...,.....
0x000055555a83c3d      ff fe ff ef 00 01 e4 2c 46 38 39 61 0a 00 03 00
.....,F89a....
0x000055555a83c4d      80 76 f4 fe fa 2c 00 47 49 46 38 39 b4 0a 80 00
.v...,.GIF89.... // <-- Here

```

This can even corrupt the heap, when parsing the remaining data as image data (for example, when attempting to reallocate image memory in gif2mem.c because “multiple” images were found, when in fact they are merely copies of the same image data):

```

152             if (ImageCount > ImageArray) {
153                 aTemp      = ImageArray;
154                 ImageArray = (GIFBYTE)((ImageArray << 1) + 1);
155                 if (!(gifImageDesc =
156                     (GIFIMAGEDESC **)realloc(gifImageDesc,
sizeof(GIFIMAGEDESC *) * ImageArray))) {
157                     printf("Out of memory!");
158                     exit(EXIT_FAILURE);
159                 }

```

The reallocation can shift the next image pointer into the middle of this palette data. From here, arbitrary memory can be freed when returning to the main function:

```
113      /* Free all buffers */
114      /* replacing int32 with long */
115      ImageCount = (long)((GifMemoryStruct.GifHeader)->ImageCount);
116
117      for (i = 0; i < ImageCount; i++) {
118          gifImageDesc = *(GifMemoryStruct.GifImageDesc[i]);
119
120          if (gifImageDesc.Image != NULL)
121              free(gifImageDesc.Image);

gef➤ p gifImageDesc.Image
$3 = (GIFBYTE *) 0x4141414141414141 <error: Cannot access memory at address
0x4141414141414141>

Program received signal SIGSEGV, Segmentation fault.
__GI___libc_free (mem=0x4141414141414141) at malloc.c:3102
3102      p = mem2chunk (mem);
```

TIMELINE

2022-03-23 - Vendor Disclosure

2022-08-16 - Public Release

CREDIT

Discovered by Dave McDaniel of Cisco Talos.

[VULNERABILITY REPORTS](#)

[PREVIOUS REPORT](#)

[NEXT REPORT](#)

[TALOS-2022-1485](#)

[TALOS-2022-1514](#)

