

Talos Vulnerability Report

TALOS-2022-1485

HDF5 Group libhdf5 gif2h5 out-of-bounds write vulnerability

AUGUST 16, 2022

CVE NUMBER

CVE-2022-25972

SUMMARY

An out-of-bounds write vulnerability exists in the gif2h5 functionality of HDF5 Group libhdf5 1.10.4. A specially-crafted GIF file can lead to code execution. An attacker can provide a malicious file to trigger this vulnerability.

CONFIRMED VULNERABLE VERSIONS

The versions below were either tested or verified to be vulnerable by Talos or confirmed to be vulnerable by the vendor.

HDF5 Group libhdf5 1.10.4

PRODUCT URLS

libhdf5 - <https://www.hdfgroup.org>

CVSSV3 SCORE

7.8 - CVSS:3.0/AV:L/AC:L/PR:N/UI:R/S:U/C:H/I:H/A:H

CWE

CWE-787 - Out-of-bounds Write

DETAILS

HDF5 is a file format that is maintained by a non-profit organization, the HDF Group. HDF5 is designed to store and organize large amounts of scientific data. It is used to exchange data structures between applications in industries (such as the GIS industry) via libraries such as GDAL, OGR or as part of software like ArcGIS.

The library that includes the gif2h5 tool is used for converting GIF data to the HDF5 file format. The vulnerability exists due to a failure to check the bounds of a heap buffer during GIF decompression while using user-provided input to calculate an offset for writing into the heap buffer.

During GIF file decompression, the code size which represents the number of bits required to represent pixel values is used to calculate various offsets and other codes, such as the clear code and end-of-file code. In the file `decompress.c`, we can see this occur:

```
/*
Example using a code size of 0x0c:
CodeSize = 0x0c
ClearCode = 0x1000
EOFCode = 0x1001
FreeCode & FirstFree = 0x1002
*/

200     CodeSize = GifImageDesc->CodeSize;    // Get CodeSize
201     ClearCode = (1 << CodeSize);           // Shift by CodeSize value
202     EOFCode   = ClearCode + 1;             // Add 1
203     FreeCode = FirstFree = ClearCode + 2;  // Add 2
```

After this information is gathered, the remaining GIF data is parsed in a loop until the EOFCode is reached. Note that there are two objects on the heap, Prefix & Suffix, that were allocated earlier in the code:

```
159     if (!(Prefix = calloc(4096, sizeof(int)))) {
160         printf("Out of memory");
161         exit(EXIT_FAILURE);
162     }
163     if (!(Suffix = calloc(4096, sizeof(int)))) {
164         printf("Out of memory");
165         exit(EXIT_FAILURE);
```

At this point, we can see the size and address of both heap objects:

```

gef> hexdump Prefix-2
0x0000555555a99e08    11 40 00 00 00 00 00 00 00 00 00 00 00 00 00 00
. @ .....
0x0000555555a99e18    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....
0x0000555555a99e28    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....
0x0000555555a99e38    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....

gef> hexdump Suffix-2
0x0000555555a9de18    11 40 00 00 00 00 00 00 00 00 00 00 00 00 00 00
. @ .....
0x0000555555a9de28    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....
0x0000555555a9de38    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....
0x0000555555a9de48    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....

```

Once the data parsing begins, the user-controlled index `FreeCode` is used to write into the heap buffers. Since our user-controlled `FreeCode` variable is `0x1002` it will index to `0x4008` when line 300 is executed. We can see this will lead to a heap-buffer overflow that will allow us to overwrite the size of the `Suffix` heap buffer, for example:

```

gef> hexdump &Prefix[FreeCode]
0x000055555a9de18    11 40 00 00 00 00 00 00 00 00 00 00 00 00 00 00
. @ .....
0x000055555a9de28    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....
0x000055555a9de38    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....
0x000055555a9de48    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....

232      /*
233      * Decompress the file, continuing until you see the GIF EOF code.  One
234      * obvious enhancement is to add checking for corrupt files here.
235      */
236
237      Code = ReadCode();
238
239      while (Code != EOFCode) {
...
296          /*
297          * Build the hash table on-the-fly. No table is stored in the
298          * file.
299          */
300          Prefix[FreeCode] = OldCode;
301          Suffix[FreeCode] = FinChar;
302          OldCode           = InCode;

```

After line 300 is executed, we can see that we have cleared the size of the Suffix buffer and can continue to corrupt the remaining heap data:

```

gef> hexdump Suffix-2
0x000055555a9de18    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....
0x000055555a9de28    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....
0x000055555a9de38    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....
0x000055555a9de48    00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
.....

```

To demonstrate triggering a crash, we will use the following values for the decompression data:

AA AA AA 04 40 47

These values will be read by the function ReadCode() called at line 237 (and later at 320), until certain codes are found. In this case, we want the EOFCode or 0x1001 to trigger the crash.

```
232     /*
233     * Decompress the file, continuing until you see the GIF EOF code. One
234     * obvious enhancement is to add checking for corrupt files here.
235     */
236
237     Code = ReadCode();
238
239     while (Code != EOFCode) {           // Loops until EOFCode is reached
...
296         /*
297         * Build the hash table on-the-fly. No table is stored in the
298         * file.
299         */
300         Prefix[FreeCode] = OldCode; // Out-of-bounds write to corrupt
Suffix
301         Suffix[FreeCode] = FinChar;
302         OldCode          = InCode;
...
318     }
319
320     Code = ReadCode();                // Call ReadCode() during each loop
iteration
321 }
```

The ReadCode() function performs a few operations against the Raster data. When the values 0x474004 are read in this function (in this case, the third time this function is called), they will evaluate to 0x1001 or the EOFCode

```

77 static int
78 ReadCode(void)
79 {
80     int RawCode, ByteOffset;
81
82     ByteOffset = BitOffset / 8; //
83     BitOffset = 0x1a, ByteOffset = 0x03
84     RawCode = Raster[ByteOffset] + (0x100 * Raster[ByteOffset + 1]); //
85     RawCode = 0x4004
86
87     if (CodeSize >= 8) //
88     CodeSize = 0x0d (Incremented during caller loop)
89     RawCode += (0x10000 * Raster[ByteOffset + 2]); //
90     RawCode = 0x474004
91
92     RawCode >>= (BitOffset % 8); //
93     RawCode = 0x11d001
94     BitOffset += (int)CodeSize; //
95     BitOffset = 0x27
96     return (RawCode & ReadMask); //
97     ReadMask = 0x1fff, (RawCode & ReadMask) = 0x1001 [This is our EOF Code]
98 }

```

Once the EOFCode is encountered, the heap buffers are freed.

```

323     free(Prefix);
324     free(Suffix);
325     free(OutCode);

```

Thus triggering the crash

```

double free or corruption (!prev)

Program received signal SIGABRT, Aborted.

```

TIMELINE

2022-03-17 - Initial Vendor Contact

2022-03-21 - Vendor Disclosure

2022-08-16 - Public Release

CREDIT

Discovered by Dave McDaniel of Cisco Talos.

VULNERABILITY REPORTS

PREVIOUS REPORT

NEXT REPORT

TALOS-2022-1487

TALOS-2022-1486
