

New issue

[Jump to bottom](#)

# AddressSanitizer: heap-buffer-overflow in function pdf\_write\_names #480

🔒 Closed

hdthky opened this issue on Mar 24 · 7 comments

Assignees



Labels

unable-to-reproduce

hdthky commented on Mar 24

## Description

Whilst experimenting with `htmldoc`, built from commit [31f7804](#), we are able to induce a vulnerability in function `pdf_write_names`, using a harness compiled from `htmldoc/htmldoc.cxx`.

Because there is no bounds checking, a heap-based out-of-bound read will be triggered when the software encounters a malformed file, result in information disclosure or denial of service.

## Proof of Concept

The POC is: [poc\\_heap\\_overflow1](#)

The command is: `./htmldoc --webpage -t pdf -f /dev/null poc_heap_overflow1`

The ASAN report is:

```
=====
==50540==ERROR: AddressSanitizer: heap-buffer-overflow on address 0x625000020f4 at pc
0x0000003cf241 bp 0x7fffffffaf90 sp 0x7fffffffaf88
READ of size 4 at 0x625000020f4 thread T0
    #0 0x3cf240 in pdf_write_names(_IO_FILE*) /work/libraries/htmldoc/htmldoc/ps-pdf.cxx:3589:39
    #1 0x3cf240 in pdf_write_document(unsigned char*, unsigned char*, unsigned char*, unsigned
char*, unsigned char*, unsigned char*, tree_str*, tree_str*) /work/libraries/htmldoc/htmldoc/ps-
pdf.cxx:2301:5
    #2 0x3cf240 in pspdf_export /work/libraries/htmldoc/htmldoc/ps-pdf.cxx:910:7
```

```
#3 0x39a254 in main /work/libraries/htmldoc/htmldoc/htmldoc.cxx:1291:3
#4 0x7ffff75070b2 in __libc_start_main (/lib/x86_64-linux-gnu/libc.so.6+0x270b2)
#5 0x2a051d in _start (/work/libraries/htmldoc/htmldoc/htmldoc+0x2a051d)
```

0x6250000020f4 is located 20 bytes to the right of 8160-byte region  
[0x625000000100,0x6250000020e0)

allocated by thread T0 here:

```
#0 0x31b6f9 in realloc (/work/libraries/htmldoc/htmldoc/htmldoc+0x31b6f9)
#1 0x3ddaf8 in check_pages(int) /work/libraries/htmldoc/htmldoc/ps-pdf.cxx:8859:24
```

SUMMARY: AddressSanitizer: heap-buffer-overflow /work/libraries/htmldoc/htmldoc/ps-pdf.cxx:3589:39  
in pdf\_write\_names(\_IO\_FILE\*)

Shadow bytes around the buggy address:

```
0x0c4a7fff83c0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0x0c4a7fff83d0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0x0c4a7fff83e0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0x0c4a7fff83f0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0x0c4a7fff8400: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
=>0x0c4a7fff8410: 00 00 00 00 00 00 00 00 00 00 00 00 fa fa[fa]fa
0x0c4a7fff8420: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
0x0c4a7fff8430: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
0x0c4a7fff8440: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
0x0c4a7fff8450: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
0x0c4a7fff8460: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
```

Shadow byte legend (one shadow byte represents 8 application bytes):

```
Addressable:           00
Partially addressable: 01 02 03 04 05 06 07
Heap left redzone:      fa
Freed heap region:      fd
Stack left redzone:     f1
Stack mid redzone:      f2
Stack right redzone:    f3
Stack after return:     f5
Stack use after scope:  f8
Global redzone:         f9
Global init order:      f6
Poisoned by user:       f7
Container overflow:     fc
Array cookie:           ac
Intra object redzone:   bb
ASan internal:          fe
Left alloca redzone:    ca
Right alloca redzone:   cb
Shadow gap:             cc
```

==50540==ABORTING

## Impact

This vulnerability is capable of inducing information disclosure or denial of service.

**michaelsweet** commented on Mar 24

Owner

OK, so I am unable to reproduce when "leak\_check\_at\_exit=false" is set. When not set I get completely different results.

 **michaelsweet** self-assigned this on Mar 24

 **michaelsweet** added the `unable-to-reproduce` label on Mar 24

**hdthky** commented on Mar 24

Author

It is a little weird. Even if I turn off `leak_check_at_exit`, I can still reproduce it.

The compile command I used is:

```
CFLAGS="-g -fsanitize=address" CXXFLAGS="-g -fsanitize=address" LDFLAGS="-fsanitize=address"
./configure
```

**hdthky** commented on Mar 24

Author

The causes of buffer overflow and memory leak are totally different. In theory, turning `leak_check_at_exit` off or not doesn't affect the result.

**michaelsweet** commented on Mar 24

Owner

@**hdthky** I am well aware of the differences. Theory doesn't matter, actual run evidence shows a difference on three different systems at my immediate disposal (iMac Pro running current macOS, Ubuntu VM on that system, and Ubuntu VM on an M1 MacBook Pro). No errors on macOS, different results on both Ubuntu VMs.

 **michaelsweet** added a commit that referenced this issue on Mar 24

 Call `check_pages` when writing links (Issue [#480](#))

✗ 46c8ec2

**michaelsweet** commented on Mar 24

Owner

@**hdthky** Assuming the cause is similar to the others, the following change should fix things for you:

[master [46c8ec2](#)] Call `check_pages` when writing links (Issue [#480](#))

**hdthky** commented on Mar 24

Author

Yes, it has been fixed now.



**michaelrsweet** closed this as completed on Mar 24

**hdthky** commented on May 19

Author

The vulnerability was found by Xingyuan Mo, Hui Lu, Zhihong Tian from Guangzhou University.

#### Assignees



**michaelrsweet**

#### Labels

unable-to-reproduce

#### Projects

None yet

#### Milestone

No milestone

#### Development

No branches or pull requests

#### 2 participants

