

f3b9bf4c3c ▾

...


tensorflow / tensorflow / core / kernels / stage_op.cc



quintinwang5 add DEVICE_DEFAULT for debug/stage ops ✖

History

10 contributors



323 lines (257 sloc) | 10.2 KB

...

```

1  /* Copyright 2017 The TensorFlow Authors. All Rights Reserved.
2
3  Licensed under the Apache License, Version 2.0 (the "License");
4  you may not use this file except in compliance with the License.
5  You may obtain a copy of the License at
6
7      http://www.apache.org/licenses/LICENSE-2.0
8
9  Unless required by applicable law or agreed to in writing, software
10 distributed under the License is distributed on an "AS IS" BASIS,
11 WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
12 See the License for the specific language governing permissions and
13 limitations under the License.
14 =====*/
15
16 #include <cstdint>
17 #include <deque>
18 #include <mutex>
19 #include <numeric>
20 #include <vector>
21
22 #include "tensorflow/core/framework/op_kernel.h"
23 #include "tensorflow/core/framework/resource_mgr.h"
24 #include "tensorflow/core/framework/tensor.h"
25 #include "tensorflow/core/framework/tensor_shape.h"
26 #include "tensorflow/core/lib/strings/strcat.h"
27 #include "tensorflow/core/platform/env.h"
28 #include "tensorflow/core/platform/mutex.h"
29

```

...

```

30 namespace tensorflow {
31 namespace {
32
33 class Buffer : public ResourceBase {
34 public:
35     using Tuple = std::vector<Tensor>;
36
37     explicit Buffer(std::size_t capacity, std::size_t memory_limit)
38         : capacity_(capacity), memory_limit_(memory_limit), current_bytes_(0) {}
39
40     // the Buffer takes ownership of the Tuple
41     Status Put(Tuple* tuple) {
42         std::unique_lock<std::mutex> lock(mu_);
43
44         std::size_t tuple_bytes = GetTupleBytes(*tuple);
45
46         // Sanity check so that we don't block for ever below
47         if (memory_limit_ > 0 && tuple_bytes > memory_limit_) {
48             return Status(
49                 errors::ResourceExhausted("Attempted to insert "
50                                         "tensors with combined size of '",
51                                         tuple_bytes,
52                                         "' bytes into "
53                                         "Staging Area with a memory limit of '",
54                                         memory_limit_, "'."));
55         }
56
57         // If buffer capacity is bounded wait until elements have been removed
58         if (IsBounded()) {
59             full_cond_var_.wait(lock, [tuple_bytes, this]() {
60                 // If there's a memory limit, check if there's space for insertion
61                 bool memory_limit_valid =
62                     memory_limit_ > 0 ? !WouldExceedMemoryLimit(tuple_bytes) : true;
63                 // If we're configured for capacity check if there's space for insertion
64                 bool capacity_valid = capacity_ > 0 ? !IsCapacityFull() : true;
65
66                 // Stop waiting upon success for both conditions
67                 return capacity_valid && memory_limit_valid;
68             });
69         }
70
71         // Update bytes in the Staging Area
72         current_bytes_ += tuple_bytes;
73
74         // Store tuple
75         buf_.push_back(std::move(*tuple));
76
77         lock.unlock();
78         // Notify all removers. Removers

```

```

79     // may be peeking at a specific element or waiting
80     // for the element at the front of the deque.
81     // As we don't know the appropriate one to wake up
82     // we should wake them all.
83     non_empty_cond_var_.notify_all();
84
85     return Status::OK();
86 }
87
88 // Get tuple at front of the buffer
89 void Get(Tuple* tuple) { // TODO(zhifengc): Support cancellation.
90     std::unique_lock<std::mutex> lock(mu_);
91
92     // Wait for data if the buffer is empty
93     non_empty_cond_var_.wait(lock, [this]() { return !buf_.empty(); });
94
95     // Move data into the output tuple
96     *tuple = std::move(buf_.front());
97     buf_.pop_front();
98
99     // Update bytes in the Staging Area
100    current_bytes_ -= GetTupleBytes(*tuple);
101
102    notify_inserters_if_bounded(&lock);
103 }
104
105 // Return tuple at index
106 Status Peek(std::size_t index, Tuple* tuple) {
107     std::unique_lock<std::mutex> lock(mu_);
108
109     // Wait if the requested index is not available
110     non_empty_cond_var_.wait(
111         lock, [index, this]() { return index < this->buf_.size(); });
112
113     // Place tensors in the output tuple
114     for (const auto& tensor : buf_[index]) {
115         tuple->push_back(tensor);
116     }
117
118     return Status::OK();
119 }
120
121 // Buffer size
122 size_t Size() {
123     std::unique_lock<std::mutex> lock(mu_);
124     return buf_.size();
125 }
126
127 void Clear() {

```

```

128     std::unique_lock<std::mutex> lock(mu_);
129     buf_.clear();
130     current_bytes_ = 0;
131
132     notify_inserters_if_bounded(&lock);
133 }
134
135 string DebugString() const override {
136     std::unique_lock<std::mutex> lock(mu_);
137     return strings::StrCat("Staging size: ", buf_.size());
138 }
139
140 private:
141     // If the buffer is configured for bounded capacity, notify
142     // waiting inserters that space is now available
143     void notify_inserters_if_bounded(std::unique_lock<std::mutex>* lock) {
144         if (IsBounded()) {
145             lock->unlock();
146             // Notify all inserters. The removal of an element
147             // may make memory available for many inserters
148             // to insert new elements
149             full_cond_var_.notify_all();
150         }
151     }
152
153     // Are there a limit number of elements or a memory limit
154     // configured on this buffer?
155     bool IsBounded() const { return capacity_ > 0 || memory_limit_ > 0; }
156
157     bool IsCapacityFull() const { return buf_.size() >= capacity_; }
158
159     bool WouldExceedMemoryLimit(std::size_t bytes) const {
160         return bytes + current_bytes_ > memory_limit_;
161     }
162
163     std::size_t GetTupleBytes(const Tuple& tuple) {
164         return std::accumulate(tuple.begin(), tuple.end(), 0,
165             [](const std::size_t& lhs, const Tensor& rhs) {
166                 return lhs + rhs.TotalBytes();
167             });
168     }
169
170     std::size_t capacity_;
171     std::size_t memory_limit_;
172     std::size_t current_bytes_;
173     mutable std::mutex mu_;
174     std::condition_variable non_empty_cond_var_;
175     std::condition_variable full_cond_var_;
176     std::deque<Tuple> buf_;

```

```

177 };
178
179 Status GetBuffer(OpKernelContext* ctx, const NodeDef& ndef, Buffer** buf) {
180     auto rm = ctx->resource_manager();
181     ContainerInfo cinfo;
182
183     // Lambda for creating the Staging Area
184     auto create_fn = [&ndef](Buffer** ret) -> Status {
185         int64_t capacity;
186         int64_t memory_limit;
187         TF_RETURN_IF_ERROR(GetNodeAttr(ndef, "capacity", &capacity));
188         TF_RETURN_IF_ERROR(GetNodeAttr(ndef, "memory_limit", &memory_limit));
189         *ret = new Buffer(capacity, memory_limit);
190         return Status::OK();
191     };
192
193     TF_RETURN_IF_ERROR(cinfo.Init(rm, ndef, true /* use name() */));
194     TF_RETURN_IF_ERROR(rm->LookupOrCreate<Buffer>(cinfo.container(), cinfo.name(),
195                                                  buf, create_fn));
196     return Status::OK();
197 }
198
199 } // namespace
200
201 class StageOp : public OpKernel {
202 public:
203     explicit StageOp(OpKernelConstruction* ctx) : OpKernel(ctx) {}
204
205     void Compute(OpKernelContext* ctx) override {
206         Buffer* buf = nullptr;
207         OP_REQUIRES_OK(ctx, GetBuffer(ctx, def(), &buf));
208         core::ScopedUnref scope(buf);
209         Buffer::Tuple tuple;
210         tuple.reserve(ctx->num_inputs());
211         for (int i = 0; i < ctx->num_inputs(); ++i) {
212             tuple.push_back(ctx->input(i));
213         }
214         OP_REQUIRES_OK(ctx, buf->Put(&tuple));
215     }
216 };
217
218 REGISTER_KERNEL_BUILDER(Name("Stage").Device(DEVICE_CPU), StageOp);
219 REGISTER_KERNEL_BUILDER(Name("Stage").Device(DEVICE_DEFAULT), StageOp);
220
221 class UnstageOp : public OpKernel {
222 public:
223     explicit UnstageOp(OpKernelConstruction* ctx) : OpKernel(ctx) {}
224
225     // Using this op in such a way that it blocks forever

```

```

226 // is an error. As such cancellation is not handled.
227 void Compute(OpKernelContext* ctx) override {
228     Buffer* buf = nullptr;
229     OP_REQUIRES_OK(ctx, GetBuffer(ctx, def(), &buf));
230     core::ScopedUnref scope(buf);
231     Buffer::Tuple tuple;
232
233     buf->Get(&tuple);
234
235     OP_REQUIRES(
236         ctx, tuple.size() == (size_t)ctx->num_outputs(),
237         errors::InvalidArgument("Mismatch stage/unstage: ", tuple.size(),
238                                 " vs. ", ctx->num_outputs()));
239
240     for (size_t i = 0; i < tuple.size(); ++i) {
241         ctx->set_output(i, tuple[i]);
242     }
243 }
244 };
245
246 REGISTER_KERNEL_BUILDER(Name("Unstage").Device(DEVICE_CPU), UnstageOp);
247 REGISTER_KERNEL_BUILDER(Name("Unstage").Device(DEVICE_DEFAULT), UnstageOp);
248
249 class StagePeekOp : public OpKernel {
250 public:
251     explicit StagePeekOp(OpKernelConstruction* ctx) : OpKernel(ctx) {}
252
253     // Using this op in such a way that it blocks forever
254     // is an error. As such cancellation is not handled.
255     void Compute(OpKernelContext* ctx) override {
256         Buffer* buf = nullptr;
257         OP_REQUIRES_OK(ctx, GetBuffer(ctx, def(), &buf));
258         core::ScopedUnref scope(buf);
259         Buffer::Tuple tuple;
260
261         std::size_t index = ctx->input(0).scalar<int>();
262
263         OP_REQUIRES_OK(ctx, buf->Peek(index, &tuple));
264
265         OP_REQUIRES(
266             ctx, tuple.size() == (size_t)ctx->num_outputs(),
267             errors::InvalidArgument("Mismatch stage/unstage: ", tuple.size(),
268                                     " vs. ", ctx->num_outputs()));
269
270         for (size_t i = 0; i < tuple.size(); ++i) {
271             ctx->set_output(i, tuple[i]);
272         }
273     }
274 };

```

```

275
276 REGISTER_KERNEL_BUILDER(Name("StagePeek").Device(DEVICE_CPU), StagePeekOp);
277 REGISTER_KERNEL_BUILDER(
278     Name("StagePeek").HostMemory("index").Device(DEVICE_DEFAULT), StagePeekOp);
279
280 class StageSizeOp : public OpKernel {
281 public:
282     explicit StageSizeOp(OpKernelConstruction* ctx) : OpKernel(ctx) {}
283
284     // Using this op in such a way that it blocks forever
285     // is an error. As such cancellation is not handled.
286     void Compute(OpKernelContext* ctx) override {
287         Buffer* buf = nullptr;
288         OP_REQUIRES_OK(ctx, GetBuffer(ctx, def(), &buf));
289         core::ScopedUnref scope(buf);
290
291         // Allocate size output tensor
292         Tensor* size = nullptr;
293         OP_REQUIRES_OK(ctx, ctx->allocate_output(0, TensorShape({}), &size));
294
295         // Set it to the actual size
296         size->scalar<int32>().setConstant(buf->Size());
297     }
298 };
299
300 REGISTER_KERNEL_BUILDER(Name("StageSize").Device(DEVICE_CPU), StageSizeOp);
301 REGISTER_KERNEL_BUILDER(
302     Name("StageSize").HostMemory("size").Device(DEVICE_DEFAULT), StageSizeOp);
303
304 class StageClearOp : public OpKernel {
305 public:
306     explicit StageClearOp(OpKernelConstruction* ctx) : OpKernel(ctx) {}
307
308     // Using this op in such a way that it blocks forever
309     // is an error. As such cancellation is not handled.
310     void Compute(OpKernelContext* ctx) override {
311         Buffer* buf = nullptr;
312         OP_REQUIRES_OK(ctx, GetBuffer(ctx, def(), &buf));
313         core::ScopedUnref scope(buf);
314
315         buf->Clear();
316     }
317 };
318
319 REGISTER_KERNEL_BUILDER(Name("StageClear").Device(DEVICE_CPU), StageClearOp);
320 REGISTER_KERNEL_BUILDER(Name("StageClear").Device(DEVICE_DEFAULT),
321     StageClearOp);
322
323 } // namespace tensorflow

```

